

# Homomorphic Commitment Schemes

Entry #:	55.86.5
Word Count:	12441 words
Reading Time:	62 minutes
Last Updated:	August 31, 2025

*"In space, no one can hear you think."*

## Table of Contents

### Contents

<b>1</b>	<b>Homomorphic Commitment Schemes</b>	<b>2</b>
1.1	Foundational Concepts & Definition . . . . .	2
1.2	Historical Development & Theoretical Roots . . . . .	3
1.3	Core Properties & Security Definitions . . . . .	5
1.4	Major Construction Paradigms . . . . .	7
1.5	Mathematical Foundations & Mechanisms . . . . .	9
1.6	Applications in Cryptography . . . . .	11
1.7	Applications Beyond Core Cryptography . . . . .	13
1.8	Implementation Challenges & Practical Considerations . . . . .	15
1.9	Security Analysis & Cryptanalytic Threats . . . . .	17
1.10	Comparisons & Trade-offs . . . . .	19
1.11	Societal Impact, Ethics & Controversies . . . . .	21
1.12	Future Directions & Open Problems . . . . .	23

# 1 Homomorphic Commitment Schemes

## 1.1 Foundational Concepts & Definition

At the heart of secure digital interaction lies a fundamental paradox: how can one party prove they possess specific information to another, without revealing the information itself? This cryptographic conundrum finds its elegant solution in the concept of a **commitment scheme**. Imagine placing a secret message into a tamper-evident, opaque envelope and handing it over. The recipient knows you've committed to *something* (the envelope is sealed), but cannot see what it is (the envelope is opaque - **hiding**). Later, when you reveal the message and the unique seal mechanism, the recipient can verify it matches the original sealed envelope, proving you couldn't have switched messages after the fact (**binding**). This digital "sealed envelope" is the bedrock upon which countless protocols are built, from simple coin flips over the phone to complex blockchain transactions. Common constructions include the remarkably simple yet powerful **Pedersen commitment**, leveraging the discrete logarithm problem within a group; **hash-based commitments** like using SHA-256 ( $c = H(\text{message} \parallel \text{random\_salt})$ ), prized for their simplicity and reliance solely on hash function security; and **RSA-based commitments**, utilizing the difficulty of factoring large integers.

However, the basic commitment envelope, while secure, is static. What if the secret inside isn't a single, fixed value, but part of a larger data set requiring computation? What if proving a complex relationship *about* committed values is needed, without opening any envelopes? This is where the transformative power of **homomorphism** enters the cryptographic stage, elevating the humble commitment into a dynamic and powerful tool. Homomorphism, borrowed from algebra (where it describes structure-preserving maps between groups, rings, or fields), in this context means that specific mathematical operations performed *directly on the commitment values* (the sealed envelopes) yield results that are *valid commitments* of the outcome of those same operations performed *on the underlying secret messages*. Symbolically, for an operation  $\square$  and a function  $f$ , if  $\text{Commit}(a) \square \text{Commit}(b)$  produces  $\text{Commit}(f(a, b))$ , the scheme is homomorphic over that operation. Consider the Pedersen commitment:  $\text{Commit}(a) = g^a * h^{r_a}$  and  $\text{Commit}(b) = g^b * h^{r_b}$  (where  $g, h$  are group generators and  $r$  is randomness). Multiplying these commitments gives  $g^{a+b} * h^{r_a + r_b}$ , which is precisely a valid Pedersen commitment to  $(a+b)$  with randomness  $(r_a + r_b)$ . This simple additive homomorphism is revolutionary; it allows computations on *sealed* data. Contrast this starkly with a standard hash-based commitment:  $H(a \parallel r_a) * H(b \parallel r_b)$  bears no computable relationship to  $H(a+b \parallel \text{something})$ , rendering such computations impossible without revealing the secrets. The algebraic structure enabling homomorphism introduces complexity, demanding careful analysis of security properties under this new computational capability, but the payoff is immense.

Why endure this complexity? The driving need stems from fundamental cryptographic ambitions that basic commitments cannot fulfill alone. The core promise of homomorphic commitment schemes (HCS) is the unprecedented ability to enforce **verifiability** of complex computations or relationships concerning sensitive data while simultaneously guaranteeing **privacy** – the raw data remains concealed. This dual capability

unlocks entire classes of advanced protocols. In **Zero-Knowledge Proofs (ZKPs)**, HCS allow a prover to demonstrate they know secrets satisfying intricate conditions (e.g., “I know  $x$  such that  $f(x) = y$  AND  $x$  is within a valid range”) by performing computations solely on commitments and generating proofs about those computations, never revealing  $x$ . The additive homomorphism of Pedersen commitments is foundational to numerous ZKP protocols like Schnorr signatures and more complex Sigma protocols. **Secure Multi-Party Computation (MPC)** benefits profoundly; multiple parties can commit to their private inputs using an HCS, then collaboratively compute a public function  $f(\text{input1}, \text{input2}, \dots, \text{inputN})$  over the *commitments*, deriving a commitment to the result. Individual inputs remain private, yet the correctness of the computation relative to the committed inputs is verifiable. Similarly, **Verifiable Computation/Delegation** leverages HCS: a client commits to its input data, sends it (and perhaps a committed program) to a powerful but untrusted server; the server computes the result, provides an output, *and* a proof demonstrating that the output is consistent with the computation applied to the committed input, all without the client needing to redo the computation or reveal the input. Within the broader cryptographic landscape, HCS occupy a crucial middle ground. They are generally far more efficient than **Fully Homomorphic Encryption (FHE)**, which allows arbitrary computations on encrypted data but often with impractical overhead. Conversely, HCS provide the essential “sealing” and verifiable computation capabilities that are fundamental building blocks *within* powerful **ZK-SNARKs and ZK-STARKs** proof systems (like the polynomial commitments pioneered by Kate, Zaverucha, and Goldberg), enabling them to handle complex statements succinctly. The journey from the simple sealed envelope to this algebraically empowered primitive represents a pivotal evolution in cryptography, setting the stage for protocols that achieve unprecedented levels of privacy coupled with ironclad verifiability. This foundational interplay between commitment, homomorphism, and verifiable secrecy forms the essential context for exploring the rich history and intricate mechanics of these schemes that follows.

## 1.2 Historical Development & Theoretical Roots

The transformative power of homomorphic commitment schemes, as established in their foundational interplay of secrecy and verifiable computation, did not emerge in a vacuum. Their evolution is deeply entwined with centuries of mathematical exploration and decades of cryptographic innovation, gradually transforming abstract algebraic concepts into practical tools for digital security. Understanding this historical trajectory reveals not only the ingenuity behind HCS but also the persistent quest to reconcile profound privacy with robust verifiability.

**The indispensable groundwork was laid by fundamental branches of mathematics.** Group theory, particularly the study of finite cyclic groups, provided the structural bedrock. The inherent difficulty of the **Discrete Logarithm Problem (DLP)** – determining the exponent  $x$  given  $g^x$  in a large multiplicative group – became a cornerstone for security. Early commitment schemes like **Pedersen commitments**, as mentioned previously, directly leveraged this asymmetry: the binding property relies on the DLP’s hardness, while hiding stems from the randomness introduced via a second generator  $h^x$ . **Finite fields**, providing well-defined arithmetic modulo a prime, offered the essential environment for these groups to operate securely. The later

advent of **elliptic curve cryptography** (ECC) in the 1980s (building on work by Koblitz and Miller) offered more efficient realizations of the same DLP hardness within smaller groups, significantly impacting commitment scheme efficiency. Simultaneously, **integer factorization**, particularly the challenge of factoring large numbers like RSA moduli ( $n = p \cdot q$ ), offered an alternative hardness foundation explored in commitments, though often yielding larger cryptographic objects. **Information-theoretic security**, pioneered by Claude Shannon, also played a crucial conceptual role, influencing definitions of perfect hiding or binding, where security holds even against computationally unbounded adversaries. These mathematical structures weren't developed for cryptography *per se*, but their inherent asymmetries and computational hardness became the raw materials cryptographers would sculpt.

**The 1980s witnessed the formalization and early exploration of commitment schemes themselves, setting the stage for homomorphism.** While the concept of a cryptographic commitment was intuitively understood earlier, rigorous definitions and constructions solidified in this era. Manuel Blum's 1981 coin-flipping protocol over the phone arguably provided one of the first concrete uses, implicitly relying on a commitment to a bit. However, it was the work of Torben Pryds Pedersen in 1991 that delivered the first explicitly recognized homomorphic commitment scheme. His elegantly simple construction within a prime-order group offered **additive homomorphism** – the multiplication of commitments corresponding to the addition of messages – as a natural consequence of the group structure. This was not merely a theoretical curiosity; Pedersen commitments became a fundamental workhorse in subsequent protocols. Crucially, David Chaum and Pedersen, in their 1992 paper on “Wallet Databases with Observers,” formalized the link between commitments and **zero-knowledge proofs (ZKPs)**, demonstrating how commitment schemes could be used within interactive proofs to show relations about committed values without revealing them. This established a vital application domain that would heavily drive HCS development. Furthermore, Gilles Brassard and Claude Crépeau's influential 1986 paper “Non-Transitive Transfer of Confidence” provided a rigorous formal framework for commitment schemes in the context of multi-party computation and ZKPs, analyzing their security properties and interactions. Concurrently, albeit less directly applicable immediately, Miklós Ajtai's groundbreaking 1996 work established the connection between worst-case and average-case hardness for lattice problems (like Shortest Vector Problem - SVP), laying the theoretical foundation for future **lattice-based cryptography** – a paradigm that would later offer post-quantum secure commitment schemes with inherent homomorphic properties. While these early schemes, particularly Pedersen's, demonstrated the power of homomorphism, they were largely confined to additive properties or specific, limited relations within interactive proofs. The quest for richer homomorphism and more efficient non-interactive proofs demanded further formalization.

**The 2000s ushered in a period of rigorous formalization and the breakthrough development of highly efficient schemes, particularly driven by the needs of non-interactive zero-knowledge proofs.** A critical step was the precise definition of **security models** for commitment schemes, distinguishing between **computational hiding/binding** (security holds only against efficient, polynomial-time adversaries) and **statistical hiding/binding** (security holds against computationally unbounded adversaries, except with negligible probability). Understanding the trade-offs – for instance, that Pedersen commitments offer computational binding but can achieve statistical hiding with appropriate parameter choices – became essential for se-

lecting schemes tailored to specific threat models. The most transformative development, however, came from **pairing-based cryptography**. The introduction of efficient bilinear pairings (functions  $e: G_1 \times G_2 \rightarrow G_T$  mapping group elements to a target group with specific algebraic properties) in the early 2000s, exemplified by Dan Boneh, Ben Lynn, and Hovav Shacham's work on short signatures (BLS signatures, 2001), opened a treasure trove of cryptographic possibilities. Pairings enabled entirely new algebraic structures where more complex computations on commitments became feasible and verifiable succinctly. This culminated in the landmark 2010 paper by Aniket Kate, Gregory Zaverucha, and Ian Goldberg introducing what became known as **KZG polynomial commitments**. KZG leveraged pairings to achieve constant-sized commitments to polynomials and constant-sized evaluation proofs. This **succinctness** was revolutionary: a prover could commit to a large polynomial (representing complex computation or state) and later prove, with a small, fixed-size proof, that the polynomial evaluates to a specific value at a given point, or that it satisfies certain divisibility conditions. This efficiency made KZG the engine powering many practical **ZK-SNARKs** (Zero-Knowledge Succinct Non-interactive Arguments of Knowledge), enabling complex private computations on blockchain platforms like Zcash and Ethereum scaling solutions. While lattice-based schemes advanced steadily, offering promising **post-quantum security** and often **transparent setups** (no trusted initialization), pairing-based constructions like KZG dominated the efficiency landscape for polynomial commitments throughout the decade, despite their significant drawback: the requirement for a **trusted setup** to generate a Structured Reference String (S

### 1.3 Core Properties & Security Definitions

The pivotal breakthroughs in pairing-based constructions like KZG polynomial commitments, while dramatically enhancing efficiency and succinctness, underscored a critical truth: the power and security of any homomorphic commitment scheme (HCS) fundamentally rests upon rigorously defined core properties. These properties govern both its basic functionality as a cryptographic commitment and its enhanced capabilities derived from homomorphism. Understanding these definitions is paramount, as they form the yardstick against which all schemes, whether rooted in discrete logs, pairings, lattices, or other foundations, must be measured. The trusted setup requirement of KZG, a potential vulnerability point highlighted at the end of our historical exploration, is just one facet of a broader security landscape defined by these essential characteristics.

**The security of any commitment scheme, homomorphic or not, hinges on a fundamental duality: the properties of Hiding and Binding.** These are not mere features but the bedrock guarantees. *Hiding* ensures the secrecy of the committed message. Formally, it requires that an adversary, given a commitment  $c$ , cannot distinguish whether  $c$  corresponds to any one of two distinct messages  $m_0$  or  $m_1$  chosen by the adversary themselves. The strength of hiding varies: **Perfect hiding** guarantees indistinguishability even against an adversary with unlimited computational power (information-theoretic security), **statistical hiding** means the advantage of any computationally unbounded adversary is negligible (vanishingly small as security parameters increase), and **computational hiding** restricts the guarantee to adversaries bounded by realistic computational limits, relying on the hardness of problems like DLP or LWE. Pedersen commitments

exemplify this nuance; with generators  $g$  and  $h$  where the discrete log of  $h$  base  $g$  is unknown, they achieve *statistical hiding* – the random blinding factor  $r$  perfectly masks  $m$  in the exponent. Conversely, *Binding* ensures the commitment’s integrity. It mandates that it is computationally infeasible for the committer to find two distinct message-opening pairs  $(m, r)$  and  $(m', r')$  that both validate for the same commitment  $c$  (i.e.,  $\text{Open}(c, m, r) = \text{true}$  and  $\text{Open}(c, m', r') = \text{true}$  with  $m \neq m'$ ). Similarly, binding can be perfect, statistical, or computational. Pedersen commitments offer *computational binding* under the Discrete Logarithm assumption – finding two openings would imply solving the DLP for  $h$  base  $g$ . A profound tension exists between these twin pillars. Achieving both perfect hiding and perfect binding simultaneously within a single scheme is generally impossible, a consequence of fundamental information theory. Schemes typically optimize one property statistically or perfectly while relying on computational hardness for the other. For example, a simple hash commitment  $c = H(m)$  offers computational hiding (assuming  $H$  is a random oracle or collision-resistant) and computational binding (via collision resistance), but lacks homomorphism entirely. The introduction of homomorphism adds layers to analyzing this duality, as operations on commitments must not inadvertently weaken hiding or create new avenues for violating binding.

**Homomorphism, the defining superpower of HCS, manifests in distinct algebraic flavors, each enabling different classes of verifiable computation.** The simplest and most common is **Additive Homomorphism**. This allows adding commitments to obtain a valid commitment of the sum of the underlying messages:  $\text{Comm}(a; r_a) + \text{Comm}(b; r_b) = \text{Comm}(a + b; r_a + r_b)$ . Pedersen commitments are the archetype:  $g^a * h^{r_a} * g^b * h^{r_b} = g^{a+b} * h^{r_a+r_b}$ . This property underpins countless protocols, from proving the sum of secret values in ZKPs to confidential balance transfers in cryptocurrencies like Mimblewimble, where additively homomorphic commitments allow verifying that transaction inputs equal outputs without revealing individual amounts. **Multiplicative Homomorphism** is rarer and more powerful, enabling the multiplication of commitments to yield a commitment to the product:  $\text{Comm}(a; r_a) * \text{Comm}(b; r_b) = \text{Comm}(a * b; r_{ab})$  (where  $r_{ab}$  is a new randomness term, often derived from  $r_a$ ,  $r_b$ , and potentially the messages). Basic RSA commitments provide a canonical example:  $\text{Comm}(m) = g^m * r^N \bmod N$  (where  $N$  is an RSA modulus). Multiplying two such commitments gives  $(g^{m_1} * r_1^N) * (g^{m_2} * r_2^N) = g^{m_1+m_2} * (r_1*r_2)^N \bmod N$ , which is *additively* homomorphic over the exponent  $m$ . True multiplicative homomorphism in the message itself often requires richer algebraic structures like groups supporting bilinear pairings. Schemes exhibiting both additive *and* multiplicative homomorphism over the same committed values come closest to mimicking Fully Homomorphic Encryption (FHE), albeit often with significant restrictions or overhead. The most expressive and impactful form for modern applications, however, is **Polynomial (or Functional) Homomorphism**. This allows evaluating an arbitrary polynomial  $f$  on committed values  $a_1, a_2, \dots, a_n$  by performing specific operations on their commitments, resulting in a commitment to  $f(a_1, a_2, \dots, a_n)$ . Crucially, schemes like KZG achieve this with remarkable succinctness. One commits not just to individual values, but to a polynomial  $\phi(x)$  representing the data or computation. The homomorphic property then enables generating a very small proof that  $\phi(z) = y$  for some point  $z$ , or that  $\phi(x)$  satisfies a certain algebraic relation (e.g.,  $\phi(x) * q(x) = t(x)$  for



some public polynomials  $q$  and  $t$ ), all while keeping  $\varphi(x)$  hidden. This polynomial homomorphism is the engine driving efficient ZK-SNARKs, enabling complex statement proofs with tiny verification footprints. The power hierarchy is clear: additive homomorphism supports linear operations, multiplicative extends to quadratic and some higher-degree forms with limitations, while polynomial homomorphism unlocks near-arbitrary computation verification, albeit often requiring more complex setup or cryptographic machinery.

**Beyond the core duality and homomorphism types, advanced security properties address the practical demands of modern cryptographic systems, particularly concerning efficiency and proof composition.**

**Succinctness** has become paramount. It refers to the size of the commitment and any associated proofs (like evaluation proofs in polynomial schemes) relative to the size of the committed data or the complexity of the verified computation. A scheme is succinct if these sizes are sublinear (ideally constant or logarithmic) in the size of the committed message/function. KZG set a high bar: the commitment to a polynomial  $\varphi(x)$  of degree  $d$  is a single group element (e.g., 48 bytes for BLS12-381), and the proof that  $\varphi(z) = y$  is also a single group element, regardless of  $d$ . This contrasts sharply with, say, committing to each coefficient individually using Pedersen, yielding a commitment size linear in  $d$ . Succinctness is essential for scalability in blockchains and efficient ZKPs. **Extractability** (also known as the “Knowledge of

## 1.4 Major Construction Paradigms

The intricate security definitions explored in Section 3 – the delicate balance of hiding and binding, the varying power of additive, multiplicative, and polynomial homomorphism, and the crucial importance of succinctness and extractability – are not abstract ideals. They are realized through concrete mathematical machinery. The diverse landscape of homomorphic commitment schemes (HCS) arises from distinct computational hardness assumptions and the algebraic structures they enable. Examining the major construction paradigms reveals how cryptographers leverage these foundations to build schemes with unique strengths and trade-offs, directly impacting their suitability for real-world applications.

**The Discrete Logarithm (DL) paradigm provides the most accessible entry point into homomorphic commitments, anchored by the enduring Pedersen scheme.** As established earlier, the Pedersen commitment  $\text{Comm}(m; r) = g^m \cdot h^r$  in a cyclic group  $G$  of prime order  $q$  (where the discrete logarithm of  $h$  base  $g$  is unknown) exemplifies elegant simplicity. Its security rests squarely on the Computational Diffie-Hellman (CDH) or Discrete Logarithm Problem (DLP) assumptions: binding requires that finding  $m \neq m'$

and  $r, r'$  such that  $g^m h^r = g^{m'} h^{r'}$  implies solving the DLP for  $h$  relative to  $g$ , while statistical hiding is guaranteed by the perfect blinding introduced by  $r$ . Its inherent **additive homomorphism** –  $\text{Comm}(a; r_a) \cdot \text{Comm}(b; r_b) = g^a h^{r_a} \cdot g^b h^{r_b} = g^{a+b} h^{r_a + r_b} = \text{Comm}(a+b; r_a + r_b)$  – emerges naturally from the multiplicative group structure. This property made it an instant classic and a foundational component in early zero-knowledge protocols like Schnorr identification and Chaum-Pedersen proofs. Practical implementations often leverage elliptic curve groups (e.g., secp256k1, Curve25519) for efficiency. Recognizing the need to commit to multiple values simultaneously, **vector Pedersen commitments** ( $\text{Comm}(\vec{m}; r) = g^r \cdot \prod_{i=1}^n h_i^{m_i}$ ) were developed, preserving additive homomorphism component-wise and finding extensive



use in confidential transaction protocols like Mimblewimble, where they allow the verification of balance (sum of inputs equals sum of outputs) without revealing individual amounts. While remarkably versatile and efficient for linear operations, DL-based schemes are fundamentally constrained. Their homomorphism is primarily additive; achieving multiplicative homomorphism within the same group is generally impossible without compromising security or efficiency. Furthermore, some advanced variants requiring specific generator relationships might necessitate a **trusted setup** to ensure the discrete log relationship between  $g$  and  $h$  remains unknown – a potential point of vulnerability, though techniques like public verifiable generation mitigate this. Despite these limitations, the simplicity, efficiency, and well-understood security of DL-based schemes like Pedersen ensure their continued, widespread deployment.

**Integer Factorization (RSA) based schemes leverage the hardness of factoring large integers  $n = p \cdot q$  or related problems like the RSA assumption.** A basic RSA commitment takes the form  $\text{Comm}(m; r) = g^m \cdot r^e \pmod n$ , where  $g$  is a public generator in  $\mathbb{Z}_n^*$ ,  $e$  is a public exponent (often a prime), and  $r$  is random blinding factor. Security relies on the difficulty of computing  $e$ -th roots modulo  $n$  (RSA assumption) for binding and the need for randomization to ensure hiding. The homomorphic properties here are interesting but require careful interpretation. Multiplying two commitments:  $\text{Comm}(a; r_a) \cdot \text{Comm}(b; r_b) = (g^a \cdot r_a^e) \cdot (g^b \cdot r_b^e) = g^{a+b} \cdot (r_a r_b)^e \pmod n = \text{Comm}(a+b; r_a \cdot r_b)$ . This demonstrates **additive homomorphism in the exponent**  $m$ , meaning the product of commitments corresponds to a commitment of the *sum* of the messages  $a+b$ . To achieve true **multiplicative homomorphism** in the message itself ( $\text{Comm}(a) \cdot \text{Comm}(b) = \text{Comm}(a \cdot b)$ ), more complex constructions are needed, often utilizing the Strong RSA assumption. These might involve commitments like  $\text{Comm}(m) = g^m h^r \pmod n$  (similar in form to Pedersen but modulo a composite) where multiplying commitments yields  $g^{m_1+m_2} h^{r_1+r_2}$ , still additive, or schemes designed specifically for multiplicative operations which can be less efficient. Fujisaki and Okamoto presented influential schemes in this vein. While RSA-based commitments can offer both additive and multiplicative capabilities under strong assumptions, they suffer from significant drawbacks. The size of RSA moduli (typically 3072+ bits for modern security) results in much larger commitments and associated proofs compared to elliptic curve DL-based schemes (256 bits). Operations modulo large composites are also computationally heavier than elliptic curve operations. Furthermore, the security landscape for factoring-based cryptography is considered less robust in the long term against quantum attacks compared to lattice problems, and the schemes themselves can be trickier to implement securely due to the nuances of composite modulus arithmetic. Consequently, while historically significant and offering a different homomorphic profile, RSA-based HCS see less practical adoption today than DL, pairing, or lattice alternatives for most applications demanding homomorphism.

**Pairing-Based Schemes represent a quantum leap in expressive power and succinctness, primarily enabling efficient polynomial homomorphism.** This paradigm leverages **bilinear pairings** – specialized functions  $e: G_1 \times G_2 \rightarrow G_T$  that map elements from two source groups (often cyclic groups of prime order) onto a target group, satisfying  $e(g^a, h^b) = e(g, h)^{a \cdot b}$  for generators  $g \in G_1, h \in G_2$ . This multiplicative property across groups is the magic ingredient. The seminal example is the **Kate-Zaverucha-Goldberg (KZG) polynomial commitment** (2010). Here, a polynomial  $\phi(x)$

$= c_0 + c_1 x + \dots + c_d x^d$  is committed using a trusted setup that generates a **Structured Reference String (SRS)** containing powers of a secret  $s$ :  $(g, g^s, g^{s^2}, \dots, g^{s^d})$  in  $G_1$ . The commitment is  $C = g^{\{\varphi(s)\}} = \prod_{i=0}^d (g^{s^i})^{c_i}$ , a single group element. To prove that  $\varphi(z) = y$ , the prover computes the quotient polynomial  $q(x) = (\varphi(x) - y) / (x - z)$  and provides the evaluation proof  $\pi = g^{\{q(s)\}}$ , another single group element. The verifier, using the SRS and pairing, checks  $e(C / g^y, h) = e(\pi, h^z / h)$  (or equivalent formulation depending on SRS structure). This provides **constant-sized commitments and proofs** for polynomial evaluation, regardless of the degree  $d$ . This **succinctness** is revolutionary. Furthermore, KZG exhibits powerful **polynomial homomorphism**: commitments to polynomials  $\varphi(x)$  and  $\psi(x)$  can be combined to get a commitment to  $\varphi(x) + \alpha \psi(x)$  for public scalar  $\alpha$ , and proofs can be constructed to show relations like  $\varphi(x) \cdot \psi(x) = \gamma(x)$  for some public  $\gamma(x)$ . This richness

## 1.5 Mathematical Foundations & Mechanisms

The exploration of major construction paradigms reveals a landscape of remarkable diversity: from the elegant simplicity of Pedersen commitments riding on the coattails of the discrete logarithm problem, to the pairing-powered succinctness of KZG unlocking polynomial homomorphism, and the lattice-based approaches promising post-quantum resilience. Yet, beneath this surface variety lies a bedrock of unifying mathematical principles. Understanding these deeper foundations – the algebraic structures enabling homomorphism, the critical role of randomness in preserving secrecy, and the precise algorithmic choreography of commitment and opening – is essential to grasp not just *how* these schemes work, but *why* they are secure and how their unique capabilities emerge. This section delves into the mathematical machinery powering homomorphic commitment schemes (HCS), illuminating the elegant interplay of algebra, number theory, and randomness that transforms abstract concepts into practical cryptographic tools.

**The remarkable properties of homomorphic commitment schemes – their ability to preserve structure while concealing secrets – are fundamentally enabled by carefully chosen algebraic structures.** Each major construction paradigm leverages a specific mathematical setting whose intrinsic properties dictate the scheme’s homomorphic capabilities and security foundations. *Finite cyclic groups* form the stage for discrete logarithm (DL)-based schemes like Pedersen commitments. Here, the multiplicative structure of the group (often an elliptic curve group) directly facilitates additive homomorphism: multiplying commitments  $(g^a * h^{r_a} * g^b * h^{r_b})$  seamlessly translates to adding the exponents  $(g^{a+b} * h^{r_a + r_b})$ . The hardness of the Discrete Logarithm Problem (DLP) within this group underpins the binding property. In contrast, *rings*, specifically the ring of integers modulo a large composite  $n = p \cdot q$  (RSA modulus), host many RSA-based commitments. The ring structure, combining addition and multiplication, allows for richer homomorphic interactions. A basic RSA commitment  $\text{Comm}(m; r) = g^m * r^e \bmod n$  exhibits additive homomorphism in the exponent ( $\text{Comm}(a) * \text{Comm}(b) = \text{Comm}(a+b)$ ) due to the multiplicative nature of the operation modulo  $n$ . Achieving true multiplicative homomorphism over the message itself often requires exploiting specific properties within this ring under stronger assumptions like Strong RSA. Pairing-based schemes, like KZG, operate within a more com-

plex ecosystem involving *finite fields* and carefully constructed *groups* ( $G_1, G_2, G_T$ ) linked by a bilinear map  $e$ . The power stems from the pairing’s ability to “multiply across groups”:  $e(g^a, h^b) = e(g, h)^{a \cdot b}$ . This property is the linchpin allowing KZG to commit to a polynomial  $\varphi(x)$  using a structured reference string ( $g^{\{s^i\}}$ ) and later generate a succinct proof  $g^{\{q(s)\}}$  that  $\varphi(z) = y$ , leveraging the homomorphism induced by the pairing to verify the polynomial relation efficiently. Finally, *modules* (generalizations of vector spaces over rings) provide the framework for lattice-based commitments, built on problems like Learning With Errors (LWE). Here, commitments often take the form  $\text{Comm}(\text{vec}\{m\}) = \mathbf{A} \cdot \text{vec}\{r\} + \text{vec}\{m\} \bmod q$ , where  $\mathbf{A}$  is a public matrix. The additive structure of the module naturally provides additive homomorphism:  $\text{Comm}(\text{vec}\{a\}) + \text{Comm}(\text{vec}\{b\}) = \mathbf{A}(\text{vec}\{r_a\} + \text{vec}\{r_b\}) + (\text{vec}\{a\} + \text{vec}\{b\}) = \text{Comm}(\text{vec}\{a\} + \text{vec}\{b\})$ . Security rests on the hardness of finding short vectors ( $\text{vec}\{r\}, \text{vec}\{m\}$ ) in high-dimensional lattices induced by  $\mathbf{A}$ . The choice of algebraic structure is thus not arbitrary; it directly determines the type of homomorphism achievable (additive, multiplicative, polynomial) and the computational hardness assumptions underpinning security.

**While the algebraic structure provides the framework for computation and homomorphism, it is the meticulous use of randomization that safeguards the fundamental promise of secrecy: the hiding property.** Imagine committing to the same confidential bid  $m$  in two separate auctions. Without randomization, the commitments would be identical ( $\text{Comm}(m)$ ), immediately revealing to an observer that the same value was committed twice, a catastrophic privacy breach. This is where the blinding factor  $r$  becomes indispensable. Virtually all secure commitment schemes, homomorphic or not, incorporate randomness into the commitment algorithm:  $\text{Comm}(m, r) \rightarrow c$ . In Pedersen, it’s  $r$  in  $g^m \cdot h^r$ ; in RSA-based schemes, it’s  $r$  in  $g^m \cdot r^e$ ; in lattice schemes, it’s the random vector  $\text{vec}\{r\}$ . This  $r$  acts as a cryptographic mask. Even for the *same* message  $m$ , choosing a fresh, random  $r$  for each commitment yields a radically different, unpredictable output  $c$ . This ensures the scheme satisfies its hiding guarantee: commitments to different messages are computationally (or statistically/perfectly) indistinguishable, and crucially, commitments to the *same* message are also indistinguishable from each other and from commitments to any other message. The quality of randomness is paramount; using a weak or predictable random number generator can fatally compromise hiding, as demonstrated in historical vulnerabilities. For homomorphic schemes, randomization plays an additional subtle role. During homomorphic operations (e.g., adding two Pedersen commitments), the randomness terms also combine ( $r_a + r_b$ ). The resulting commitment  $\text{Comm}(a+b; r_a + r_b)$  remains properly randomized because the sum of random values (assuming they are drawn from a large enough space) is itself random. Schemes must be designed to ensure this “homomorphism of randomness” preserves the hiding property and doesn’t leak information through the combined randomness term. Managing this randomness – its generation, size, domain, and combination during homomorphic operations – is a critical aspect of the scheme’s design and implementation, often requiring careful analysis to prevent subtle attacks exploiting randomness correlation.

**The theoretical elegance of algebraic structures and randomization crystallizes into practical functionality through precisely defined commitment and opening algorithms.** These algorithms embody the core interaction between the committer and verifier. The `Commit` algorithm formalizes the act of sealing the

secret. It takes the message  $m$  (and potentially auxiliary data), the randomness  $r$ , and public parameters ( $pp$  – e.g., the group description  $(G, q, g, h)$  for Pedersen, the SRS for KZG, the matrix  $\mathbf{A}$  and modulus  $q$  for lattices), and outputs the commitment  $c$ . Crucially, `Commit` must be efficient (polynomial time) and must satisfy the **correctness** property: honestly generated commitments must always open successfully later. For homomorphic schemes, the `Commit` algorithm is intrinsically linked to the homomorphism; its output structure ( $c$ ) must belong to a domain where the desired homomorphic operations (addition, multiplication, function evaluation) are well-defined and yield valid outputs that themselves are commitments to the correct derived value. The `Open` algorithm (sometimes called `Verify` or `OpenVerify`) is the counterpart for revelation and verification. It takes the commitment  $c$ , the claimed message  $m$ , the randomness  $r$  used to create it, and the public parameters  $pp$ . Its output is a binary decision: `accept` (if  $c$  is indeed a valid commitment to  $m$  using  $r$ ) or `reject`. The ‘

## 1.6 Applications in Cryptography

The intricate mathematical machinery of homomorphic commitment schemes (HCS), encompassing diverse algebraic structures, precise randomization techniques, and rigorously defined commitment and opening algorithms, is not developed for its own abstract beauty. Rather, it springs to life as a fundamental enabler within advanced cryptographic protocols, solving real-world problems of privacy, verifiability, and distributed trust that were previously intractable. The unique combination of binding secrecy and structured computability inherent in HCS makes them indispensable building blocks in several groundbreaking cryptographic domains, transforming theoretical potential into practical power.

**In the realm of Zero-Knowledge Proofs (ZKPs), homomorphic commitment schemes act as the indispensable cryptographic glue, allowing provers to demonstrate complex knowledge about secrets without revealing them.** Consider the foundational **Schnorr identification protocol**, a precursor to modern ZKPs. A prover, Peggy, wants to convince verifier Victor she knows the discrete logarithm  $x$  of a public value  $y = g^x$ , without disclosing  $x$ . Peggy commits to a random secret  $k$  (the “blinding factor”) using a Pedersen commitment or simply  $r = g^k$ . Victor sends a challenge  $c$ . Peggy then computes a response  $s = k + c \cdot x \bmod q$ . Victor accepts if  $g^s = r \cdot y^c$ . The core magic enabling Victor’s verification without learning  $x$  lies in the homomorphism:  $g^s = g^{k + c \cdot x} = g^k \cdot (g^x)^c = r \cdot y^c$ . The Pedersen commitment  $r$  (or  $g^k$ ) and the public key  $y$  are combined homomorphically (multiplication in the group corresponds to addition in the exponents) to verify the linear relation  $s = k + c \cdot x$  holds, proving knowledge of  $x$  relative to  $y$ . This principle scales dramatically. In more complex **Sigma protocols** proving statements like “I know  $x$  such that  $g^x = y$  AND  $h^x = z$ ” (a discrete log equivalence proof), Pedersen commitments are used to commit to intermediate values and proofs, leveraging their additive homomorphism to allow Victor to check relations between committed exponents. The revolution, however, came with succinct non-interactive proofs (**ZK-SNARKs** and **ZK-STARKs**). Here, the computation or statement to be proven is often encoded as a polynomial. **KZG polynomial commitments** become the workhorse. The prover commits to the polynomial  $\phi(x)$  representing the witness (the secret data satisfying the statement) using the KZG `Commit` algorithm. To prove  $\phi$  satisfies certain constraints (e.g.,  $\phi(z)$

$= y$  for public  $z$ ,  $y$ , or that  $\varphi(x)$  vanishes on a specific set), the prover leverages KZG's homomorphic properties to compute a commitment to a quotient polynomial and generates a constant-sized proof  $\pi$  using the Open algorithm specialized for evaluations. The verifier checks this proof using the pairing, confirming the polynomial relation holds relative to the committed  $\varphi(x)$ , without ever seeing the witness. This is the engine powering the privacy in Zcash (ZeroCash protocol) and the scalability in Ethereum rollups like zkSync and StarkNet, where massive batches of transactions are proven correct via a single succinct proof verifying computations over committed state. HCS are not merely *used in* ZKPs; they are foundational to enforcing the consistency of the prover's hidden computations within the proof structure itself.

**Secure Multi-Party Computation (MPC) leverages homomorphic commitment schemes to ensure integrity and fairness among mutually distrusting parties collaborating on a joint computation with private inputs.** The core challenge in MPC is preventing cheating: a participant might provide incorrect input shares or deviate during the computation phase to bias the result. HCS provide powerful mechanisms to bind parties to their inputs and intermediate states. A primary application is in **Verifiable Secret Sharing (VSS)**, a crucial MPC building block. In Shamir's secret sharing, a dealer splits a secret  $s$  into shares distributed among  $n$  parties. Basic VSS ensures the dealer commits to a consistent polynomial  $\varphi(x)$  (where  $\varphi(0) = s$ ) before distributing shares. Pedersen commitments are ideal for this. The dealer broadcasts a Pedersen vector commitment to the coefficients of  $\varphi(x)$ :  $C_j = g^{a_j} h^{r_j}$  for each coefficient  $a_j$ . Parties receive their share  $\varphi(i)$  and can locally verify its consistency against these public commitments using the homomorphism:  $g^{\{\varphi(i)\}}$  should equal  $\prod_j (C_j)^{i^j} = \prod_j (g^{a_j} h^{r_j})^{i^j} = g^{\sum_j a_j i^j} h^{\sum_j r_j i^j} = g^{\{\varphi(i)\}} h^{r'}$ . They cannot recover  $s$ , but the binding property ensures the dealer is committed to a single polynomial, guaranteeing consistent shares and enabling recovery if the dealer later disappears. Beyond VSS, HCS enforce honest behavior during the computation phase itself. Consider parties holding private inputs  $x_i$  wanting to compute  $f(x_1, \dots, x_n)$ . They first commit to their inputs using an HCS, say  $C_i = \text{Commit}(x_i; r_i)$ . During the MPC protocol (often based on secret sharing), parties perform computations on shares. At critical points, or at the end, parties can be required to reveal openings of commitments to intermediate or final results. Crucially, the homomorphism allows verifiable computation *on the commitments*. If  $f$  is linear (e.g., sum, weighted average), parties can homomorphically compute  $C_f = f(C_1, \dots, C_n)$ , which should be  $\text{Commit}(f(x_1, \dots, x_n); r_f)$ . They can then open  $C_f$  to reveal the result, and anyone can verify the opening and that  $C_f$  was correctly derived from the input commitments, ensuring participants used the inputs they originally committed to and followed the protocol for computing  $f$ . While non-linear computations are more complex, often involving interactive protocols or ZKPs *within* the MPC, the use of HCS for input commitment and linear step verification remains a cornerstone for ensuring robustness against malicious adversaries in MPC frameworks like SPDZ and its variants.

**Verifiable Computation and Delegation harness homomorphic commitment schemes to allow resource-constrained clients to securely outsource expensive computations to powerful but untrusted servers, with cryptographic guarantees of correctness.** Imagine a mobile device needing the result of a complex machine learning model inference on private user data, or a lightweight blockchain node needing to verify



the validity of a block containing thousands of transactions. Performing the computation locally is infeasible. The client commits to its private input data  $x$  (and potentially the function  $f$  itself, if proprietary or large) using an HCS, sending the commitment  $C_x$  (and  $C_f$ , if needed) to the server. The server performs the computation  $y = f(x)$  and returns  $y$  along with a proof  $\pi$ . The proof  $\pi$  demonstrates that  $y$  is indeed the correct output of  $f$  applied to the data committed in  $C_x$ . The power of HCS shines in generating and verifying this proof efficiently. For polynomial computations, KZG commitments are transformative. The client commits to

## 1.7 Applications Beyond Core Cryptography

The transformative power of homomorphic commitment schemes (HCS), previously demonstrated as the cryptographic engine driving verifiable computation, zero-knowledge proofs, and secure multi-party protocols, extends far beyond these core cryptographic constructs. Their unique ability to seal data while enabling verifiable computations on that sealed state unlocks profound applications across diverse domains, fundamentally reshaping how trust, privacy, and verifiability are engineered into real-world systems. From securing financial transactions on global ledgers to safeguarding the sanctity of democratic votes, and enabling collaborative analysis of sensitive datasets, HCS are becoming indispensable tools for a privacy-centric digital future.

**Within the volatile and transparent world of cryptocurrencies and blockchain technology, HCS provide essential mechanisms for reconciling the often-conflicting demands of privacy, scalability, and verifiability.** The most direct application lies in **confidential transactions**. Protocols like **Mimblewimble** (implemented in cryptocurrencies like Grin and Beam) leverage the **additive homomorphism** of Pedersen commitments to revolutionary effect. Here, transaction amounts are hidden within commitments. Crucially, the homomorphic property allows the network to publicly verify that the sum of inputs equals the sum of outputs plus fees ( $\text{Comm}(\text{input\_total}) = \text{Comm}(\text{output\_total} + \text{fees})$ ), ensuring no coins are created from thin air, *without* revealing any individual amounts. This provides strong financial privacy akin to cash. **Zcash**, utilizing the zk-SNARK-based Zerocash protocol, takes this further. While its core privacy relies on zk-SNARKs, the underlying proof system heavily depends on polynomial commitments (originally based on inner-product arguments, evolving towards more efficient schemes) to commit to the transaction state and witness data, enabling the generation of succinct proofs validating complex transaction logic while shielding sender, receiver, and amount. The quest for **scalability** also heavily leverages HCS, particularly through **ZK-Rollups**. Systems like zkSync, StarkNet, Polygon zkEVM, and Scroll use powerful polynomial commitment schemes (often KZG or FRI-based STARKs) as core components. In a ZK-Rollup, thousands of transactions are batched off-chain. The rollup operator generates a succinct zero-knowledge proof (SNARK or STARK) demonstrating the validity of all transactions in the batch relative to the previous blockchain state. Crucially, the new state root (a commitment to the entire state of accounts and balances) is computed as part of this proof. The homomorphic properties of the underlying commitment scheme (within the ZKP system) are essential for efficiently committing to the massive state and proving correct transitions. This single proof and new state commitment are then posted on-chain, enabling **light**

**clients** to verify the entire batch’s validity with minimal computation by simply checking the ZKP against the compact state commitment, drastically reducing the on-chain footprint and enabling blockchain scaling by orders of magnitude.

**The integrity and secrecy of democratic processes find a powerful ally in homomorphic commitment schemes, enabling the construction of verifiable voting systems where votes remain private yet the tally is provably correct.** The core challenge is allowing public verification of the election outcome while preventing coercion or vote-buying by keeping individual votes secret. HCS provide an elegant solution. Each voter encrypts their vote  $v_i$  (e.g.,  $v_i = g^{\text{candidate}}$ ) and commits to this encrypted vote using a homomorphic commitment scheme, generating  $C_i = \text{Comm}(\text{Enc}(v_i); r_i)$ . This commitment is posted to a public bulletin board. Crucially, the encryption scheme itself is often additively homomorphic (e.g., using Paillier or ElGamal variants), and the commitment scheme’s homomorphism is chosen to align with this. After the election closes, the encrypted votes are homomorphically tallied:  $\text{Tally} = \prod \text{Enc}(v_i) = \text{Enc}(\sum v_i)$ . Due to the homomorphism of both the encryption and the commitment scheme, authorities (or any verifier) can compute the combined commitment to the tally  $C_{\text{tally}}$  directly from the individual vote commitments  $C_i$  on the bulletin board, without decrypting any votes. The decrypted result ( $\sum v_i$ ) can then be published. Verifiers can check two things: 1) That the published result correctly decrypts  $\text{Tally}$ , and 2) Crucially, that  $C_{\text{tally}}$  derived from the bulletin board commitments matches the commitment expected for  $\text{Tally}$ . This provides **universal verifiability** – anyone can mathematically verify the tally corresponds to the committed encrypted votes. Simultaneously, **individual verifiability** is often provided: voters receive a receipt allowing them to later check their specific encrypted vote  $\text{Enc}(v_i)$  was correctly included in the bulletin board under commitment  $C_i$ , without revealing *how* they voted to anyone else. Real-world implementations like **SwissPost’s system** (developed with experts including Rolf Haenni and Reto Koenig) and elements within **Estonia’s i-voting system** have explored variations of this approach. The non-profit **Open Vote Network** leverages Pedersen commitments specifically for verifiable multi-party computation of the tally in boardroom-style elections. While adoption faces hurdles related to usability and trust in the encryption/decryption authorities, HCS offer the strongest cryptographic guarantees for end-to-end verifiable (E2E-V) voting systems that truly protect ballot secrecy.

**The burgeoning fields of data analysis and machine learning face immense pressure to utilize sensitive information while respecting privacy regulations and ethical boundaries. Homomorphic commitment schemes offer a pathway towards privacy-preserving collaboration and verifiable outsourced computation.** Consider a consortium of hospitals wanting to train a model on patient data without sharing the raw, sensitive records. Participants can commit to their local datasets or model parameters using HCS. Leveraging the homomorphic properties, they can then collaboratively compute aggregate statistics (sums, averages) or even perform federated learning updates directly on the commitments. The final aggregate or global model parameters can be revealed, along with proofs derived from the homomorphic computations demonstrating that the result was correctly computed from the committed inputs, all while the individual participant data remains confidential. This provides **verifiable data sharing** – participants know the computation used their sealed data correctly, but competitors or malicious actors cannot reconstruct sensitive details. Similarly, for **verifiable outsourced machine learning**, a data owner can commit to their dataset  $D$  and a specific ML



model  $M$  (or its architecture and hyperparameters) using a suitable HCS. They send these commitments to a powerful cloud service. The service trains the model  $M$  on  $D$ , producing the trained model weights  $W$  and predictions  $P$ . Crucially, it also generates a proof  $\pi$  using the homomorphic properties of the commitment scheme (often integrated within a larger ZKP system) demonstrating that  $W$  and  $P$  were correctly derived by executing the committed model  $M$  on the committed data  $D$ . The client can verify this proof efficiently. This ensures the integrity of the outsourced computation – the client knows the expensive training was performed faithfully on their data – while the cloud provider never gains access to the raw sensitive  $D$ . Projects like **zkML** (Zero-Knowledge Machine Learning) are actively exploring the use of zk-SNARKs/STARKs, which inherently rely on polynomial commitments, to achieve this verifiability for complex models. Significant challenges remain, particularly regarding the **efficiency** of generating proofs for very large models and datasets and the **expressiveness** of the functions that can be efficiently verified compared to Fully Homomorphic Encryption (FHE). However, HCS provide a crucial building block for balancing the immense potential of data-driven insights with the fundamental right to privacy.

\*\*Even established

## 1.8 Implementation Challenges & Practical Considerations

The transformative potential of homomorphic commitment schemes (HCS) across domains like confidential blockchain transactions, verifiable voting, and privacy-preserving machine learning, as explored in the previous section, represents a compelling vision. However, bridging the gap between elegant cryptographic theory and robust, real-world deployment introduces significant practical hurdles. Implementing HCS efficiently and securely demands careful navigation of computational constraints, precise parameterization, complex trust dynamics, and evolving software ecosystems. These implementation challenges shape the viability and adoption of HCS-powered technologies.

**The computational complexity inherent in different homomorphic commitment paradigms directly impacts performance and scalability, often dictating their suitability for specific applications.** Pairing-based schemes like KZG offer revolutionary succinctness for polynomial commitments, enabling tiny proofs crucial for ZK-Rollups. However, this comes at the cost of computationally intensive operations. Generating a KZG commitment to a large polynomial involves multi-exponentiations scaling linearly with the polynomial's degree, while pairing operations themselves, though constant-time for verification, are orders of magnitude slower than simple elliptic curve operations. Verifying a single pairing on modern curves like BLS12-381 can take several milliseconds, whereas verifying a Pedersen commitment opening might take microseconds. This asymmetry becomes critical in high-throughput systems like blockchains processing thousands of transactions per second. Conversely, lattice-based schemes promise post-quantum security but currently impose a substantial computational burden. Commitments often involve large matrix-vector multiplications over high-dimensional integer lattices (e.g., dimensions  $1024 \times 1024$  or larger), generating significantly larger objects (kilobytes versus bytes for elliptic curve commitments) and requiring complex operations like rejection sampling. Discrete Logarithm (DL) schemes like Pedersen offer the best raw performance for basic additive operations but lack the expressive polynomial homomorphism needed for complex

ZKPs. Developers employ various optimizations to mitigate these bottlenecks. **Batching** combines multiple operations (e.g., verifying many KZG evaluation proofs simultaneously) to amortize the cost of expensive pairings. **Precomputation** allows generating parts of commitments or proofs ahead of time. Projects like Zcash (Sapling upgrade) invested heavily in optimized pairing libraries and circuit design to make zk-SNARKs practically usable on consumer hardware. Nevertheless, the computational overhead remains a primary factor influencing paradigm choice; applications demanding high throughput for linear operations might favor Pedersen, while those needing succinct proofs for complex logic may accept pairing costs, and those prioritizing post-quantum security brace for current lattice inefficiencies, hoping for future algorithmic improvements.

**Closely tied to performance is the critical task of parameter selection, where mathematical security assumptions must be translated into concrete, secure bit-lengths and dimensions, balancing efficiency against robust security margins.** Choosing insecure parameters can render an otherwise sound scheme completely broken, while overly conservative choices cripple performance. For DL-based schemes like Pedersen, the security level (e.g., 128-bit) dictates the required size of the underlying elliptic curve group or finite field. A 128-bit security level typically requires a 256-bit prime for elliptic curves (like secp256k1 or BLS12-381) or a 3072-bit RSA modulus. However, this isn't static; cryptanalytic advances constantly threaten these estimates. The Logjam attack demonstrated how vulnerabilities in the TLS protocol could downgrade Diffie-Hellman security by exploiting weak 512-bit groups, highlighting the dangers of outdated parameter choices. Lattice-based schemes face even greater complexity. Security relies on the hardness of problems like Learning With Errors (LWE) or Ring-LWE, parameterized by the lattice dimension  $n$ , the modulus  $q$ , and the error distribution  $\chi$ . Selecting these involves navigating complex trade-offs: larger  $n$  improves security but increases key/commitment sizes and slows operations; larger  $q$  can sometimes allow smaller  $n$  but impacts security reduction tightness and noise management. The NIST Post-Quantum Cryptography Standardization project provides vital guidance, but concrete parameter sets for lattice-based commitments are still evolving as research refines security estimates against both classical and quantum attacks. Furthermore, the **randomness quality** used for blinding factors must be cryptographically secure; weak entropy sources can compromise hiding, as seen in historical vulnerabilities in embedded systems. Consequently, parameter selection is not a one-time task but an ongoing process requiring vigilance against cryptanalytic progress and adherence to best practices from standardization bodies, balancing provable security guarantees with the practical need for usable systems.

**The requirement for a trusted setup in certain powerful schemes, notably pairing-based polynomial commitments like KZG, introduces a significant procedural and trust challenge unique within the deployment lifecycle.** The KZG scheme relies on a **Structured Reference String (SRS)**, typically generated as powers of a secret scalar  $s$  ( $[g, g^s, g^{s^2}, \dots, g^{s^d}]$ ). Knowledge of  $s$  – the **toxic waste** – allows forging commitments and proofs. This creates a critical point of vulnerability: if the entity generating the SRS retains or leaks  $s$ , the entire system's security collapses. Relying on a single trusted party is antithetical to the decentralized ethos of many applications, particularly blockchains. To mitigate this, the cryptographic community developed **trusted setup ceremonies**, also known as **Power of Tau ceremonies** or MPC (Multi-Party Computation) setups. In these elaborate protocols, multiple participants sequentially

contribute to generating the SRS. Each participant  $i$  receives the current SRS state, samples a fresh random secret  $\tau_i$ , updates the SRS by exponentiating all elements by  $\tau_i$ , and then destroys  $\tau_i$ . Crucially, they also provide a proof (often a ZKP) that they performed the exponentiation correctly without leaking  $\tau_i$ . As long as at least one participant is honest and destroys their secret, the final SRS is secure, even if all other participants are malicious and collude. The security rests on the difficulty of inverting the product  $\tau = \tau_1 * \tau_2 * \dots * \tau_n$  even knowing the individual contributions to the SRS. Real-world examples demonstrate both the feasibility and complexity. The Zcash “Sprout” setup involved six participants in 2016, while later ceremonies like “Sapling” and the massive Ethereum-oriented **Perpetual Powers of Tau** (with over 70 participants globally, including Vitalik Buterin) scaled significantly. These ceremonies involve meticulous procedures, specialized software, and physical security measures (filmed destruction of hardware). While MPC setups greatly reduce trust, they are complex logistical undertakings and require careful verification of participant proofs. This inherent complexity drives interest in **transparent (trustless setup)** alternatives, such as lattice-based schemes or hash-based constructions like FRI (used in ZK-STARKs), which require only public randomness. However, these often come with their own trade-offs in proof size or computational efficiency compared to KZG. The trusted setup remains a distinctive and challenging aspect of deploying the most succinct polynomial commitment schemes.

\*\*The maturation and wider adoption of hom

## 1.9 Security Analysis & Cryptanalytic Threats

The practical deployment hurdles detailed in Section 8—computational overhead, parameter selection pitfalls, the logistical complexity of trusted setups, and nascent library support—underscore a fundamental reality: cryptographic elegance alone is insufficient. Theoretical security guarantees must withstand relentless assault in the real world. This brings us to the critical evaluation of homomorphic commitment schemes (HCS) against known and emerging cryptanalytic threats. The very properties that empower HCS—algebraic structure enabling homomorphism, reliance on computational hardness assumptions, and the intricate dance of hiding and binding—also create unique vectors for attack. Understanding these vulnerabilities is paramount for assessing the robustness of specific schemes and guiding secure implementation choices.

**The foundational security guarantees of any HCS, its hiding and binding properties, remain primary targets for cryptanalysts, with threats arising from both algorithmic advances and implementation flaws.** A core vulnerability lies in **weak randomness generation**. The blinding factor  $r$  is the linchpin of hiding; if  $r$  is predictable, reused, or drawn from an insufficiently large space, secrecy collapses catastrophically. The infamous 2010 breach of the Sony PlayStation 3’s ECDSA implementation serves as a stark historical parallel: reuse of the nonce  $k$  (analogous to a weak blinding factor) allowed attackers to extract the master private signing key, compromising the entire system. For HCS like Pedersen, if the same  $r$  is reused for different messages  $m_1$  and  $m_2$ , the commitments become  $g^{m_1} h^r$  and  $g^{m_2} h^r$ ; an attacker can compute  $\text{Comm}(m_1) / \text{Comm}(m_2) = g^{m_1 - m_2}$ , directly revealing the difference  $m_1 - m_2$ . If one message is known or guessed, the other is exposed. Even partial leakage of  $r$  bits can significantly weaken hiding. Beyond randomness failures, **cryptanalytic advances against the underly-**

**ing mathematical problems** pose existential threats. Shor’s algorithm renders schemes based on Integer Factorization (RSA) or Discrete Logarithms (Pedersen, basic pairing-based schemes) completely insecure against a sufficiently large, fault-tolerant quantum computer, as it efficiently solves these problems. While practical quantum computers of this scale remain years away, the threat necessitates post-quantum migration planning. For classical adversaries, improvements in algorithms for the DLP (e.g., the ongoing refinement of index calculus and Pollard’s rho methods) or for lattice problems (like the Shortest Vector Problem - SVP, targeted by sieving or basis reduction algorithms such as BKZ 2.0) constantly erode concrete security margins. A scheme deemed secure with 128-bit parameters today might offer only 100-bit security tomorrow due to algorithmic breakthroughs, emphasizing the need for conservative parameter selection and agility. Furthermore, **exploiting statistical deviations** can break hiding. While Pedersen offers statistical hiding under DLP, poorly chosen parameters (e.g., a group order  $q$  where  $m$  has limited possibilities) or biased message distributions might allow distinguishing attacks. For lattice-based schemes, the precise shape of the error distribution  $\chi$  is crucial; deviations or correlations in sampled errors can leak information about the committed message  $\text{vec}\{m\}$ . Ensuring the hardness assumptions hold *under homomorphic operations* is also critical; multiplying many RSA commitments might inadvertently create structures vulnerable to factoring attacks, or repeated operations on lattice commitments could amplify noise in ways exploitable by advanced lattice reduction techniques. Vigilance against these evolving attacks on the core assumptions is non-negotiable for long-term security.

**The defining feature of HCS—their homomorphic capabilities—introduces novel attack surfaces distinct from basic commitments. Malicious actors can exploit the algebraic structure to perform unintended computations or forge openings.** A prevalent class is **misuse attacks**, where an adversary leverages homomorphism to compute a valid commitment  $C'$  that corresponds not to a meaningful or intended message, but to a maliciously crafted result. For instance, in a confidential transaction system using additive Pedersen commitments, an attacker might homomorphically combine their own valid commitment  $\text{Comm}(\text{amount\_A})$  with a stolen commitment  $\text{Comm}(\text{amount\_B})$  (perhaps intercepted from another transaction) to create  $\text{Comm}(\text{amount\_A} + \text{amount\_B})$ , effectively attempting to steal funds. While higher-level protocol rules should prevent such invalid combinations from being accepted (e.g., by verifying input/output ownership), the homomorphic property *enables* the algebraic manipulation that must then be detected and rejected by the application logic. Failure to implement such checks creates a vulnerability. More insidiously, **forging openings through algebraic manipulation** attacks target the binding property under homomorphic operations. An adversary might attempt to find two distinct pairs  $(m, r)$  and  $(m', r')$  for a commitment  $C$  derived homomorphically from other commitments, violating binding for the derived value. In pairing-based schemes like KZG, while the scheme itself may be binding, vulnerabilities in the specific way the polynomial or its evaluations are used within a larger protocol (e.g., a specific ZK-SNARK construction) might allow an attacker to exploit linear dependencies in the polynomial coefficients or evaluation points to create a forged evaluation proof  $\pi$  for an incorrect value  $y' \neq \phi(z)$ , if the verifier’s checks aren’t sufficiently rigorous. **Distinguishing attacks** represent another threat vector amplified by homomorphism. If the homomorphic operation creates a commitment  $C_{\text{result}}$  whose structure or distribution differs statistically from a commitment to a “benign” result, an observer might distinguish be-

tween valid and invalid computations or even infer properties about the underlying secrets. For example, in an additively homomorphic scheme, the sum of many small-value commitments might have a different randomness distribution (e.g., smaller variance in the combined  $r\_sum$ ) compared to a commitment to a single large value, potentially leaking information about the nature of the computation or inputs. The algebraic richness enabling polynomial homomorphism in KZG also creates complexity; subtle interactions between the committed polynomial, the evaluation point  $z$ , and the SRS might theoretically enable attacks if the parameters or usage deviates from strict security proofs. Defending against these attacks requires not only the inherent security of the commitment scheme but also careful protocol design that anticipates and constrains how homomorphism can be applied.

**Even when the underlying mathematical scheme is theoretically sound, real-world implementations are vulnerable to a plethora of side-channel and software-based attacks that bypass cryptographic hardness entirely, exploiting physical artifacts or coding errors. Timing attacks** are particularly pernicious. If the time taken to execute the `Commit`, `Open`, or homomorphic operation algorithms varies depending on the secret bits (e.g., the message  $m$  or the blinding factor  $r$ ), an attacker measuring execution times can glean information. A classic example involves early modular exponentiation implementations (relevant to RSA-based commitments or pairing computations in KZG) where branches or lookups dependent on secret exponents created timing variances sufficient to recover keys. Similarly, operations in lattice-based schemes involving rejection sampling or conditional loops based on secret values are potential timing leak vectors. **Power analysis attacks** take this further. By monitoring the fluctuating power consumption of a device (like a hardware wallet or TEE) while it performs cryptographic operations, sophisticated adversaries can correlate power traces with secret data. Simple Power Analysis (SPA) might visually identify operations like point doubling vs. point addition in elliptic curve Pedersen commitments, revealing bits of the exponent. Differential Power Analysis (DPA) uses statistical methods on numerous traces to extract secrets even with significant noise. **Fault injection attacks** introduce deliberate perturbations—via voltage glitching, clock manipulation, or targeted radiation—to cause computational errors during critical operations. For instance, inducing a fault during the computation of a KZG opening proof might result in an invalid

## 1.10 Comparisons & Trade-offs

The relentless scrutiny of security threats against homomorphic commitment schemes (HCS) – from cryptanalytic assaults on core assumptions to sophisticated side-channel exploits – underscores a critical reality: no single scheme offers a universal solution. As these powerful primitives migrate from theoretical constructs into real-world systems, practitioners face a complex landscape of trade-offs. Selecting the optimal HCS demands careful evaluation of mathematical foundations, performance profiles, security horizons, and alignment with application-specific requirements. This comparative analysis provides a structured framework for navigating these choices, contrasting the dominant construction paradigms, differentiating HCS from related technologies, and illustrating decision pathways through concrete use cases.

**Comparing the major construction paradigms reveals distinct constellations of strengths and limitations, directly influencing their suitability across the cryptographic ecosystem.** Discrete Logarithm



(DL)-based schemes, epitomized by Pedersen commitments, offer compelling simplicity and efficiency. Leveraging well-understood hardness assumptions like the DLP in elliptic curve groups, they achieve statistical hiding and computational binding with minimal computational overhead. Their additive homomorphism enables elegant solutions for confidential transactions (Mimblewimble) and basic zero-knowledge proofs (Schnorr). However, they are fundamentally constrained to linear operations and lack succinctness for complex computations. Crucially, they are vulnerable to Shor's algorithm, offering no post-quantum (PQ) security. Pairing-based schemes, particularly KZG polynomial commitments, represent the pinnacle of expressive power and succinctness. Their ability to commit to entire polynomials and generate constant-sized evaluation proofs revolutionized ZK-SNARKs, enabling scalable blockchain rollups (zkSync, Scroll). The trade-off is substantial computational cost (complex pairings, large multi-exponentiations) and the persistent specter of the trusted setup ceremony required to generate the Structured Reference String (SRS), introducing logistical complexity and a single point of failure if compromised. Like DL schemes, they are not PQ-secure. Lattice-based schemes, founded on Learning With Errors (LWE) or Ring-LWE, offer a promising PQ-secure alternative. They provide inherent additive homomorphism and, increasingly, support for more complex operations without trusted setups, appealing to decentralized ethos. Projects like Filecoin have explored lattice-based VDFs and commitments. However, current implementations suffer from significant inefficiencies: large key/commitment sizes (kilobytes vs. bytes for elliptic curves) and slower operations due to matrix multiplications and noise management, hindering adoption in high-throughput systems. RSA-based schemes, while historically significant for potential multiplicative homomorphism under the Strong RSA assumption, are largely sidelined today. Their large parameter sizes (3072+ bit moduli) make them inefficient, and their vulnerability to both classical cryptanalytic advances and quantum factorization relegates them to niche roles. Thus, the paradigm choice often reduces to a trilemma: prioritize efficiency and simplicity (DL, but quantum-vulnerable), embrace maximal expressiveness and succinctness despite setup and cost (Pairing, also quantum-vulnerable), or opt for future-proof PQ security at the expense of current performance (Lattice).

**Understanding the unique niche of homomorphic commitment schemes requires contrasting them with related cryptographic technologies, particularly Fully Homomorphic Encryption (FHE) and Zero-Knowledge Proofs (ZKPs), clarifying their complementary roles.** FHE allows arbitrary computations on *encrypted* data, returning encrypted results that only the secret key holder can decrypt. This provides unparalleled flexibility for privacy-preserving outsourcing (e.g., private cloud computation). However, FHE incurs immense computational overhead (orders of magnitude slower than HCS), large ciphertext expansion, and complex key management. HCS, in contrast, focus on *verifiable computation over committed data*. They prove that a specific computation on hidden inputs yields a specific (potentially revealed) output, or that the hidden data satisfies certain properties. This is achieved with far greater efficiency, especially for linear operations or when paired with succinct proofs (KZG), but lacks FHE's generality – HCS cannot directly compute arbitrary functions on sealed data *and* keep the result sealed without additional proof machinery. The relationship with ZKPs is symbiotic yet distinct. HCS are fundamental building blocks *within* ZKP systems, not direct competitors. Polynomial commitments like KZG are the engine enabling efficient ZK-SNARKs to prove complex statements about committed polynomials. Basic commitments (Pedersen)

are essential components in Sigma protocols. ZKPs are the broader framework for proving arbitrary NP statements in zero-knowledge; HCS provide the mechanism to seal the witness (prover's secret data) and enforce consistency during the proof generation. When simpler hiding suffices without needing verifiable computation *on* the sealed data, basic non-homomorphic commitments (e.g., hash-based) remain sufficient and more efficient. Therefore, the choice hinges on the core need: FHE for private computation on encrypted data, basic commitments for simple sealing, ZKPs for proving complex knowledge about secrets, and HCS specifically for verifiable computation on committed secrets – often acting as a critical component within ZKPs themselves.

**Selecting the optimal homomorphic commitment scheme for a specific application necessitates weighing five pivotal factors: the required homomorphic capability, post-quantum security needs, setup constraints, performance demands, and succinctness requirements.** Consider the contrasting demands of three archetypal use cases. For **confidential blockchain transactions** (e.g., designing a privacy coin), the primary need is efficient additive homomorphism to verify balance ( $\sum \text{inputs} = \sum \text{outputs} + \text{fees}$ ) without revealing amounts. Performance is paramount for network scalability. Trusted setups are undesirable in decentralized environments. Post-quantum security, while desirable, might be a secondary concern if migration paths exist. Here, **Pedersen commitments (DL-based)** are the pragmatic choice, proven in Mimblewimble implementations. Their blinding efficiency and simple additive properties perfectly match the requirement. If PQ security is non-negotiable, **lattice-based additive schemes** emerge as the alternative, accepting current performance penalties for future-proofing. Conversely, the **backend for a ZK-Rollup** demands polynomial homomorphism and

## 1.11 Societal Impact, Ethics & Controversies

The intricate technical trade-offs explored in Section 10 – balancing efficiency, expressiveness, quantum vulnerability, and trust assumptions – ultimately serve a purpose beyond cryptographic elegance. Homomorphic commitment schemes (HCS) are not merely abstract mathematical constructs; they are foundational enablers of technologies that profoundly reshape societal interactions with privacy, trust, and authority. As these schemes migrate from academic papers and niche protocols into real-world systems governing finance, identity, and democratic participation, their deployment sparks significant ethical debates, regulatory challenges, and controversies that touch the core of individual rights and collective security. Examining these broader implications reveals the complex interplay between cryptographic potential and societal values.

**The core promise of HCS – enabling verifiable computation on sealed data – offers unprecedented tools for enhancing individual privacy in an increasingly surveilled digital landscape, yet simultaneously fuels concerns about enabling illicit activities beyond the reach of oversight.** Technologies powered by HCS provide tangible privacy benefits. Blockchain platforms like Zcash and Monero utilize HCS-based zero-knowledge proofs to shield transaction details, offering financial confidentiality akin to physical cash for law-abiding citizens seeking protection from predatory surveillance, corporate profiling, or oppressive regimes. Verifiable e-voting systems employing homomorphic tallying on committed ballots (e.g., SwissPost's pilot) aim to preserve ballot secrecy while guaranteeing auditability, a cornerstone of



democratic integrity often eroded in opaque digital systems. Privacy-preserving data analysis using HCS allows collaborative research on sensitive medical or financial datasets without exposing individual records, potentially accelerating scientific breakthroughs while respecting confidentiality. However, this powerful privacy enhancement cuts both ways. Regulatory bodies and law enforcement agencies worldwide express deep concern that the very cryptographic guarantees protecting legitimate users also provide cover for money laundering, terrorist financing, and other illicit transactions. The 2022 sanctioning of the Ethereum-based mixer Tornado Cash (which relied on cryptographic techniques including commitments and ZKPs) by the US Treasury Department’s Office of Foreign Assets Control (OFAC) exemplifies this tension. Authorities argued Tornado Cash materially assisted North Korean hackers (Lazarus Group) in laundering stolen cryptocurrency, effectively anonymizing transactions in a way traditional financial surveillance could not penetrate. This incident ignited fierce debate within the crypto and privacy communities, pitting arguments for financial privacy and permissionless innovation against legitimate national security imperatives and regulatory mandates like the Financial Action Task Force’s (FATF) “Travel Rule” (Recommendation 16), which requires Virtual Asset Service Providers (VASPs) to share sender/receiver information – a requirement fundamentally at odds with strong on-chain privacy protocols built using HCS. This duality manifests starkly: HCS empower individuals against unwarranted scrutiny but can also obstruct lawful investigation, forcing societies to grapple with the difficult equilibrium between fundamental privacy rights and collective security needs in the digital age. Calls for “backdoors” or exceptional access mechanisms tailored to HCS-based systems routinely surface, but cryptographers universally warn these fatally compromise the fundamental security properties, echoing the historical “Crypto Wars” debates over encryption.

**HCS inherently mediate questions of trust and transparency, offering powerful mechanisms to engineer verifiability into opaque processes while simultaneously introducing new, sometimes paradoxical, trust dependencies.** A key societal benefit lies in their ability to foster **algorithmic accountability**. ZK-Rollups on blockchains like Ethereum leverage HCS (e.g., KZG commitments) to generate succinct proofs that massive batches of off-chain transactions were processed correctly relative to a committed state. This allows anyone, even resource-constrained users, to cryptographically verify the integrity of the entire system’s operation without trusting the rollup operator – a paradigm shift towards transparent and auditable infrastructure. Similarly, end-to-end verifiable voting systems built with HCS enable voters to independently confirm their ballot was included correctly and auditors to verify the entire tally matches the cast votes, all while preserving ballot secrecy. This reduces reliance on trusting central authorities or proprietary software. However, this engineered verifiability often clashes with **practical obscurity**. The mathematical complexity underlying HCS, particularly succinct zero-knowledge proofs utilizing polynomial commitments, creates a significant barrier to public understanding. While the outputs (a valid proof) are verifiable, the process itself remains a “black box” for non-experts. This creates a paradox: the system is *mathematically* transparent and verifiable, yet *pragmatically* opaque to the average user, potentially undermining public trust if the verification process feels inaccessible or mystical. Furthermore, the **trusted setup problem**, inherent to some of the most efficient schemes like KZG, introduces a critical societal wrinkle. While multi-party ceremonies (like Ethereum’s Perpetual Powers of Tau) significantly mitigate the risk by distributing trust among many participants (ideally destroying their individual secrets), the necessity of *any* initial trust ritual remains a point of

contention, especially in communities valuing radical decentralization. The 2016 Zcash “Sprout” ceremony, involving six participants, faced intense scrutiny precisely because of this inherent philosophical tension – could the system be truly trustless if its birth required trusted actors? This highlights how cryptographic trust models intersect with societal values, where mathematical security must navigate real-world human constraints and the perception of trustworthiness. Explaining the nuanced security guarantees of HCS (e.g., computational binding vs. statistical hiding) and the meaning of a verified ZK proof to policymakers and the public remains a significant communication challenge, essential for fostering informed societal acceptance.

**The deployment of HCS-enabled technologies occurs within a rapidly evolving and fragmented global regulatory landscape, marked by divergent approaches to privacy-enhancing technologies (PETs) and intense debates within standardization bodies about future-proofing cryptographic infrastructure.** Regulatory attitudes towards PETs incorporating HCS vary dramatically. The European Union, through frameworks like the General Data Protection Regulation (GDPR), implicitly encourages PETs as a means to achieve “data protection by design and by default” (Article 25), recognizing technologies like secure computation on committed data as potential compliance tools. Conversely, some national regulators, particularly in the financial sector influenced by FATF recommendations, view strong cryptographic privacy in digital assets with deep suspicion, leading to proposals for strict controls or outright bans on privacy coins like Monero or Zcash in jurisdictions like Japan and South Korea. This regulatory patchwork creates significant uncertainty for developers and businesses seeking to leverage HCS. Within **technical standardization bodies**, crucial debates rage over the path forward. Organizations like the Internet Engineering Task Force (IETF) and the National Institute of Standards and Technology (NIST) are grappling with how to standardize HCS, particularly in the context of the looming quantum threat. NIST’s ongoing Post-Quantum Cryptography (PQC) standardization project focuses primarily on digital signatures and KEMs, leaving the standardization of post-quantum secure HCS (primarily lattice-based) less mature. This creates a dilemma: should current standards prioritize the well-understood efficiency of pairing-based schemes like KZG (despite lacking PQ security and requiring trusted setup) or push for nascent, less efficient lattice-based alternatives offering PQ security and transparency? Furthermore, regulating open-source cryptographic protocols themselves, rather than just the entities using them, presents a novel challenge, as evidenced by the Tornado Cash sanctions. Can code be responsible? How does regulation interact with decentralized, permissionless systems whose core protocols embed privacy guarantees via HCS? These questions remain largely unresolved, fueling controversy and highlighting the tension between fostering innovation in privacy technology and establishing frameworks to prevent abuse. The trajectory of HCS adoption will be significantly shaped by the outcomes of these regulatory and standardization debates, determining whether these powerful tools become widely accessible instruments for empowering individual privacy and verifiability or become constrained by compliance burdens and

## 1.12 Future Directions & Open Problems

The regulatory fragmentation and intense standardization debates surrounding homomorphic commitment schemes (HCS), as discussed in the context of societal impact, underscore a critical reality: the field is not

static. It is propelled forward by a vibrant research community tackling profound technical challenges and exploring transformative new frontiers. As the limitations of current paradigms become clearer under the pressures of quantum threats, efficiency demands, and novel application requirements, several key trajectories define the cutting edge of HCS research, alongside stubbornly persistent open problems that continue to shape the landscape.

**The specter of large-scale quantum computing casts a long shadow, driving intense focus on developing practical post-quantum secure HCS.** Lattice-based cryptography, primarily founded on the Learning With Errors (LWE) and Module Short Integer Solution (Module-SIS) problems, remains the most promising avenue. Recent research focuses on enhancing the efficiency and expressiveness of lattice-based commitments. Projects like the **Brakerski-Lombardi-Polychroniadou (BLP)** commitment scheme demonstrate progress towards achieving functionalities akin to KZG polynomial commitments but without trusted setups and with quantum resistance. Techniques involve intricate use of lattice trapdoors and rejection sampling to enable commitment opening and evaluation proofs. However, significant hurdles remain. The computational overhead and large parameter sizes (commitments often measured in kilobytes rather than bytes) are primary bottlenecks. The **NIST Post-Quantum Cryptography standardization process**, while establishing lattice-based signatures and KEMs, has yet to converge on standardized HCS primitives, creating uncertainty. Beyond lattices, researchers explore alternative PQ foundations. **Isogeny-based cryptography**, leveraging the hardness of computing isogenies between supersingular elliptic curves, offers potential for compact constructions. Schemes like **SeaSign** inspired ideas for commitments, though efficient homomorphic properties remain challenging to achieve. **Code-based cryptography**, relying on the syndrome decoding problem, provides another PQ candidate. Schemes like the **Wave signature protocol** incorporate commitment-like structures, and research into adapting code-based techniques specifically for succinct HCS is ongoing. **Hash-based constructions**, while inherently PQ-secure and transparent, typically lack efficient homomorphism beyond simple binding. Integrating them with protocols like FRI (used in transparent ZK-STARKs) offers a path for verifiable computation without algebraic homomorphism but often at the cost of larger proof sizes compared to KZG. The quest is not merely for PQ security but for PQ schemes achieving performance and functionality parity with current non-PQ leaders like KZG – a goal still distant but actively pursued through initiatives like the **PQ-CRYPTO** project and research consortia.

**Alongside the quantum imperative, relentless pressure exists to enhance the efficiency of existing schemes and expand their expressive capabilities.** For pairing-based constructions, particularly KZG, research focuses on **algorithmic optimization** and **hardware acceleration**. Improving multi-scalar multiplication algorithms for commitment generation and optimizing pairing implementations for specific curves (e.g., BLS12-381, BN254) are crucial for real-world systems like ZK-Rollups. Projects like **Arkworks** and **Halo2** showcase significant performance gains through such optimizations. Simultaneously, **transparent alternatives** to KZG are advancing rapidly. **FRI (Fast Reed-Solomon IOPP)** forms the backbone of ZK-STARKs and inherently provides a commitment mechanism for large polynomials through Merkle trees built over FRI codewords. While transparent and PQ-secure, FRI proofs grow poly-logarithmically with the circuit size, making them larger than KZG proofs for complex computations, though often faster to generate. Schemes like **RedShift** (from Polygon Zero) and **ethSTARK** exemplify this approach. **DARK (Diophan-**

**tine Arguments of Knowledge**) commitments, based on groups of unknown order like RSA groups or class groups, offer another transparent path. While potentially PQ-secure if using class groups, DARK suffers from larger proof sizes and slower verification than KZG. **Vector commitments** (e.g., **Verkle trees**, using polynomial commitments for internal nodes) represent a powerful application-focused efficiency gain, dramatically reducing proof sizes for sparse state trees in blockchains like Ethereum's upcoming upgrade. Beyond efficiency, expanding **expressiveness** is crucial. Can HCS support richer operations – division, comparisons, or even arbitrary circuits – with efficiency surpassing Fully Homomorphic Encryption (FHE)? Research explores techniques like combining different commitment types (e.g., additive Pedersen for some inputs with KZG for polynomials) or using HCS within recursive proof compositions to handle complex computations in manageable chunks. Projects like **Nova** (using folding schemes) and **HyperNova** explore these frontiers, aiming for broader functionality while maintaining practical verifier costs.

**The unique properties of HCS – binding secrecy coupled with verifiable computation – continue to inspire novel applications and integration with emerging paradigms. Decentralized Identity (DID) and Verifiable Credentials (VCs)** stand to benefit immensely. HCS can enable the issuance of credentials where claims are committed, allowing selective disclosure of specific attributes via zero-knowledge proofs derived from the commitment, while enabling verifiers to cryptographically confirm the credential's integrity and issuer without accessing all data. The **W3C Verifiable Credentials Data Model** community actively explores such integrations. In **privacy-preserving machine learning**, HCS offer a middle ground between pure MPC and FHE. Federated learning can leverage HCS for participants to commit to their model updates; aggregators can homomorphically combine these commitments and prove the correctness of the aggregation without accessing individual updates, enhancing privacy and verifiability in collaborative training. The intersection with **confidential computing** using **Trusted Execution Environments (TEEs)** like Intel SGX or AMD SEV presents intriguing synergies. While TEEs aim to shield computation, they require remote attestation to prove correct initialization. HCS could commit to the initial state (enclave hash, public keys) and attestation reports within a blockchain or public ledger, providing tamper-evident verification of the TEE's state before sensitive data is submitted. Furthermore, the computation results generated within the TEE could be verifiably linked back to the committed initial state using HCS and ZKPs, creating a chain of cryptographic trust augmenting the hardware root. **Token binding and privacy-preserving access control** are other emerging areas, where HCS commit to access tokens or policies, enabling efficient verification of complex authorization logic without revealing the underlying credentials or user attributes prematurely. The potential for HCS to act as cryptographic glue connecting disparate privacy-enhancing technologies is a fertile ground for innovation.

**Despite significant progress, several fundamental open challenges persistently resist elegant solutions, shaping the long-term research agenda. Achieving truly practical transparent polynomial commitments** remains elusive. While FRI and DARK offer transparency, their proof sizes or verification costs often exceed those of pairing-based KZG, limiting adoption in performance-critical settings like high-throughput blockchains. Closing this gap without compromising security or introducing new assumptions is a major goal. Bridging the **chasm between theoretical security proofs and practical implementation security** is another enduring struggle. Formal methods and verified implementations (e.g., using tools like **EasyCrypt**

or **HACL\***) are crucial but complex for sophisticated HCS. Side-channel resistance remains challenging, especially for lattice-based schemes involving complex sampling. Ensuring that abstract security models accurately capture real-world adversary capabilities, including novel quantum attacks or sophisticated side-channels, requires constant vigilance. Furthermore, the field grapples with **cryptographic agility** – the ability for systems to transition smoothly between different