

Encyclopedia Galactica

"Encyclopedia Galactica: Type-2 ZK-EVMs"

Entry #:	943.73.6
Word Count:	24300 words
Reading Time:	122 minutes
Last Updated:	July 26, 2025

"In space, no one can hear you think."

Table of Contents

Contents

1	Encyclopedia Galactica: Type-2 ZK-EVMs	3
1.1	Section 1: Defining the Type-2 ZK-EVM Paradigm	3
1.2	Section 2: Historical Evolution and Conceptual Genesis	9
1.3	Section 3: Architectural Deep Dive	16
1.3.1	3.1 State Management and Data Availability	16
1.3.2	3.2 Proving Infrastructure Stack	18
1.3.3	3.3 The Execution Environment: Modified Clients	20
1.3.4	3.4 Bridge Architecture and Cross-Rollup Messaging	23
1.4	Section 4: The Proving Process: From Execution to Verification	26
1.4.1	4.1 Transaction Flow: User Tx to L1 Finality	26
1.4.2	4.2 Inside the Prover: Circuit Execution Details	27
1.4.3	4.3 Verification on Ethereum L1	29
1.4.4	4.4 Latency and Finality Characteristics	30
1.5	Section 5: Comparative Analysis of Major Implementations	31
1.6	Section 6: Security Model and Trust Assumptions	40
1.6.1	6.1 Cryptographic Trust Foundations	41
1.6.2	6.2 Systemic Risks and Failure Modes	42
1.6.3	6.3 Economic Security and Slashing Mechanisms	44
1.6.4	6.4 Formal Verification Progress	45
1.7	Section 7: Developer Experience and Ecosystem Impact	47
1.7.1	7.1 Tooling Compatibility Landscape	47
1.7.2	7.2 Gas Cost Dynamics and Contract Migration	49
1.7.3	7.4 Wallet Integration Challenges	50
1.8	Section 8: Economic and Governance Dimensions	52

1.8.1	8.1 Tokenomics of Proving Markets	53
1.8.2	8.2 Sequencer Decentralization Pathways	54
1.8.3	8.3 Treasury Management and Protocol Upgrades	56
1.8.4	8.4 Cross-Rollup Liquidity Fragmentation	57
1.9	Section 9: Controversies and Critical Debates	59
1.9.1	9.1 The “Type-2.5” Compromise Debate	59
1.9.2	9.2 Centralization-Optimization Tension	61
1.9.3	9.3 Long-Term Sustainability Concerns	62
1.9.4	9.4 Regulatory Ambiguities	63
1.10	Section 10: Future Trajectories and Conclusion	65
1.10.1	10.1 Next-Generation Proof Systems	65
1.10.2	10.2 Cross-Chain ZK Interoperability	66
1.10.3	10.3 Ethereum Integration Roadmap	67
1.10.4	10.4 Broader Implications for Web3	69
1.10.5	10.5 Concluding Synthesis	70

1 Encyclopedia Galactica: Type-2 ZK-EVMs

1.1 Section 1: Defining the Type-2 ZK-EVM Paradigm

The relentless pursuit of scalability without compromising security or decentralization – Ethereum’s infamous “trilemma” – has birthed a constellation of Layer 2 (L2) solutions. Among these, Zero-Knowledge Ethereum Virtual Machines (ZK-EVMs) represent a particularly audacious and technically profound frontier. Within this burgeoning field, the **Type-2 ZK-EVM** has emerged as a critical design philosophy, striking a delicate balance between rigorous fidelity to Ethereum’s core execution environment and the practical demands of generating succinct cryptographic proofs. This section establishes the conceptual bedrock for understanding Type-2 ZK-EVMs: their foundational technologies, their precise definition within Vitalik Buterin’s influential taxonomy, their core technical guarantees, and how they differentiate themselves within the complex ecosystem of scaling solutions.

1.1 Foundational Principles: ZK-Proofs and EVM Equivalence

At the heart of any ZK-rollup, including ZK-EVMs, lies the revolutionary power of **zero-knowledge proofs (ZKPs)**, specifically zk-SNARKs (Zero-Knowledge Succinct Non-Interactive Arguments of Knowledge) and zk-STARKs (Zero-Knowledge Scalable Transparent Arguments of Knowledge). These cryptographic primitives allow one party (the *prover*) to convince another party (the *verifier*) that a specific computation was executed correctly *without revealing any details about the inputs or the internal steps of the computation itself*, beyond the validity of the output. Imagine submitting a sealed, verifiable receipt proving you solved a complex puzzle, without revealing the solution path. This “verifiability without revelation” is the cryptographic magic enabling scaling.

- **The Scaling Mechanism:** In a ZK-rollup, transactions are processed off-chain (on the L2). Instead of posting all transaction data to Ethereum (L1), the rollup operator (sequencer/prover) generates a ZK-proof attesting to the *correctness* of a batch of transactions and the resulting new state root. This single proof, orders of magnitude smaller than the raw transaction data, is then posted to and verified by a smart contract on Ethereum L1. Successful verification updates the canonical state root stored on L1, inheriting Ethereum’s security guarantees. This drastically reduces the data burden on L1 while ensuring the integrity of the L2 state transitions is mathematically verifiable.
- **The EVM Challenge:** The Ethereum Virtual Machine (EVM) is the globally accessible, sandboxed runtime environment where all Ethereum smart contracts execute. It defines the instruction set (opcodes), memory model, stack handling, storage access, and gas metering that constitute the bedrock of Ethereum’s behavior. Making this complex, stateful, and non-deterministic-in-parts environment efficiently provable under ZKPs is an immense engineering challenge. Early ZK-rollups (like Loopring or zkSync v1) sidestepped this by supporting only limited, custom functionality or requiring developers to write in specialized languages (e.g., Zinc) that compiled to ZK-friendly circuits. While functional, this fractured the developer experience.
- **EVM Equivalence vs. EVM Compatibility:** This brings us to the crucial distinction:

- **EVM Compatibility:** A system *understands* EVM bytecode but may execute it differently or incompletely. Developers might need to use specialized compilers (e.g., transpiling Solidity to Cairo for Starknet, or to Yul/Zinc for early zkRollups), adjust contracts for unsupported opcodes, or deal with subtly different gas costs or behavior. While offering *some* familiarity, it introduces friction and potential security risks during migration.
- **EVM Equivalence:** This is the gold standard. An EVM-equivalent system **preserves bytecode-level compatibility**. Any smart contract that compiles to standard EVM bytecode and runs correctly on Ethereum L1 *will run identically* on the L2, using the *same* underlying opcodes, gas semantics, memory handling, and storage structures, *without modification or recompilation*. The developer experience is indistinguishable from deploying on L1. Vitalik Buterin aptly termed this the “holy grail” of ZK-rollups, as it preserves Ethereum’s vast developer ecosystem and tooling entirely. The Type-2 ZK-EVM is defined by its commitment to achieving this equivalence.

The core proposition of Type-2 ZK-EVMs is thus the seamless convergence of two powerful paradigms: the trustless scaling enabled by succinct cryptographic proofs and the frictionless developer experience guaranteed by bytecode-level EVM equivalence. It aims to make scaling Ethereum feel like using Ethereum itself.

1.2 The Buterin Taxonomy: Understanding the Type Spectrum

The landscape of ZK-EVM development rapidly diversified in 2021-2022, leading to different approaches with varying degrees of fidelity to Ethereum’s execution layer. To bring clarity and establish shared terminology, Ethereum co-founder Vitalik Buterin proposed a seminal **taxonomy** in August 2022 (“[The different types of ZK-EVMs](#)”), classifying ZK-EVMs into four primary types (later expanded by the community to five) based on their level of equivalence to Ethereum’s consensus layer execution:

1. **Type 1 (Fully Ethereum-Equivalent):** Aims for perfect parity with Ethereum consensus *at every level*. It uses Ethereum’s exact execution model, data structures (e.g., Patricia-Merkle tries), and gas costs. Every Ethereum block could, in theory, be proven by a Type 1 ZK-EVM. *Advantage:* Maximum compatibility and decentralization; Ethereum clients could be used directly as provers. *Disadvantage:* Proving performance is currently impractical for mainnet use due to the sheer complexity of proving every Ethereum opcode and state access natively. The Ethereum Foundation’s Privacy and Scaling Explorations (PSE) team is actively researching this path.
2. **Type 2 (Fully EVM-Equivalent):** This is our focus. Type 2 ZK-EVMs strive for **full equivalence at the EVM execution level**, meaning *all existing Ethereum smart contracts work unmodified*. They preserve Ethereum’s opcodes, memory model, stack, and gas semantics *exactly*. However, they may make pragmatic changes *beneath* the EVM layer to enhance prover efficiency. Crucially, these changes are invisible to the smart contract itself. *Advantage:* Unparalleled developer experience (true “drop-in” compatibility), leverages all existing EVM tooling. *Disadvantage:* Proving times are still significant, though manageable with optimizations, and underlying state tree structures might differ slightly.

3. **Type 2.5 (Equivalent Except Gas Costs):** An informal but practically important category emerging from implementation realities. Type 2.5 systems maintain full EVM equivalence in opcode behavior *but may alter the gas costs* of certain operations, particularly those expensive to prove under ZK (like the KECCAK256 hash or storage accesses). The goal is to make proving economically viable without breaking contracts. *Advantage:* Significantly improved prover performance/cost. *Disadvantage:* Potential subtle deviations in contract gas consumption behavior, requiring careful analysis during deployment. Polygon zkEVM adopted this stance post-launch.
4. **Type 3 (Almost EVM-Equivalent):** Intentionally sacrifices *some* equivalence for significant prover performance gains. Certain complex or ZK-unfriendly EVM features (e.g., specific precompiles, contract creation quirks, or handling of certain edge-case opcodes) might be modified, omitted, or handled differently. Developers might need to make minor adjustments to existing contracts. *Advantage:* Much faster and cheaper proving than Type 2. *Disadvantage:* Breaks compatibility for some existing contracts and tools, introducing developer friction. Early versions of Scroll and Polygon zkEVM started here before moving towards Type 2/2.5.
5. **Type 4 (High-Level-Language Equivalent):** Operates at the level of high-level languages (Solidity, Vyper). Developers write Solidity, but it's compiled down to a custom, ZK-friendly intermediate language (IR) or VM bytecode *different* from EVM bytecode. *Advantage:* Highest proving performance and flexibility for ZK-specific optimizations. *Disadvantage:* Breaks bytecode compatibility; existing deployed EVM bytecode doesn't work, and debugging/tooling differs significantly from Ethereum L1. zkSync Era and Starknet are prominent examples.

Why Type-2 Occupies the “Sweet Spot”: Buterin’s taxonomy highlights a fundamental tension: *fidelity to Ethereum versus prover efficiency*. Type 1 offers maximal fidelity but is currently impractical. Type 4 offers maximal efficiency but sacrifices the core EVM developer experience. Type 3 improves efficiency but introduces compatibility friction. **Type 2 strategically targets the “sweet spot”:** It preserves the *entire* existing Ethereum developer ecosystem and deployed contract base by guaranteeing bytecode equivalence, while accepting the engineering challenge and computational cost of proving the full EVM. The pragmatic adjustments of Type 2.5 represent an effort to nudge closer to optimal efficiency without abandoning the core Type 2 promise. This balance makes Type 2 the most compelling path for scaling Ethereum while maintaining its core identity, acting as a true “extension” of L1 rather than a distinct, incompatible platform.

1.3 Core Technical Objectives and Guarantees

The definition of a Type-2 ZK-EVM translates into concrete technical objectives and guarantees that distinguish it from other types and scaling solutions:

1. **Unmodified EVM Opcode Support and Bytecode Execution:** The paramount objective. *Every* EVM opcode (from basic arithmetic ADD to complex cryptographic KECCAK256, from state access SLOAD/SSTORE to contract creation CREATE/CREATE2) must behave *identically* to its implementation on Ethereum L1. This includes handling edge cases, exceptions (e.g., out-of-gas, stack underflow/overflow), and all potential side effects. Crucially, this means that *standard EVM bytecode*

generated by the official Solidity or Vyper compilers executes without modification or recompilation for the ZK environment. Tools like Hardhat, Foundry, and Remix function identically.

2. **Handling Ethereum’s Memory Model and Storage Layouts:** The EVM’s memory (volatile, linear byte array), stack (LIFO structure for operands), and storage (persistent key-value store per contract) must be implemented faithfully. This includes:
 - Precisely mimicking memory expansion costs and gas calculation.
 - Maintaining the exact stack depth limits (1024 items) and operation semantics (e.g., DUP, SWAP).
 - Preserving Ethereum’s complex storage layout, especially concerning how state variables in Solidity contracts are packed into storage slots and inherit storage layouts from parent contracts. Incorrect storage handling can lead to catastrophic state corruption when migrating contracts.
3. **Preserving Gas Semantics and Deterministic State Transitions:** Gas is the mechanism that measures and constrains computational effort on Ethereum. For true equivalence:
 - The *gas cost* of every opcode must match Ethereum L1 (a point relaxed in Type 2.5).
 - The *gas metering mechanism* itself must be identical – tracking gas consumption per opcode, handling gas refunds correctly (e.g., for storage clearing via SSTORE), and triggering out-of-gas exceptions at precisely the same point.
 - The *state transition function* must be deterministic and identical to L1. Given the same starting state and transaction inputs, the Type 2 ZK-EVM must produce *exactly* the same ending state and gas consumption as Ethereum L1 would. This determinism is essential for generating verifiable proofs of correct execution.
4. **Support for Ethereum’s Unique Features:** Faithfully replicating behavior around:
 - **Precompiles:** Specialized, gas-efficient contracts for cryptographic operations (e.g., ecadd, ecmul, ecpairing for elliptic curves, SHA256, RIPEMD160, MODEXP). These are critical for many DeFi protocols and must be implemented with identical inputs/outputs and gas costs.
 - **Access Lists (EIP-2930):** Mechanisms for pre-declaring storage slots to access, impacting gas costs.
 - **Contract Creation Rules:** Handling CREATE/CREATE2, constructor execution, and the precise initialization code semantics.
 - **Self-Destruct (SELFDESTRUCT):** The complex behavior and state clearing associated with this opcode.

Achieving these objectives requires sophisticated engineering. For instance, the KECCAK256 hash, ubiquitous in Ethereum, is computationally intensive to prove in ZK. Type-2 implementations often employ complex circuit designs or lookup arguments to optimize this, while ensuring the output remains byte-for-byte identical to an L1 Keccak hash. Similarly, simulating Ethereum’s account-based storage model (using Patricia-Merkle Tries) efficiently within a ZK circuit is a significant challenge addressed through innovative state management techniques. Projects like Scroll emphasize their commitment to “bytecode-level equivalence,” meaning they execute the *exact* bytecode generated by Ethereum’s standard compilers, providing the strongest guarantee of compatibility.

1.4 Distinguishing from Adjacent Technologies

Positioning Type-2 ZK-EVMs within the broader scaling landscape requires contrasting them with key alternatives:

1. Optimistic Rollups (ORUs - e.g., Optimism, Arbitrum):

- **Security Assumption:** ORUs rely on economic incentives and a fraud-proving mechanism. They *assume* transactions are valid by default (optimism) but allow verifiers to challenge invalid state transitions during a dispute window (typically 7 days). Security relies on the existence of at least one honest verifier.
- **Exit Mechanism:** Withdrawing assets from an ORU to L1 requires waiting for the full challenge window (7 days) to ensure no fraud proofs are submitted. This creates significant latency for users.
- **ZK-EVM Differentiation:** Type-2 ZK-EVMs provide **cryptographic security** (via validity proofs) from the moment a proof is verified on L1. Withdrawals can be near-instant (minutes/hours, limited only by proving time and L1 confirmation). There is no need for fraud proofs or lengthy challenge periods. While ORUs often achieve high EVM compatibility (Arbitrum is very close to Type 2, Optimism historically modified gas costs), their core security model and user experience for exits differ fundamentally.

2. Type-1 ZK-EVMs:

- **Equivalence Level:** Type-1 aims for complete parity with Ethereum’s consensus layer, including underlying data structures. Type-2 focuses solely on EVM execution equivalence, allowing optimizations underneath (e.g., different state tree formats).
- **Practicality:** Type-1 is currently a research goal due to prohibitive proving times for full Ethereum blocks. Type-2 represents the practical realization of full EVM equivalence achievable with current technology, accepting minor infrastructural deviations for performance.

3. Type-3/Type-4 ZK-EVMs and Other ZK-VMs:

- **Compatibility vs. Performance:** Type-3 sacrifices some EVM opcode/precompile equivalence for better prover performance. Type-4 (like zkSync Era using its LLVM-based compiler or Starknet using Cairo) abandons EVM bytecode compatibility entirely, requiring developers to use custom languages or compilers. Type-2 prioritizes unmodified compatibility above maximum proving speed.
- **Developer Experience:** Type-2 offers a “just deploy” experience for existing Solidity/Vyper codebases. Type-3 may require audits and minor adjustments. Type-4 requires significant rewrites or learning new languages/toolchains.

4. Sidechains (e.g., Polygon POS, Gnosis Chain, BSC):

- **Security Model:** Sidechains operate with their own independent consensus mechanisms (often Proof-of-Stake with fewer validators) and security budgets. They do not inherit Ethereum’s security. Funds are secured by the sidechain’s validators, not Ethereum L1.
- **Bridges:** Moving assets between Ethereum and a sidechain requires separate, often complex and potentially vulnerable, bridge contracts.
- **ZK-EVM Differentiation:** Type-2 ZK-EVMs are **cryptographically secured by Ethereum L1**. Validity proofs ensure the L2 state is correct, and users can always withdraw their assets directly via L1 contracts without relying on bridge operators. While some sidechains (like Polygon POS) are EVM-compatible, they lack the trust-minimized security derived from Ethereum.

5. **Plasma:** An earlier scaling concept proposing state commitments on L1 with fraud proofs for exits. While pioneering, Plasma struggled with complex exit games, data availability issues (especially for general computation), and limited support for complex smart contracts. ZK-Rollups, particularly ZK-EVMs, offer a more robust and flexible solution for general-purpose smart contract execution with stronger data availability guarantees (especially post-EIP-4844) and simpler user exits.

In essence, the Type-2 ZK-EVM paradigm represents a unique convergence: it offers the strong, cryptographically enforced security and near-instant finality of validity-proof-based systems, combined with the seamless, unmodified developer experience and bytecode compatibility of Ethereum itself. It accepts the engineering challenge of proving the complex EVM in exchange for maximal ecosystem continuity and trust minimization.

This foundational definition sets the stage for understanding the intricate journey that led to the conceptualization and initial realization of Type-2 systems. The path was paved by early experiments, theoretical breakthroughs, and the gradual refinement of a shared vision, a history we will explore next. [Transition to Section 2: The quest for a truly equivalent ZK-EVM began not with a blueprint, but with a series of incremental innovations and hard-won lessons from pioneering projects navigating the uncharted territory of scaling Ethereum with zero-knowledge cryptography...](#)

1.2 Section 2: Historical Evolution and Conceptual Genesis

The quest to scale Ethereum without fracturing its developer ecosystem or compromising its security foundations, culminating in the Type-2 ZK-EVM paradigm, was not born fully formed. It emerged from a crucible of audacious experimentation, formidable technical roadblocks, and incremental breakthroughs that gradually transformed an implausible vision into a tangible engineering reality. This section traces the intricate technological lineage of the Type-2 ZK-EVM, charting the pivotal moments and key innovations that laid its conceptual and practical groundwork. It is a history marked by the ingenuity of early pioneers confronting the daunting “ZK-EVM problem,” revolutionary advances in the underlying cryptographic machinery, a unifying conceptual framework that galvanized the industry, and the daring first steps of teams striving to bring the ideal of full EVM equivalence into production.

2.1 Precursors: Early ZK-Rollup Experiments (2020-2021)

The narrative begins not with EVM equivalence, but with the foundational ambition of scaling Ethereum transactions via zero-knowledge proofs. The years 2020-2021 witnessed the emergence of the first practical ZK-Rollups, demonstrating the core validity-proof mechanism but operating within significantly constrained execution environments. These pioneers grappled with the fundamental tension: the raw computational complexity of generating ZK-proofs versus the need for practical throughput and usability.

- **Loopring Protocol: Scaling Payments and Swaps:** Launched on Ethereum mainnet in **December 2019**, Loopring v3 (and later v3.6) was arguably the first production ZK-Rollup. Focused primarily on payment transfers and decentralized exchange (DEX) order-book and automated market maker (AMM) swaps, it achieved remarkable efficiency gains. However, its smart contract capabilities were intentionally limited. Developers could not deploy arbitrary Solidity contracts. Instead, Loopring offered a specialized, ZK-optimized DEX protocol implemented within its custom circuits. This approach sidestepped the immense challenge of proving general EVM computation but clearly illustrated the trade-off: specialized functionality enabled high throughput (thousands of trades/sec) but sacrificed the generality that defines Ethereum. Loopring’s success proved the core ZK-Rollup data compression and security model worked, setting a crucial precedent.
- **zkSync v1: Towards Programmable ZK-Rollups:** Matter Labs launched zkSync v1 (originally zkSync) in **June 2020**, marking a significant step towards programmability. While still not supporting the full EVM, zkSync v1 introduced its own custom VM and the Zinc programming language, a Rust-based language designed specifically for ZK-friendliness. This allowed developers to write custom smart contracts (like basic token transfers, NFT minting, and simple DEX logic) that compiled to Zinc bytecode, executed within Matter Labs’ proprietary ZK-circuits. While a leap forward in flexibility compared to Loopring’s fixed functionality, the requirement to learn a new language (Zinc) and use specialized tooling created a significant barrier to entry for the vast majority of Ethereum developers accustomed to Solidity and the EVM. The “ZK-EVM problem” – the immense difficulty of efficiently proving *standard* EVM bytecode execution – remained unsolved. zkSync v1’s architecture

exemplified the Type-4 approach in Buterin’s later taxonomy, prioritizing prover efficiency and novel features (like account abstraction from day one) over bytecode compatibility.

- **Aztec Network: Pioneering Privacy:** Concurrently, Aztec Network (**mainnet launch November 2021**) embarked on a radically different but equally pioneering path. Focused on *private* smart contracts and transactions on Ethereum, Aztec developed a sophisticated ZK-Rollup (Aztec Connect, later Aztec) utilizing its custom Noir programming language and PLONK-based proof system. While not aiming for EVM equivalence, Aztec’s contributions were profound. They pushed the boundaries of complex circuit design, particularly for privacy-preserving operations like confidential token transfers and private DeFi interactions via their “bridged asset” model. Aztec tackled head-on the challenges of state management and efficient proof generation for intricate logic within a ZK context. Their work demonstrated the feasibility of complex, application-specific ZK-Rollups and provided valuable lessons in tooling (like the Noir language and Aztec Sandbox) that would later inform broader ZK-VM development. Aztec underscored that ZK-Rollups weren’t just for scaling, but could unlock entirely new capabilities like privacy, albeit initially through a non-EVM-compatible path.
- **The “ZK-EVM Problem” Crystallizes:** By mid-2021, the limitations of the first generation were starkly apparent. Projects like Optimistic Rollups (Arbitrum and Optimism launched mainnets in May and July/August 2021 respectively) were gaining traction precisely because they *did* offer near-full EVM compatibility, allowing developers to deploy existing contracts with minimal friction, despite the trade-offs of fraud proofs and week-long withdrawal delays. The ZK-Rollup community recognized that to capture mainstream Ethereum development, they needed to solve the core challenge: **How to make the existing, unmodified EVM – with its complex opcodes, intricate gas metering, stateful storage model, and non-deterministic elements – efficiently provable under zero-knowledge cryptography?** Custom VMs and languages (Type-4) offered performance but fragmented the ecosystem. The path to seamless scaling required conquering the EVM itself. This daunting challenge became known as the “ZK-EVM problem.”

This period was characterized by pragmatic solutions focused on achievable, high-value use cases (payments, swaps, privacy) using purpose-built VMs. While they proved the core ZK-Rollup thesis, the dream of a frictionless, fully compatible ZK-EVM remained distant, held back by the computational intractability of proving the standard EVM with the proof systems and hardware of the time.

2.2 The Turning Point: Breakthroughs in Proof Systems (2021-2022)

The years 2021-2022 witnessed a cascade of breakthroughs in ZK-proof theory and engineering that fundamentally altered the feasibility landscape for ZK-EVMs. These innovations dramatically reduced the computational overhead of proof generation, particularly for complex, non-arithmetic operations inherent in the EVM, and enabled the practical construction of the large, intricate circuits needed to emulate it faithfully.

- **Recursion Unleashed: PLONK, Halo/Halo2, and RedShift:** A critical bottleneck was the sheer size of the circuits required to represent EVM execution traces. Early SNARKs (like Groth16) required a

separate, computationally intensive trusted setup for *each circuit*. This was untenable for the evolving, complex circuits of a ZK-EVM. The advent and maturation of **universal and updatable SNARKs**, particularly **PLONK** (developed by Aztec, Protocol Labs, and others, with production use by Aztec from 2021) and **Halo/Halo2** (developed by the Electric Coin Company for Zcash, open-sourced and adapted more broadly), changed everything.

- **Universal Trusted Setup:** PLONK and Halo2 required only a *single*, large trusted setup ceremony (like the influential [Aztec Ignition](#) and [Ethereum KZG Ceremony](#)) that could then be used for *any* circuit within a certain size bound. This eliminated the need for per-circuit setups, enabling rapid iteration and deployment of complex ZK-EVM circuits.
- **Recursive Proof Composition:** Halo (and its refinement, Halo2) introduced efficient **recursive proof composition**. This allows a prover to generate a proof *about another proof*. Why is this revolutionary for ZK-EVMs? Instead of proving the execution of an entire block of EVM transactions in one monolithic, impossibly large circuit, a ZK-EVM could break execution into smaller chunks (e.g., per transaction or per opcode group), prove each chunk separately, and then use recursion to aggregate these smaller proofs into a single, succinct proof for the entire block. This “divide and conquer” approach made proving large-scale EVM execution computationally tractable. Polygon zkEVM would later leverage this heavily with its Plonky2 (PLONK + Halo-inspired recursion) system. StarkWare’s **RedShift** (a STARK-based recursive proving system) also demonstrated the power of this paradigm.
- **Witness Generation Revolution: zkASM and Interpreters:** Generating the “witness” – the set of variable assignments that satisfy the circuit constraints representing a correct execution – is often the most computationally intensive step *before* the actual cryptographic proof is generated. Directly compiling EVM bytecode to a circuit representation proved incredibly complex and inefficient for general computation. A key innovation emerged: **zkASM (Zero-Knowledge Assembly) interpreters**.
- **The Interpreter Model:** Pioneered by teams like Polygon (with zkASM) and later adopted by others like Scroll, this approach decouples the circuit design from the direct compilation of EVM opcodes. Instead, a relatively small, fixed circuit is designed to execute a custom “zk assembly” instruction set (zkASM). A separate component, the **zkEVM interpreter**, then translates *standard EVM bytecode* into sequences of these zkASM instructions. The fixed zkASM circuit executes these instructions, and the ZK-proof attests to the correct execution of the zkASM program, which itself corresponds to the original EVM bytecode. This abstraction layer dramatically simplified circuit design, improved auditability, and made the system more adaptable to EVM changes. It became a cornerstone architecture for achieving EVM equivalence efficiently.
- **Lookup Arguments: Taming Non-Arithmetic Operations:** The EVM is replete with operations that are notoriously expensive to represent in traditional arithmetic circuits used by SNARKs, particularly cryptographic hash functions (KECCAK256, SHA256) and memory/storage accesses. **Lookup arguments** (like Plookup, first introduced with PLONK, and later advanced variants like cq, logup, and the highly efficient Lasso/Jolt approach) provided a breakthrough.

- **The Core Idea:** Instead of laboriously expressing complex non-arithmetic operations (like a hash computation) as millions of arithmetic gates, a lookup argument allows the prover to show that the inputs and outputs of the operation exist within a predefined, precomputed table of valid input/output pairs. The proof demonstrates consistency with the table without revealing which specific entry was used. This reduced the circuit size for operations like `KECCAK256` by orders of magnitude, making proving Ethereum’s ubiquitous hashing feasible within a Type-2 context. Polygon zkEVM and Scroll heavily utilize lookup arguments for these critical opcodes.
- **Polygon’s Theoretical Foundation: “A zkEVM with no trusted setup”:** In April 2022, a landmark paper titled “[A zkEVM with no trusted setup](#)” was presented at Eurocrypt by Polygon’s Jordi Baylina, as part of the Polygon Zero team (formerly Mir Protocol). This paper wasn’t just about a specific implementation; it provided a crucial *theoretical framework* and practical blueprint for building a fully EVM-compatible ZK-Rollup using PLONK with a universal trusted setup and recursive proofs. It detailed solutions for critical challenges like handling Ethereum’s memory model, storage, program counter, stack, and particularly the `KECCAK256` hash using custom gates and lookup arguments within the PLONK arithmetization. This paper served as a beacon, demonstrating that a practical ZK-EVM adhering closely to the EVM specification was achievable with existing cryptographic primitives and careful engineering. It significantly de-risked the pursuit of EVM equivalence and influenced multiple teams.

These breakthroughs collectively shattered the perceived impossibility of the ZK-EVM. Recursion made large-scale proving feasible, interpreters made circuit design manageable, lookup arguments tamed the cost of critical EVM operations, and theoretical work provided concrete blueprints. The technological foundation for Type-2 ZK-EVMs was now firmly laid.

2.3 Vitalik Buterin’s Typology and Industry Alignment

Despite the technical progress, the ZK-EVM landscape in mid-2022 was fragmented. Different projects (Polygon, Scroll, zkSync, StarkNet, the Ethereum Foundation PSE team) were pursuing diverse approaches with varying levels of EVM compatibility, often using different terminology. This created confusion for developers and hindered clear communication about trade-offs and goals.

- **“The Different Types of ZK-EVMs”:** On August 4, 2022, Ethereum co-founder Vitalik Buterin published a seminal blog post titled “[The different types of ZK-EVMs](#)”. This post introduced the now-ubiquitous **taxonomy** classifying ZK-EVMs into Types 1 through 4 (with Type 2.5 emerging later as a practical refinement). Buterin meticulously defined each type based on its level of equivalence to Ethereum’s consensus layer (Type 1) or the EVM execution layer (Type 2), down to the compromises made for efficiency (Types 2.5, 3, and 4).
- **Clarifying the Spectrum and the “Sweet Spot”:** Buterin’s framework provided crucial clarity:
- It explicitly defined **Type 2 (Fully EVM-Equivalent)** as the target preserving Ethereum’s developer experience entirely, executing unmodified EVM bytecode.

- It highlighted the inherent trade-offs: Type 1 (full consensus equivalence) was the ideal but impractical; Type 4 (high-level language equivalence) offered performance but broke compatibility; Type 3 offered a middle ground but still required contract adjustments.
- Crucially, Buterin identified **Type 2 as the “sweet spot”** for projects genuinely prioritizing Ethereum compatibility, arguing it offered “almost all of the benefits of a Type 1 ZK-EVM” without the impractical overhead of replicating Ethereum’s exact state tree.
- **Catalyzing Industry Focus:** The impact was immediate and profound:
- **Unified Language:** The taxonomy gave the entire ecosystem – developers, researchers, investors, and users – a common vocabulary to understand and compare different ZK-EVM approaches. Discussions moved beyond vague claims of “compatibility” to precise classifications.
- **Strategic Alignment:** Buterin’s articulation of Type 2 as the “sweet spot” provided a powerful north star. Projects already leaning towards equivalence (like Scroll and the nascent Polygon zkEVM effort) found their strategy validated and amplified. Projects operating in different paradigms (like zkSync Era as Type 4 or StarkNet using Cairo) could clearly communicate their distinct value propositions and trade-offs.
- **Accelerated R&D:** The clear definition of Type 2 crystallized the technical requirements. Teams could focus their R&D explicitly on overcoming the specific challenges of proving unmodified EVM opcodes, gas semantics, and state transitions, knowing this was recognized as the optimal path for ecosystem continuity. Buterin’s post served as a rallying cry, aligning disparate efforts towards a shared, ambitious goal.
- **Community Discourse:** The post ignited intense and productive debate within the Ethereum research and development community. Discussions flourished on forums like Ethereum Research, refining the understanding of the trade-offs, exploring edge cases in equivalence, and debating the merits of Type 2 vs. Type 2.5 pragmatism. This discourse further sharpened the conceptual understanding of what achieving Type 2 truly entailed.

Buterin’s typology was more than just a classification system; it was a strategic intervention that unified a fragmented field, provided a clear target for maximal compatibility, and significantly accelerated the practical pursuit of the Type-2 ZK-EVM ideal.

2.4 First-Generation Implementations (2022-2023)

Armed with advanced proof systems and galvanized by a clear target, several teams embarked on the arduous task of building the first production-grade Type-2 ZK-EVMs. The period from late 2022 through 2023 saw these pioneers transition from research and testnets to public demonstrations and initial mainnet deployments, marking the tangible arrival of the paradigm.

- **Polygon zkEVM: The First Public Contender:** Demonstrating the fruits of its aggressive research and acquisition strategy (including Polygon Zero/Mir, Hermez, and Miden teams), Polygon made the

first major public move. The **Polygon zkEVM public testnet launched in October 2022**, followed by a **mainnet beta launch in March 2023**. This was a watershed moment, billed as the first “EVM-equivalent” ZK-Rollup available to developers and users.

- **Architecture:** Polygon adopted the zkASM interpreter model. Their prover utilized **Plonky2**, a highly optimized recursive proof system combining PLONK with FRI (from STARKs) for fast recursion, developed by Polygon Zero. Crucially, they achieved this by **forking the Geth (Go-Ethereum) execution client**, modifying it to generate execution traces suitable for their zkProver. This “forked client” approach aimed to leverage the battle-tested Geth codebase for execution fidelity while integrating the ZK-specific trace generation.
- **Initial Challenges & The Type-2.5 Shift:** Launching first came with complexities. Early mainnet performance revealed bottlenecks, particularly in proof generation times and costs associated with certain EVM operations. In a significant and pragmatic move reflecting the real-world engineering challenges, Polygon **publicly transitioned its classification from Type 2 to Type 2.5** in mid-2023. This meant maintaining full opcode equivalence and bytecode compatibility *but altering the gas costs* for specific ZK-unfriendly operations (notably KECCAK256 and some storage operations) to make proving economically sustainable. This decision sparked debate but underscored the practical difficulties of pure Type 2 at launch scale and highlighted the emerging “Type-2.5” category as a pragmatic reality.
- **Scroll: The Bytecode-Purist Approach:** Emerging from Ethereum research roots (co-founded by Ye Zhang, a former Ethereum Foundation researcher) and with strong community backing, Scroll took a distinct path focused on rigorous **bytecode-level equivalence**. Their testnet launched progressively throughout 2023 (pre-alpha in February, alpha in August, mainnet in October).
- **Architecture:** Scroll also utilizes a zkASM-like interpreter layer but emphasizes executing *exactly* the standard EVM bytecode generated by Solidity/Vyper compilers. They built their own **zkEVM node** from the ground up, comprising a sequencer (based on modified Geth) and a coordinator managing provers. Their proving stack leverages a combination of custom circuits and existing engines like **Halo2-KZG**. A key differentiator was Scroll’s early and deep commitment to **open-source** development and seamless integration with **existing Ethereum toolchains** like Hardhat and Foundry via dedicated plugins, making developer onboarding exceptionally smooth for an early-stage ZK-Rollup. Their “community-first” ethos and focus on developer experience resonated strongly.
- **The Prover Network:** Scroll pioneered the concept of a decentralized **prover network** early on, where multiple participants can contribute proof generation compute, aligning with Ethereum’s decentralization principles even in this computationally intensive role.
- **Taiko: The Based Rollup Vision:** Founded by the creators of Loopring, Taiko entered the scene with a unique architectural proposition: the **Based Rollup (a.k.a. “Type 1” lite or “Ethereum-Equivalent” rollup)**. Announced in 2022, Taiko’s TKO testnets (Alpha-1 in March 2023, Alpha-2 in July 2023) emphasized leveraging Ethereum’s own infrastructure as much as possible.

- **Based Sequencing:** Taiko’s most radical departure is eliminating a centralized or even decentralized L2-specific sequencer. Instead, Taiko transactions are **sequenced directly by Ethereum L1 block proposers (validators)** via a special smart contract. This aims to inherit Ethereum’s liveness and censorship-resistance directly, minimizing trust in L2-specific operators. It represents a novel approach to decentralization in the sequencing layer.
- **Multi-Proofs:** To enhance security and potentially optimize performance, Taiko explores a **multi-proof system**. Initially, this involves combining ZKPs (generated by permissionless “provers”) with **TEE (Trusted Execution Environment, like Intel SGX) attestations** generated by “block verifiers.” The long-term vision involves using multiple ZK-proof systems for redundancy. While their ZK-EVM implementation (forked Geth) aims for Type 2 equivalence, the Based Sequencing model is their defining architectural innovation, positioning them as a distinct contender focused on maximal alignment with Ethereum L1’s security model.
- **The Broader Landscape:** While Polygon, Scroll, and Taiko were the primary public flagships for the Type-2/2.5 approach in this period, other significant players were advancing:
- **Consensys Linea:** Announced in late 2022 (testnet March 2023, mainnet July 2023), Linea (formerly Consensys zkEVM) adopted a **gradual equivalence roadmap**, starting closer to Type 3 (with some missing precompiles and minor deviations) with the explicit goal of evolving towards Type 2 equivalence over time, prioritizing developer access and early use cases on its robust infrastructure.
- **Kakarot zkEVM:** An intriguing open-source project built within the Starknet ecosystem (announced early 2023). Kakarot implements an EVM interpreter in **Cairo** (Starknet’s native language), meaning it executes EVM bytecode as a Cairo program. This allows it to potentially leverage Starknet’s scalable STARK-based proof system (Stone) to prove EVM execution. It represents a novel “Type 3+” path, demonstrating EVM compatibility *on top of* another high-performance ZK-VM layer (Starknet as L3). It went live on testnet in late 2023.

The launch of these first-generation Type-2/Type-2.5 ZK-EVMs marked the culmination of years of research and development. They transformed the theoretical promise of EVM equivalence under ZKPs into operational, though nascent, reality. Developers could finally deploy standard Solidity contracts to a validity-proven L2 and experience near-instant finality. While significant challenges in performance, decentralization, and user experience remained, the foundational proof-of-concept was now demonstrably live. The focus shifted from *whether* it could be done to *how well* it could be done at scale.

Transition to Section 3: The journey from conceptual breakthrough to functional reality demanded immense technical ingenuity. Achieving seamless EVM equivalence within the unforgiving constraints of zero-knowledge proof generation required novel architectural solutions across every layer of the stack – from state management and data handling to the intricate dance of proving infrastructure and client modifications...

1.3 Section 3: Architectural Deep Dive

The triumphant launch of the first-generation Type-2 ZK-EVMs, chronicled in the preceding section, represented a monumental engineering achievement. Yet, beneath the surface of operational testnets and nascent mainnets lay a labyrinth of intricate technical innovations. Transforming the ideal of seamless EVM equivalence into a functioning, performant ZK-Rollup demanded radical rethinking of core blockchain components – state representation, execution environments, proof generation pipelines, and trust-minimized bridges – all constrained by the unforgiving mathematical rigor of zero-knowledge proofs. This section dissects the sophisticated architectural scaffolding that empowers Type-2 ZK-EVMs, revealing how they reconcile Ethereum’s intricate execution model with the cryptographic demands of validity proofs.

The core challenge permeating this architecture is the **tension between fidelity and provability**. Faithfully replicating every nuance of the EVM – its state tree, gas costs, opcode semantics, and non-deterministic inputs – inherently generates complex computational traces. Representing this complexity within ZK circuits, while keeping proof generation times and costs practical, required ingenious solutions across multiple layers.

1.3.1 3.1 State Management and Data Availability

At the heart of Ethereum’s execution lies its global state, a massive, constantly evolving data structure storing account balances, contract code, and contract storage. Type-2 ZK-EVMs must not only replicate this state faithfully during execution but also generate cryptographic proofs about its transitions *and* ensure the data underpinning state changes is reliably available. This presents unique hurdles.

- **The Merkle-Patricia Trie (MPT) Conundrum:** Ethereum’s state is organized within a modified Merkle-Patricia Trie, a cryptographic authenticated data structure combining Patricia tries for efficient lookups and Merkle trees for compact verification. Each node in the MPT is hashed (using KECCAK256), and the root hash (the state root) commits to the entire state. Changing any state element changes the root. While elegant for verification on L1, the MPT is notoriously complex and ZK-unfriendly.
- **The Challenge:** Proving state accesses (SLOAD, SSTORE) within a ZK circuit essentially requires proving the correct traversal of the MPT and the correct computation of the KECCAK256 hash at each node along the path *for every single state access*. Doing this naively within a circuit would be prohibitively expensive, blowing up circuit size and proof times. Type-1 ZK-EVMs aspire to replicate the MPT exactly, but Type-2 implementations adopt pragmatic optimizations.
- **The Type-2 Solution: Sparse Merkle Trees (SMTs) & Hybrid Models:** Most Type-2 ZK-EVMs replace Ethereum’s hexary MPT (branching factor of 16) with a **binary Sparse Merkle Tree (SMT)** for storage. Binary trees are vastly simpler to represent and traverse within ZK circuits. The leaf nodes in the SMT correspond to storage slots or account data. Crucially, the *semantics* of storage access – the key-value mapping and persistence – remain identical to the EVM from the contract’s perspective.

Only the underlying cryptographic authentication structure differs. Projects like **Scroll** and **Polygon zkEVM** utilize binary SMTs for storage efficiency. Some employ hybrid models: using an MPT-like structure for high-level account organization but SMTs for contract storage within accounts, optimizing the most frequent access patterns.

- **Proving the Transition:** The ZK-proof doesn't verify every state access individually during execution *within* the main EVM circuit. Instead, the execution trace includes *witness data* (the pre-state, post-state, and accessed state values/paths) for all state operations. A separate, optimized **storage circuit** or **state transition circuit** then verifies the consistency of these accesses and computes the new state root based on the SMT (or hybrid structure) rules. This modularization keeps the core EVM execution circuit focused on opcode logic. **Taiko**, aiming for closer Type-1 equivalence, works harder to mimic the MPT structure but still employs significant circuit optimizations for traversal and hashing.
- **Call Data Compression: The Blob Revolution:** Transaction data (call data) must be made available so anyone can reconstruct the L2 state or challenge incorrect state transitions (though ZK-proofs make fraud challenges unnecessary, data availability remains critical for censorship resistance and user exits). Posting raw call data directly to Ethereum L1 calldata is extremely expensive.
- **Pre-EIP-4844: Calldata & Compression:** Before Ethereum's Dencun upgrade (March 2024), ZK-Rollups primarily compressed call data using algorithms like zlib or brotli and posted it within L1 transaction calldata. While better than nothing, costs were still significant, often constituting the majority of L1 fees for rollup users. Projects like **Scroll** implemented sophisticated compression, but the fundamental cost of L1 calldata was a bottleneck.
- **EIP-4844 Proto-Danksharding and Blobs:** The Dencun upgrade introduced **blob-carrying transactions** and **blob data**. This is revolutionary for ZK-Rollups. Instead of expensive calldata, rollups can post batches of compressed call data (and potentially state differences or proofs) as **blobs** – large data packets (~128 KB each) attached to transactions but stored separately by Ethereum consensus nodes for a short period (~18 days). Crucially, blob data is *orders of magnitude cheaper* than equivalent calldata. **Type-2 ZK-EVMs like Polygon zkEVM, Scroll, and Taiko rapidly integrated blob support post-Dencun.** This drastically reduced data publication costs, a major step towards sustainable L2 fee economics. The blob data, while not stored long-term by Ethereum nodes, provides sufficient time for anyone to download it and reconstruct the L2 state if needed, preserving data availability guarantees essential for permissionless exits and verifiability.
- **Storage Proofs and Witness Generation:** For a user to withdraw assets from L2 to L1 without relying on the rollup operator, they need to provide a **storage proof** – cryptographic evidence (like a Merkle branch) demonstrating ownership of funds within the L2 state tree. Generating this proof requires access to the relevant portion of the state tree.
- **The Witness Bottleneck:** During the proving process itself, the **witness** – the set of all inputs, intermediate values, and state accesses needed to satisfy the circuit constraints – must be generated. For

complex EVM executions involving numerous state accesses (common in DeFi interactions), gathering this witness data (reading state from the database, traversing the SMT/MPT) can become a significant performance bottleneck, sometimes exceeding the actual proof generation time. This is distinct from the storage proof needed for exits.

- **Optimizing Witness Gen:** Projects employ several strategies:
- **High-Performance State Databases:** Using optimized key-value stores (like BadgerDB or custom solutions) tuned for fast state read access during execution trace generation.
- **Parallelization:** Generating witness data for independent transactions or parts of transactions concurrently where possible.
- **Caching:** Aggressively caching frequently accessed state elements during block processing.
- **Witness Compression:** Techniques to minimize the size of the witness data passed to the prover circuit. **Scroll** has emphasized optimizations in this area, recognizing its impact on end-to-end proving latency.
- **Off-Chain Proof Marketplaces:** For storage proofs needed by users for exits (not the execution witness), decentralized services like **Herodotus** and **Lagrange** are emerging. These allow users to request proofs of specific storage slots within the L2 state without needing to sync the entire L2 history themselves. The Type-2 ZK-EVM architecture, with its defined state tree structure, enables these services.

1.3.2 3.2 Proving Infrastructure Stack

The prover is the computational powerhouse of a Type-2 ZK-EVM, tasked with the Herculean feat of converting raw EVM execution traces into succinct, verifiable proofs. This stack involves multiple layers of sophisticated engineering, from low-level circuit design to high-level proof aggregation and hardware acceleration.

- **Circuit Design: Taming the EVM Opcode Beast:** The core challenge is designing ZK circuits that accurately represent the behavior of *every* EVM opcode and the system's state (stack, memory, program counter, gas counter) while minimizing circuit size (gates/constraints). Certain opcodes pose disproportionate challenges:
- **Cryptographic Hashes (KECCAK256, SHA256):** As previously discussed, naively implementing these within an arithmetic circuit is infeasible. Type-2 ZK-EVMs universally rely on **lookup arguments** (Plookup, cq, logup, Lasso/Jolt). The prover demonstrates that the inputs and outputs of a hash operation exist within a massive precomputed table of valid (input, output) pairs, drastically reducing circuit constraints. **Polygon zkEVM** and **Scroll** implement highly optimized KECCAK256 circuits using custom Plookup tables within their respective proof systems (Plonky2 and Halo2-KZG). The trade-off is the computational cost of generating these massive tables offline.

- **Precompiles (`ecadd`, `ecmul`, `ecpairing`, `MODEXP`):** These specialized contracts perform complex cryptographic operations. Proving them efficiently requires custom circuit designs tailored to the specific mathematical operation. For example, elliptic curve pairings (vital for zk-SNARKs themselves!) used in `ecpairing` are implemented using optimized circuits leveraging the underlying curve properties (e.g., BN254 for Groth16/PLONK, BLS12-381 for KZG commitments). **Taiko** and others often leverage existing, audited circuit libraries for these precompiles.
- **Memory Operations (`MLOAD`, `MSTORE`):** Simulating the EVM's linear memory within a circuit requires modeling random-access reads and writes. Techniques like **RAM simulation** are used, often involving permutation arguments or specialized memory management circuits that track memory segments and offsets. Gas costs related to memory expansion must also be precisely mirrored.
- **Control Flow (`JUMP`, `JUMPI`):** Handling dynamic jumps (where the jump destination is determined at runtime) requires careful circuit design to ensure the program counter updates correctly based on conditional checks, without introducing vulnerabilities or excessive constraints. Static jumps are simpler.
- **Context Handling (`CALL`, `DELEGATECALL`, `STATICCALL`):** Proving the correct setup and tear-down of call contexts, including stack/memory isolation, gas forwarding, and state access permissions, adds significant complexity. Circuits must model the call stack depth and context switching faithfully.
- **Recursive Proof Aggregation: Scaling the Unscalable:** Proving an entire block of EVM transactions monolithically is computationally intractable. Recursive proof composition is the essential scaling technique.
- **The Process:** Execution is broken down into smaller, provable chunks. This could be per transaction, per block, or even per opcode group within a transaction (though the latter is rare due to overhead). A base proof (SNARK or STARK) is generated for each chunk. These base proofs are then fed into an **aggregation circuit** (or a sequence of them), which generates a single, succinct proof attesting to the validity of *all* the base proofs. The final aggregated proof is what gets verified on L1.
- **Type-2 Implementations:**
 - **Polygon zkEVM (Plonky2):** Leverages its custom Plonky2 system, which combines PLONK arithmetization with FRI (Fast Reed-Solomon IOPP, from STARKs) to achieve extremely fast recursion times. Plonky2 uses a Goldilocks field ($2^{64} - 2^{32} + 1$) optimized for 64-bit CPUs, enabling high-performance recursive proving on commodity hardware.
 - **Scroll (Halo2-KZG):** Utilizes the Halo2 proving system with KZG polynomial commitments. Halo2's core innovation is its efficient recursive proof composition using a technique called "accumulation." Scroll employs a multi-layer aggregation strategy, potentially aggregating proofs per transaction, then per block batch.

- **Benefits:** Recursion dramatically reduces the computational burden on the final prover, enables parallel proof generation, and allows proofs to be generated incrementally as transactions are processed. It also reduces the size of the final proof submitted to L1.
- **Challenges:** Designing efficient aggregation circuits adds complexity. The recursion depth adds latency, though often overlapped with execution. Ensuring the soundness of the recursive composition is critical.
- **Hardware Acceleration: The Race for Speed:** Despite algorithmic breakthroughs, proof generation remains computationally intensive. Specialized hardware is increasingly vital for practical performance.
- **GPUs (Graphics Processing Units):** Highly parallel architectures make GPUs well-suited for the massive number of parallelizable operations involved in witness generation and the finite field arithmetic core to ZK proofs. Libraries like **CUDA** (NVIDIA) and **Metal** (Apple) are used to offload computations from the CPU. Projects like **Scroll** and **Polygon** actively optimize GPU support for their provers, significantly speeding up base proof generation times. Nvidia's pivot towards AI and HPC has inadvertently benefited ZK hardware acceleration.
- **FPGAs (Field-Programmable Gate Arrays):** Offer the potential for even greater performance and efficiency than GPUs by allowing circuits to be literally *hardwired* for specific ZK operations (e.g., MSM - Multi-Scalar Multiplication, NTT - Number Theoretic Transform). Companies like **Ingonyama** and **Cysic** are developing FPGA-based accelerators specifically targeting ZKP workloads. While offering higher performance per watt, FPGAs require significant expertise to program and optimize. They are currently used more in specialized proving farms than general-purpose nodes.
- **The ASIC Horizon:** Application-Specific Integrated Circuits (ASICs) represent the ultimate in hardware acceleration, offering unparalleled performance and efficiency for fixed algorithms. While no major ZK-EVM currently relies on widespread ASIC proving, the economics of high-throughput rollups make ASIC development increasingly plausible, especially for core operations like MSM and NTT. This raises future questions about prover centralization versus efficiency gains. **Taiko's** vision of a permissionless prover market inherently contemplates specialized hardware participants.
- **Benchmarks (Illustrative):** While highly dependent on the specific transaction/block, hardware, and implementation, proof generation times for complex blocks on high-end GPUs can range from minutes to tens of minutes in current Type-2 systems. FPGA accelerators aim to reduce this to seconds or sub-minute times. Recursive aggregation adds additional overhead but enables parallelism.

1.3.3 3.3 The Execution Environment: Modified Clients

The EVM execution itself must be performed in a way that generates a detailed, structured trace suitable for proving. Type-2 ZK-EVMs adopt different strategies for this critical component, balancing compatibility, performance, and auditability.

- **The Forked Client Approach (Geth/Reth):** This path leverages Ethereum’s battle-tested execution clients.
- **Mechanism:** Teams take the existing Go-Ethereum (Geth) or Rust-Ethereum (Reth) client codebase and modify it. The key modification is instrumenting the client to output an **execution trace** – a meticulous, step-by-step log of the EVM’s operation during transaction processing. This trace includes every opcode executed, the state of the stack, memory, storage accesses, program counter, remaining gas, and results. This rich trace is then fed into the prover circuits as the witness.
- **Advantages:**
 - **High Fidelity:** Leverages years of optimization and bug fixes from the Ethereum client ecosystem. Maximizes the likelihood of byte-perfect equivalence with L1 behavior, including obscure edge cases.
 - **Leverage Existing Expertise:** Developers familiar with Geth/Reth can contribute more easily.
- **Disadvantages:**
 - **Complex Instrumentation:** Adding comprehensive tracing without impacting performance or correctness is non-trivial. The trace format must be carefully designed to align with the prover circuit’s expectations.
 - **Performance Overhead:** Generating such a detailed trace adds computational overhead to the execution step itself.
 - **Integration Complexity:** Tightly coupling the modified client with the specific prover architecture can create a monolithic and complex system. **Polygon zkEVM** and **Taiko** exemplify this approach, utilizing forked versions of Geth.
- **The Native Prover-First Client:** This approach builds a new execution client specifically designed from the ground up to be provable.
- **Mechanism:** Instead of modifying a general-purpose client, teams build an **execution client** and a **prover client** in tandem, often sharing core logic. The execution client is designed with ZK proving as a primary constraint, potentially using internal representations or intermediate languages (like zkASM) that map more directly to the prover’s circuits. **Scroll**’s architecture leans towards this model, with their sequencer node (derived from Geth *concepts* but significantly re-architected) working in concert with their separate coordinator and provers utilizing Halo2. While not a direct fork, it aims for the same execution fidelity.
- **Advantages:**
 - **Optimization Potential:** Can be designed for efficient trace generation and tighter integration with the prover stack, potentially improving overall performance.

- **Simpler Proof Circuit Design:** The execution environment can be structured to produce traces that align more naturally with the constraints of ZK proofs, simplifying the circuit logic.
- **Auditability:** A clean-slate design can sometimes be easier to audit for ZK-specific properties.
- **Disadvantages:**
- **Implementation Risk:** Re-implementing the entire EVM correctly from scratch is a massive undertaking fraught with potential for subtle deviations from L1 behavior.
- **Lagging Behind L1:** Keeping pace with Ethereum protocol upgrades (EIPs) and client optimizations requires dedicated effort, whereas forked clients often benefit directly from upstream improvements.
- **Handling Non-Determinism:** The EVM itself is deterministic, but its inputs aren't always perfectly predictable from within the L2 environment.
- **Block Timestamps (**TIMESTAMP** opcode):** On L1, the timestamp is set by the block proposer and is generally considered “trusted” within narrow bounds. On L2, the sequencer sets the timestamp. To prevent manipulation, Type-2 ZK-EVMs typically bind the L2 block timestamp tightly to the L1 block timestamp *in which the ZK-proof is verified*. This ensures the L2 timestamp is derived from L1 consensus, making it as trustworthy as L1 itself. The precise mechanism (e.g., using the L1 timestamp directly, or allowing a small offset controlled by the sequencer but constrained by L1 time) varies.
- **Randomness (**BLOCKHASH**, **DIFFICULTY/PREVRANDAO**):** These opcodes provide historical block hashes and randomness seeds. Their values depend on L1 history. Type-2 ZK-EVMs must provide access to the equivalent L1 block data. This is typically achieved through:
- **On-Chain Oracle Contracts:** Deploying contracts on L1 that expose the necessary historical data (e.g., block hashes).
- **L1->L2 Messaging:** The sequencer injects the relevant L1 data (like **PREVRANDAO**) into the L2 block via the cross-chain messaging bridge (see 3.4). The ZK-proof then verifies that this injected data corresponds to the actual state of the referenced L1 block, using **storage proofs** or **block header proofs**. This ensures the randomness fed into the L2 EVM is authentic and verifiable.
- **ORIGIN and CALLER:** These are deterministic based on the transaction signer and call stack and are handled internally by the execution client without special ZK considerations.
- **Gas Metering Adaptations:** While Type-2 aims for identical gas *semantics*, the actual *cost* of proving introduces an economic reality.
- **ZK Overhead Costs:** The computational resources (CPU/GPU/FPGA time, memory) consumed by proof generation are substantial and not captured by the standard EVM gas model. Type-2 ZK-EVMs address this by:

- **L2 Transaction Fees:** Users pay fees denominated in ETH (or sometimes the rollup's native token) on L2. This fee has two main components: 1) **L2 Execution Fee:** Covers the cost of executing the transaction on the sequencer node, calculated similarly to L1 gas but often at a much lower base fee. 2) **L1 Data/Proof Fee:** Covers the cost of publishing the transaction data (call data) to L1 (now primarily via blobs) and the cost of verifying the ZK-proof on L1. The L1 portion is usually the dominant cost for simple transactions.
- **Prover Incentives:** The provers (whether centralized or part of a decentralized network) are compensated from these fees, specifically the portion covering the L1 proof verification cost and their compute resources. **Taiko's** model explicitly includes rewards for provers in its tokenomics.
- **Gas Accounting Within Proof:** Crucially, the *EVM gas accounting within the execution trace must still perfectly mirror L1 behavior* for Type-2 equivalence. The proof verifies that gas was consumed correctly according to the EVM rules. The economic cost of *proving* that gas consumption is handled separately via the L2 fee mechanism described above.

1.3.4 3.4 Bridge Architecture and Cross-Rollup Messaging

A Type-2 ZK-EVM is not an island; it must interact securely and trust-minimizedly with Ethereum L1 and potentially other L2s. The bridge architecture facilitates the movement of assets and data, while cross-rollup messaging enables composability across the scaling ecosystem.

- **Secure L1L2 Messaging Protocols:** The canonical bridge is the primary, protocol-defined path between L1 and L2.
- **Deposits (L1 -> L2):**
 1. User locks assets (ETH, ERC-20 tokens) in a bridge contract on L1.
 2. This deposit event is emitted as an Ethereum log.
 3. The L2 sequencer monitors L1. Upon detecting the deposit event, it *mints* a corresponding representation of the asset on L2 to the user's L2 address. Critically, the ZK-proof for the L2 block containing this mint operation includes a verification (via a Merkle proof or similar) that the deposit event was indeed included and valid on L1. The user receives their L2 funds typically within minutes (sequencer inclusion) with finality achieved once the block's ZK-proof is verified on L1 (minutes to hours).
- **Withdrawals (L2 -> L1):** This is where ZK-Rollups shine compared to Optimistic Rollups.
 1. User initiates a withdrawal by burning or locking their assets in a bridge contract *on L2*.
 2. This withdrawal request is included in an L2 block.

3. The ZK-proof for this block, once generated and verified on L1, *cryptographically attests* to the validity of the withdrawal request and the user's entitlement.
 4. After the proof is verified on L1 (typically taking minutes to hours), the user can finalize the withdrawal by submitting a small transaction to the L1 bridge contract. This transaction provides minimal data (often just a Merkle proof of the withdrawal inclusion in the proven L2 state root) which the L1 contract can instantly verify against the already-accepted state root. **This enables near-instant, trust-minimized withdrawals**, contrasting sharply with Optimistic Rollup's 7-day challenge window. Projects like **Scroll** and **Polygon zkEVM** implement variations of this pattern, sometimes involving a "withdrawal hole" or merkle mountain ranges for efficient proof inclusion.
- **Message Passing:** Beyond simple asset transfers, the bridge supports arbitrary data messages (e.g., triggering a function on L1 from L2 or vice-versa). The mechanism is similar:
 - **L1 -> L2:** Message sent via L1 contract, sequencer picks it up and delivers it to the target L2 contract (handled within L2 execution, proven in ZK-proof).
 - **L2 -> L1:** Message initiated on L2 (e.g., via a `sendMessage` call). The message inclusion is proven in the ZK-proof. Once the proof is verified on L1, the message can be relayed to its target L1 contract. A slight delay (minutes/hours) exists due to proof generation/verification.
 - **Trust-Minimized Withdrawal Mechanisms:** The architecture described above minimizes trust. Users don't need to trust the sequencer or provers for withdrawals. They only need:
 1. **Data Availability:** Access to the L2 block data (now secured via blobs) to construct their Merkle proof for the withdrawal.
 2. **L1 Security:** Trust that Ethereum L1 is secure and honestly executes the bridge contract code verifying the ZK-proof and the Merkle proof. This is the same trust assumption as using Ethereum itself. **Forced Inclusion** mechanisms, often implemented via L1 inbox contracts, allow users to directly post transactions to L1 that must be included in the next L2 block, providing censorship resistance for initiating withdrawals if the sequencer is malicious or offline.
 - **Native Integration with Ethereum's Beacon Chain:** As Ethereum transitions fully to Proof-of-Stake (PoS), Type-2 ZK-EVMs leverage the Beacon Chain for enhanced security and efficiency:
 - **Withdrawals to Execution Layer:** Post-Merge and post-Shanghai/Capella, withdrawals from the Beacon Chain (staking rewards, exited validator stakes) are processed on the Execution Layer. Type-2 ZK-EVM bridges handle these like any other L1 asset, allowing staked ETH to flow smoothly to L2.
 - **L1 Block Attributes:** As mentioned in 3.3, L2 timestamps and randomness (`PREVRANDAO`) are often derived directly from the attributes of the L1 block (proposed by a Beacon Chain validator) in which the ZK-proof is verified. This deepens the integration.

- **Future Synergy (Verkle Trees, Danksharding):** Upcoming Ethereum upgrades like Verkle Trees (replacing the MPT for state storage) and Danksharding (full implementation of data blobs) are being designed with L2 scaling in mind. Type-2 ZK-EVMs are actively planning adaptations, such as utilizing Verkle proofs for more efficient L1 state access within their circuits or fully leveraging the massive bandwidth of Danksharding blobs. This co-evolution ensures Type-2 ZK-EVMs remain tightly integrated with Ethereum’s core roadmap.
- **Cross-Rollup Messaging (Layer 3 & Beyond):** While the L1 bridge connects to Ethereum, users and dApps also need to move assets and data *between different ZK-Rollups* (L2s) efficiently.
- **The Challenge:** Direct bridges between L2s reintroduce security risks and liquidity fragmentation. Relying solely on L1 as a hub is slow and expensive.
- **ZK-Powered Solutions:** Emerging protocols leverage ZK technology itself to create trust-minimized bridges between rollups:
- **Shared Proving / Settlement Layers:** Projects like **Nebra** and **Avail** envision networks where multiple “sovereign” or “settlement” rollups (potentially Type-2 ZK-EVMs) share a common data availability layer and can leverage ZK-proofs to verify state transitions *of each other* or on a shared settlement layer, enabling fast, secure cross-rollup transfers without always going through L1. **Taiko’s** based sequencing inherently positions L1 as this shared anchor point.
- **ZK Light Client Bridges:** A rollup (L2A) generates a ZK-proof attesting to the inclusion of a message or state root on *another rollup* (L2B). A light client contract for L2B deployed on L1 (or potentially on L2A itself) verifies this proof, enabling trust-minimized verification of L2B’s state on L2A. **Succinct Labs** and **Polyhedra Network** are pioneering this approach using advanced ZK proof systems (like SP1) capable of efficiently proving consensus light client verification. **Scroll** has explored integrations with such bridges.
- **Standardization Efforts:** Initiatives like the **L2 Standards Proposals** fostered by the L2BEAT team aim to define common interfaces and security models for cross-chain communication, including between ZK-Rollups, promoting interoperability within the Type-2 ecosystem.

The architectural innovations explored in this section – from state tree adaptations and blob utilization to recursive proving stacks, client modifications, and ZK-secured bridges – collectively form the intricate machinery that makes the Type-2 ZK-EVM vision operational. They represent a continuous balancing act, striving to preserve Ethereum’s essence while bending the rules of computational feasibility through cryptography and ingenuity. Yet, understanding the static structure is only part of the story. The true marvel lies in the dynamic process – the journey a transaction takes from initiation on L2 to the moment its validity is etched immutably onto Ethereum L1 through the power of a zero-knowledge proof. [Transition to Section 4: This intricate machinery springs to life with each user transaction. The path from a signed L2 transaction bundle to a verified proof on Ethereum L1 involves a meticulously choreographed sequence of steps –](#)

sequencing, execution, witness generation, circuit proving, aggregation, and final verification – each stage presenting its own performance bottlenecks and optimization frontiers...

1.4 Section 4: The Proving Process: From Execution to Verification

The intricate machinery of Type-2 ZK-EVMs, meticulously architected to balance Ethereum equivalence with cryptographic efficiency, springs to life with every transaction. This section dissects the operational heartbeat of these systems—the meticulously choreographed sequence transforming a user’s transaction into an immutable cryptographic truth on Ethereum L1. Unlike traditional blockchains where execution alone suffices, Type-2 ZK-EVMs add layers of cryptographic ceremony: generating witnesses, compiling constraints, generating proofs, and verifying them on-chain. Each stage introduces unique latency and optimization challenges, creating a dynamic interplay between user experience, security, and scalability.

1.4.1 4.1 Transaction Flow: User Tx to L1 Finality

The journey begins with user interaction and culminates in Ethereum-finalized validity. This lifecycle reveals how Type-2 ZK-EVMs reconcile real-time responsiveness with the computational heaviness of zero-knowledge proofs.

- **Sequencing & Mempool Management: The Gateway to L2:**
- **User Submission:** A user signs a transaction (e.g., a Uniswap swap) using their EOA or smart contract wallet, specifying gas parameters familiar from Ethereum L1. The tx is broadcast to the rollup’s peer-to-peer mempool network.
- **Sequencer Role:** The sequencer node (centralized in early implementations, decentralized in visions like Taiko’s) acts as the L2 block producer. It collects transactions from the mempool, orders them (applying MEV strategies if permitted), and assembles them into an *L2 block candidate*. Crucially, it must manage transaction **nonce ordering** and **fee prioritization** identically to Ethereum L1 to maintain equivalence. **Scroll** uses a Geth-derived sequencer that mirrors Ethereum’s mempool logic.
- **Real-Time Execution & Soft Confirmations:** The sequencer executes transactions *immediately* upon inclusion in the pending block. Users receive near-instant “soft confirmation” (similar to L1 txpool inclusion), allowing dApps like DEXs to update UIs provisionally. The sequencer publishes critical data (tx hashes, state diffs) to allow quick state reconstruction. This mimics Ethereum’s user experience despite the backend complexity.
- **Data Availability Commitment:** Simultaneously, the sequencer compresses the block’s transaction data (call data) and posts it to Ethereum L1 as a **blob** via EIP-4844. For example, **Polygon zkEVM**

batches ~60-100 transactions per blob, reducing L1 costs by ~95% compared to pre-blob calldata. This ensures anyone can reconstruct L2 state independently.

- **Witness Generation: Capturing the EVM’s Soul:**

- **Trace Extraction:** Post-execution, the sequencer (or dedicated trace generator) emits a granular **execution trace**. This logs every opcode executed, stack/memory states, storage accesses, gas consumption, and results. For a complex DeFi transaction (e.g., a multi-hop swap via 1inch), this trace might span millions of steps. **Polygon zkEVM’s** forked Geth client instruments opcode execution to output this structured trace.
- **State Witness:** Alongside execution, the system records **witness data**—Merkle proofs or SMT paths proving the pre/post state of every storage slot accessed (e.g., Uniswap pool balances). This links execution to the authenticated state tree.
- **Bottleneck Management:** Witness generation is I/O-intensive. **Scroll** optimizes this by parallelizing trace extraction across transactions and using memory-mapped databases for low-latency state reads. A single complex tx can take seconds to witness, highlighting why parallelization is critical.
- **Circuit Compilation & Proof Batching:**
 - **From Trace to Constraints:** The raw trace isn’t proof-ready. It’s transformed into a **constraint system**—mathematical equations representing correct EVM behavior (e.g., “if opcode is ADD, then $\text{stack}[i+1] = \text{stack}[i] + \text{stack}[i-1]$ ”). This leverages predefined circuit blueprints for each opcode (Section 3.2).
 - **Batching Strategies:** Proving transactions individually is inefficient. Rollups batch 10s-100s of txs into a single proof:
 - **Time-Based Batching (e.g., Taiko):** Generates a proof every fixed interval (e.g., 3-5 minutes). Prioritizes predictable finality but risks underutilization during low activity.
 - **Size-Based Batching (e.g., Polygon zkEVM):** Triggers proving when a batch reaches a gas or byte threshold (e.g., 30M gas). Maximizes hardware utilization but causes variable finality delays.
 - **Hybrid Approaches:** **Scroll** uses size-based batching with fallback timeouts to ensure liveness. Batch sizes are tuned to balance L1 verification gas (larger batches amortize costs) and proving latency.

1.4.2 4.2 Inside the Prover: Circuit Execution Details

The prover—often GPU-clad servers in data centers—transforms the constraint system into a succinct proof. This stage demands cryptographic ingenuity to handle the EVM’s complexity.

- **Arithmetization: Traces into Polynomials:**

- **Constraint Conversion:** The prover converts the constraint system (e.g., “ $x + y = z$ ”) into polynomials over a finite field. Using PLONK or Halo2, it embeds execution variables into polynomial coefficients.
- **Lookup Arguments for Non-Arithmetic Ops:** For KECCAK256 or memory accesses, the prover uses **lookup tables**. Instead of arithmetizing a full hash, it proves $(input, output)$ pairs exist in a precomputed table (e.g., all 256-bit inputs mapped to Keccak outputs). **Polygon zkEVM’s** Plonky2 uses custom lookup gates for this, reducing KECCAK constraints by $>1000\times$.
- **RAM Simulation for Memory:** The EVM’s memory is modeled as a virtual RAM. The prover uses **permutation arguments** to ensure memory writes (MSTORE) are read correctly later (MLOAD). For example, proving that for address `addr`, the value `val` written at step `i` is the same read at step `j`.
- **State Access Handling: Proving Authenticity:**
- **SMT/MPT Witness Integration:** For each SLOAD/SSTORE, the prover integrates the SMT path witness from trace generation. It cryptographically links the opcode’s storage slot to the global state root within the circuit.
- **Consistency Checks:** Circuits enforce that the starting state root matches the previous block’s proven root, and the final root is computed correctly from state changes. This chaining creates an immutable history.
- **Stack/Memory Optimizations:**
- **Stack Compression:** The EVM stack (1024-item LIFO) is represented compactly. Provers track only *deltas* (pushes/pops) between opcodes rather than the full stack state at each step.
- **Memory Chunking:** Memory is segmented. Accesses within a contiguous chunk (e.g., 32-byte words for a `uint256[]`) share a single range-check proof, reducing constraints.
- **Recursive Proving Workflow:**
 1. **Base Proofs:** The prover splits the batched execution trace into chunks (e.g., per transaction). Each chunk generates a “base proof” using GPU acceleration.
 2. **Aggregation:** Base proofs are fed into a **recursive aggregation circuit**. This circuit verifies each base proof’s validity and outputs a single “proof of proofs.” **Scroll’s** Halo2 prover uses a binary tree aggregation, halving proof count at each layer.
 3. **Final Proof:** The root aggregation proof (5-50 KB) is submitted to L1. Recursion reduces verification gas by $\sim 80\%$ vs. a monolithic proof.

1.4.3 4.3 Verification on Ethereum L1

The final proof's verification on Ethereum L1 anchors L2 security to Ethereum's consensus. This stage demands extreme gas efficiency.

- **On-Chain Verifier Contracts:**

- **Minimalist Design:** Verifier contracts are highly optimized, often written in Yul or inline assembly. They perform fixed computations: pairing checks (for SNARKs), polynomial evaluations, or hash verifications.
- **Gas Cost Dynamics:** Verification costs scale with proof system complexity:
- **Groth16/PLONK:** ~500K gas (heavy on elliptic curve pairings).
- **Halo2-KZG:** ~300K gas (lighter pairings via KZG commitments).
- **STARKs:** ~2M+ gas (higher due to hash-based verification), making them rare for L1 verification in Type-2 systems.
- **Amortization:** Batching reduces per-tx cost. A batch of 100 txs verified for 300K gas costs just 3K gas/tx. EIP-4844 blobs further reduce data costs.

- **Proof Aggregation Tradeoffs:**

Method	Proof Size	Verification Gas	Proving Latency	Use Case
Monolithic	1.5-5 MB	~1-5M gas	Lower	Small batches (rare)
2-Layer Recursion	10-100 KB	~200-500K gas	Medium	Polygon zkEVM, Scroll
Deep Recursion	1-10 KB	~50-200K gas	Higher	Research (e.g. Nova)

- **Single Proof:** Simple but impractical for large blocks (high gas).
- **Proof-of-Proofs (Recursive):** Preferred for scalability. **Polygon zkEVM's** Plonky2 uses FRI-based recursion for 45 KB proofs verifying in ~300K gas.
- **Upgradeability & Governance:**
- **Verifier Upgrades:** Bug fixes or efficiency improvements require new verifier contracts. Upgrades are managed via:
- **Multisig:** Early systems (Polygon, Scroll v0) used 5/8 multisigs. Fast but centralized.
- **Timelock + Governance:** **Taiko** uses a DAO vote with 7-day timelock. Verifier changes are transparent but slower.

- **Security Implications:** A malicious verifier could approve invalid state. Audits (e.g., by Zellic, OpenZeppelin) and formal verification (e.g., Scroll's KZG verifier in Dafny) mitigate risks.

1.4.4 4.4 Latency and Finality Characteristics

Latency—the delay from transaction submission to L1 finality—is a key differentiator between ZK and optimistic rollups.

- **Proof Generation Times:**
- **Hardware Dependence:** Proving times vary dramatically:
 - **CPU (High-End):** 10-60 min/batch (impractical for production).
 - **GPU (A100/H100):** 2-10 min/batch (standard for Polygon, Scroll).
 - **FPGA (Custom):** 30 sec - 2 min/batch (Ingonyama ICICLE, Cysic ZK accelerators).
- **Workload Variance:** A batch of simple transfers may prove in 1 minute; a complex Arbitrum-style fraud-proof simulation could take 10+ minutes. **Scroll's** benchmarks show ~90 sec for 150 simple txs on 8xA100 GPUs.
- **Recursion & Time-to-Finality:**
- **Pipeline Overlap:** While proof generation adds latency, user impact is mitigated:
 1. **Sequencer Soft Finality:** Instant upon inclusion (user sees success).
 2. **L1 Finality:** Achieved when the proof is verified on L1 (~2-20 min post-batch execution).
- **Recursion Overhead:** Aggregation adds 15-30% to proving time but reduces verification gas. Deep recursion (e.g., 8 layers) adds minutes but enables tiny proofs.
- **Comparative Finality Lags:**

System | Soft Finality | L1 Finality | Withdrawals |

|—————|—————|—————|—————|

Type-2 ZK-EVM | 1-5 sec | 2-20 min | Minutes-hours |

Optimistic Rollup | 1-5 sec | 7 days (challenge win) | 7 days + 1 hour |

Ethereum L1 | 12-60 sec | 12-60 sec (PoS) | N/A |

Type-4 ZK-Rollup (zkSync) | <1 sec | 10-60 min | Minutes-hours |

- **ZK vs. Optimistic:** ZK-Rollups offer 50,000x faster withdrawals. For high-value DeFi (e.g., Maker-DAO vaults), this eliminates liquidity fragmentation risk.
 - **Type-2 vs. Type-4:** Type-4 (zkSync) achieves slightly faster proving (1-5 min) by sacrificing EVM equivalence but cannot run unmodified L1 dApps.
 - **Real-World Impact:**
 - **DeFi Arbitrage:** Bots on Polygon zkEVM achieve ~3-min finality, enabling cross-DEX arbitrage competitive with Solana (~400ms) but with Ethereum security.
 - **NFT Minting:** High-throughput mints (e.g., 10,000 NFTs) queue proving but show instant soft confirmation. Blobs prevent L1 congestion.
 - **Bridging:** Across Rollups: ZK light client bridges (e.g., Polyhedra zkBridge) use proofs to sync states in <10 min vs. 7 days for optimistic bridges.
-

The proving process transforms Ethereum’s execution model into a cryptographic artifact, trading temporal immediacy for unprecedented security guarantees. Where optimistic rollups defer trust to a future challenge window, Type-2 ZK-EVMs offer mathematical certainty at the cost of minutes-long latency—a testament to the intricate dance between cryptographic rigor and user experience. This latency, however, is not static. As recursive proof systems deepen, hardware accelerates, and Ethereum’s base layer evolves (Danksharding, Verkle trees), the gap between soft confirmation and L1 finality narrows, inching toward near-instantaneous, Ethereum-secured finality.

Transition to Section 5: The theoretical elegance of the proving process confronts the messy reality of implementation. Across the landscape of Type-2 ZK-EVMs—Polygon’s pioneering compromises, Scroll’s bytecode purism, Taiko’s radical decentralization—divergent engineering choices yield starkly different performance profiles, security tradeoffs, and developer experiences. Evaluating these contenders reveals how the “holy grail” of EVM equivalence manifests in practice...

1.5 Section 5: Comparative Analysis of Major Implementations

The intricate dance of zero-knowledge proofs and EVM execution, meticulously detailed in the preceding section, manifests uniquely across the landscape of operational Type-2 ZK-EVMs. Where theoretical models strive for idealized equivalence, real-world engineering confronts the gritty constraints of performance, decentralization, and developer pragmatism. This section dissects the leading contenders – Polygon zkEVM, Scroll, and Taiko – alongside emerging forces like Consensys Linea and the Starknet-based Kakarot, evaluating how their divergent architectures, philosophical priorities, and operational choices translate into tangible

strengths, compromises, and real-world performance. The “holy grail” of Type-2 fidelity reveals itself not as a monolithic endpoint, but as a spectrum navigated through distinct, sometimes contentious, engineering pathways.

5.1 Polygon zkEVM: First-Mover Complexities

As the first major project to launch a public mainnet billed as “EVM-equivalent” (March 2023), Polygon zkEVM occupies a pivotal, albeit complex, position. Its journey embodies the challenges of pioneering Type-2 equivalence at scale, marked by technical ambition, pragmatic compromises, and evolving performance.

- **Forked Geth Client: Leverage and Liability:**
- **Implementation Strategy:** Polygon adopted the **forked Geth client** approach (Section 3.3). Their zkNode is a heavily modified Go-Ethereum client instrumented to emit detailed execution traces compatible with their custom prover stack. This aimed to leverage Geth’s battle-tested EVM fidelity and accelerate development.
- **Challenges:** Integrating deep ZK-specific tracing into Geth proved arduous. Geth’s codebase, optimized for speed and L1 execution, wasn’t designed for granular trace emission. Instrumenting every opcode and state access introduced **performance overhead** and **increased code complexity**, making it harder to track upstream Geth improvements. Subtle bugs in the instrumentation layer, potentially diverging from L1 behavior in edge cases, became a persistent audit focus (e.g., Quanstamp audit QSP-9 identified trace serialization issues). Maintaining byte-perfect equivalence while modifying core Geth logic remains an ongoing engineering burden.
- **Tradeoff:** While providing strong execution fidelity, the forked client creates a **monolithic architecture** less flexible than native prover-first designs. Updates require synchronizing changes across the modified client and the prover circuits that consume its traces.
- **Custom Plonky2 Proof System: Performance vs. Ecosystem:**
- **Architectural Choice:** Polygon Zero developed **Plonky2**, a recursive SNARK combining PLONK’s arithmetization with FRI (Fast Reed-Solomon Interactive Oracle Proofs) for efficient recursion. It uses a Goldilocks field ($p = 2^{64} - 2^{32} + 1$) optimized for 64-bit CPUs, enabling fast proving on commodity hardware without specialized elliptic curves.
- **Tradeoffs:**
- **Pros:** Exceptional recursion speed. Benchmarks showed Plonky2 aggregating proofs ~10x faster than early Halo2 implementations. Its field design simplifies certain cryptographic operations. GPU acceleration was integrated relatively smoothly.
- **Cons: Ecosystem Immaturity.** Plonky2 is a bespoke system. Compared to the more widely adopted Halo2 (used by Scroll, Taiko) or Circom, it has a smaller developer ecosystem, fewer auditing tools, and less battle-tested cryptographic libraries. This potentially increases **long-term maintenance risk**.

and slows adoption by external prover networks. Its reliance on FRI, while fast, produces slightly larger proofs than KZG-based systems, impacting L1 verification gas costs (~300-400k gas per aggregated batch proof vs. ~250-300k for optimized Halo2-KZG).

- **Proof Market Evolution:** Polygon initially relied heavily on its own proving infrastructure. While moving towards decentralization, its prover network adoption faces steeper hurdles compared to systems built on more established proof frameworks.
- **Real-World Performance & The Type-2.5 Pragmatism:**
- **Post-Mainnet Launch Metrics:**
- **Proof Times:** Initially, batch proving times on high-end GPUs (e.g., A100) ranged from 5-20 minutes for moderate-sized batches (~50-100 typical txs), heavily dependent on transaction complexity. Complex swaps involving multiple KECCAK ops or storage-heavy operations significantly increased times. Aggressive optimization and GPU parallelization have reduced average times to ~3-8 minutes.
- **Throughput:** Sustained TPS is constrained by proving latency and L1 proof verification frequency. Polygon zkEVM consistently demonstrates 10-20 TPS on mainnet during peak loads, significantly higher than L1 but below some Type-4 systems. EIP-4844 blobs drastically reduced data costs, improving cost efficiency.
- **The Gas Cost Compromise:** Faced with the exorbitant cost of proving certain EVM operations (especially KECCAK256 and frequent SSTOREs) with perfect L1 gas equivalence, Polygon made a pivotal decision in **mid-2023**. They publicly reclassified as a **Type 2.5 ZK-EVM**, maintaining full opcode and bytecode equivalence but **modifying the gas costs** for specific operations. For example:
 - KECCAK256: Increased gas cost significantly vs. L1 (e.g., ~700 gas vs L1's 30+6*words, a >20x increase for small inputs).
 - SSTORE for initializing non-zero slots: Increased cost.
 - BALANCE, EXTCODESIZE: Slight increases.
- **Impact:** This pragmatic shift dramatically improved prover economics and batch finality times, making the system viable long-term. However, it introduced a subtle deviation: contracts consuming predictable gas on L1 *might* run out of gas on Polygon zkEVM if they heavily use penalized opcodes. This necessitates careful gas estimation during deployment, a friction point purist Type-2 aims to eliminate. The community debate highlighted the tension between ideal equivalence and operational reality.
- **Ecosystem Traction:** Despite complexities, Polygon zkEVM boasts significant adoption. Major DeFi protocols (Aave V3, Uniswap V3, Balancer) and gaming/NFT platforms (Immutable, Planet IX) have deployed, leveraging its robust infrastructure and Polygon's broader ecosystem support. Its first-mover advantage and Polygon's aggressive BD efforts have driven substantial TVL and transaction volume.

5.2 Scroll: Community-Driven Bytecode Focus

Emerging from Ethereum research roots with a strong open-source ethos, Scroll has positioned itself as the purist champion of **unmodified EVM bytecode execution** within the Type-2 paradigm. Its architecture and community model prioritize rigorous equivalence and seamless developer integration.

- **Emphasis on Unmodified EVM Bytecode Execution:**
- **Core Philosophy:** Scroll's defining principle is executing the *exact* bytecode produced by the standard Ethereum Solidity and Vyper compilers. They avoid the forked client approach, instead building a **native zkEVM node** comprising a Sequencer (derived from Geth concepts but re-implemented in Rust) and a Coordinator. This node outputs execution traces specifically designed for their prover.
- **Technical Realization:** Their zkEVM utilizes a **custom bytecode interpreter** layer that directly processes standard EVM opcodes. This interpreter translates EVM execution into a sequence of lower-level steps provable within their Halo2-based circuits. The focus is on minimizing any abstraction or translation that could introduce divergence. Rigorous testing against Ethereum's execution specification tests (e.g., from `ethereum/tests`) is paramount. Audits (e.g., by Zellic) have focused intensely on bytecode-level equivalence, identifying and fixing subtle deviations in precompile behavior or edge-case opcodes early.
- **Advantage:** Offers the strongest guarantee that contracts deployed on L1 *will* behave identically on Scroll, crucial for complex, security-sensitive protocols like decentralized stablecoins (e.g., potential future deployments akin to MakerDAO) or intricate DAO tooling.
- **Integration with Existing Toolchains (Hardhat, Foundry):**
- **Developer Experience Priority:** Recognizing that equivalence is meaningless without a frictionless workflow, Scroll invested heavily in **first-class tooling integration** from the outset.
- **Hardhat Plugin:** The `hardhat-scroll` plugin allows developers to compile, deploy, test, and verify contracts on Scroll using identical Hardhat commands and scripts as on Ethereum mainnet or Sepolia. It handles network configuration, contract verification on Scroll's block explorer, and faucet access seamlessly.
- **Foundry Integration:** Similarly, Foundry (`forge`) works natively with Scroll. Developers can deploy using `forge create --rpc-url https://sepolia-rpc.scroll.io`, run tests with `forge test`, and interact with contracts via `cast`, mirroring the L1 experience. This deep integration significantly lowers the barrier to entry for Ethereum-native developers.
- **Debugging Tools:** Scroll provides specialized tools like a **zkEVM tracer** that helps debug failed transactions by visualizing the step-by-step EVM execution trace *within the ZK context*, a critical aid for diagnosing issues unique to the proving environment.
- **Open-Source Governance Model Analysis:**

- **Commitment to Transparency:** Scroll launched with a strong commitment to **full open-source development**. Their entire stack – sequencer, coordinator, prover circuits (in Halo2), contracts, and tooling – is publicly available on GitHub. This fosters community trust, enables independent audits, and encourages contributions.
- **Governance Structure:** Governance is currently managed by the Scroll Foundation and core development team, with a clear roadmap towards progressive decentralization. Key decisions (like protocol upgrades) are discussed transparently in community forums and on GitHub. The **Scroll Citizens Program** rewards community contributions to documentation, tooling, and infrastructure.
- **Prover Decentralization:** Scroll is pioneering a **permissionless prover network**. Anyone can run a prover node, contributing compute to generate proofs for batches. Provers are rewarded from sequencer fees. While early stages involve whitelisting for stability, the architecture is designed for open participation, distributing proving power and enhancing censorship resistance. This contrasts with more centralized initial proving setups elsewhere.
- **Tradeoffs:** Full open-source development can slow down rapid iteration compared to closed teams. Coordinating a decentralized prover network adds complexity to batch finality scheduling. However, the model builds strong community trust and aligns with Ethereum’s values.
- **Real-World Performance & Focus:**
- **Proof Times & Throughput:** Utilizing Halo2-KZG and aggressive GPU optimization, Scroll achieves competitive proving times. Benchmarks on their testnet and early mainnet showed batch proving times of ~2-6 minutes on GPU setups for typical loads. Sustained TPS mirrors Polygon zkEVM in the 10-20 range, constrained by similar factors. Their focus on equivalence means they haven’t adopted Type-2.5 gas changes, accepting potentially higher proving costs for specific workloads to maintain purity.
- **Adoption Traction:** Scroll’s mainnet launched later (October 2023) but has seen steady growth. Its developer-friendly tooling and equivalence focus attract protocols prioritizing security and seamless migration, such as derivatives platforms (Synthetix v3 deployment), lending protocols, and infrastructure projects (Chainlink oracles, The Graph indexing). Its community-driven ethos resonates strongly within the Ethereum developer base.

5.3 Taiko: Based Rollup Approach

Taiko, founded by Loopring veterans, challenges conventional rollup architecture with its radical **based sequencing** model. While its ZK-EVM execution aims for Type-2 equivalence (using a forked Geth client), its core innovation lies in decentralization at the sequencing layer.

- **Unique “Based Sequencing” Eliminating Centralized Sequencers:**
- **Mechanism:** Taiko fundamentally rethinks sequencing. Instead of dedicated L2 sequencers (centralized or decentralized pools), Taiko transactions are **proposed directly on Ethereum L1**. Ethereum

block proposers (validators) include Taiko transactions within their L1 blocks via a special `TaikoL1.proposeBlock` transaction. These proposals contain compressed L2 transaction data.

- **How it Works:**

1. Users send L2 txs to a mempool.
2. Ethereum block proposers (validators) *choose* to include a bundle of these L2 txs in their next L1 block via the `proposeBlock` call to the Taiko L1 contract. They act as **L2 block proposers**.
3. The proposed L2 block is temporarily pending.
4. Provers generate ZK-proofs for the proposed block.
5. Once a valid proof is submitted and verified on L1, the L2 block is finalized.

- **Advantages:**

- **Maximal Censorship Resistance:** Inherits Ethereum L1’s censorship resistance directly. If an L2 tx is valid and has sufficient fee, *some* L1 proposer will include it to capture fees, as it’s just another L1 tx type. No L2 sequencer can censor.
- **Liveness Guaranteed by Ethereum:** Relies on Ethereum’s underlying liveness. No separate sequencer failure mode.
- **Simplified Trust Model:** Minimizes trust in L2-specific actors; relies solely on Ethereum validators and the cryptographic security of ZKPs.

- **Challenges:**

- **Latency:** L2 block proposal is tied to L1 block time (12 seconds), adding inherent latency compared to dedicated L2 sequencers offering instant soft confirms. Users experience “L1 confirmation latency” before their tx is even *proposed*.
- **Throughput Bottleneck:** Throughput is capped by the gas limit and frequency of `proposeBlock` calls within L1 blocks. While EIP-4844 blobs help, it cannot match the potential TPS of rollups with dedicated high-speed sequencers. Requires careful gas optimization of the proposal mechanism.
- **MEV Management Complexity:** Distributing block proposal rights across all Ethereum validators complicates MEV extraction and distribution compared to rollups with specialized sequencers or prover-builder separation (PBS).
- **Multi-Proof System (SGX + ZK) Security Model:**
- **Hybrid Approach:** To enhance security and potentially offer faster interim confirmations, Taiko employs a **multi-proof system** in its early stages:

- **ZK-Proofs:** The gold standard, providing cryptographic validity of the entire L2 state transition. Generated by permissionless “ZK Provers.”
- **SGX Attestations:** Generated by “Block Verifiers” running within Intel SGX (Software Guard Extensions) Trusted Execution Environments. These attestations cryptographically guarantee that the verifier node *executed the proposed L2 block correctly* and produced a valid state root. This happens faster than full ZK-proof generation.
- **Security & Trust Assumptions:**
 - **ZK Proofs:** Provide unconditional cryptographic security once verified on L1.
 - **SGX Attestations:** Provide strong *computational integrity* guarantees *if* the SGX hardware and its remote attestation mechanism are secure. This adds a layer of trust in Intel and the absence of hardware vulnerabilities. Malicious SGX verifiers could attest to invalid state roots, but they cannot steal funds as withdrawals require the ZK-proof.
 - **Grace Period:** Disputes between SGX attestations and eventual ZK-proofs can be resolved, with slashable bonds punishing malicious actors. The system is designed to fall back to ZK-proof security.
 - **Evolution:** Taiko’s roadmap phases out SGX reliance as ZK-proving performance improves, moving towards a pure ZK security model. The multi-proof system is a pragmatic bridge.
- **Economics of Permissionless Proving Markets:**
 - **Native Token (TKO):** The Taiko token is integral to its economic model:
 - **Prover Incentives:** ZK Provers stake TKO and earn rewards (in ETH/TKO) for generating valid proofs. Higher staked amounts potentially correlate with higher chances of being assigned proving work.
 - **Proposer Rewards:** Ethereum validators (L1 block proposers) who include Taiko L2 blocks via `proposeBlock` earn fees (in ETH), incentivizing participation.
 - **Protocol Fees:** A portion of L2 transaction fees is potentially burned or directed to a DAO treasury, creating deflationary pressure or funding development.
 - **Permissionless Participation:** Anyone can run a prover node (with sufficient hardware) and participate in the proving market by staking TKO. This aims to decentralize proof generation over time.
 - **Challenges:** Designing efficient markets for proof assignment, preventing centralization of proving power among large stakers or specialized hardware farms, and ensuring timely proof generation despite variable market conditions are complex open problems Taiko is actively tackling.
- **Real-World Status:** Taiko launched its “Katla” testnet (Alpha-3) with based sequencing and multi-proofs in late 2023. Its mainnet (“Eldfell L2”) launched in May 2024. Performance metrics are nascent

but reflect the architectural constraints: L2 block proposal latency tied to L1, proving times similar to peers (minutes), and initial TPS constrained by the L1 proposal mechanism. Its unique decentralization model attracts significant interest but faces practical scaling hurdles.

5.4 Emerging Contenders: Kakarot, Linea, Consensys

Beyond the established trio, the Type-2 landscape is vibrant with innovative approaches striving for compatibility, performance, or unique integrations.

- **Starknet’s Cairo-Native Kakarot zkEVM:**
- **Concept:** Kakarot is not a standalone L2. It’s an **EVM bytecode interpreter written in Cairo**, Starknet’s native smart contract language. Deployed as a smart contract *on Starknet* (operating as an L3 or even within an appchain), Kakarot executes standard EVM bytecode.
- **Mechanism:** Developers deploy standard Solidity contracts. The Kakarot contract interprets the EVM bytecode instruction-by-instruction, executing it within the Cairo VM. The entire execution trace of the Kakarot contract is then proven using Starknet’s STARK-based proof system (Stone Prover).
- **Advantages:** Leverages Starknet’s high-throughput, low-cost proving infrastructure (STARKs scale better than SNARKs for complex computations). Allows EVM compatibility *on top of* Starknet’s ecosystem and tooling (e.g., native account abstraction). Offers a path for Ethereum dApps to tap into Starknet’s scalability without full migration.
- **Tradeoffs: Type 3+.** While aiming for EVM equivalence, Kakarot operates within the Cairo VM environment. Subtle differences in gas metering, memory handling, or access to Cairo-specific features could cause deviations. Proof finality inherits Starknet’s latency (~hours). Security depends on Starknet’s validity proofs and decentralized sequencer. Kakarot zkEVM launched on testnet in late 2023, showcasing working Uniswap V2 deployments. Its success hinges on seamless equivalence and bridging liquidity between Ethereum/Starknet/Kakarot environments.
- **Consensys Linea: Gradual Equivalence Roadmap:**
- **Strategic Approach:** Linea, developed by Consensys (creators of MetaMask, Infura), launched its mainnet in July 2023 with a pragmatic **gradual evolution** strategy. It started as a **Type 3 ZK-EVM** (almost equivalent) with explicit plans to progress towards Type 2.
- **Initial Compromises (Type 3):** At launch, Linea lacked support for certain precompiles (e.g., the elliptic curve pairing `ecPairing` vital for advanced ZK-circuits like zk-SNARK verifiers themselves!) and had minor deviations in gas costs and edge-case opcode behavior. This allowed faster time-to-market and optimized initial performance.
- **Tooling & Ecosystem Strength:** Leveraging Consensys’ vast ecosystem, Linea prioritized **developer accessibility**:

- **MetaMask Integration:** Native support in the world’s most popular wallet via the Linea network RPC.
- **Infura & Truffle Support:** Seamless deployment and management via Consensys’ industry-standard infrastructure.
- **Foundry/Hardhat Plugins:** Robust toolchain integration comparable to Scroll.
- **Roadmap to Type 2:** Linea has systematically implemented missing features. Key milestones included adding `ecPairing` support and refining gas costs. Ongoing work focuses on eliminating remaining deviations and achieving full bytecode equivalence. Consensys’ resources and focus make Linea a formidable contender, particularly attractive to projects already within its ecosystem orbit. Its TVL and transaction volume grew rapidly post-EIP-4844, demonstrating the power of its integrated tooling despite not yet being fully Type-2.
- **Performance:** Linea benefits from Consensys’ engineering scale. It demonstrates competitive proving times and high TPS, leveraging its initial Type-3 optimizations while incrementally adding equivalence.
- **Benchmarking Progress Toward Full Type-2 Compliance:**
- **Evaluation Framework:** Assessing “Type-2 compliance” involves rigorous testing:
 1. **Ethereum Execution Spec Tests:** Passing the full suite of official Ethereum tests is a baseline. Polygon zkEVM, Scroll, and Linea report high pass rates (>95%), with failures often highlighting edge cases under active development.
 2. **Differential Fuzzing:** Tools like **HEVM** or custom frameworks replay vast numbers of transactions with random inputs simultaneously on L1 and the ZK-EVM, comparing state roots and gas usage. Discrepancies reveal subtle equivalence bugs. All major projects employ this.
 3. **Real Contract Deployment:** Testing complex, mainstream contracts (e.g., Uniswap V3, Aave V3) and verifying identical behavior under high load and edge conditions.
 4. **Gas Cost Parity:** Measuring deviation from L1 gas consumption across a wide range of operations. Polygon 2.5 explicitly diverges; others strive for minimal deviation.
- **Current State (Mid-2024):**
 - **Scroll:** Closest to pure Type-2 in bytecode execution and avoiding gas changes. Focuses on eliminating last edge cases.
 - **Polygon zkEVM:** Stable Type-2.5 with known gas deviations but robust bytecode support.
 - **Taiko:** Aims for Type-2 equivalence in execution; based sequencing is its differentiation.

- **Linea:** Progressing rapidly from Type-3, nearing full equivalence in execution semantics, gas parity remains a focus.
- **Kakarot:** Type-3+ (EVM equivalence within Cairo VM constraints), equivalence verification is ongoing.
- **L2BEAT Verification:** Independent tracker L2BEAT provides a “Stage” rating assessing security and decentralization. As of mid-2024, Polygon zkEVM, Scroll, Taiko, and Linea all hold a “Stage 0” rating, indicating their security still relies heavily on centralized components (like upgradable multisigs controlling bridges/verifiers), a common stage for young L2s. Achieving higher stages (full decentralization) is a next frontier for all contenders.

The landscape of Type-2 ZK-EVMs reveals a dynamic tension between the uncompromising ideal of perfect equivalence and the practical demands of performance, decentralization, and developer adoption. Polygon zkEVM, the pioneer, demonstrates the necessary compromises inherent in scaling first. Scroll champions purity and open-source ethos, betting that seamless equivalence will win developers long-term. Taiko challenges the very notion of centralized sequencing, prioritizing alignment with Ethereum’s trust model at the cost of inherent L1 latency. Linea leverages vast resources to bridge the gap from pragmatic Type-3 to full equivalence, while Kakarot explores a novel path via Starknet’s scaling infrastructure. Each implementation embodies a distinct interpretation of the Type-2 promise, their successes and struggles collectively defining the evolving frontier of trustless Ethereum scaling. Yet, the cryptographic bedrock upon which they all rely demands rigorous scrutiny – how secure are these complex systems in practice, and what vulnerabilities lurk beneath the surface of their mathematical guarantees? [Transition to Section 6: The elegance of zero-knowledge proofs offers profound security benefits, but the intricate machinery of Type-2 ZK-EVMs introduces novel cryptographic dependencies, systemic risks, and economic attack surfaces that demand critical examination...](#)

1.6 Section 6: Security Model and Trust Assumptions

The elegant mathematics of zero-knowledge proofs promise near-perfect security—a cryptographic guarantee that state transitions are valid, eliminating the need for fraud proofs or optimistic delays. Yet the intricate machinery of Type-2 ZK-EVMs introduces layers of complexity where theoretical ideals confront engineering realities. Beneath the veneer of cryptographic certainty lie nuanced trust assumptions, systemic vulnerabilities, and economic attack surfaces that demand rigorous scrutiny. This section dissects the security foundations of Type-2 ZK-EVMs, exposing the delicate balance between mathematical perfection and operational pragmatism that defines their trust model. From the cryptographic bedrock to the human governance of upgrade keys, we critically examine where these systems derive their security—and where hidden fractures might emerge.

1.6.1 6.1 Cryptographic Trust Foundations

The security of Type-2 ZK-EVMs rests on cryptographic primitives whose assumptions must hold. While proofs themselves are information-theoretically sound, their practical implementation introduces subtle trust dependencies.

- **Setup Ceremonies and Toxic Waste Risks:**
- **The Peril of “Toxic Waste”:** SNARKs like Groth16, PLONK, and Marlin require a *trusted setup ceremony* to generate public parameters (Common Reference String - CRS). This process involves participants collaboratively generating secret random values (“toxic waste”) that *must be destroyed*. If any participant retains these values, they could forge fake proofs that pass verification. The infamous **Zcash “Ceremony” flaw (2018)**—where a developer accidentally saved toxic waste to a public AWS bucket for weeks—illustrates the catastrophic consequences of failure.
- **Mitigation Strategies:** Modern Type-2 projects employ MPC-based ceremonies:
- **Polygon zkEVM:** Used the **Hermes 1 Ceremony (2021)** with 72 participants (including Vitalik Buterin, Daniel Lubarovich). The ceremony utilized a multi-party computation (MPC) protocol where each participant added entropy, making it statistically improbable for even a majority of colluding participants to reconstruct the toxic waste.
- **Scroll:** Leveraged the **Ethereum KZG Ceremony (2023)**, the largest trusted setup in history with 141,416 participants. Its “Powers of Tau” structure ensures that as long as *one* participant was honest and destroyed their entropy, the entire CRS remains secure.
- **Persistent Risks:** While massively scaled ceremonies mitigate risks, they introduce new concerns. The complexity of MPC protocols increases the chance of implementation bugs (e.g., the 2022 **Tornado Cash ceremony bug** that forced a restart). Projects like **Taiko** explore non-trusted-setup options like **STARKs** (though currently impractical for full EVM equivalence) to eliminate this vector entirely.
- **Post-Quantum Considerations:**
- **SNARKs vs. STARKs Vulnerability:** Most Type-2 ZK-EVMs use pairing-based SNARKs (Groth16, PLONK) vulnerable to **quantum attacks** via Shor’s algorithm, which could break elliptic curve discrete logarithm problems. A sufficiently powerful quantum computer could forge proofs or extract private keys. In contrast, **STARKs** (used in Kakarot on Starknet) rely on hash-based cryptography (e.g., SHA-3) considered quantum-resistant.
- **Practical Timeline vs. Theoretical Risk:** While large-scale quantum computers are likely decades away, the long-lived nature of blockchain systems demands foresight. **Scroll’s Halo2-KZG** implementation includes a **recursive proof upgrade path**, allowing future migration to quantum-resistant proofs without breaking historical state. **Polygon zkEVM’s Plonky2** uses FRI (STARK component), creating a hybrid foundation for eventual transition. Nevertheless, the *current* security of all major Type-2 systems assumes the non-existence of practical quantum adversaries.

- **Verifier Contract Vulnerabilities:**
- **The Single Point of Failure:** The on-chain verifier smart contract is the ultimate arbiter of truth. A bug here could accept invalid proofs, corrupting the L2 state. Historical incidents underscore the danger:
- The **zkSync Lite verifier bug (2021)** forced a 72-hour shutdown after a logic error allowed invalid proofs to pass. No funds were lost, but the system’s liveness failed.
- **Polygon zkEVM’s** initial verifier (audited by Hexens) required emergency patching when a rounding error in a field operation was discovered during internal fuzzing.
- **Audit Rigor and Formal Methods:** Leading projects employ layered defenses:
- **Scroll:** Implemented its KZG verifier in **Dafny**, a verification-aware language, generating mathematical proofs of correctness before compiling to EVM bytecode. Audited by Zellic and Veridise.
- **Consensys Linea:** Used **Certora’s** formal verification tools to specify and prove properties of its verifier contract, catching 3 critical bugs pre-launch.
- **Industry Benchmark:** As of 2024, all major Type-2 verifiers have undergone ≥ 3 independent audits (e.g., OpenZeppelin, Trail of Bits, Spearbit). However, the complexity of pairing operations (e.g., 600 lines of Yul in Polygon’s verifier) leaves room for subtle errors.

1.6.2 6.2 Systemic Risks and Failure Modes

Beyond cryptography, Type-2 ZK-EVMs inherit systemic risks from their reliance on centralized components, upgrade mechanisms, and data availability layers.

- **Sequencer Centralization Risks:**
- **Censorship and MEV Extraction:** In early implementations (Polygon, Scroll, Linea), a single entity operates the sequencer. This creates risks:
- **Transaction Censorship:** Sequencers can exclude specific addresses or transactions (e.g., OFAC-sanctioned Tornado Cash interactions). Polygon faced allegations in 2023 after delaying transactions to a privacy tool.
- **MEV Centralization:** Sequencers can front-run, back-run, or sandwich user transactions. **Scroll’s** decentralized prover network still relies on a centralized sequencer for ordering, creating a MEV goldmine. In Q1 2024, $>60\%$ of Polygon zkEVM’s MEV revenue came from just two searcher bots colluding with the sequencer operator.
- **Mitigations:** Solutions are emerging but immature:

- **Taiko's Based Sequencing:** Eliminates dedicated sequencers by using Ethereum validators, inheriting L1's censorship resistance. However, L1 proposers can still censor L2 txs by excluding `proposeBlock` calls.
- **PBS (Proposer-Builder Separation):** Scroll and Polygon are testing MEV-Boost-like architectures where specialized “builders” construct blocks and “proposers” (sequencers) select the highest-bidder bundle. This distributes MEV profits but doesn't prevent censorship.
- **Force Inclusion Mechanisms:** All major Type-2 L2s implement L1 “inbox” contracts allowing users to bypass sequencers by submitting txs directly to L1. These txs must be included in the next L2 block but cost 10-100× typical L2 fees, making them impractical for regular use.
- **Upgrade Key Control Mechanisms:**
 - **The Multisig Problem:** Admin keys controlling upgradable contracts (bridges, verifiers, sequencer selection) pose extreme centralization risks. As of mid-2024:
 - **Polygon zkEVM:** 5/8 multisig held by Polygon Labs engineers.
 - **Scroll:** 4/7 multisig with keys in hardware wallets managed by the foundation.
 - **Consensys Linea:** 3/5 multisig controlled by Consensys.
 - **Exploit Scenarios:** A compromised key could:
 1. Upgrade the verifier to accept fake proofs (stealing all bridge funds).
 2. Halt the bridge (freezing assets).
 3. Drain insurance funds.

The **Nomad Bridge hack (2022)**, where a single upgradeable contract lost \$190M, exemplifies this risk.

- **Pathways to Decentralization:** Projects are slowly transitioning:
- **Timelocks:** Taiko uses 7-day timelocks for upgrades, allowing users to exit if malicious changes are proposed.
- **DAO Governance:** Scroll's roadmap delegates upgrade control to a token-based DAO by 2025. However, token-based voting introduces plutocracy risks.
- **Escape Hatches:** Linea implements a “security council” with veto power over malicious upgrades, but council members are still trusted entities.
- **Data Availability Failures:**

- **Blob Reliance and Risks:** EIP-4844 blobs reduced costs but introduced new risks. Blobs are pruned by Ethereum nodes after ~18 days. If rollup operators fail to persist this data externally:
- **State Reconstruction Failure:** Users cannot compute Merkle proofs for withdrawals after the pruning period.
- **Forced Transaction Inoperability:** Users relying on L1 force-inclusion mechanisms cannot prove their tx was ignored by the sequencer without blob data.
- **Solutions and Gaps:**
 - **Decentralized Storage:** Polygon uses Celestia and Arweave for long-term blob storage. Scroll relies on its community-run “Guardian” nodes.
 - **Proof-of-Custody:** None of the major Type-2 ZK-EVMs yet implement Ethereum’s proposed proof-of-custody games, where validators cryptographically attest to storing blob data. This remains a critical vulnerability for long-term recoverability.

1.6.3 6.3 Economic Security and Slashing Mechanisms

ZK-Rollups replace fraud proofs with cryptographic validity, but their economic security models protect against liveness failures and governance attacks.

- **Bond Requirements for Sequencers/Provers:**
 - **Sequencer Bonds:** To deter malicious behavior (e.g., censoring force-inclusion txs), sequencers post bonds (e.g., 500 ETH in Polygon zkEVM). If provably malicious, bonds are slashed. However, proving censorship cryptographically remains impractical—only verifiable downtime or incorrect state transitions trigger slashing.
 - **Prover Bonds:** In decentralized networks (Scroll, Taiko), provers stake tokens to participate. **Scroll’s** design slashes stakes for:
 - **Unresponsiveness:** Failing to generate proofs when selected.
 - **Invalid Proofs:** Submitting proofs that fail verification (theoretically impossible if cryptography holds, but possible via implementation bugs).
 - **Example:** In Scroll’s testnet, a prover simulating a “lazy actor” had its entire 50k TKO testnet stake slashed after missing 3 consecutive proof assignments.
- **Fraud Detection Challenges:**
 - **The ZK Advantage:** Unlike Optimistic Rollups (ORUs), Type-2 ZK-EVMs have no “fraud proof window.” Validity is proven instantly. This eliminates:

- **Reorg Attacks:** Adversaries cannot temporarily fork the chain during a dispute window (e.g., the **Optimism “fault proof” delay exploit (2022)**).
- **Withdrawal Delays:** Users withdraw in minutes, not days.
- **Residual Risks:** ZK systems can still suffer *liveness* failures:
- **Prover Collusion:** If all provers in a permissionless network halt (e.g., due to unprofitability), the chain stalls. **Taiko’s** TKO tokenomics include **prover subsidies** to prevent this.
- **Data Withholding:** A sequencer could withhold transaction data (blobs) while still generating valid proofs of state transitions. Users couldn’t verify their balances or initiate withdrawals. This requires robust data availability sampling (not yet implemented).
- **Insurance Fund Designs:**
 - **Purpose:** To cover losses from undiscovered bugs (e.g., a circuit flaw allowing invalid withdrawals). Funds are sourced from sequencer fees or token inflation.
 - **Implementation Variance:**
 - **Polygon zkEVM:** Maintains a 10,000 ETH fund managed by the Polygon Foundation. Claims require manual review.
 - **Scroll:** Uses a decentralized model where token holders vote on bug bounty payouts from a community treasury.
 - **Coverage Gaps:** No fund covers 100% of TVL. Polygon’s fund covered <2% of its \$120M peak TVL in 2023. A catastrophic bug could still leave users uncompensated.

1.6.4 6.4 Formal Verification Progress

Formal verification (FV)—mathematically proving the correctness of code—offers the highest assurance against implementation bugs. Its adoption in Type-2 ZK-EVMs is nascent but accelerating.

- **Verified Circuits: Current State:**
 - **Circuit Complexity:** Proving an entire EVM circuit (millions of gates) is currently infeasible. Projects focus on critical components:
 - **Scroll:** Formally verified its **Keccak-256 circuit** using the **Coq** proof assistant, proving its equivalence to the Ethereum Yellow Paper specification.
 - **Polygon zkEVM:** Used **Giza** (custom tool) to verify the arithmetic correctness of its **PLONK constraint system compiler**.

- **Taiko:** Partnered with **Veridise** to verify the **elliptic curve precompile circuits** using symbolic execution.
- **Limitations:** FV tools struggle with non-arithmetic operations (e.g., memory access patterns) and recursive aggregation layers. Most verified components are isolated modules, not the full proving stack.
- **Runtime Verification vs. Full Protocol Verification:**
- **Runtime Verification (RV):** Tools like **K framework** execute EVM bytecode symbolically to detect deviations between L1 and L2:
- **Success:** RV caught a critical **gas metering divergence** in an early Scroll build where `SSTORE` re-funds were miscalculated.
- **Scope:** RV verifies *execution equivalence* but not the ZK circuits themselves.
- **Full Protocol Verification:** Aims to prove the *entire system* (client + prover + verifier) adheres to a formal specification. This remains aspirational:
- **Ethereum Foundation’s “Type-1” Effort:** The PSE team is building a **formally verified client in Rust** (Reth-based) with a matching **Lean-proven circuit**. If successful, this could set a benchmark for Type-2 systems.
- **Barriers:** Requires formalizing the EVM specification itself (an ongoing effort via **Ethereum Execution Specification (EELS)**) and bridging the gap between high-level specs and low-level circuit code.
- **Leading Projects and Tools:**
- **Certora:** Used by Consensys Linea and Polygon to verify critical invariants (e.g., “no invalid block can be finalized”).
- **Veridise:** Audited Taiko and Scroll circuits using **concolic execution**, hybridizing concrete and symbolic testing.
- **OtterSec:** Focused on **ZK-specific vulnerabilities** (e.g., under-constrained circuits) in Scroll’s Halo2 implementation.
- **Community Efforts:** The **Zellic FV team** open-sourced **zkLLVM**, a tool for generating formally verifiable circuits from LLVM IR, potentially used in future Type-2 iterations.

The security model of Type-2 ZK-EVMs resembles a layered fortress: the impenetrable walls of zero-knowledge cryptography defend the core, but these are buttressed by human-governed upgrade mechanisms,

economically secured sequencers, and manually administered insurance funds—each layer introducing its own vulnerabilities. While mathematically superior to optimistic approaches in preventing invalid state transitions, they trade fraud proof risks for more nuanced threats: trusted setup ceremonies, quantum fragility, verifier bugs, and the persistent specter of centralized control. As formal verification slowly expands its reach—from isolated circuits to entire protocol stacks—the aspiration shifts from “trust, but verify” to “verify, then trust.” Yet, for now, the security of even the most advanced Type-2 ZK-EVM rests as much on institutional credibility (Polygon Labs, Consensys, Ethereum Foundation) and bug bounty programs as it does on abstract algebra. The journey toward truly trustless scaling remains a work in progress, where cryptographic ideals are tempered by the messy realities of implementation and governance.

Transition to Section 7: Beyond the abstract realm of security models and cryptographic guarantees lies the tangible arena of developer adoption. The promise of “EVM equivalence” means little if developers face fractured tooling, unpredictable gas costs, or debugging nightmares when migrating their dApps. The ultimate test of Type-2 ZK-EVMs unfolds not in the prover’s data center, but at the developer’s workstation...

1.7 Section 7: Developer Experience and Ecosystem Impact

The cryptographic elegance and security guarantees of Type-2 ZK-EVMs, dissected in the preceding section, ultimately face their most critical validation not in academic papers or audit reports, but in the trenches of developer workflows and real-world dApp deployment. The core promise – *unmodified* EVM equivalence – was conceived precisely to eliminate friction, promising Ethereum developers a frictionless scaling path: “compile once, deploy anywhere.” Yet, the journey from this ideal to practical reality reveals a landscape marked by nuanced tooling gaps, unpredictable cost dynamics under ZK constraints, subtle performance bottlenecks, and fragmented user experience. This section critically assesses the tangible developer experience and ecosystem impact of Type-2 ZK-EVMs, exploring where the “drop-in compatibility” dream meets the gritty reality of building and deploying scalable dApps on these nascent validity-proven layers. Success hinges not just on cryptographic soundness, but on conquering the mundane yet crucial hurdles of debuggers, gas estimators, and wallet integrations.

1.7.1 7.1 Tooling Compatibility Landscape

The initial allure of Type-2 ZK-EVMs lies in leveraging Ethereum’s mature toolchain. While significant strides have been made, the ZK context introduces unique complexities that challenge seamless compatibility.

- **Hardhat/Foundry Plugin Support Status: The Foundation:**
- **Maturity Gradient:** Support for the dominant Ethereum development frameworks (Hardhat, Foundry) is now table stakes. Projects compete on plugin sophistication:

- **Scroll:** Offers arguably the most polished integration. Its `hardhat-scroll` plugin (v0.6.0+) provides near-identical workflows: `npm run hardhat compile`, `npm run hardhat test`, `npm run hardhat run scripts/deploy.js --network scrollSepolia`. Foundry (`forge test`, `forge create`) works seamlessly via RPC configuration. The plugin handles faucet access, block explorer verification, and network management transparently. This reflects Scroll's core ethos of minimal developer friction.
- **Polygon zkEVM:** Provides robust `hardhat-zkevm` (v1.0.0+) and Foundry support. Deployment and testing are smooth, though developers must consciously account for its Type-2.5 gas deviations during testing (see 7.2). Its documentation explicitly flags KECCAK256 cost differences.
- **Taiko:** `hardhat-taiko` (v0.1.4) and Foundry integration are functional but less mature. Developers face minor quirks related to its based sequencing model, like slightly longer initial tx confirmation times waiting for L1 proposal. Its unique architecture demands more bespoke tooling adjustments.
- **Consensys Linea:** Leverages Consensys' infrastructure strength. Tight integration with **Truffle** and **Hardhat** via `@consensys/linea-hardhat-plugin` is a major draw, complemented by native **Infura** RPC endpoints and **MetaMask** support. This ecosystem cohesion lowers barriers significantly.
- **Remaining Gaps:** While core compilation and deployment work, advanced plugin features sometimes lag:
- **Forking Mainnet State:** Crucial for testing complex interactions with live protocols (e.g., simulating interactions with L1 Aave). Scroll and Polygon support forking via their RPCs (`hardhat node --fork`), but performance and state syncing lag behind Optimism/Arbitrum forks due to proving overhead during fork simulation. Taiko's fork support is experimental.
- **Gas Reporting:** Hardhat/Foundry's detailed gas reports may not perfectly reflect *actual* L2 execution + proof costs, especially for operations penalized in Type-2.5 systems. Plugins are evolving to provide more accurate L2-specific cost breakdowns.
- **Debugging Challenges in ZK Environments: Stepping into the Unknown:**
- **The Black Box Problem:** Traditional EVM debuggers (Hardhat's `console.log`, Foundry's `forge test -vvvv`, Remix debugger) operate by inspecting EVM execution traces *after* execution. In a ZK-EVM, the execution trace is processed *offline* by the prover. Debugging a failed transaction becomes non-trivial:
- **Trace Accessibility:** Developers need access to the granular execution trace *before* proof generation to pinpoint failures (e.g., an out-of-gas error at opcode 1,234,567). This trace isn't readily available on standard RPC nodes.
- **Prover-Specific Errors:** Failures might occur not in EVM execution itself, but during witness generation or proving (e.g., an unconstrained circuit path, a lookup argument failure). These errors are opaque and often reported as generic "execution reverted" or "proof generation failed."

- **Emerging Solutions:**
 - **Scroll zkEVM Tracer:** Scroll’s standout tool is its custom tracer. When a transaction fails, developers can fetch a detailed trace via a dedicated RPC call or explorer interface. This trace mirrors Hardhat/Foundry’s output, showing step-by-step opcode execution, stack, memory, and storage changes *within the ZK context*, highlighting the exact opcode where failure occurred (e.g., `SSTORE` at depth 43 with insufficient gas). This significantly demystifies ZK debugging.
 - **Polygon zkEVM Debug Bridge:** Polygon offers a more limited “`debug_traceTransaction`” RPC method, providing basic execution traces. However, correlating this with potential prover errors requires deeper infrastructure access or reliance on support channels.
 - **Local Prover/Node Setups:** Projects encourage running local testnet nodes with proving capabilities (`zkevm-node`). This allows developers to generate and inspect full traces locally but demands significant computational resources (high-RAM servers, often with GPUs), creating a high barrier for individual developers. **Linea’s** well-documented local setup with Docker mitigates this somewhat.
- **Tracer Tools for Failed Transactions: Bridging the Gap:**
 - **Block Explorer Integration:** Mature explorers (Blockscout forks for Polygon, Scroll, Taiko; Consensus’ Linea Explorer) are integrating transaction tracing features. While not as granular as dedicated tracer tools, they visualize contract calls, internal transactions, and event logs effectively. Identifying *which* internal call reverted is standard; pinpointing the *exact opcode* within that call is still evolving.
 - **Third-Party Tools:** Services like **Tenderly** are adding support for major ZK-EVMs. Tenderly’s visual debugger for Polygon zkEVM allows stepping through transactions, inspecting state, and seeing gas usage per call, providing a significant UX improvement over raw RPC traces. Similar integrations for Scroll and Linea are in progress. However, deep ZK-specific error decoding (prover circuit failures) remains outside their scope.

The tooling landscape is rapidly evolving from “it compiles” to “it debugs.” Scroll leads in transparency, Consensus Linea in integrated ecosystem smoothness, while Polygon and Taiko balance maturity with their unique architectural quirks. The goal remains enabling developers to spend time on logic, not layer idiosyncrasies.

1.7.2 7.2 Gas Cost Dynamics and Contract Migration

While Type-2 ZK-EVMs preserve EVM *opcode semantics*, the *economic cost* of operations diverges significantly from L1 due to the immense overhead of ZK-proof generation. Understanding and navigating these dynamics is crucial for successful contract migration and efficient dApp operation.

- **Comparative L1 vs L2 Gas Profiles: Beyond Execution Cost:**

- **Breaking Down L2 Gas Fees:** An L2 transaction fee has distinct components:

1. **L2 Execution Fee:** The cost of running the transaction on the sequencer's node, calculated using L1-equivalent gas costs *or* the rollup's specific Type-2.5 costs. This is usually minimal (e.g., 0.00001 ETH for a simple transfer).
2. **L2 Storage Write Fee:** Specific to writes (SSTORE creating new slots). While L1 charges 20,000 gas for init, ZK-EVMs often charge more to account for the proving cost of complex state updates. Polygon zkEVM charges ~100,000 gas equivalent.
3. **L1 Data Fee (Blob Cost):** The cost of publishing the transaction's compressed call data via an EIP-4844 blob. This dominates simple transactions (transfers, swaps). Cost scales with calldata size and Ethereum blob base fee. (~0.0001 - 0.001 ETH).
4. **L1 Proof Verification Fee:** The amortized cost of verifying the ZK-proof for the batch containing the transaction on L1. This is usually small per tx (L2 messaging inherit the proof finality latency. A price update posted on L1 takes minutes to be proven and accepted on L2. dApps must be designed to tolerate this ~5-15 minute stale data window or use native L2 oracle solutions (like RedStone or Supra on Scroll) that post data directly to L2 with faster soft confirms, though potentially with different security assumptions.

- **Randomness (PREVRANDAO/BLOCKHASH):** As detailed in Section 3.3, these opcodes derive values from L1. Using `block.prevrandao` on L2 involves:

1. Sequencer injecting the L1 block's `prevrandao` value into the L2 block.
2. The ZK-proof verifying this value matches the actual L1 block via a storage proof.

This introduces a 1-2 L1 block delay (~12-24 seconds) plus proving time before the L2 `prevrandao` value is finalized. Applications needing on-chain randomness must account for this delay or use commit-reveal schemes or VRF oracles designed for L2 latency.

1.7.3 7.4 Wallet Integration Challenges

The end-user experience hinges on seamless wallet interaction. Type-2 ZK-EVMs, while EVM-equivalent, introduce novel complexities at the wallet layer.

- **Account Abstraction (AA) Support Variations:**

- **Native vs. Emulated AA:** Ethereum L1 relies on Externally Owned Accounts (EOAs). Account Abstraction (ERC-4337) allows smart contract wallets. ZK-EVMs vary in support:

- **Native AA (e.g., zkSync Era, Starknet):** Not Type-2. These have AA baked into their core protocol, offering superior UX (gas sponsorship, social recovery).
- **Type-2 ZK-EVMs:** Primarily support **ERC-4337 emulated AA**. This means deploying standard ERC-4337 Bundler, Paymaster, and EntryPoint contracts *on the ZK-rollup*. Functionality is similar to L1 but with rollup-specific latency.
- **Implementation Maturity:**
- **Polygon zkEVM:** Supports ERC-4337 via community bundlers (e.g., Biconomy, Stackup). Paymaster integrations allow gas sponsorship in ETH or stablecoins.
- **Scroll:** Active development of ERC-4337 infrastructure; testnet bundlers operational. Strong focus on compatibility with existing AA providers.
- **Taiko/Linea:** ERC-4337 support is present but less battle-tested. Linea benefits from Consensys' "Session Keys" research for improved gaming/DeFi AA UX.
- **User Experience:** While possible, AA on Type-2 ZK-EVMs currently lacks the seamless "out-of-box" feel of native AA chains. Users often need to actively seek out AA-enabled dApps or configure bundler RPCs.
- **Signature Scheme Incompatibilities:**
- **The secp256r1 Opportunity (and Challenge):** Ethereum EOAs use the secp256k1 curve. The rise of passkeys and WebAuthn uses the secp256r1 (P-256) curve. While not an incompatibility per se, Type-2 ZK-EVMs present an opportunity to natively support r1 signatures for seamless Web2-like logins via AA wallets. However:
- **Precompile Absence:** Ethereum L1 lacks a secp256r1 precompile. Verifying r1 signatures requires writing a Solidity contract, costing ~1M+ gas on L1, making it impractical. On L2, cheaper gas *could* make this feasible, but the absence of a standardized precompile means each dApp must implement its own verifier, fragmenting support.
- **ZK-Proving Cost:** Verifying r1 signatures in-circuit is computationally heavy. No major Type-2 ZK-EVM currently offers a native r1 precompile equivalent, pushing this functionality into expensive contract calls. **Projects like Pimlico are experimenting** with off-chain proof generation for r1 within AA, but native ZK-EVM support remains future work.
- **Multi-Chain Wallet UX Friction:**
- **Network Discovery and Configuration:** While MetaMask and Rabby support major ZK-EVMs, users still need to manually add the network RPC (chain ID, symbol, explorer URL). Non-technical users find this confusing. Solutions like **Chainlist** help, but friction remains compared to established L1s.

- **Chain Switching Delays:** Switching between L1 and an L2 ZK-EVM (or between two ZK-EVMs) in wallets involves RPC reconfiguration. This breaks flow compared to single-chain interactions. AA and smart wallets promise session-based multi-chain interactions but aren't widespread.
- **Bridging UX:** Initiating a bridge deposit is usually smooth (built into explorers/wallets like Socket). *Withdrawing* funds, while cryptographically secure, involves multiple steps: 1) Initiate withdrawal on L2, 2) Wait for proof finality (minutes), 3) Claim funds on L1 via a separate transaction. Wallet UIs are improving to track this multi-step process (e.g., showing “Withdrawal Initiated” -> “Ready to Claim”), but it's inherently more complex than an L1 transfer. **Scroll's** explorer provides a unified withdrawal status view.

The developer and user experience on Type-2 ZK-EVMs is a story of remarkable progress shadowed by persistent friction. The “unmodified deployment” promise largely holds for contract *logic*, but navigating the altered economic landscape, debugging the ZK stack, and smoothing the end-user wallet journey demand ongoing adaptation. Success hinges on the ecosystem maturing – from tracer tools becoming standard to wallets abstracting multi-chain complexity and gas cost profiles stabilizing. While not yet frictionless, the trajectory points towards a future where the cryptographic magic of ZK-proofs fades into the infrastructure, leaving developers free to build and users free to interact, unburdened by the complexities of the layer beneath.

Transition to Section 8: The viability of this developer-centric future, however, rests not just on tooling and UX, but on sustainable economic models. How do Type-2 ZK-EVMs incentivize decentralized sequencers and provers? What tokenomics govern fee markets and MEV distribution? And how do treasuries fund the ongoing evolution of this complex infrastructure? The economic and governance dimensions become the critical enablers – or bottlenecks – for long-term ecosystem health...

1.8 Section 8: Economic and Governance Dimensions

The developer experience and technical performance of Type-2 ZK-EVMs, while crucial, ultimately rest upon sustainable economic foundations and robust governance mechanisms. The intricate dance of zero-knowledge proofs and EVM execution demands carefully calibrated incentive structures to ensure decentralized participation, protocol longevity, and ecosystem health. This section examines the economic engines powering these systems – from token-driven proving markets to sequencer decentralization pathways – and explores how treasuries manage billions in value while navigating protocol upgrades. It further confronts the emerging challenge of cross-rollup liquidity fragmentation, where the very success of multiple ZK-EVMs threatens to Balkanize Ethereum's unified liquidity pool. The viability of the Type-2 paradigm hinges not just on cryptographic elegance, but on creating incentive-compatible systems that align diverse actors toward shared security and growth.

1.8.1 8.1 Tokenomics of Proving Markets

The computationally intensive nature of ZK-proof generation necessitates sophisticated incentive structures. Native tokens often serve as coordination mechanisms, though their design varies significantly across implementations.

- **Native Token Utilities: Beyond Simple Gas:**
- **Sequencing Rights:** In decentralized sequencing models, tokens confer the right to propose blocks. **Taiko's TKO token** is staked by Ethereum validators who wish to include Taiko blocks via L1 `proposeBlock` transactions. Validators with higher TKO stakes receive priority in block proposal opportunities, earning ETH fees from included transactions. This creates a direct link between token stake and sequencing revenue.
- **Prover Staking:** Decentralized prover networks require sybil resistance. **Scroll's** emerging prover market requires participants to stake **ETH** (not a native token) as a bond. However, projects like **Risc Zero** (influencing ZK-EVM designs) utilize native tokens (e.g., **RISC0**) for staking, where token-holders earn fees for proof generation and face slashing for malfeasance. The trade-off: ETH stakes leverage existing value but offer less protocol control; native tokens enable tighter economic policy but face bootstrapping challenges.
- **Governance Rights:** Tokens often govern protocol upgrades, parameter tuning, and treasury allocation. **Polygon's POL token** (multichain token) will eventually govern the Polygon 2.0 ecosystem, including its Type-2.5 zkEVM, enabling token-weighted voting on sequencer selection rules or fee market parameters.
- **Prover Subsidy Mechanisms and Fee Markets:**
- **The Cost-Recovery Challenge:** L2 transaction fees (dominated by L1 data/proof costs) often fall short of covering the true compute cost of proof generation, especially during early adoption. Projects employ creative subsidy models:
- **Sequencer Subsidies (Scroll):** A portion of sequencer fees (from L2 execution) is redirected to provers. During low activity, the treasury can top up subsidies to ensure liveness.
- **Token Emission (Taiko):** Newly minted TKO tokens supplement proving rewards during the bootstrapping phase, gradually transitioning to fee-only rewards as network activity grows. This mirrors Bitcoin's block subsidy model but for ZK compute.
- **Priority Fee Auctions (Polygon zkEVM):** Users bidding higher priority fees incentivize sequencers to include their transactions in blocks with faster proving schedules. Provers receive a share of these fees proportional to the computational intensity of the batch they prove.

- **Prover Markets in Practice: Ingonyama’s IGNY Pool** demonstrates an externalized model – a marketplace where rollups (including experimental Type-2s) auction proof generation tasks to competitive prover nodes. Provers stake IGNY tokens for reputation and earn fees paid in the rollup’s native token or ETH. This commoditizes proving, potentially lowering costs but introducing coordination overhead.
- **MEV Distribution Frameworks:**
- **The Centralization Risk:** Without intervention, MEV naturally flows to sequencers who control transaction ordering. In early centralized sequencers (Polygon, Scroll pre-decentralization), >80% of MEV was captured by the sequencer operator or privileged bots.
- **ZK-Rollup PBS (Proposer-Builder Separation):** Inspired by Ethereum’s MEV-Boost, this decouples block *building* (ordering txs for MEV extraction) from block *proposal* (selecting the best block):
- **Builders:** Specialized actors (e.g., **Flashbots SUAVE**, **BloXroute**) compete to create MEV-optimized blocks, submitting bids (containing the block + a fee) to the proposer.
- **Proposers (Sequencers):** Select the highest-bidding valid block. Revenue is split between the proposer (sequencer), builder, and potentially the protocol treasury or a public goods fund.
- **Implementation Status:**
- **Polygon zkEVM:** Piloting PBS via a modified MEV-Boost integration. Builders submit bids containing pre-confirmed blocks and ZK-proof commitments. The sequencer acts as proposer, selecting bids based on fee value and proof validity promises.
- **Scroll:** Designing a PBS variant where builders must also stake ETH, ensuring they complete proof generation for their blocks or face slashing. This ties MEV extraction to proving responsibility.
- **Taiko:** Its based sequencing complicates PBS. Ethereum validators (proposers) receive MEV bids directly when including `proposeBlock` txs. MEV revenue flows primarily to L1 validators rather than a dedicated L2 ecosystem.

1.8.2 8.2 Sequencer Decentralization Pathways

Centralized sequencers represent a critical point of failure and censorship. Decentralizing this function is paramount for credible neutrality and resilience.

- **PoS Sequencing Pools: The Established Path:**
- **Mechanism:** Sequencers stake tokens (ETH or native) to join a permissionless pool. A leader selection algorithm (e.g., pseudo-random selection weighted by stake) chooses the sequencer for each slot. **Polygon zkEVM’s** roadmap transitions its single sequencer to a **PoS pool** staking **POL tokens**. Slashing punishes liveness failures (missing blocks) or submitting invalid blocks (detectable via ZK-proofs).

- **Challenges:** Achieving low-latency communication between decentralized sequencers is complex. **Fast Finality Gadgets:** Projects like **Astria** (shared sequencer network) offer solutions where a decentralized sequencer set reaches rapid consensus on block ordering before execution and proving, minimizing soft confirmation delays.
- **DVT (Distributed Validator Technology) Integration:** Inspired by Ethereum staking pools, DVT splits the sequencer key among multiple operators using threshold cryptography. **Obol Network's Charon** is being adapted by **Scroll** for its sequencer pool. No single operator holds the full key, mitigating the risk of a single point of compromise or censorship. A majority (e.g., 4-of-7) must collaborate to sign blocks.
- **Slashing for Liveness Failures:**
- **Detectable Offenses:** ZK-proofs make invalid state transitions impossible to finalize. Slashing therefore focuses on:
 - **Liveness Failure:** Sequencer fails to propose a block within its allocated slot (e.g., within 12 seconds). Easily verifiable on-chain.
 - **Censorship:** Provably ignoring valid “forced inclusion” transactions submitted via L1. Requires cryptographic proof that the transaction met fee requirements and was excluded. **Polygon's** implementation uses a verifiable delay function (VDF) to timestamp L1 force-include txs, allowing demonstrable exclusion.
 - **Slashing Mechanics:** A portion of the sequencer's stake (e.g., 1-10% for liveness, up to 100% for provable censorship) is burned or redirected to a treasury/insurance fund. **Taiko's** based model slashes the ETH stake of the L1 validator who proposed a malicious L2 block via their `proposeBlock` call.
- **Anti-Censorship Guarantees via Forced Inclusion:**
- **The Safety Valve:** All major Type-2 ZK-EVMs implement **L1 Inbox Contracts** (e.g., `Inbox.sol` on Ethereum). Users can submit transactions directly to this contract by paying a premium L1 gas fee (~5-10x typical L2 fees). These transactions *must* be included in the next L2 block by the sequencer.
- **Cryptographic Enforcement:** The ZK-proof for the L2 block *must* include and correctly process these forced transactions. If omitted, the proof verification fails on L1, preventing block finalization and potentially triggering sequencer slashing. This creates a costly but cryptographically guaranteed escape hatch.
- **Cost as a Deterrent:** The high fee prevents routine use but ensures censorship is economically impractical for the sequencer and technically impossible to finalize. **Scroll's** implementation caps forced inclusion fees during congestion, preventing predatory pricing.

1.8.3 8.3 Treasury Management and Protocol Upgrades

Sustaining protocol development, security, and growth requires robust treasury management and transparent upgrade mechanisms.

- **Fee Revenue Distribution Models:**
- **Sources:** Treasuries are funded from:
 - **Sequencer Fee Cuts:** A percentage (e.g., 10-20%) of L2 execution fees.
 - **MEV Revenue Shares:** A portion of PBS auction proceeds directed to the treasury.
 - **Token Inflation (Taiko, Polygon):** New token issuance funds early development and subsidies.
 - **Bridge Fees:** Optional small fees on asset bridges (less common to avoid disincentivizing usage).
- **Allocation:** Funds typically support:
 - **Core Development:** Salaries for protocol engineers and researchers (e.g., Polygon Labs, Scroll core team).
 - **Grants:** Ecosystem incentives (e.g., **Scroll's Grants Program** funding novel dApps, developer tools).
 - **Security:** Audits, bug bounties (e.g., **Immunefi** programs with million-dollar pots), and formal verification efforts.
 - **Prover Subsidies:** Ensuring proof generation remains profitable during low-usage periods.
 - **Public Goods:** Funding infrastructure like RPC nodes, block explorers, and educational content (e.g., **Taiko's allocation to Bitcoin matching rounds**).
 - **Transparency:** Leading projects publish quarterly treasury reports. **Consensys Linea** leverages its parent company's resources but also maintains a transparent **Ecosystem Fund** dashboard showing grant distributions.
- **On-Chain Governance vs. Off-Chain Coordination:**
- **Off-Chain Foundations (Early Stage):** Polygon (Polygon Labs), Scroll (Scroll Foundation), and Taiko (Taiko Labs) currently rely on off-chain foundations or core teams for strategic direction and major upgrades. Decision-making involves community forums (Discourse, Discord), community calls, and technical working groups. This allows agility but lacks formal accountability.
- **On-Chain DAOs (Emerging):** Transitioning to token-based on-chain governance is a common roadmap item:
- **Proposal Types:** Upgrading bridge contracts, adjusting fee parameters (e.g., Type-2.5 gas costs), allocating treasury funds, electing security councils.

- **Voting Mechanisms:** **Optimism’s Collective** uses token-weighted voting for major decisions and citizen (NFT-based) voting for grants, a model influencing ZK-EVMs. **Scroll’s** DAO design emphasizes quadratic voting to mitigate plutocracy.
- **Risks:** Low voter turnout (common in DAOs) can lead to governance capture by whales or activist funds. **Compound’s Proposal 64** demonstrated how a small, motivated group can pass controversial measures with minimal overall participation.
- **Upgrade Veto Mechanisms and Timelocks:**
- **The Centralized Upgrade Risk:** Upgradable contracts controlled by multisigs are the largest security vulnerability (see Section 6). Mitigations include:
 - **Timelocks:** Mandatory delays (e.g., **Taiko’s 7-day timelock**) between a governance vote approving an upgrade and its execution. This allows users to withdraw funds if malicious changes are proposed. Vital for bridge and verifier contracts.
 - **Security Councils:** **Polygon’s** “System of Smart Contracts” includes a 12-member Security Council with veto power over upgrades during the timelock period. Members are reputable entities (e.g., Coinbase, Chainlink Labs). A supermajority (8/12) can veto a potentially malicious upgrade. This balances agility with emergency response.
 - **Escalation Games:** Inspired by Optimism, proposals for critical upgrades (like verifier changes) could be subject to a multi-stage challenge period where users can contest the upgrade’s safety via cryptographic proofs before it activates.

1.8.4 8.4 Cross-Rollup Liquidity Fragmentation

The proliferation of Type-2 ZK-EVMs, while scaling Ethereum, threatens to splinter liquidity across isolated domains, increasing friction and reducing capital efficiency.

- **Bridging Costs and UX Friction:**
- **The Multi-Hop Problem:** Moving assets from **Polygon zkEVM** -> **Scroll** -> **Arbitrum** requires:
 1. Bridge from Polygon zkEVM to Ethereum L1 (minutes, \$0.10-\$1.00).
 2. Wait for Ethereum L1 finality (12 sec).
 3. Bridge from L1 to Scroll (minutes, \$0.10-\$1.00).
 4. Bridge from Scroll to Arbitrum via L1 again (minutes, \$0.10-\$1.00 + 7-day challenge period).

Total cost: \$0.30-\$3.00+, time: 10+ minutes (excluding Optimistic delay). This cripples cross-rollup arbitrage, lending, and user experience.

- **Trust and Security Tradeoffs:** Native bridges offer security but high latency/cost. Third-party “fast bridges” (e.g., **Across**, **Socket**) use liquidity pools and optimistic oracles for instant transfers but introduce custodial risk or require overcollateralization. The **Wormhole exploit (\$325M loss)** highlights the systemic risk of bridge vulnerabilities proliferating across L2s.
- **Shared Liquidity Pools Initiatives:**
- **Omnichain Liquidity Networks:** Protocols build unified liquidity layers spanning multiple L2s:
- **Chainlink CCIP:** Enables smart contracts to send messages and tokens across chains via a decentralized oracle network. A Uniswap V4 pool on Polygon zkEVM could use CCIP to access liquidity locked in a Scroll pool, creating a virtual shared pool. Early integration by **Synthetix**.
- **LayerZero’s Stargate:** Creates unified liquidity pools (e.g., USDC) across supported chains. Users deposit into a pool on one chain; Stargate’s protocol manages the underlying assets across chains via its “Delta Algorithm” and LayerZero’s ultra-light nodes. Significant TVL but introduces dependency on LayerZero’s security model.
- **Native Cross-Rollup AMMs:** Projects like **Swaap** build AMMs deployed natively on multiple ZK-EVMs (Polygon, Scroll) with synchronous state updates via ZK-light client messages, allowing near-instant arbitrage and shared fee generation.
- **ZK-Powered Synchronization: Succinct Labs’ Telepathy** uses zkSNARKs to prove the state of one rollup (e.g., Polygon zkEVM) to another (e.g., Scroll) via an on-chain light client contract. This allows trust-minimized cross-rollup token transfers without routing through L1, reducing latency and cost. **Scroll is an early adopter** for its cross-rollup messaging layer.
- **Standardization Efforts (L2 Standards Proposals):**
- **L2Beat’s L2SP Initiative:** The team behind the L2 tracking site L2Beat spearheads **L2 Standards Proposals (L2SP)**, defining common interfaces and security expectations:
- **L2SP-1: Standard Bridge Interface:** Specifies functions for deposits, withdrawals, and finalizing withdrawals. Ensures wallets and explorers can integrate bridges uniformly. Adopted by Polygon, Arbitrum, Optimism, and emerging ZK-EVMs.
- **L2SP-2: Data Format Standardization:** Defining how rollups format batch data posted to L1 (especially blobs), enabling universal parsers and shared data availability layers like **EigenDA** or **Avail**.
- **L2SP-5: Prover Standardization (Draft):** Exploring common interfaces for proof submission and verification, facilitating shared proving networks and cross-rollup state proofs. Early feedback from **Scroll** and **Taiko** teams.
- **Ethereum Foundation’s ERC-7281 (xERC-20):** Standardizes “bridging-minimized” tokens. Token contracts on each L2 hold their own mint/burn logic and communicate via a standardized registry on

L1. This allows token issuers (like **USDC’s Circle**) to control canonical bridges and permissions, reducing fragmentation and improving security compared to wrapped assets. Adoption is nascent but critical for stablecoin liquidity across ZK-EVMs.

The economic and governance architecture of Type-2 ZK-EVMs represents a complex balancing act. Tokenomics must incentivize decentralized participation in computationally intensive roles without fostering centralization through economies of scale in specialized hardware. Sequencer decentralization pathways must navigate the trade-offs between latency, censorship resistance, and MEV democratization. Treasuries, fueled by fees and potentially inflation, must sustainably fund security and innovation while resisting capture. And the looming specter of liquidity fragmentation demands innovative technical solutions and robust standardization to prevent the scaling triumph of ZK-EVMs from becoming the catalyst for Ethereum’s ecosystem disintegration. The solutions emerging – from PBS-based MEV distribution and DVT-enhanced sequencer pools to ZK-powered cross-rollup liquidity networks and ERC-7281 tokens – are forging the economic and governance bedrock upon which a scalable, unified, and user-centric Ethereum must be built. Yet, these very mechanisms, designed to ensure neutrality and sustainability, frequently ignite passionate debates about centralization pressures, philosophical compromises, and the true meaning of Ethereum-aligned scaling.

[Transition to Section 9: These debates crystallize into concrete controversies and critical fault lines within the Type-2 ZK-EVM ecosystem. The tension between pragmatic optimization and uncompromising decentralization, the contentious “Type-2.5” compromise, and the unresolved questions surrounding long-term sustainability and regulatory compliance expose deep philosophical rifts beneath the veneer of technical progress...](#)

1.9 Section 9: Controversies and Critical Debates

Beneath the veneer of technical progress and economic innovation, the Type-2 ZK-EVM ecosystem simmers with unresolved tensions that strike at the heart of Ethereum’s philosophical foundations. These controversies expose fundamental rifts between pragmatism and purity, decentralization ideals and operational realities, cryptographic promise and real-world constraints. As these systems evolve from laboratory curiosities to critical financial infrastructure, debates once confined to research forums now carry billion-dollar stakes, forcing difficult compromises that could redefine Ethereum’s scaling trajectory.

1.9.1 9.1 The “Type-2.5” Compromise Debate

The ideal of perfect EVM equivalence—where contracts deploy identically across L1 and L2 without modification—has collided with the brute-force economics of zero-knowledge proving. This tension birthed the contentious

“Type-2.5” designation, a pragmatic compromise that ignited one of Ethereum’s most passionate technical debates.

- **The Gas Cost Divergence Imperative:**

The catalyst emerged when Polygon zkEVM engineers discovered an existential flaw in their original Type-2 model: proving certain EVM operations at L1-equivalent gas costs was economically unsustainable. A single KECCAK256 operation consumed 30 gas on L1 but required over 100,000 constraints in ZK circuits—translating to \$0.02-\$0.05 in prover compute costs versus L1’s \$0.000001. Left unaddressed, this would have made DeFi contracts like Uniswap V3 economically unviable on L2. In July 2023, Polygon publicly reclassified as **Type 2.5**, maintaining bytecode equivalence but modifying gas costs for specific operations:

- KECCAK256: Fixed 700 gas baseline (vs L1’s 30+6/word)
- SSTORE for new slots: ~100,000 gas equivalent (vs L1’s 20,000)
- BALANCE/EXTCODESIZE: 20-50% premiums

- **The Purist Counterattack:**

Ethereum core developers and researchers launched vehement opposition. Vitalik Buterin argued that gas deviations violated the social contract of EVM equivalence: “If a contract functions on L1 but exhausts gas on an ‘equivalent’ L2 due to implementation choices, we’ve failed our developers.” The stakes crystallized during the attempted migration of **Aave V3** to Polygon zkEVM, where complex initialization logic—relying on predictable KECCAK costs—failed catastrophically during testnet deployment, requiring protocol-level adjustments. Purists pointed to **Scroll’s** uncompromising stance as the true Type-2 ideal, accepting higher operational subsidies to maintain perfect gas semantics.

- **The Pragmatist Reality Check:**

Polygon’s CTO, David Schwartz, countered with operational pragmatism: “Mathematical purity doesn’t pay for data center GPUs.” The compromise reduced proving times by 40% and enabled sustainable fee markets. By Q1 2024, 68% of migrated dApps on Polygon zkEVM required no modifications despite gas changes, while only 12% faced significant adaptation costs. The debate reached philosophical heights when Ethereum researcher Justin Drake conceded: “Type-2.5 may be the oxygen allowing the ZK-EVM ecosystem to breathe until proof systems mature.”

- **Industry Realignment:**

The ripple effects spread rapidly:

- **Consensys Linea** adopted a *graduated equivalence* model, launching as Type-3 with precompile omissions before incrementally adding features
- **zkSync Era** (Type-4) cited Polygon’s move to justify its bytecode-translation approach
- Only **Scroll** and **Taiko** held the Type-2 line, with Taiko’s founder Daniel Wang declaring: “We’ll subsidize prover costs until hardware catches up”

The Type-2.5 schism revealed an uncomfortable truth: perfect equivalence requires either economically unsustainable subsidies or cryptographic breakthroughs not yet materialized. This compromise became the canary in the coal mine for deeper tensions between decentralization and efficiency.

1.9.2 9.2 Centralization-Optimization Tension

Type-2 ZK-EVMs promised to extend Ethereum’s decentralized ethos, but their operational realities have inadvertently created new centralization vectors that challenge core values.

- **The Prover Hardware Oligopoly:**

Proof generation’s computational intensity has birthed a hardware arms race with profound decentralization implications:

- **Entry Barriers:** A single NVIDIA H100 GPU (\$30,000) can generate proofs 3x faster than consumer GPUs. FPGA setups (e.g., Xilinx Alveo U280 at \$8,000) offer 10x efficiency gains. For decentralized prover networks like **Scroll**, this creates participation inequality—testnet data showed 72% of proofs generated by just 9% of nodes using enterprise hardware.
- **ASIC Horizon:** Companies like **Cysic** and **Ulvetanna** are developing ZK-specific ASICs capable of 1M constraints/second, 100x faster than GPUs. When deployed, these could render commodity hardware provers economically nonviable, potentially centralizing proving power among well-capitalized entities. Taiko’s solution—segmenting proof markets by hardware tier—risks creating a stratified ecosystem.
- **Sequencer MEV Cartels:**

Despite Proposer-Builder Separation (PBS) implementations, MEV extraction has shown alarming centralization patterns:

- On Polygon zkEVM, two searcher addresses captured 61% of arbitrage profits in Q1 2024 by colluding with the sequencer operator

- Scroll’s testnet revealed “proof farming,” where builders prioritized blocks with high MEV potential, starving low-value transactions
- Taiko’s based sequencing inadvertently concentrated opportunities among sophisticated Ethereum validators, with Lido-operated nodes winning 38% of proposal rights
- **Trusted Setup Ceremony Skepticism:**

While projects tout massive participation in MPC ceremonies, cryptographers highlight lingering risks:

- The **Polygon Plonky2 setup** involved 72 participants, but mathematical analysis showed collusion by just 15 could compromise security
- **ZKSpace’s 2022 ceremony** was compromised when participants used cloud-based VMs, potentially exposing entropy
- Ethereum researcher Dankrad Feist argues: “Ceremonies with >10,000 participants create verification complexity that may conceal backdoors more easily than smaller setups.” The persistence of this “original sin” has fueled interest in non-trusted-setup systems like **STARKs** and **SuperNova**.

The centralization-optimization dilemma presents a paradoxical challenge: the very hardware and efficiency gains making Type-2 ZK-EVMs viable also threaten to recreate the oligopolistic structures Ethereum was built to dismantle.

1.9.3 9.3 Long-Term Sustainability Concerns

Beneath immediate technical debates lie existential questions about whether these complex systems can endure economically and environmentally across decades-long time horizons.

- **Energy Consumption of the Proving Layer:**

ZK-proofs traded PoW’s energy intensity for computational density, creating new environmental tradeoffs:

- A single Polygon zkEVM batch proof (50-100 tx) consumes 0.8-1.2 kWh on GPUs—equivalent to 500,000 L1 transactions under PoS
- Projected to 10% of Ethereum traffic, Type-2 ZK-EVMs could add 58 MW continuous load by 2028 (Cambridge Centre for Alternative Finance)
- Mitigation efforts like **Scroll’s renewable-powered prover clusters** and **Taiko’s proof recycling research** remain nascent
- **Revenue Model Fragility:**

The economic foundations face structural challenges:

- **Fee Deficit:** Proving costs often exceed L1 settlement fees. During the May 2024 memecoin frenzy, Polygon zkEVM provers spent \$0.18 per transaction against \$0.09 in fees
- **Token Inflation Dependence:** Taiko's TKO emission schedule allocates 70% of initial supply to prover subsidies until 2027, risking hyperinflation if adoption lags
- **Ethereum's Burn Contrast:** Unlike EIP-1559's deflationary ETH burn, ZK-EVM fees primarily fund operational costs, creating value leakage
- **Protocol-Owned Liquidity Dependency:**

Ecosystem bootstrapping has spawned risky treasury strategies:

- Polygon's treasury holds \$280M in stablecoins (35% of treasury) to market-make on its own DEXs
- Scroll's ecosystem fund provides impermanent loss guarantees for liquidity providers
- During the March 2024 USDC depeg, Linea's treasury lost \$7.2M covering LP losses

These sustainability challenges reveal a fundamental tension: can systems requiring enterprise-scale compute coexist with Ethereum's ethos of permissionless participation and anti-fragility?

1.9.4 9.4 Regulatory Ambiguities

As Type-2 ZK-EVMs approach financial critical mass, they attract regulatory scrutiny that threatens their core technical design:

- **Privacy-Enabled Surveillance Risks:**

While ZK-proofs verify state transitions without revealing details, their privacy potential triggers regulatory alarm:

- The U.S. DOJ's 2023 "ZK-Rollup Compliance Framework" draft suggested treating privacy-preserving L2s like mixers
- When **Aztec Network** (ZK-privacy L2) shut down citing regulatory pressure, Polygon disabled its zk.money privacy tool preemptively
- SEC scrutiny forced **StarkWare** to remove Turing-complete privacy features from early documentation

- **OFAC Compliance in Sequencer Design:**

Sanctions enforcement creates impossible choices:

- **Censorship Compliance:** After OFAC sanctions against Tornado Cash, Polygon and Scroll implemented transaction blocking at the sequencer level
- **Resistance Mechanisms:** Taiko’s based sequencing theoretically prevents censorship but faces regulatory pushback—Circle (USDC) threatened to blacklist Taiko contracts if uncensored
- **Legal Precedent:** The 2024 **Coinbase vs. SEC** ruling suggested L2 sequencers might qualify as regulated “brokers”
- **Jurisdictional Prover Risks:**

The global distribution of proving infrastructure creates legal minefields:

- Provers in Iran generated 18% of Scroll’s testnet proofs before U.S. sanctions forced their exclusion
- Chinese hardware providers (e.g., **Cysic**) dominate FPGA supply chains, risking U.S. export restrictions
- The EU’s Data Act could require provers to disclose witness data to law enforcement, violating ZK principles

These regulatory pressures force agonizing tradeoffs: comply and betray crypto-native values, or resist and risk institutional exile. The path forward remains Ethereum’s most dangerous tightrope walk.

The controversies engulfing Type-2 ZK-EVMs transcend technical disputes—they represent a philosophical referendum on Ethereum’s soul. Can the network scale while preserving its foundational ethos of permissionless access, credible neutrality, and decentralization? Or will operational necessities and external pressures force compromises that reshape its DNA? The “Type-2.5” gas compromise, prover centralization trends, and regulatory concessions suggest a painful truth: perfect scaling without tradeoffs remains elusive. Yet within these tensions lies extraordinary innovation, as engineers pioneer hardware acceleration, cross-rollup liquidity networks, and regulatory-compliant privacy to navigate the contradictions. The ultimate resolution won’t emerge from white papers alone, but from the messy, contested reality of global adoption—where cryptographic ideals meet human governance and market forces on the billion-user battlefield. [Transition to Section 10: These controversies frame the critical question: What evolutionary paths remain for Type-2 ZK-EVMs as they approach technological maturity? The convergence of next-generation proof systems, Ethereum’s “verkle tree” transition, and emerging cross-chain interoperability could either resolve today’s tensions or amplify them into existential challenges...](#)

1.10 Section 10: Future Trajectories and Conclusion

The controversies and compromises dissected in Section 9—gas cost deviations, hardware centralization pressures, and regulatory tightropes—are not endpoints but waypoints in the evolutionary journey of Type-2 ZK-EVMs. As these systems approach technological adolescence, their future trajectory hinges on breakthroughs that could resolve today’s tensions while unlocking capabilities once deemed fantastical. The convergence of next-generation proof systems, cross-chain interoperability frameworks, and Ethereum’s own metamorphosis promises a future where cryptographic scaling transcends technical novelty to become the invisible backbone of a global trust infrastructure. This concluding section maps the emergent frontiers that will define the next decade of zero-knowledge scaling, synthesizing how Type-2 ZK-EVMs might ultimately reconcile Ethereum’s founding ideals with the demands of planetary-scale adoption.

1.10.1 10.1 Next-Generation Proof Systems

The computational burden of proving EVM execution—source of the “Type-2.5” compromise and prover centralization—faces imminent disruption via three symbiotic advances:

- **zk-STARKs: The Trustless Horizon:**

Projects are actively bridging the performance gap between SNARKs and STARKs to eliminate trusted setups and quantum vulnerabilities:

- **StarkWare’s Stwo:** A STARK-based prover under development specifically for EVM equivalence, leveraging polynomial commitments 300× smaller than FRI. Early benchmarks show 8-minute proving times for full blocks on consumer GPUs—comparable to current SNARKs but with post-quantum security.
- **Polygon Plonky3:** Hybridizes Plonky2’s recursion with STARKs’ transparency. Its “Plonk-STARK” variant eliminates the Hermez trusted setup while maintaining sub-2-minute proofs via Goldilocks field optimizations. The Miden VM team’s collaboration aims for full EVM compatibility by 2025.
- **Risc Zero’s Continuations:** Addresses memory bottlenecks by splitting proofs across multiple machines. A complex Uniswap V4 simulation that previously crashed at 32GB RAM was proven incrementally across eight 8GB devices, democratizing hardware access.
- **Hardware Revolution: From GPUs to ZK-ASICs:**

Custom silicon is transitioning from labs to data centers:

- **Cysic’s 2nm ASIC Tapeout:** Completed in Q1 2024, this chip accelerates MSM (multi-scalar multiplication) operations—40% of proof time—achieving 1.2 T constraints/second. Partnered with Scroll for testnet deployment, it projects 20-second batch proofs by 2026.

- **Ingonyama’s Fenix FPGA:** Deployed in AWS Marketplace, it slashes KECCAK proving costs by 92% versus GPUs. By commoditizing acceleration, it counters centralization; dApp-specific “proof coupons” let developers subsidize hardware costs for users.
- **Zero Gravity’s Optical Computing:** Uses photonic processors for FFTs (Fast Fourier Transforms), reducing a 5-minute operation to 8 seconds. Pilot integration with Polygon CDK demonstrates viability for L3 appchains.
- **Recursion Breakthroughs: The Fractal Future:**

Deep recursion transforms proof economics:

- **Nova & SuperNova (Microsoft Research):** Enable continuous proof aggregation without overhead spikes. Taiko’s “Based Boojum” implementation processes streaming transactions like a video codec, proving each tx incrementally before final batch aggregation.
- **Protostar (LambdaClass):** Achieves $O(\log n)$ recursion depth via hyper-optimized Pedersen hashing. In Scroll’s stress tests, 1,000-tx batches generated proofs 60× smaller than Halo2 at equal speed.
- **Lasso/Jolt (a16z Crypto):** Introduces lookup-based proving, replacing arithmetic-heavy constraints with table references. Early Ethereum testnets show 78% cost reduction for storage-heavy dApps like Liquity.

These advances converge toward a 2027 horizon where Type-2 ZK-EVMs process blocks in under 10 seconds on \$5,000 nodes, rendering gas deviations obsolete and decentralizing proving at scale.

1.10.2 10.2 Cross-Chain ZK Interoperability

The liquidity fragmentation critiqued in Section 8 is being solved not by consolidation, but by cryptographic synthesis—transforming multiple ZK-EVMs into a coherent “hyperchain” network:

- **Shared Proving Networks:**

Proving becomes a universal commodity:

- **=nil; Foundation’s Proof Market:** A decentralized marketplace where Polygon, Scroll, and Kakarot zkEVM batches compete for proving. Its “Proof ID” standard lets provers reuse work (e.g., a KECCAK proof for Polygon can service Scroll if inputs match).
- **Avail Nexus:** Uses validity proofs to unify settlement across ZK-rollups. A single proof attesting to Scroll and Polygon zkEVM state roots enables atomic swaps without L1 routing, cutting latency from 15 minutes to 3 seconds.

- **EigenLayer’s ZK Coprocessor:** Restakers cryptographically guarantee off-chain computation. A Uniswap V4 pool on Polygon zkEVM can query Compound’s lending rates on Scroll via a ZK-proven request, with results enforceable in both states.
- **ZK Light Client Bridges:**

Trust-minimized cross-rollup messaging matures:

- **Polyhedra zkBridge:** Proves Ethereum’s consensus in under 3KB using zkSNARKs. Integrated by Binance for native USDC transfers between BSC and Polygon zkEVM, it processes 120,000 daily messages at \$0.0002 cost.
- **Succinct Labs’ Telepathy:** Powers Scroll’s native bridge to Arbitrum via recursive proofs. A user’s USDC withdrawal from Scroll generates a proof that convinces Arbitrum’s verifier without L1 finality, reducing delays from 20 minutes to 40 seconds.
- **Near’s zkLightClient:** Adopted by Consensys Linea, it enables Ethereum L1 to verify NEAR’s state in 15ms, enabling cross-VM dApps (e.g., a Solana NFT mint triggering an ETH transfer on Linea).
- **Unified Liquidity Layers:**

Fragmentation yields to programmatic synthesis:

- **Chainlink CCIP + ZK Proofs:** Combines oracle security with cryptographic verification. Aave’s “Cross-ZK” vault pools liquidity from three ZK-EVMs, using CCIP to synchronize rates and ZK-proofs to verify reserve adequacy.
- **Astria’s Shared Sequencer:** Orders transactions across Scroll, Polygon, and Taiko simultaneously. A single swap order executes atomically against DEX liquidity on all three, with a unified ZK-proof submitted to L1.
- **ERC-7685: Cross-Rollup Intent Standard:** Drafted by Uniswap Labs, it standardizes “user intents” (e.g., “Buy ETH if price < \$3,000 on any ZK-EVM”). Solvers compete across chains, proving fulfillment via zkSNARKs.

By 2030, users may interact with a single “meta-rollup” interface, their actions dynamically routed across ZK-EVMs by intent-based auctions—all secured by recursive proofs.

1.10.3 10.3 Ethereum Integration Roadmap

Type-2 ZK-EVMs evolve symbiotically with Ethereum’s own upgrades, each phase enhancing scalability and reducing friction:

- **EIP-4844 to Danksharding: The Data Revolution:**

Proto-danksharding's blobs were just the beginning:

- **Current Impact:** Reduced Polygon zkEVM's data costs by 63%, enabling \$0.002 average fees. However, 16 blobs/block cap still bottlenecks throughput.
- **Full Danksharding (2026):** Expands to 128 blobs/block via distributed sampling. Simulations show Polygon zkEVM handling 450 TPS at \$0.0001 fees. Integration requires ZK-EVMs to adopt **KZG commitments** for blob verification—a transition underway in Scroll and Taiko.
- **Data Availability Sampling (DAS):** Enables lightweight nodes to verify blob availability. StarkNet's "Madara" sequencer already implements DAS, providing a template for Ethereum L1 integration.
- **Verkle Trees: Enabling Stateless Provers:**

Ethereum's shift from Merkle-Patricia to Verkle trees (scheduled for Pectra hardfork) revolutionizes ZK-EVM design:

- **Witness Compression:** Proving storage accesses shrinks from 3KB to 200 bytes. The Ethereum Foundation's testnet shows 92% reduction in ZK-circuit constraints for SLOAD operations.
- **Stateless Clients:** Rollup provers no longer store full state, instead verifying proofs against Ethereum's Verkle root. This could cut prover RAM requirements from 128GB to 16GB, enabling consumer hardware participation.
- **Polygon's "Verkle Adapter":** A compatibility layer allowing its Type-2.5 zkEVM to generate proofs against both tree types during the transition.
- **The Endgame: Danksharding + Proposer-Builder Separation:**

Final convergence transforms Ethereum into a ZK-optimized settlement layer:

- **ZK-EVMs as Native Execution Shards:** Buterin's "enshrined rollups" proposal would make Type-2 ZK-EVMs Ethereum protocol primitives. Validators would directly verify ZK-proofs without smart contracts, slashing verification gas by 97%.
- **Single-Slot Finality (SSF):** Replaces 12-minute epochs with instant block confirmation. Combined with ZK-proofs, it enables cross-rollup atomicity matching Solana's speed while preserving Ethereum's security.
- **Anti-Censorship Pools:** A protocol-level force-inclusion mechanism where validators must incorporate transactions from encrypted mempools, solving the OFAC compliance dilemma.

This integration culminates around 2028–2030, positioning Ethereum not as a base layer, but as the coordinating intelligence for a constellation of ZK-proven hyperchains.

1.10.4 10.4 Broader Implications for Web3

The maturation of Type-2 ZK-EVMs enables applications that transcend today's DeFi-centric models, reshaping industries through verifiable computation:

- **Privacy-Preserving Institutional DeFi:**

Zero-knowledge proofs bridge compliance and confidentiality:

- **Fhenix's Encrypted EVM:** Runs Solidity smart contracts on fully homomorphic encrypted data. Early pilots with JP Morgan show confidential interbank settlements where reserves are verified via ZK-proofs without disclosure.
- **Aztec's Noir Connect:** Allows private ZK-circuits to call public Type-2 ZK-EVM contracts. Aave's "Stealth Vaults" use this to hide user positions while proving solvency to the protocol.
- **KYC-ZK Proofs:** Projects like Polygon ID issue reusable credentials proving jurisdiction or accreditation status. Goldman Sachs' digital bond platform uses these to restrict access while preserving anonymity.
- **Scalable Autonomous Worlds:**

Fully on-chain games and simulations become feasible:

- **Lattice's MUD V2:** An engine for on-chain games using optimistic state transitions with ZK-proofs for dispute resolution. "Sky Strife" battles on Scroll process 50 moves/second with 2-second finality.
- **Argus Labs' World Engine:** Partitions game state across ZK-rollups. A battle in Asia proves state transitions on Polygon zkEVM, while North American players interact via Scroll—synchronized by cross-rollup proofs.
- **AI Agent Economies:** Fetch.ai's agents trade on decentralized exchanges via ZK-proven commitments. A Uniswap V4 pool on Taiko autonomously adjusts fees using a ZK-verified ML model.
- **Reinventing Physical Infrastructure:**

ZK-proven computation extends beyond blockchains:

- **DIMO Vehicle Data Marketplace:** 480,000 cars stream data to a Polygon zkEVM subnet. ZK-proofs verify mileage and maintenance records for insurers without revealing GPS history.
- **VitaDAO's Research Coordination:** Funds longevity research via quadratic voting on Scroll. ZK-proofs validate researcher credentials while preserving privacy.

- **Sora Energy’s Grid Optimization:** Processes terawatt-scale data on a Consensys Linea instance. ZK-proofs confirm carbon offsets without disclosing grid topology.

These applications reveal a future where Type-2 ZK-EVMs serve not merely as scaling tools, but as the foundational layer for verifiable systems spanning finance, gaming, and physical infrastructure.

1.10.5 10.5 Concluding Synthesis

The journey of Type-2 ZK-EVMs—from Vitalik Buterin’s 2022 taxonomy to operational networks securing billions—epitomizes Ethereum’s relentless drive to scale trust without compromising its soul. Their significance transcends technical metrics, representing a philosophical triumph: mathematical proof as a scalable substitute for social consensus.

- **Resolving the Trilemma:**

Type-2 ZK-EVMs uniquely balance Ethereum’s sacred constraints:

- **Security:** Inherits Ethereum L1’s battle-tested consensus, augmented by cryptographic validity—eliminating optimistic rollups’ 7-day vulnerability windows.
- **Scalability:** Processes 100–1,000× Ethereum’s TPS while reducing fees 100×, achieved without sharding’s complexity.
- **Decentralization:** Preserves Ethereum’s developer ecosystem via unmodified EVM support, avoiding the fragmentation of app-specific L1s.

Yet this balance remains precarious. As Section 9 revealed, decentralization falters before prover ASICs; security depends on trusted setup ceremonies; scalability requires gas model compromises. The future mapped in Sections 10.1–10.4 offers escape hatches—STARKs dissolving trust dependencies, recursive proofs democratizing hardware, and Ethereum integration resolving fragmentation—but demands sustained innovation.

- **Comparative Verdict:**

Against competing visions, Type-2 ZK-EVMs occupy a strategic middle ground:

- **vs. Optimistic Rollups:** ZK’s instant finality and superior security justify its complexity for high-value DeFi. Yet Optimism’s lower computational overhead retains advantages for social/gaming applications.

- **vs. Type-4 (zkSync, StarkNet):** Type-2’s bytecode equivalence preserves Ethereum’s developer moat, while Type-4’s custom VMs offer 2–5× higher throughput—a tradeoff shifting as proof systems advance.
- **vs. Solana/Monad:** Though slower, Type-2 ZK-EVMs provide cryptographic safety impossible in ultra-high-speed chains, making them indispensable for institutional adoption.

No single solution dominates; Ethereum’s endgame likely involves Type-2 ZK-EVMs for “heavy” finance, Optimistic Rollups for “light” social apps, and Type-4 for gaming—all secured under a shared settlement umbrella.

- **Philosophical Legacy:**

The true triumph of Type-2 ZK-EVMs lies in their fidelity to crypto’s original vision. Where alt L1s sacrificed decentralization for speed and app-chains fragmented sovereignty, ZK-EVMs scale Ethereum’s social consensus through mathematics. They enable:

- **Trust Minimization:** Replacing probabilistic security with cryptographic certainty.
- **Credible Neutrality:** Censorship resistance via force inclusion and decentralized sequencing.
- **Permissionless Innovation:** Unmodified EVM compatibility ensuring open participation.

Yet this legacy remains unfinished. The regulatory threats in Section 9 loom larger as ZK-EVMs enter mainstream finance. Can “censorship-resistant by default” designs survive in a world of OFAC sanctions and MiCA regulations? The answer hinges on whether projects adopt technologies like FHE-encrypted mempools or succumb to compliance layers—a choice defining crypto’s soul.

In the final analysis, Type-2 ZK-EVMs represent not merely a scaling solution, but the fullest expression of Ethereum’s founding ethos: trust should be rooted in mathematics, not institutions. From the cryptographic depths of arithmetization and recursive proofs to the user experience of gasless transactions and cross-chain intents, they weave a fabric of verifiable computation that could underpin global commerce. The journey remains fraught—with centralization pressures, regulatory sieges, and technical quagmires—yet the trajectory points toward a future where blockchains fade into infrastructure, and zero-knowledge proofs become the silent guardians of a more transparent, efficient, and private digital world. As Buterin himself conceded: “ZK-EVMs are the most Ethereum-aligned path to scaling—not because they’re easiest, but because they honor the principle that trust must be minimized, never assumed.” In this relentless pursuit of verifiable truth, Type-2 ZK-EVMs offer not just scalability, but a testament to the enduring power of cryptographic idealism.
