

Encyclopedia Galactica

# "Encyclopedia Galactica: Verifiable Delay Functions"

Entry #:	473.1.9
Word Count:	10521 words
Reading Time:	53 minutes
Last Updated:	July 25, 2025

*"In space, no one can hear you think."*

## Table of Contents

### Contents

<b>1</b>	<b>Encyclopedia Galactica: Verifiable Delay Functions</b>	<b>3</b>
1.1	Section 1: The Problem of Trusted Time: Why Verifiable Delay Functions Exist . . . . .	3
1.1.1	1.1 The Tyranny of Parallelism and the Need for Sequentiality .	3
1.1.2	1.2 Defining the Gap: Trusted Timestamps vs. Decentralized Time	5
1.1.3	1.3 Core Intuition: Delay + Verifiability = Trust . . . . .	7
1.2	Section 2: Genesis and Evolution: The Historical Path to VDFs . . . .	9
1.2.1	2.1 Precursors: Time-Lock Puzzles and Early Concepts . . . . .	9
1.2.2	2.2 The Catalytic Moment: Boneh, Bonneau, Bünz, and Fisch (2018) . . . . .	11
1.2.3	2.3 Rapid Expansion: Pietrzak, Wesolowski, and Beyond . . . .	12
1.2.4	2.4 Standardization and Community Efforts . . . . .	15
1.3	Section 3: Mathematical Underpinnings and Core Properties . . . . .	16
1.3.1	3.1 The Formal VDF Definition: A Triple of Algorithms . . . . .	17
1.3.2	3.2 The Holy Trinity of Properties . . . . .	19
1.3.3	3.3 Complexity Assumptions: The Bedrock of Security . . . . .	22
1.4	Section 4: Constructing VDFs: Major Schemes and Mechanisms . . .	25
1.4.1	4.1 Wesolowski's Scheme: Efficient Verification via SNARGs . .	25
1.4.2	4.3 The Crucial Role of Groups of Unknown Order (GUOs) . . .	28
1.5	Section 5: Security Analysis: Attacks, Assumptions, and Limitations .	30
1.5.1	5.1 Cryptanalytic Attacks on Core Assumptions . . . . .	30
1.5.2	5.2 The Trusted Setup Problem . . . . .	32
1.5.3	5.3 Hardware Advantages and the ASIC Threat . . . . .	34
1.5.4	5.4 Amortization and Precomputation Attacks . . . . .	36
1.6	Section 6: VDFs in Action: Core Applications and Use Cases . . . . .	39

1.6.1	6.1 Randomness Beacons: Unpredictable and Unbiasable . . .	39
1.6.2	6.2 Enhancing Proof-of-Stake (PoS) Security . . . . .	41
1.6.3	6.3 Proofs of Space-Time (PoST) and Sustainable Consensus .	43
1.6.4	6.4 Mitigating Miner Extractable Value (MEV) . . . . .	45
1.6.5	6.5 Beyond Blockchain: Timed Releases, Auctions, and More .	47
1.7	Section 7: Implementation Landscape: Hardware, Software, and Op- timization . . . . .	49
1.7.1	7.1 The Computational Core: Modular Exponentiation and Be- yond . . . . .	50
1.7.2	7.2 Hardware Acceleration: ASICs and FPGAs . . . . .	52
1.7.3	7.3 Software Implementations and Libraries . . . . .	55
1.7.4	7.4 Optimizations for Verification and Proof Generation . . . . .	56
1.7.5	7.5 Benchmarking, Testing, and Standardization Efforts . . . . .	58
1.8	Section 8: Societal and Economic Implications . . . . .	60
1.8.1	8.1 Decentralization Revisited: Power Dynamics and Access . .	61
1.8.2	8.2 Environmental Footprint: Energy vs. Storage vs. Time . . .	62
1.8.3	8.3 Economic Models and Incentives . . . . .	64
1.8.4	8.4 Accessibility, Open Source, and Geopolitics . . . . .	66
1.9	Section 9: Frontiers of Research and Open Problems . . . . .	67
1.9.1	9.1 Post-Quantum Secure VDFs . . . . .	68
1.9.2	9.2 Continuous VDFs and Adaptive Security . . . . .	69
1.9.3	9.3 Distributed VDFs and MPC Protocols . . . . .	70
1.9.4	9.4 Improved Constructions and Proof Systems . . . . .	71
1.9.5	9.5 Major Unsolved Problems . . . . .	72
1.10	Section 10: Synthesis and Future Trajectories . . . . .	73
1.10.1	10.1 VDFs as Foundational Cryptographic Primitives . . . . .	73
1.10.2	10.2 Impact Assessment: Promises Fulfilled and Challenges Ahead . . . . .	74
1.10.3	10.3 The Broader Vision: VDFs in the Fabric of Future Systems	75
1.10.4	10.4 Ethical Considerations and Responsible Development . . .	76
1.10.5	10.5 Concluding Thoughts: The Enduring Quest for Trust in Time	77

# 1 Encyclopedia Galactica: Verifiable Delay Functions

## 1.1 Section 1: The Problem of Trusted Time: Why Verifiable Delay Functions Exist

Time is the invisible scaffolding upon which human coordination, fairness, and trust are built. From auction deadlines and bond maturation to embargoed news releases and cryptographic key rotation, the reliable passage of time underpins countless critical processes. Yet, in the digital realm, where physical clocks can be manipulated, networks experience unpredictable delays, and adversaries operate at the speed of light across continents, establishing a *provable, unforgeable* measure of elapsed time has long been a profound challenge. Verifiable Delay Functions (VDFs) emerged not merely as an interesting cryptographic curiosity, but as a direct response to this fundamental deficit in our digital infrastructure: the need to create **trustless, publicly verifiable proof that a specific, inherently sequential amount of computation time has elapsed**. This section delves into the historical and conceptual crucible from which VDFs were forged, exploring the limitations of existing solutions and articulating the precise problem they uniquely solve.

### 1.1.1 1.1 The Tyranny of Parallelism and the Need for Sequentiality

The digital age is synonymous with parallel processing. From multi-core CPUs to vast distributed computing grids and specialized ASICs (Application-Specific Integrated Circuits), throwing more hardware at a problem typically yields faster results. This is the engine of progress for most computational tasks. However, this very power becomes a liability when the goal is not merely to compute *a result*, but to *prove that a minimum amount of sequential time has been spent computing it*. Many real-world scenarios demand precisely this:

- **Timed Releases & “Crypto Time Capsules”:** Imagine wanting to encrypt a message so it can only be decrypted after a specific date in the future, perhaps to enforce an embargo, reveal a will, or schedule software updates. Simply encrypting with a key held by a trusted party reintroduces centralization and trust. Ron Rivest, Adi Shamir, and David Wagner foresaw this need in 1996, proposing the concept of “Time-Lock Puzzles.” Their elegant solution for the “MIT Time Capsule” involved creating a computational puzzle based on repeated squaring modulo a large composite number, designed to take approximately 35 years of continuous computation to solve on the hardware of the time. While a brilliant early conceptualization, it lacked *public verifiability* – anyone could solve it faster with better hardware or parallelization, and crucially, only the solver knew when it was done; there was no proof for others.
- **Fair Auctions and Preventing Last-Second Sniping:** Online auctions often suffer from “sniping,” where bidders wait until the literal last second to submit a winning bid, denying others a fair chance to respond. A naive solution might involve a trusted server enforcing a strict deadline. But in decentralized settings, how can participants *prove* that their bid was submitted *before* the deadline without revealing it prematurely? A mechanism proving that a bidder genuinely spent time (e.g., solving a puzzle) *before* the deadline could deter sniping by making instantaneous last-second bids impossible. Traditional cryptographic commitments don’t inherently encode time spent.

- **Mitigating Brute-Force Attacks and Spam:** Proof-of-Work (PoW), famously used in Bitcoin, forces participants to expend computational effort (finding a hash below a target) to access resources (e.g., sending email, minting a block). While effective in deterring spam and Sybil attacks by imposing cost, PoW fundamentally measures *work*, not *time*. An adversary with vast parallel resources (like a botnet or ASIC farm) can solve many PoW puzzles *in parallel*, effectively compressing the perceived time required for an individual action. This parallelism makes PoW unsuitable for applications requiring a *guaranteed minimum time delay* between cause and effect.
- **Generating Unpredictable Randomness:** Creating unbiased, unpredictable public randomness in decentralized systems is notoriously difficult. A common vulnerability is the “last-revealer bias,” where the final participant contributing to a randomness seed can manipulate the outcome based on knowledge of previous contributions. If participants were forced to commit to their contribution and then wait a *fixed, verifiable minimum time* before revealing it, this bias could be mitigated, as the last participant couldn’t compute an advantageous value instantly upon seeing others.

The core limitation exposed by these examples is the **parallelizability** of traditional cryptographic primitives and naive time-delay mechanisms.

- **Hashing (e.g., SHA-256):** Finding a preimage or collision is highly parallelizable. Throw more cores at it, and the solution time decreases proportionally.
- **Asymmetric Cryptography (e.g., RSA, ECC):** While signing or decrypting a *single* message might be sequential, verifying a signature or encrypting a message is fast. Crucially, processing *multiple independent inputs* (like PoW puzzles or time-lock puzzles) can be done completely in parallel. An adversary with  $N$  processors can solve  $N$  puzzles roughly  $N$  times faster than a single processor.
- **Simple Computational Puzzles:** Puzzles lacking a rigorously sequential structure are vulnerable to parallelization and hardware acceleration.

**The Challenge Defined:** The fundamental problem, therefore, is to design a *function* that is:

1. **Sequentially Slow:** Evaluating the function on a specific input  $x$  *must* require a minimum number of *sequential* computational steps,  $T$ , even for an adversary possessing arbitrary amounts of parallel computing power (e.g., polynomially bounded in the security parameter, but vastly exceeding honest participants).
2. **Efficiently Verifiable:** Given the output  $y$  and a small proof  $\pi$ , anyone must be able to verify *extremely quickly* (ideally in time logarithmic in  $T$ ,  $O(\text{poly}(\lambda, \log T))$ ) that  $y$  is indeed the correct output of the function applied to  $x$  after *approximately*  $T$  sequential steps.
3. **Unforgeable:** It should be computationally infeasible to find a valid  $(y, \pi)$  for a given  $x$  without actually performing the sequential computation, or to find two valid outputs  $(y, \pi)$  and  $(y', \pi')$  with  $y' \neq y$  for the same  $x$ .

Creating such a function means constructing computational quicksand: a task that stubbornly resists the crushing weight of parallelism, forcing even the most powerful adversary down a narrow, sequential path, while allowing anyone else to effortlessly verify the path was indeed traversed.

### 1.1.2 1.2 Defining the Gap: Trusted Timestamps vs. Decentralized Time

Before the advent of VDFs, the primary mechanisms for attesting to the passage of time relied heavily on trust in centralized authorities or were fundamentally unsuited for measuring pure, sequential time in trustless environments.

- **The Perils of Centralized Timestamping:**

- **Vulnerability:** A centralized Timestamping Authority (TSA) cryptographically signs a document hash and a claimed time. While technically sound if implemented correctly (e.g., RFC 3161), the TSA itself becomes a single point of failure and attack. Compromise of its signing key allows an attacker to backdate or forward-date any document at will.
- **Collusion:** The authority could collude with a specific party to issue fraudulent timestamps, providing false alibis or antedating contracts.
- **Availability:** Reliance on a central service introduces a dependency; if the TSA is offline or censors requests, the timestamping service becomes unavailable.
- **Scalability & Cost:** High-volume timestamping can become expensive and create bottlenecks. While systems like Certificate Transparency logs add public verifiability to the issuance of digital certificates, they still rely on trusted logs and monitors and don't inherently prove *sequential time elapsed* between events. The trust model remains fundamentally centralized.
- **Proof-of-Work: A Measure of Energy, Not Time:**

Bitcoin's Proof-of-Work (PoW) demonstrated a revolutionary way to achieve Byzantine fault tolerance in a decentralized network without a trusted timestamp server. Miners race to solve computationally difficult puzzles (hash preimages). The longest chain, representing the most cumulative work, is considered valid. While PoW incorporates *clock time* through difficulty adjustment (aiming for ~10 min/block), its core function is to measure *economic cost* (energy expended), not *sequential time*.

- **Parallelization Exploit:** The key weakness for *time measurement* is that PoW puzzles are embarrassingly parallel. If you double the hashing power (e.g., by adding more ASICs), you double the probability of finding the solution in any given time interval. There is no *inherent* minimum time enforced between blocks; a sudden massive influx of hashing power could, theoretically, solve many blocks very rapidly. This makes PoW unsuitable for applications requiring a guaranteed minimum delay (e.g., enforcing a cooldown period between actions).

- **Energy Inefficiency:** The massive energy consumption of PoW systems like Bitcoin is a direct consequence of using easily parallelizable computations to impose cost. This is environmentally unsustainable and economically wasteful if the *only* goal is to prove time elapsed, rather than securing a multi-trillion dollar ledger.
- **Subjectivity:** The “time” measured by PoW is probabilistic and relative to the current global hashrate. It doesn’t provide a concrete proof that “at least X seconds of sequential computation occurred” on a specific input.
- **Proof-of-Stake and the Missing Time Dimension:** Proof-of-Stake (PoS) consensus mechanisms replace computational work with economic stake. Validators are chosen to propose and attest to blocks based on the amount of cryptocurrency they “stake” as collateral. While vastly more energy-efficient than PoW, early PoS designs faced significant security challenges:
- **Nothing at Stake:** In a fork (multiple competing chains), a rational validator has nothing to lose by voting on *all* forks to maximize reward chances, potentially hindering consensus finality.
- **Long-Range Attacks:** An attacker acquiring old private keys (even if the associated coins were later spent or slashed) could theoretically rewrite history from a point far in the past, as creating blocks in PoS has negligible computational cost compared to PoW. Defenses like “weak subjectivity” require new nodes to trust recent checkpoints, reintroducing a form of trust.
- **The Need for Physical Time:** These vulnerabilities stem partly from the lack of a robust, decentralized link to *physical time*. PoS provides cryptographic security based on economic incentives but lacks a mechanism to prove that real, sequential time has passed between critical events, making certain attacks feasible.

**Articulating the Precise Need:** The limitations of trusted authorities and existing decentralized mechanisms like PoW and PoS highlight the stark gap: the absence of a cryptographic primitive capable of generating **unforgeable, publicly verifiable evidence of sequential computation time elapsed**. What is required is a function that:

- Binds a specific input  $x$  to an output  $y$  through a computation that *cannot be shortcut* by any amount of parallel processing.
- Produces a succinct proof  $\pi$  allowing anyone to instantly confirm that  $y$  is the correct output of  $T$  sequential steps on  $x$ .
- Operates without any trusted third party.
- Is efficient enough to be practical for real-world applications.

This is the void Verifiable Delay Functions were conceived to fill.

### 1.1.3 1.3 Core Intuition: Delay + Verifiability = Trust

The essence of a Verifiable Delay Function can be grasped intuitively: **It is a mathematical function that is intentionally slow to compute in a sequential manner, yet fast for anyone to verify.** Formally, a VDF is defined by a triple of algorithms:

1. **Setup( $\lambda, T$ )  $\rightarrow$  pp:** Generates public parameters pp based on a security parameter  $\lambda$  (governing the difficulty of breaking cryptographic assumptions) and the desired time delay  $T$ .
2. **Eval(pp, x)  $\rightarrow$  (y,  $\pi$ ):** Takes the public parameters pp and an input x, performs a sequential computation requiring approximately  $T$  steps, and outputs a value y and a (typically small) proof  $\pi$ .
3. **Verify(pp, x, y,  $\pi$ )  $\rightarrow$  {Accept, Reject}:** Takes pp, x, y, and  $\pi$ , and quickly verifies (in time much less than  $T$ , ideally  $O(\text{poly}(\lambda, \log T))$ ) whether y is indeed the correct output of Eval on x.

The magic lies in the properties enforced by the underlying mathematics:

- **Sequentiality:** The evaluation function Eval must be inherently sequential. No matter how many parallel processors an adversary throws at the problem (up to a limit defined by  $\lambda$ ), they cannot compute (y,  $\pi$ ) significantly faster than  $T$  sequential steps. The computation acts like a single-file path; adding more people doesn't make the path wider or traversable faster in parallel.
- **$\delta$ -Efficiency (Fast Verification):** Verification must be exponentially faster than evaluation. If evaluation takes time  $T$ , verification should take time logarithmic in  $T$  (e.g.,  $O(\log T)$ ), or at worst, polynomial in the security parameter but independent of  $T$ . This ensures that verifying the proof of time elapsed is trivial compared to generating it.
- **Uniqueness (or Soundness):** For a given (pp, x), it should be computationally infeasible to find valid (y,  $\pi$ ) and (y',  $\pi'$ ) where  $y' \neq y$ . This guarantees that the output y is uniquely determined by x and pp, preventing an adversary from creating multiple “valid” proofs for different outcomes. Some definitions emphasize soundness, meaning it's hard to find *any* valid (y,  $\pi$ ) for an incorrect y.

### Distinguishing VDFs from Proof-of-Work and Proof-of-Stake:

- **VDF vs. PoW:** While both impose computational effort, their goals and mechanisms differ fundamentally.
- **Goal:** PoW measures *total work* (easily parallelizable) to impose economic cost and secure consensus. VDFs measure *sequential time elapsed* (resistant to parallelization) to create verifiable delays.
- **Parallelism:** PoW thrives on parallelism; more hash power solves puzzles faster. VDFs *resist* parallelism; more processors don't speed up a single evaluation.



- **Verification:** PoW verification is fast (hashing is quick), similar to VDFs. However, PoW verification confirms *work done*, while VDF verification confirms *sequential time elapsed*.
- **Resource:** PoW consumes significant energy proportional to the work done. VDF evaluation consumes energy proportional to the time delay  $T$ , but crucially, this energy cost is *fixed* for a given  $T$  and hardware efficiency; throwing more hardware at it doesn't reduce the time or the *per-evaluation* energy cost (though it allows more evaluations *concurrently*).
- **VDF vs. PoS:** PoS is resource-based, relying on staked capital. VDFs are resource-agnostic in principle; they require computation, but the key property is the *sequential nature* of that computation, not the economic value staked. VDFs provide a *physical time anchor* that PoS inherently lacks.

### The Promise: Trustless Coordination Based on Elapsed Time

The advent of VDFs unlocks transformative possibilities for decentralized systems:

- **Fair Randomness Beacons:** By requiring participants to commit to a random seed and then wait a fixed, verifiable delay enforced by a VDF before revealing it, the ability of the last revealer to manipulate the final result is neutralized (e.g., Ethereum 2.0's RANDAO/VDF hybrid).
- **Securing Proof-of-Stake:** Incorporating VDFs ("Proof-of-Sequential-Time") into PoS protocols can mitigate long-range attacks and the nothing-at-stake problem by adding a physical time dimension. Validators might need to complete a VDF after being selected before they can propose a block, making rapid chain reorganization infeasible.
- **Proofs of Space-Time (PoST):** Combining Proofs-of-Space (demonstrating allocated storage) with VDFs ensures that storage proofs cannot be generated instantly on demand. Miners must prove they stored the data for the duration of the VDF delay, enabling sustainable consensus mechanisms (e.g., Chia, Spacemesh).
- **Mitigating Miner Extractable Value (MEV):** Enforcing a VDF delay between transaction ordering and block building can prevent front-running by allowing decentralized actors time to detect and potentially counter manipulative ordering.
- **Enhanced Timed Releases and Auctions:** Truly decentralized and verifiable time-lock encryption and auction deadlines become feasible.
- **Resource Pricing and Spam Control:** Revisiting early concepts like "pricing via processing," VDFs offer a way to impose mandatory, verifiable delays as a spam deterrent that is resistant to parallelization by botnets, unlike simple PoW.

VDFs thus emerge as a foundational primitive, bridging the gap between the physical reality of time and the digital need for verifiable proof of its passage. They offer a way to build *trust* in the progression of events within decentralized networks, not through centralized authorities or probabilistic economic mechanisms

alone, but through the unforgeable constraints of sequential computation itself. The quest to formalize this powerful intuition, to find mathematical functions possessing these remarkable properties, would become a central pursuit in theoretical and applied cryptography, driven by the urgent needs of a rapidly decentralizing world.

The conceptual groundwork laid here – the tyranny of parallelism, the inadequacy of trusted authorities and traditional mechanisms like PoW for measuring pure sequential time, and the core intuition of delay plus verifiability equating trust – sets the stage for exploring the fascinating intellectual journey that led to the concrete realization of VDFs. We now turn to the **Genesis and Evolution** of these functions, tracing the path from early precursors like Rivest’s time-lock puzzle to the seminal formalization by Boneh, Bonneau, Bünz, and Fisch, and the subsequent explosion of research and constructions that followed.

(Word Count: ~1,980)

---

## 1.2 Section 2: Genesis and Evolution: The Historical Path to VDFs

The conceptual imperative outlined in Section 1 – the urgent need for a primitive capable of generating trust-less, publicly verifiable proofs of sequential time elapsed – did not emerge in a vacuum, nor was its solution immediate. The formalization of Verifiable Delay Functions (VDFs) represents the culmination of decades of cryptographic exploration, driven by evolving needs and punctuated by key intellectual breakthroughs. This section traces that fascinating journey, from early, partial solutions grappling with the core challenge to the catalytic moment of formal definition and the subsequent explosion of research and practical constructions that solidified VDFs as a fundamental cryptographic tool.

The concluding insight of Section 1 – that VDFs bridge the gap between physical time and digital verification through the unforgeable constraints of sequential computation – sets the stage for understanding the historical struggle to capture this elusive property. Prior to 2018, cryptographers devised ingenious mechanisms that hinted at the possibility but fell short of the complete, robust solution VDFs would provide. The path to VDFs is a testament to incremental progress, where each step illuminated a facet of the problem or offered a building block, waiting for the right catalyst to assemble them into a coherent whole.

### 1.2.1 2.1 Precursors: Time-Lock Puzzles and Early Concepts

Long before the term “Verifiable Delay Function” was coined, cryptographers recognized the need to enforce computational delays and grappled with the challenge of sequentiality.

- **Rivest, Shamir, and Wagner’s Time-Lock Puzzle (1996):** This seminal work, explicitly motivated by the desire for “sending messages into the future” (like the famous **MIT Time Capsule Crypto-Puzzle** celebrating 35 years of MIT’s Laboratory for Computer Science), provided the first clear cryptographic mechanism for creating a computational delay. Their elegant solution relied on **repeated**

**squaring modulo a large composite number  $N = pq$ .** To encrypt a message  $M$  for time  $T$ , the sender:

1. Generates  $N = pq$  (keeping  $p$  and  $q$  secret).
2. Computes  $\phi(N) = (p-1)(q-1)$ .
3. Computes  $t = T * K$  (where  $K$  is an ops-per-second estimate for the target hardware).
4. Chooses a random key  $K_s$  for a symmetric cipher.
5. Computes  $C = (2^{(2^t \bmod \phi(N))} \bmod N) \bmod N$  *using the knowledge of  $\phi(N)$*  (via Euler's theorem:  $2^{(k \bmod \phi(N))} \equiv 2^k \bmod N$ ). This allows fast computation *for the sender*.
6. Sets the puzzle as  $(N, t, E_{K_s}(M), C)$ .

The receiver must compute  $S = 2^{(2^t)} \bmod N$  through  $t$  sequential squarings to recover  $K_s = S \bmod N$  and decrypt  $M$ . While revolutionary, this scheme had critical limitations: **Lack of Public Verifiability.** Only the solver knows when the puzzle is solved; there is no succinct proof  $\pi$  they can show others to *prove* they performed  $t$  steps and recovered the correct  $K_s$ . **Trusted Setup.** The security relies entirely on the sender generating  $N$  correctly and destroying  $p$  and  $q$ . If  $p$  and  $q$  are leaked, the puzzle can be solved instantly. **Parallelization Vulnerability (Theoretical).** While inherently sequential *per puzzle*, an adversary could solve *many independent puzzles* in parallel. Rivest et al. acknowledged this, suggesting  $t$  be set large enough to deter such attacks, but it lacked the formal sequentiality guarantee against polynomially bounded parallel adversaries that defines VDFs. Nevertheless, the core idea of using inherently sequential computations (like iterated squaring in a group) as a time delay anchor was foundational.

- **Mahmoody, Moran, and Vadhan's Publicly Verifiable Proofs of Sequential Work (PVPoSW) (2013):** This work made a crucial leap towards VDFs by explicitly formalizing the notion of proving sequential computation. They constructed a protocol where a prover, given a statement  $x$  and a time parameter  $T$ , could compute a proof demonstrating they spent  $T$  sequential steps starting from  $x$ . Crucially, this proof was *publicly verifiable* in time much less than  $T$ . Their construction was complex, relying on depth-robust graphs and Merkle tree commitments to force sequential computation. While achieving the core goals of sequentiality and verifiability, it suffered from significant practical drawbacks: **Large Proof Size.** Proofs were linear in  $T$  (e.g., hundreds of MB or even GB for large  $T$ ), making them impractical for many applications. **Complex Verification.** While asymptotically faster than  $T$ , verification was still relatively expensive ( $O(T / \log T)$ ). **No Uniqueness.** Their construction didn't guarantee a unique output for a given input. Despite these limitations, PVPoSW provided the first rigorous cryptographic definition and proof for the concept of verifiable sequential computation, demonstrating its feasibility and laying vital theoretical groundwork. It explicitly framed the problem as a cryptographic primitive distinct from Proof-of-Work.

- **Dwork and Naor’s Pricing via Processing (1992) and the Spam Battle:** Motivated by combating junk email, Cynthia Dwork and Moni Naor proposed a scheme requiring senders to solve a moderately hard, but not necessarily sequential, computational puzzle for each email. The cost imposed by this “pricing via processing” was intended to deter mass spamming. While not focused on *sequential* time per se, this work was an early influential example of using computational effort as a sybil-resistance mechanism in decentralized settings. It highlighted the need for functions that were inherently costly (in time or resources) to compute but cheap to verify, a core characteristic shared with VDFs, though VDFs specifically target *sequential* cost resistant to parallelization. Later proposals explored more sequential puzzles for spam, but often lacked formal guarantees or efficient public verifiability.

These precursors illuminated different facets of the problem: Rivest et al. demonstrated the power of sequential squaring for enforced delays but lacked verifiability. Mahmoody et al. achieved rigorous sequentiality proofs and verifiability but with impractical overhead. Dwork and Naor highlighted the application space requiring cost functions. The stage was set for a synthesis – a primitive combining the practical sequentiality of time-lock puzzles with the robust public verifiability of PVPoS, all within an efficient and uniquely defined framework.

### 1.2.2 2.2 The Catalytic Moment: Boneh, Boneau, Bünz, and Fisch (2018)

The year 2018 marked the pivotal moment when the scattered concepts coalesced into the modern definition of Verifiable Delay Functions. The catalyst was the rapidly evolving landscape of blockchain technology, particularly the drive towards **Proof-of-Stake (PoS) consensus** and the need for **unbiasable randomness beacons**.

- **The Seminal Paper:** Dan Boneh, Joseph Boneau, Benedikt Bünz, and Ben Fisch published “[Verifiable Delay Functions](#)” in May 2018 (with a major update in October 2018). This paper did more than just propose a new construction; it **formally defined the VDF primitive** and its three core properties: **Sequentiality**,  **$\delta$ -Efficiency (Fast Verification)**, and **Uniqueness**. This clear, rigorous definition provided the essential vocabulary and framework for the field. They articulated the necessity for each property in the context of applications like secure leader election in PoS and randomness generation.
- **Initial Constructions:** The paper proposed two main candidate constructions:
  1. **Injective Rational Maps:** Based on the sequentiality of computing isogenies of large degree between supersingular elliptic curves. While theoretically promising, especially for post-quantum security, this construction was complex and lacked efficient implementations at the time.
  2. **Repeated Squaring in Groups of Unknown Order (GUOs):** Building directly on Rivest’s time-lock puzzle concept, but crucially adding mechanisms for *verifiability*. Their primary candidate used **RSA groups** ( $\mathbb{Z}_{N^*}$  for  $N = pq$ ). They proposed an interactive protocol (based on the Fiat-Shamir heuristic to make it non-interactive) where the prover performs  $T$  sequential squarings  $y = x^{(2^T)}$

$\text{mod } N$  and generates a proof  $\pi$  by interacting with a verifier (or simulating the interaction via a hash function) to prove correctness without revealing  $T$  intermediate steps. This addressed the critical verifiability gap in Rivest’s puzzle.

- **The Blockchain Catalyst:** The timing was not coincidental. Ethereum, spearheaded by Vitalik Buterin and researchers like Justin Drake, was actively exploring the transition from Proof-of-Work (PoW) to PoS (Ethereum 2.0, now the Ethereum consensus layer). A major challenge was generating unpredictable, unbiased randomness for validator selection (the “randomness beacon” problem). Existing solutions like RANDAO (where validators collectively generate randomness by revealing hashes) were vulnerable to last-revealer manipulation. VDFs offered an elegant solution: feed the RANDAO output into a VDF. The enforced sequential delay prevents the last revealer from predicting and manipulating the final output, as they cannot compute the VDF output faster than anyone else. VDFs were also seen as a potential defense against long-range attacks in PoS by adding a physical time cost to block production or validation. The paper explicitly cited these blockchain applications as primary motivators.
- **Immediate Impact:** The Boneh et al. paper sent ripples through both the cryptography and blockchain communities. It provided a formal foundation and a practical path forward for solving critical problems in decentralized systems. Cryptographers recognized it as defining a fundamental new primitive with wide-ranging implications beyond blockchains. Blockchain developers saw it as a potential keystone for building more secure, efficient, and fair protocols. Research groups at institutions like the Ethereum Foundation, Stanford, and EPFL, along with nascent blockchain projects like Chia and Dfinity, quickly began exploring implementations and further refinements. The hunt for simpler, more efficient, and setup-minimizing constructions intensified immediately.

The Boneh-Bonneau-Bünz-Fisch paper crystallized years of conceptual exploration. By naming the primitive, rigorously defining its properties, proposing concrete constructions rooted in established concepts like time-lock puzzles, and linking them directly to the pressing needs of emerging decentralized technologies, they ignited the field of VDF research.

### 1.2.3 2.3 Rapid Expansion: Pietrzak, Wesolowski, and Beyond

Following the formal definition, progress on VDF constructions accelerated at a remarkable pace. Within months, simpler and more efficient schemes emerged, addressing limitations of the initial proposals and solidifying groups of unknown order as the leading paradigm.

- **Krzysztof Pietrzak’s Elegant Recursive Scheme (2018):** Pietrzak, building on concepts from his earlier work on Proofs of Sequential Work, proposed a beautifully simple and efficient VDF based on **repeated squaring in RSA groups**. His key innovation was a **recursive proof composition** strategy:

1. **Evaluation:** The prover computes  $y = x^{(2^T)} \text{ mod } N$ .

2. **Proof Generation:** The prover doesn't just output  $y$ ; they compute intermediate points (e.g.,  $\mu = x^{(2^{\{T/2\}})} \bmod N$ ). The proof  $\pi$  consists of  $\mu$  and proofs for the two halves: that  $\mu$  is  $x^{(2^{\{T/2\}})}$  and that  $y$  is  $\mu^{(2^{\{T/2\}})}$ . These sub-proofs are themselves generated recursively. This structure allows the verifier to check consistency between the root ( $y$ ), the midpoint ( $\mu$ ), and the input ( $x$ ).
  3. **Verification:** The verifier checks the recursive structure. Crucially, the proof size is  $O(\log T)$  (the depth of the recursion tree) and verification time is also  $O(\log T)$ , dominated by a few modular exponentiations. The proof leverages the homomorphic property of exponentiation ( $(x^a)^b = x^{a*b}$ ) and the Fiat-Shamir transform to make the challenge for the midpoint non-interactive. Pietrzak's scheme stood out for its conceptual clarity and relatively straightforward implementation compared to the interactive protocols in the original Boneh et al. paper. It became a primary reference construction.
- **Benjamin Wesolowski's Efficient Verification Scheme (2018):** Wesolowski, independently and nearly simultaneously, proposed another major breakthrough based on **class groups of imaginary quadratic fields**, although his scheme also works with RSA groups. His scheme achieved **constant-size proofs** and **constant-time verification** (with respect to  $T$ ), a significant efficiency leap:

1. **Evaluation:** Compute  $y = x^{(2^T)} \bmod N$  (RSA) or in a class group.
2. **Proof Generation:**
  - The verifier (or a hash function, via Fiat-Shamir) generates a small prime  $\ell$  based on  $x, y, T$ .
  - The prover computes  $q$  and  $r$  such that  $2^T = q * \ell + r$  (with  $0 \leq r < \ell$ ).
  - The prover computes  $\pi = x^q \bmod N$  (or in the class group). The proof is simply  $\pi$ .
3. **Verification:**
  - The verifier recomputes  $\ell$  (via Fiat-Shamir).
  - Computes  $r = 2^T \bmod \ell$ .
  - Checks that  $y \stackrel{?}{=} \pi^\ell * x^r \bmod N$ . This holds because  $\pi^\ell * x^r = (x^q)^\ell * x^r = x^{q\ell + r} = x^{2^T} = y$ .

Verification requires essentially one full exponentiation ( $\pi^\ell$ ) and one small exponentiation ( $x^r$ ), independent of  $T$ . The proof is just one group element. This made Wesolowski's scheme exceptionally attractive for blockchain applications where small proof size and fast verification are paramount. The use of class groups offered a major additional advantage: **no trusted setup**. Unlike RSA groups, which require someone to

generate the modulus  $N = pq$  and then (ideally) destroy  $p$  and  $q$  – a significant point of trust – class groups of imaginary quadratic fields are defined by a publicly chosen discriminant  $-D$  (where  $-D$  is a negative prime congruent to 1 mod 4, or a product of such primes). There is no trapdoor; the group structure is inherently “unknown order” based on the difficulty of computing the class number. This aligned perfectly with the decentralization ethos.

- **Subsequent Refinements and Alternative Paths:** The Pietrzak and Wesolowski schemes became the dominant practical VDF paradigms (RSA and Class Group based). Research immediately focused on:
- **Optimizations:** Improving the efficiency of the core squaring operation, proof generation, and verification, especially for class group arithmetic which is more complex than RSA modular arithmetic.
- **Security Proofs:** Strengthening the security arguments, particularly formalizing the Sequentiality Assumption within the Algebraic Group Model (AGM) or similar frameworks.
- **Groups of Unknown Order (GUOs):** Deepening the understanding of RSA groups vs. class groups – their relative security (factoring vs. class group discrete logarithm/order computation), performance trade-offs, and suitability for different applications. Exploring other potential GUO candidates like Jacobians of hyperelliptic curves.
- **Isogeny-Based VDFs:** Revisiting Boneh et al.’s initial isogeny proposal. Researchers explored constructions based on the sequentiality of computing large-degree isogenies between supersingular elliptic curves (e.g., using CSIDH or SQIsign primitives). These hold promise for **post-quantum security** as they don’t rely on factoring or discrete log, but currently suffer from much slower evaluation and verification times and larger parameters compared to GUO-based VDFs. Lattice-based VDFs were also investigated but faced significant hurdles in achieving competitive sequentiality guarantees.
- **Incremental Verification and SNARKs:** Exploring whether techniques from the realm of succinct non-interactive arguments of knowledge (SNARKs) and Incrementally Verifiable Computation (IVC) could be adapted to construct VDFs, potentially offering even smaller proofs or adaptive security, though often at the cost of heavier cryptographic machinery and trusted setups. **Hash-based VDFs (e.g., Sloth)** were proposed as simple, quantum-safe alternatives relying on the sequentiality of hash chains or graph pebbling, but typically offer weaker sequentiality guarantees than number-theoretic VDFs, as parallelism can offer some speedup, especially with precomputation.

The period immediately following the 2018 definition was characterized by intense creativity and rapid convergence. Pietrzak and Wesolowski provided the elegant, efficient, and relatively simple constructions that made practical deployment conceivable, while the broader cryptographic community explored the boundaries of what was possible, seeking alternatives and improvements.



### 1.2.4 2.4 Standardization and Community Efforts

The surge of theoretical interest and the clear demand from practical applications, particularly within major blockchain ecosystems, necessitated a shift towards standardization, implementation, and collaborative research to ensure robustness, interoperability, and security.

- **The IETF VDF Working Group (2019-2022):** Recognizing the potential for VDFs to become critical internet infrastructure, the Internet Engineering Task Force (IETF) established the [VDF Working Group](#) in 2019. Its mission was to “standardize one or a small number of VDF constructions along with their parameters to provide interoperability and security.” The group engaged cryptographers, blockchain developers, and hardware engineers. It focused intensely on evaluating the security and practicality of the leading candidates (primarily Wesolowski and Pietrzak schemes using RSA and Class Groups), analyzing implementation considerations, discussing parameter selection, and exploring the trusted setup problem for RSA-based VDFs. While the working group concluded in 2022 without publishing a formal standard RFC, it produced valuable internet-drafts and fostered significant collaboration and understanding. The work highlighted the complexities involved, especially concerning trusted setups and the desire for post-quantum options, leaving the door open for future standardization efforts as the field matures.
- **The Ethereum Foundation’s VDF Initiative:** As a primary driver of VDF adoption for its PoS beacon chain randomness, the Ethereum Foundation invested heavily in VDF research and development. This included:
  - **VDF Research:** Funding academic teams (like those at Stanford, UIUC, and EPFL) to analyze security, optimize constructions, and explore alternatives.
  - **RSA ASIC Challenge:** Launching a high-profile competition to design and build **specialized ASICs** for accelerating the repeated squaring core of RSA-based VDFs (the Wesolowski/Pietrzak model). The goal was to democratize access to efficient VDF computation, preventing centralization by entities who could afford custom hardware. Teams from Supranational, Synopsys, and others participated, leading to significant open-source hardware designs (e.g., [EPFL’s VDF ASIC](#)). While Ethereum ultimately shifted its near-term randomness beacon design away from VDFs due to complexity and hardware dependency concerns (opting for a simpler RANDAO+attester shuffling approach), the competition dramatically advanced the state-of-the-art in VDF hardware acceleration and demonstrated the feasibility of ASIC implementations achieving the necessary performance.
  - **Software Implementations:** Developing and maintaining open-source VDF libraries (like [Chia VDF](#) - though primarily driven by Chia, Ethereum contributed significantly, and research codebases) for experimentation and integration.
- **Collaborative Research Hubs and Ongoing Exploration:** Beyond formal standardization and Ethereum, a vibrant research community continues to explore VDFs. Academic conferences (CRYPTO, EURO-CRYPT, TCC) regularly feature new VDF results. Projects like Chia Network (using VDFs heavily



in its Proofs of Space and Time consensus) and Dfinity/Internet Computer have become significant drivers of applied research and development, pushing optimizations for class groups and exploring integration challenges. Open-source libraries like [Chia’s VDF](#) (supporting class groups) and [Filecoin’s rust-fil-proofs](#) (incorporating VDF elements) provide critical tools for the community. Forums like the [VDF Research Discord](#) (initially driven by Ethereum efforts) serve as hubs for discussion among researchers and engineers.

The journey from Rivest’s 1996 puzzle to the vibrant ecosystem of VDF research and development circa 2024 illustrates the dynamic interplay between theoretical insight and practical necessity. The formal definition by Boneh, Bonneau, Bünz, and Fisch acted as the catalyst, crystallizing decades of conceptual work. Pietrzak and Wesolowski provided the elegant, efficient constructions that unlocked practical potential. And the collaborative efforts of standardization bodies, blockchain foundations, and the open-source research community are now forging these ideas into robust, implementable tools. VDFs transitioned from an intriguing theoretical possibility to a concrete primitive actively shaping the design of decentralized systems.

This historical evolution sets the stage for a deeper dive into the **Mathematical Underpinnings and Core Properties** that make VDFs possible. Understanding the formal definitions, the precise security properties, and the complexity-theoretic assumptions they rely upon is essential for appreciating their power, their limitations, and the ongoing quest to improve them.

(Word Count: ~2,020)

---

### 1.3 Section 3: Mathematical Underpinnings and Core Properties

The historical journey from Rivest’s time-lock puzzle to the catalytic formalization by Boneh et al. and subsequent breakthroughs by Pietrzak and Wesolowski reveals a profound truth: Verifiable Delay Functions (VDFs) are not mere clever engineering, but deep mathematical constructs resting on rigorous foundations. This section dissects the formal anatomy of VDFs, laying bare the precise definitions, the indispensable properties that define their essence, and the complex computational assumptions that underpin their security. Understanding this bedrock is crucial, for it illuminates both the remarkable power and the inherent limitations of these cryptographic primitives that bind digital trust to physical time.

Following the explosive innovation chronicled in Section 2, where elegant constructions like Pietrzak’s recursive proofs and Wesolowski’s constant-verification scheme emerged, the cryptographic community faced a critical task: establishing a rigorous, shared understanding of what constitutes a *true* VDF. This demanded moving beyond specific implementations to define the abstract mathematical skeleton – the algorithms, properties, and assumptions – that any valid VDF must satisfy. It is this formal framework that enables security proofs, meaningful comparisons between schemes, and confidence in their deployment within critical systems like Ethereum’s beacon chain or Chia’s consensus protocol.

### 1.3.1 3.1 The Formal VDF Definition: A Triple of Algorithms

At its core, a Verifiable Delay Function is formally defined as a triple of probabilistic polynomial-time algorithms: ( $\text{Setup}$ ,  $\text{Eval}$ ,  $\text{Verify}$ ). This structure encapsulates the lifecycle of a VDF instance, from initialization to proof generation and validation. Each algorithm plays a specific role, governed by security ( $\lambda$ ) and time ( $T$ ) parameters:

#### 1. $\text{Setup}(\lambda, T) \rightarrow \text{pp}$ (Public Parameter Generation):

- **Inputs:** A **security parameter**  $\lambda$  (typically expressed in bits, e.g.,  $\lambda = 128$  or  $\lambda = 256$ ) governing the cryptographic strength and the difficulty of breaking the underlying assumptions, and a **delay parameter**  $T$  specifying the minimum number of sequential computational steps required for evaluation.
- **Output: Public parameters**  $\text{pp}$ . These define the specific instance of the VDF. Crucially,  $\text{pp}$  must be generated *independently* of any specific input  $x$ .
- **Function:** This algorithm sets the stage. For constructions based on **Groups of Unknown Order (GUOs)**:
  - *RSA-based (e.g., Pietrzak, Wesolowski):*  $\text{Setup}$  generates a large RSA modulus  $N = pq$ , where  $p$  and  $q$  are distinct safe primes. The public parameters  $\text{pp}$  include  $N$  and often the generator  $g$  (usually 2 for simplicity). The critical point is that the factorization of  $N$  ( $p$  and  $q$ ) must be securely discarded or managed via MPC (see Section 3.3).  $T$  is incorporated implicitly by the evaluator performing  $T$  squarings.
  - *Class Group-based (e.g., Wesolowski):*  $\text{Setup}$  selects a suitable discriminant  $-D$  for an imaginary quadratic field, defining the class group  $\text{Cl}(-D)$ . The discriminant  $-D$  becomes part of  $\text{pp}$ . No secrets need discarding; the group order is inherently unknown based on the difficulty of computing the class number.  $T$  is again handled during evaluation.
- **Example:** Consider an RSA-based VDF targeting 128-bit security ( $\lambda=128$ ) and a 1-minute delay ( $T \approx 10^9$  sequential steps on target hardware).  $\text{Setup}$  would generate a 3072-bit or 4096-bit RSA modulus  $N$  (depending on best practices against factoring attacks), outputting  $\text{pp} = (N, g=2)$ . The secrecy of  $p$  and  $q$  is paramount; their compromise allows instant VDF evaluation using Euler's theorem, as in Rivest's original time-lock puzzle.

#### 2. $\text{Eval}(\text{pp}, x) \rightarrow (y, \pi)$ (Evaluation):

- **Inputs:** The public parameters  $\text{pp}$  and an **input**  $x$  (an arbitrary bit-string, often a hash or commitment in applications).
- **Output:** An **output**  $y$  and a **proof**  $\pi$ .

- **Function:** This is the computationally intensive heart of the VDF. Given  $pp$  and  $x$ ,  $\text{Eval}$  performs a computation designed to take *at least*  $T$  sequential steps, even given arbitrary polynomial (in  $\lambda$ ) parallelism. The output  $y$  is deterministically derived from  $pp$  and  $x$  via this computation. The proof  $\pi$  is auxiliary data enabling fast verification.
- **Mechanism (Illustrative - RSA Repeated Squaring):** For  $pp = (N, g)$  and input  $x$  treated as an element in  $\mathbb{Z}_N^*$ ,  $\text{Eval}$  computes:

$$y = g^{(2^T)} \bmod N$$

This requires performing  $T$  sequential modular squarings: start with  $a_0 = g$ , compute  $a_1 = a_0^2 \bmod N$ ,  $a_2 = a_1^2 \bmod N$ , ...,  $a_T = a_{T-1}^2 \bmod N$ , so  $y = a_T$ . Crucially, while each squaring is fast, the process *cannot* be parallelized; each step depends on the result of the previous one. Pietrzak's scheme generates a recursive proof  $\pi$  involving intermediate values, while Wesolowski's generates a compact proof  $\pi$  based on a challenge derived from  $x, y, T$ .

### 3. **Verify**( $pp, x, y, \pi$ ) $\rightarrow$ {**Accept**, **Reject**} (Verification):

- **Inputs:** Public parameters  $pp$ , input  $x$ , claimed output  $y$ , and proof  $\pi$ .
- **Output:** A binary decision: **Accept** if  $y$  and  $\pi$  constitute valid proof that  $y$  is the correct output of  $\text{Eval}(pp, x)$  after  $\sim T$  sequential steps; **Reject** otherwise.
- **Function:** This algorithm must be extremely efficient, running in time significantly less than  $T$  – ideally, polynomial in  $\lambda$  and *logarithmic* in  $T$  ( $O(\text{poly}(\lambda, \log T))$ ).
- **Mechanism (Illustrative - Wesolowski Verification):** Given  $pp = (N, g)$ ,  $x$  (as  $g$  in this context),  $y$ , and  $\pi$  (a single group element):
  - Derive a small prime  $\ell$  using a hash function (Fiat-Shamir) applied to  $(pp, x, y, T)$ .
  - Compute  $r = 2^T \bmod \ell$  (fast modular exponentiation).
  - Check if  $y \stackrel{?}{=} \pi^\ell * x^r \bmod N$ .
- **Efficiency:** The dominant cost is computing  $\pi^\ell \bmod N$ , which depends on the bit-length of  $\ell$  (e.g., 128-256 bits), *not* on  $T$ . This is  $O(\lambda)$  time, exponentially faster than the  $O(T)$  evaluation time.

### The Role of $\lambda$ and $T$ :

- **$\lambda$  (Security Parameter):** Governs the cryptographic hardness. Larger  $\lambda$  means larger groups (e.g., larger  $N$  or larger  $|D|$ ), making underlying problems like factoring or computing class group structure exponentially harder. It bounds the computational power of any efficient adversary (modeled as probabilistic polynomial-time in  $\lambda$ ).

- **$T$  (Delay Parameter):** Governs the minimum *sequential* time required for honest evaluation. It is typically set based on the desired real-world time delay (e.g., 1 minute, 1 hour) and the known speed of the fastest sequential implementation (often on specialized hardware like ASICs). Critically,  $T$  is independent of  $\lambda$ ; one can have a very secure VDF ( $\lambda=256$ ) with a very short delay ( $T=1000$ ), or a moderately secure one ( $\lambda=128$ ) with a very long delay ( $T=10^{\{15\}}$ ).

This formal triple  $(\text{Setup}, \text{Eval}, \text{Verify})$  provides the structural blueprint. However, the true power and security of a VDF stem from the rigorous properties this triple must guarantee.

### 1.3.2 3.2 The Holy Trinity of Properties

The utility and security of a VDF hinge on three non-negotiable properties: **Sequentiality**, **Uniqueness**, and  **$\delta$ -Efficiency**. These properties distinguish VDFs from weaker constructs and ensure they fulfill their promise of verifiable elapsed time.

#### 1. Sequentiality: The Heartbeat of Delay

- **Definition:** Informally, no adversary, even equipped with a polynomial (in  $\lambda$ ) number of parallel processors, can compute  $(y, \pi) = \text{Eval}(pp, x)$  for a randomly chosen input  $x$  significantly faster than  $T$  sequential steps, except with negligible probability (in  $\lambda$ ). Formally, for all probabilistic parallel algorithms  $A$  running in parallel time  $\sigma(T)$  with  $\text{poly}(\lambda)$  processors, and for all large enough  $\lambda$ :

$$\Pr[A(pp, x) = (y, \pi) \wedge \text{Verify}(pp, x, y, \pi) = \text{Accept}] \leq \text{negl}(\lambda)$$

where  $\sigma(T) \neq 0$  (adversary saves a constant fraction of time), or sometimes defined with  $\sigma(T) = o(T)$  (asymptotically slower than linear).

- **Significance:** This is the core property that enforces the minimum passage of *sequential* computational time. It ensures that parallelism provides no substantial shortcut. An adversary with a million processors cannot compute the VDF output a million times faster; they are forced onto a sequential path.
- **Computational Model:** Sequentiality is defined relative to a specific model of computation. The most common is the **Parallel Random-Access Machine (PRAM)** model with polynomially many processors. Crucially, the sequentiality assumption posits that the *inherent depth* (parallel time complexity) of the `Eval` function is  $\Omega(T)$ . For repeated squaring in a GUO, this depth is exactly  $T$ , as each squaring depends directly on the previous result. No amount of parallelism can reduce the depth below  $T$ .

- **Example & Limitation:** Consider Pietrzak’s RSA-based VDF ( $\text{Eval} = T$  sequential squarings). The best-known attack relies on parallelizing the computation *if* the group order  $\phi(N)$  is known (via Euler’s theorem). However, under the assumption that factoring  $N$  (or computing  $\phi(N)$ ) is hard, and that repeated squaring is inherently sequential when the order is unknown, sequentiality holds. A critical caveat is **hardware speedup**: While parallelism doesn’t help, *faster sequential hardware* (like ASICs) *does* reduce the real-world time for  $T$  steps. Sequentiality guarantees that reducing the time *below*  $t_{\min} = (T * \text{time\_per\_step\_on\_best\_hardware})$  is infeasible, but  $t_{\min}$  decreases as hardware improves. Setting  $T$  requires anticipating future hardware advances (see Section 5.3).

## 2. Uniqueness: Preventing Equivocation

- **Definition:** For a given setup  $pp$  and input  $x$ , it is computationally infeasible for any efficient adversary to find two *different* valid outputs with their corresponding proofs. Formally, for all probabilistic polynomial-time (PPT) adversaries  $A$  and for all large enough  $\lambda$ :

$$\Pr[(y, \pi, y', \pi') \leftarrow A(pp, x) : y \neq y' \wedge \text{Verify}(pp, x, y, \pi) = \text{Accept} \wedge \text{Verify}(pp, x, y', \pi') = \text{Accept}] \leq \text{negl}(\lambda)$$

This is also sometimes called **Soundness**, emphasizing that it’s hard to find *any* valid proof for an *incorrect* output  $y'$ .

- **Significance:** Uniqueness ensures the VDF output is *binding* and *unambiguous*. For a given  $x$ , there is only one  $y$  that will pass verification with a valid proof. This is critical for applications:
- **Randomness Beacons:** If uniqueness fails, an adversary could generate multiple valid outputs  $y_1, y_2, \dots$  for the same committed seed  $x$  *after* seeing other participants’ reveals, allowing them to choose the  $y$  that biases the final randomness in their favor, completely undermining the beacon’s fairness.
- **Consensus:** In PoS chains using VDFs for leader election, an adversary could propose blocks with different VDF outputs for the same slot, causing forks and instability.
- **Timed Releases:** An attacker could potentially “open” a time-lock to multiple different messages.
- **Mechanism & Challenges:** Uniqueness typically relies on strong cryptographic assumptions within the group used. In RSA-based VDFs (Pietrzak, Wesolowski), uniqueness relies on the **Low Order Assumption (LOA)** or the **Adaptive Root Assumption (ARA)**. Essentially, these assume it’s hard to find elements of small order or to find  $x$  and  $y$  satisfying the verification equation unless  $y$  is the genuine  $g^{(2^T)} \bmod N$ . Class group-based VDFs rely on similar assumptions adapted to the structure of class groups. Proving uniqueness is often more challenging than sequentiality and sometimes requires modeling the group as a *Generic Group of Unknown Order (GGUO)* or working within the *Algebraic Group Model (AGM)* to obtain rigorous security proofs.

### 3. $\delta$ -Efficiency (Fast Verification): The Keystone of Practicality

- **Definition:** The verification algorithm  $\text{Verify}(\text{pp}, x, y, \pi)$  must run in time polynomial in the security parameter  $\lambda$  and *logarithmic* in the delay parameter  $T$  ( $O(\text{poly}(\lambda, \log T))$ ), while evaluation  $\text{Eval}(\text{pp}, x)$  inherently requires  $\Omega(T)$  sequential time. The  $\delta$  represents the *gap* between evaluation and verification time.
- **Significance:** This property is what makes VDFs *scalable* and *practical*. If verification took time proportional to  $T$ , the entire advantage over simply re-running the computation vanishes, especially for large  $T$  (minutes, hours, days). Fast verification allows lightweight clients (e.g., mobile phones) to trustlessly confirm proofs of significant sequential work performed by powerful servers or specialized hardware.
- **Achievement in Practice:** Wesolowski’s scheme exemplifies  $\delta$ -Efficiency. Verification requires only a few group operations (like  $\pi^\ell$  and  $x^r$ ) whose cost depends on the bit-length of the exponents ( $O(\lambda)$ ), completely independent of  $T$ . Pietrzak’s scheme achieves  $O(\log T)$  verification time, as the recursive proof has depth  $\log T$  and each step involves a constant number of group operations. This logarithmic growth is still exceptionally efficient even for astronomically large  $T$  (e.g.,  $\log_2(10^{15}) \approx 50$ ).

#### The Interplay and Implications:

These three properties are deeply intertwined and collectively define a VDF:

- **Sequentiality +  $\delta$ -Efficiency:** Creates the essential tension: generating the proof takes enforced sequential time, but verifying its correctness is almost instantaneous. This asymmetry is the source of the VDF’s power.
- **Uniqueness + Sequentiality:** Ensures that the fast verification actually corresponds to a specific, genuinely delayed outcome. Without uniqueness, sequentiality alone could allow adversaries to generate multiple “valid” delayed outputs, breaking applications.
- **$\delta$ -Efficiency + Uniqueness:** Makes the verification of the *single* correct outcome practical and accessible.

The “Holy Trinity” forms an indivisible set of requirements. A construction lacking any one of these properties fails to be a true VDF. Pietrzak’s and Wesolowski’s schemes achieve all three under specific cryptographic assumptions, which brings us to the bedrock upon which they stand: complexity-theoretic conjectures.

### 1.3.3 3.3 Complexity Assumptions: The Bedrock of Security

The security of practical VDFs, particularly those based on groups like RSA or class groups, rests on computational hardness assumptions. These are conjectures, not proven facts, but they are widely believed within the cryptographic community based on extensive research and failed attempts to break them. The security of the VDF properties hinges on these assumptions holding against even powerful adversaries.

#### 1. The Sequentiality Assumption: Modeling Computational Depth

- **Core Conjecture:** For specific functions  $f$  (like iterated squaring modulo  $N$ ) and computational models (like PRAM), the parallel time complexity of computing  $f^{\{T\}}(x)$  (applying  $f$   $T$  times) is  $\Omega(T)$ , even for adversaries with  $\text{poly}(\lambda)$  processors. This is a *fine-grained* assumption about the *inherent depth* of the computation.
- **Relation to Other Assumptions:** For the prevalent repeated squaring VDFs ( $y = x^{(2^T \bmod N)} \bmod N$ ), sequentiality relies indirectly on the hardness of the underlying group problem (factoring for RSA, class group discrete log for class groups). If an adversary could compute the group order  $\text{ord}(g)$  (e.g.,  $\phi(N)$  for RSA), they could compute  $y = g^{\{2^T \bmod \text{ord}(g)\}} \bmod N$  using fast modular exponentiation (like Euler's theorem), bypassing the sequential squaring entirely. Therefore, the sequentiality assumption for squaring VDFs implies that computing the group order is hard, but it is *stronger*: it posits that even *without* knowing the order, there is no parallel algorithm significantly faster than  $T$  sequential squarings. This is believed to hold for well-chosen RSA moduli and class group discriminants.
- **Theoretical Support:** While not proven unconditionally, the sequentiality of iterated squaring in GUOs has strong intuitive and theoretical backing. No non-trivial parallel algorithms are known, and it seems to require a sequential dependency chain of length  $T$ .

#### 2. The Centrality of Groups of Unknown Order (GUOs):

GUOs provide the algebraic structure where sequentiality naturally arises. The two dominant families are:

- **RSA Groups ( $\mathbb{Z}_N^*$  for  $N = pq$ ):**
- **Assumption:** The **RSA Assumption** (hard to compute  $y = x^e \bmod N$  given  $x, e, N$  for random  $x$  and large  $e$ ) and the **Factoring Assumption** (hard to factor  $N = pq$ ).
- **Sequentiality Basis:** Computing  $g^{(2^T \bmod N)}$  is believed to require  $T$  sequential squarings if the factors of  $N$  are unknown. Knowledge of  $\phi(N) = (p-1)(q-1)$  allows computation in  $O(\log T)$  time via  $g^{\{2^T \bmod \phi(N)\}} \bmod N$ .



- **The “Nothing Up My Sleeve” Controversy:** Generating  $N = pq$  requires someone (or some process) to know  $p$  and  $q$ . If this entity is malicious and *keeps* this knowledge, they possess a **trapdoor** allowing them to compute VDF outputs instantly. This introduces a **trusted setup** requirement. Mitigations involve complex **Multi-Party Computation (MPC) ceremonies** (see Section 5.2), where multiple parties collaboratively generate  $N$  such that *no single party* (and ideally, no coalition below a threshold) knows the factorization. The Ethereum Foundation, for example, ran a complex MPC ceremony for generating an RSA modulus for potential VDF use. While MPC significantly reduces trust, it remains a point of scrutiny and potential vulnerability compared to setup-free alternatives.
- **Ideal Class Groups of Imaginary Quadratic Fields ( $Cl(-D)$ ):**
- **Assumption:** The **Class Group Discrete Logarithm Problem (CGDLP)** or the **Computation of the Class Number  $h(-D)$**  is hard. Essentially, computing the structure (group order and discrete logs) of the class group defined by a discriminant  $-D$  is computationally difficult.
- **Sequentiality Basis:** Computing  $[g]^{2^T}$  (repeated squaring of an ideal class  $[g]$ ) is believed to require  $T$  sequential steps when the class number  $h(-D)$  (the group order) is unknown. Knowledge of  $h(-D)$  would again allow shortcutting via exponentiation modulo  $h(-D)$ .
- **Setup-Free Advantage:** The discriminant  $-D$  can be generated *publicly* and transparently using a “nothing-up-my-sleeve” number (e.g., derived from the digits of  $\pi$  or the output of a public hash function) or via a simple random process. *No secrets* need discarding, eliminating the trusted setup risk inherent in RSA groups. This makes class groups highly attractive for decentralized systems prioritizing maximal trust minimization.
- **Trade-offs:** While class groups offer superior trust minimization, **class group arithmetic is significantly more complex and slower** than RSA modular arithmetic. Operations like ideal multiplication and reduction require more computational steps. This impacts both evaluation and (to a lesser extent) verification speed. RSA groups benefit from decades of hardware optimization for modular arithmetic.

### 3. Isogeny-Based Assumptions: A Post-Quantum Horizon?

- **Potential:** Supersingular elliptic curve isogenies offer a potential path to VDFs secure against quantum computers. The core idea, reminiscent of Boneh et al.’s initial proposal, leverages the conjectured sequentiality of computing large-degree isogenies between elliptic curves.
- **Assumptions:** Security would rely on variants of the **Sequential Isogeny Assumption** – that computing a long chain of isogenies requires inherently sequential steps. Candidate constructions often build on primitives like **CSIDH** or **SQIsign**.
- **Current Status (as of 2024):** Isogeny-based VDFs remain primarily theoretical. Key challenges include:



- **Slower Operations:** Isogeny computations are orders of magnitude slower than group operations in RSA or class groups.
- **Larger Proofs/Parameters:** Keys and proofs tend to be larger.
- **Less Mature Assumptions:** The sequentiality of isogeny computation and the underlying security assumptions are less studied and scrutinized than factoring or class group problems. The recent break of the SIDH isogeny problem (2022) highlights the evolving nature of isogeny-based cryptography.
- **Outlook:** While promising for long-term post-quantum security, isogeny-based VDFs are not yet practical competitors to GUO-based schemes. Significant algorithmic and implementation breakthroughs are needed.

#### 4. The Nature of VDF Assumptions:

VDF assumptions, particularly sequentiality, are qualitatively different from standard cryptographic assumptions like factoring hardness:

- **Non-Falsifiable:** Proving that *no* faster parallel algorithm exists is typically beyond current complexity theory. Sequentiality is assumed based on the absence of known attacks and the inherent structure of the computation.
- **Fine-Grained:** They concern the *exact* parallel time complexity ( $\Omega(T)$ ), not just polynomial vs. exponential hardness.
- **Implementation-Dependent:** The real-world sequentiality depends on the concrete implementation of the group operation (squaring). While modular squaring is believed inherently sequential, highly optimized ASICs might reveal subtle parallelism or algorithmic improvements, effectively reducing the perceived  $T$  (see Section 5.3).

The mathematical edifice of VDFs – defined by the algorithmic triple (*Setup*, *Eval*, *Verify*), fortified by the Holy Trinity of properties (Sequentiality, Uniqueness,  $\delta$ -Efficiency), and grounded in the bedrock of complexity assumptions (primarily the hardness of problems in Groups of Unknown Order) – provides a rigorous framework for understanding their capabilities and limitations. This framework transforms the compelling intuition of verifiable delay into a concrete, analyzable cryptographic primitive.

The elegance of Pietrzak’s recursion and Wesolowski’s compact verification, glimpsed historically in Section 2, now emerges as direct manifestations of this mathematical structure. However, understanding *how* these properties are concretely achieved within specific schemes requires delving into their construction. This leads naturally to an exploration of the **Major Schemes and Mechanisms** that bring the abstract VDF definition to life, examining the trade-offs and implementation nuances that govern their real-world application.

(Word Count: ~2,050)

## 1.4 Section 4: Constructing VDFs: Major Schemes and Mechanisms

The rigorous mathematical framework established in Section 3 – the algorithmic triple, the Holy Trinity of properties, and the reliance on complexity assumptions – transforms the abstract concept of verifiable delay into a blueprint for real-world implementation. Yet, the journey from mathematical elegance to practical cryptographic tool hinges on concrete constructions that realize this blueprint efficiently and securely. This section dissects the primary mechanisms that bring Verifiable Delay Functions to life, focusing on the two dominant paradigms that emerged from the post-2018 explosion of research: Wesolowski’s compact proof scheme and Pietrzak’s recursive composition. We explore their operational principles, the pivotal role of Groups of Unknown Order (GUOs) as their algebraic engines, and the trade-offs inherent in alternative approaches. Understanding these constructions reveals how the unforgeable arrow of sequential time is encoded within modular arithmetic and group theory.

Following the formalization of VDF properties, the cryptographic community converged on constructions leveraging the inherent sequentiality of repeated squaring in algebraic structures where the group order remains secret. Pietrzak and Wesolowski, building on Boneh et al.’s foundation and Rivest’s original time-lock insight, provided the elegant, efficient solutions that dominate practical applications. Their schemes, while sharing a common computational core, diverge dramatically in how they achieve the crucial feat of fast verification, embodying distinct design philosophies with profound implications for implementation.

### 1.4.1 4.1 Wesolowski’s Scheme: Efficient Verification via SNARGs

Benjamin Wesolowski’s 2018 scheme represents a paradigm shift in VDF verification efficiency. Its brilliance lies in achieving **constant-size proofs** and **constant-time verification** (relative to the delay parameter  $T$ ), making it exceptionally suitable for bandwidth-constrained environments like blockchain networks. The scheme leverages the structure of **Groups of Unknown Order (GUOs)** – RSA groups ( $\mathbb{Z}_N^*$ ) or class groups ( $\text{Cl}(-D)$ ) – and employs a clever challenge-response mechanism inspired by the Fiat-Shamir heuristic and succinct non-interactive arguments (SNARGs).

#### Core Computational Engine:

Regardless of the group ( $G$ ), the evaluation step is fundamentally the same as the time-lock puzzle:

1. **Input:** Public parameters  $pp$  (defining the group  $G$  and generator  $g$ ; for RSA,  $pp = (N, g)$ ; for class groups,  $pp = (\text{discriminant } -D, [g]$  where  $[g]$  is an ideal class) and input  $x$  (embedded into  $G$ ; often  $x$  is used directly as the base).
2. **Evaluation:** Compute the output  $y$  via  $T$  sequential squarings:

$y = g^{(2^T)}$  (in multiplicative notation, e.g.,  $g^{(2^T)} \bmod N$  for RSA, or  $[g]^{2^T}$  in the class group).

#### The Magic: Proof Generation and Verification:

The innovation lies in generating a tiny proof  $\pi$  that enables verification exponentially faster than  $T$ . Wesolowski achieves this through an interactive protocol made non-interactive using the Fiat-Shamir transform:

### 1. Proof Generation (**Eval** outputting $\pi$ ):

- After computing  $y = g^{(2^T)}$ , the prover:
- Computes a **challenge prime**  $\ell$ : This is derived by hashing the public parameters, input  $x$ , output  $y$ , and  $T$  using a cryptographic hash function  $H$  (e.g., SHA-256), interpreted as a prime number within a range determined by the security parameter  $\lambda$  (e.g.,  $\ell$  is the first prime larger than  $H(\dots)$  or chosen via a deterministic process from the hash output). Formally:  $\ell = \text{get\_prime}(H(pp \parallel x \parallel y \parallel T))$ .
- Computes the **quotient**  $q$  and **remainder**  $r$  of  $2^T$  divided by  $\ell$ :

$2^T = q * \ell + r, \text{ where } 0 \leq r < \ell$ :

- Compute the midpoint value  $\mu$  at step  $T/2$  (assuming  $T$  is a power of 2 for simplicity; generalizations exist):

$\mu = g^{(2^{T/2})}$  (Requires  $T/2$  sequential squarings starting from  $g$ ).

- **Recursively compute proofs** for the two halves:
- Proof  $\pi_L$  that  $\mu$  is the correct output starting from  $g$  after  $T/2$  steps:  $(\mu, \pi_L) = \text{Eval}(pp, g, T/2)$ .
- Proof  $\pi_R$  that  $y$  is the correct output starting from  $\mu$  after  $T/2$  steps:  $(y, \pi_R) = \text{Eval}(pp, \mu, T/2)$ . Note: Computing this requires another  $T/2$  squarings starting from  $\mu$ , but the prover already knows  $y = g^{(2^T)} = (g^{(2^{T/2})})^{2^{T/2}} = \mu^{2^{T/2}}$ , so they can compute  $\pi_R$  without redoing the squarings if they store intermediates or recompute  $\mu^{2^{T/2}}$  efficiently (which is possible as they know  $\mu$ ).
- The proof  $\pi$  for the full computation  $(g, T)$  is the tuple  $(\mu, \pi_L, \pi_R)$ .
- **Output:**  $(y, \pi = (\mu, \pi_L, \pi_R))$ .
- **Structure:** The proof  $\pi$  forms a binary tree (a Merkle tree of computation). The root corresponds to the full computation  $(g, T) \rightarrow y$ . Each internal node corresponds to a midpoint  $\mu_i$  at depth  $i$ , and its children correspond to proofs for the left half  $(g, T/2^i) \rightarrow \mu_i$  and the right half  $(\mu_i, T/2^i) \rightarrow \dots$ . The leaves correspond to small computations (e.g.,  $T=1$ ).

### 2. Verification (**Verify**):

Verification is also recursive, checking consistency at each level of the tree:

- **Input:**  $pp, x = g, y, \pi = (\mu, \pi_L, \pi_R), T$ .
- **Procedure:**
  1. Verify the proof  $\pi_L$  for the left half:  $\text{Verify}(pp, g, \mu, \pi_L, T/2) = \text{Accept?}$
  2. Verify the proof  $\pi_R$  for the right half:  $\text{Verify}(pp, \mu, y, \pi_R, T/2) = \text{Accept?}$
  3. **Output:**  $\text{Accept}$  only if both sub-proofs verify correctly.
- **Base Case ( $T=1$ ):** Check  $y \stackrel{?}{=} g^2$  directly (one squaring).
- **Efficiency:** Each recursive step involves two sub-verifications for size  $T/2$  and a constant amount of work (processing the tuple  $(\mu, \pi_L, \pi_R)$ ). The recursion depth is  $\log T$ . Assuming the work per non-leaf node is constant  $c$ , the total verification time is  $O(c * \log T)$ . The dominant cost typically involves a few group operations per level (e.g., storing and comparing group elements). Proof size is also  $O(\log T)$  group elements (one  $\mu$  per level of the tree).

### Why it Works (Intuition):

The security hinges on the **sequentiality of the computation chain** and the **binding nature of the commitments**. If the prover attempts to cheat by providing an incorrect midpoint  $\mu'$ , they must then provide valid proofs  $\pi_L'$  for  $(g, T/2) \rightarrow \mu'$  and  $\pi_R'$  for  $(\mu', T/2) \rightarrow y$ . However:

- If  $\mu'$  is not the true midpoint  $\mu = g^{(2^{T/2})}$ , then the statement  $(g, T/2) \rightarrow \mu'$  is false. Proving it requires breaking the sequentiality or soundness of the VDF at the lower level  $T/2$ .
- Even if they somehow forged  $\pi_L'$ , the statement  $(\mu', T/2) \rightarrow y$  would likely be false unless  $\mu'$  was chosen maliciously to make  $(\mu')^{(2^{T/2})} = y$ , but this contradicts the true computation  $\mu^{(2^{T/2})} = y$  and the uniqueness property (assuming it holds at each level).

Recursion amplifies security: forging the proof for a large  $T$  requires forging proofs at multiple smaller levels, which is assumed to be computationally infeasible under the VDF's sequentiality and uniqueness properties. The Fiat-Shamir transform is often incorporated implicitly by deriving any necessary challenges within the proof structure itself from the input data, making the entire proof non-interactive.

### Trade-offs and Nuances:

- **Advantages:** Simpler and often faster proof generation than Wesolowski's scheme, as it avoids expensive root computations. Primarily involves computing intermediate points ( $\mu$ ) and concatenating sub-proofs. Potential for **parallel proof generation**: the proofs  $\pi_L$  and  $\pi_R$  for the two halves can be generated concurrently once  $\mu$  is computed. The recursive structure is conceptually clear and aligns well with techniques like Merkle trees.

- **Disadvantages:** Proof size and verification time grow logarithmically with  $T$  ( $O(\log T)$ ). For very large  $T$  (e.g.,  $T = 2^{40}$ ), the proof might contain 40 group elements (e.g.,  $40 * 256 \text{ bytes} = 10 \text{ KB}$  for class groups), and verification might take 40 times a few milliseconds. While manageable, this is less efficient than Wesolowski’s constant terms for massive  $T$ . Requires careful handling if  $T$  is not a power of two.
- **Uniqueness:** Proving uniqueness for Pietrzak’s scheme is more complex than for Wesolowski’s and often requires stronger assumptions or modeling (e.g., the Algebraic Group Model - AGM). The recursive composition itself doesn’t inherently guarantee global uniqueness without additional cryptographic properties of the group.

Pietrzak’s scheme demonstrates the power of recursive decomposition. By breaking down the monolithic  $T$ -step computation into smaller, verifiable chunks and proving their consistency, it achieves efficient verification with a structure that is both elegant and amenable to optimization.

### 1.4.2 4.3 The Crucial Role of Groups of Unknown Order (GUOs)

Both Wesolowski’s and Pietrzak’s schemes, along with Boneh et al.’s initial proposals, rely fundamentally on **Groups of Unknown Order (GUOs)**. These algebraic structures provide the fertile ground where the sequentiality of repeated squaring can flourish. Understanding the characteristics, trade-offs, and specific implementations of these groups is paramount for deploying secure and efficient VDFs.

#### Why GUOs are Essential:

The security of the sequential squaring operation  $y = g^{(2^T)}$  hinges on two properties provided by GUOs:

1. **Hardness of Order Computation:** Computing the order (number of elements) of the group,  $|G|$ , must be computationally infeasible. If  $|G|$  (or  $\text{ord}(g)$ , the order of the generator) is known, then  $y = g^{(2^T \bmod \text{ord}(g))}$  can be computed using fast modular exponentiation (like `pow(g, pow(2, T, ord(g)), |G|)`), bypassing the  $T$  sequential squarings entirely. This breaks sequentiality.
2. **Absence of Shortcuts:** There should be no known algorithms to compute  $g^{(2^T)}$  significantly faster than  $T$  sequential squarings, even without knowing  $|G|$ . This is the core sequentiality assumption discussed in Section 3.3.

#### The Two Dominant GUO Families:

##### 1. RSA Groups ( $\mathbb{Z}_{N^*}$ ):

- **Structure:** The multiplicative group of integers modulo  $N$ , where  $N = pq$  is a product of two large distinct primes. The group order is  $\phi(N) = (p-1)(q-1)$ .

- **Security Assumptions:** Security relies on the **RSA Assumption** (hard to compute  $e$ -th roots modulo  $N$  for  $e > 1$ ) and the **Factoring Assumption** (hard to factor  $N$ ). Knowledge of  $\phi(N)$  allows instant computation of  $g^{(2^T)} = g^{\{ (2^T \bmod \phi(N)) \}} \bmod N$ .
- **Pros:**
- **Well-Understood:** Decades of cryptanalysis provide high confidence in the underlying assumptions.
- **Efficient Arithmetic:** Modular multiplication and squaring modulo  $N$  are highly optimized operations on CPUs, GPUs, and especially ASICs. Hardware support is mature.
- **Small Element Size:** Group elements (numbers mod  $N$ ) have compact representation (e.g., 3072 bits for  $\sim 128$ -bit security).
- **Cons:**
- **Trusted Setup Required:** Generating  $N = pq$  requires someone to know  $p$  and  $q$ . If this secret is not securely destroyed or distributed (e.g., via MPC), it creates a **trapdoor** allowing instant VDF evaluation. MPC ceremonies mitigate this but add complexity and potential risk (see Section 5.2).
- **Vulnerability to Quantum Computers:** Shor's algorithm efficiently factors  $N$  and computes  $\phi(N)$  on a sufficiently large quantum computer, breaking RSA-based VDFs.

## 2. Ideal Class Groups of Imaginary Quadratic Fields ( $\text{C1}(-D)$ ):

- **Structure:** The group of equivalence classes (ideals modulo principal ideals) in the ring of integers of the imaginary quadratic field  $\mathbb{Q}(\sqrt{-D})$ , where  $-D$  Pietrzak » Hash-based  $\approx$  IVC? (if efficient) » Isogeny/Lattice?.
- **Evaluation Speed (for given  $T$ ):** RSA Group > Class Group » Isogeny/Lattice.
- **Sequentiality Strength:** GUO-based (RSA/Class) > Isogeny? > Hash-based/Lattice?.
- **Trust Minimization:** Class Group > Isogeny/Lattice? > RSA Group (requires MPC).
- **Post-Quantum Potential:** Isogeny?/Lattice?/Hash-based > Class Group? > RSA Group (broken by Shor).
- **Implementation Maturity:** RSA Group > Class Group » Isogeny/Lattice/IVC.

This landscape underscores why Wesolowski (Class or RSA) and Pietrzak (RSA) dominate practical deployments: they offer the best balance of strong sequentiality, efficient verification, and (for class groups) trust minimization, despite the GUO trade-offs. Hash-based and isogeny approaches remain intriguing for niche applications or future horizons.

The elegant machinery of Wesolowski’s compact proofs and Pietrzak’s recursive trees, built upon the sequential bedrock of Groups of Unknown Order, provides the robust foundation for practical VDF deployment. Yet, the very mechanisms that grant their power – complex arithmetic, trusted setups, and reliance on specific computational assumptions – also introduce potential vulnerabilities and limitations. Having explored how VDFs are built, we must now rigorously examine their **Security Analysis: Attacks, Assumptions, and Limitations**, scrutinizing the boundaries of their guarantees and the practical threats they face in adversarial environments.

---

## 1.5 Section 5: Security Analysis: Attacks, Assumptions, and Limitations

The elegant machinery of Wesolowski’s compact proofs and Pietrzak’s recursive trees, built upon the sequential bedrock of Groups of Unknown Order (GUOs), provides the robust foundation for practical Verifiable Delay Function (VDF) deployment. Yet, the very mechanisms that grant their power – complex arithmetic, trusted setups, and reliance on specific computational assumptions – also introduce potential vulnerabilities and limitations. **No cryptographic primitive is invulnerable, and VDFs exist within a dynamic landscape of evolving hardware, sophisticated adversaries, and unforeseen attack vectors.** This section critically examines the security boundaries of VDFs, dissecting known cryptanalytic threats, the treacherous terrain of trusted setups, the relentless pressure of hardware acceleration, and subtle attacks exploiting amortization and precomputation. Understanding these limitations is not a sign of weakness but a prerequisite for deploying VDFs with realistic expectations and robust mitigations in high-stakes decentralized systems.

The mathematical elegance of repeated squaring in GUOs, as explored in Section 4, offers compelling sequentiality guarantees *under specific assumptions*. However, the transition from abstract model to concrete implementation introduces cracks where adversaries can pry. Cryptanalysis probes the bedrock assumptions, hardware advancements shrink the perceived delay  $T$ , trusted setup ceremonies become high-value targets, and clever attackers seek ways to sidestep the sequentiality requirement through parallelism across multiple inputs. The security of VDFs is thus a multi-faceted challenge demanding constant vigilance and adaptation.

### 1.5.1 5.1 Cryptanalytic Attacks on Core Assumptions

The security of the dominant GUO-based VDFs (Pietrzak, Wesolowski) hinges on the hardness of specific mathematical problems in RSA and class groups. Cryptanalytic breakthroughs targeting these problems could catastrophically undermine VDF sequentiality or uniqueness. While no such breaks have occurred for properly parameterized groups, the threat landscape is dynamic, and theoretical vulnerabilities highlight the fragility of the underlying conjectures.

- **Attacks on the Sequentiality Assumption:**



- The Parallelization Mirage:** The core promise of VDFs is resistance to parallelization. However, this relies on the conjecture that *no* non-trivial parallel algorithm exists for computing  $g^{(2^T)}$  in a GUO without knowing the group order. While no such algorithm is known, its non-existence is not proven. Theoretical models exploring circuit complexity or parallel RAM lower bounds provide some justification, but a breakthrough proving significant parallelism possible would collapse the sequentiality guarantee. For example, discovering an algorithm evaluating  $k$  squarings in depth  $\mathcal{O}(k)$  (e.g.,  $\mathcal{O}(\log k)$ ) using  $\text{poly}(k)$  processors would allow computing  $g^{(2^T)}$  in parallel time  $\mathcal{O}(T)$ , violating sequentiality.
- Exploiting Known Order (The Trapdoor):** As established, knowledge of the group order  $\text{ord}(G)$  (e.g.,  $\phi(N)$  for RSA,  $h(-D)$  for class groups) allows instant computation of  $g^{(2^T)} = g^{\{2^T \bmod \text{ord}(G)\}}$  using fast exponentiation. This isn't an attack per se but underscores the criticality of the GUO property. Any cryptanalytic advance making  $\text{ord}(G)$  computable efficiently breaks *all* sequentiality for that group instance. For RSA, this means progress in **factoring algorithms**. While the general Number Field Sieve (NFS) complexity remains sub-exponential, constant-factor improvements or specialized hardware (like TWIRL or SHARK) could force larger modulus sizes. The 2023 [Lattice Sieve improvement](#) reducing the NFS complexity constant is a reminder that factoring is not static. For class groups, advances in **computing class numbers** or solving the **Principal Ideal Problem (PIP)** pose analogous threats. The 2014 breakthrough computing a record class number for a 180-digit discriminant using massive parallelism showed the problem is parallelizable *if* the discriminant is known, emphasizing the need for sufficiently large  $|D|$ .
- Fault Attacks and Side-Channels:** Implementation flaws can leak information or allow manipulation bypassing the sequential computation. **Fault attacks** aim to induce errors during the long squaring process (e.g., via voltage glitching or laser injection) to reveal information about internal states or the modulus  $N$ . **Side-channel attacks** (timing, power analysis, electromagnetic emanation) could potentially leak bits of the exponent or intermediate values during squaring or proof generation, especially in Wesolowski's scheme where the  $\ell$ -th root computation is sensitive. A successful attack recovering even part of the group order or internal state could dramatically reduce the effective  $T$ . Constant-time implementations and robust fault detection are essential countermeasures.
- Breaking Uniqueness:**
  - The Adaptive Root Assumption (ARA) Under Fire:** Uniqueness in Wesolowski's scheme relies heavily on the **Adaptive Root Assumption (ARA)** (or variants like the Low Order Assumption - LOA). This assumes that for a random  $g$  in a GUO, it's hard to find *both* an exponent  $e > 1$  and an  $h$  such that  $h^e = g$ , *even after* seeing other group elements and performing operations. A break of this assumption would allow an adversary, given a challenge  $\ell$  in the verification, to find a  $\pi'$  and  $y' \neq y$  satisfying  $(\pi')^\ell * g^\tau = y'$  for the computed  $\tau$ , forging a valid proof for an incorrect output. While the ARA holds in the Generic Group Model (GGM), real groups have extra structure.
 

**Potential ARA Violations:** Research has explored potential weaknesses. The 2020 paper "[On the Security of Time-Lock Puzzles and Timed Commitments](#)" by Malavolta and Thyagarajan identified a



vulnerability in a *related* timed commitment scheme using RSA groups, exploiting properties of small subgroups if the exponent  $e$  (analogous to  $\ell$ ) is smooth. While not directly breaking Wesolowski's VDF, it highlighted the importance of careful parameter choice. Ensuring  $\ell$  is a sufficiently large prime ( $\geq 128$  bits) mitigates known attacks by making the subgroup order large and exploiting smoothness infeasible. Similar scrutiny applies to Pietrzak's scheme, where uniqueness relies on the difficulty of finding collisions in the recursive Merkle tree of computation under the group's algebraic properties.

- **Consequences of Broken Uniqueness:** The failure of uniqueness is catastrophic for core applications:
- **Randomness Beacons:** An adversary controlling the last revealer could compute multiple valid  $(y_i, \pi_i)$  pairs for the same committed seed  $x$ . After seeing other participants' reveals, they choose the  $y_i$  that biases the final randomness output maximally in their favor. This completely undermines the beacon's unpredictability and fairness.
- **Proof-of-Stake Leader Election:** A malicious validator could generate multiple valid VDF outputs for the same slot, enabling them to propose multiple conflicting blocks, causing forks and consensus instability.
- **Mitigation Strategies:** Rigorous security proofs within models like the Algebraic Group Model (AGM), careful parameter selection (large prime  $\ell$  in Wesolowski), and ongoing cryptanalysis are essential. Diversifying VDF constructions or combining them with other primitives (like signatures) might add layers of defense.

Cryptanalysis is an ongoing arms race. While current GUO-based VDFs with conservative parameters (large moduli/discriminants, large challenge primes) appear secure, the assumptions they rely on are not invincible. Constant monitoring for algorithmic advances and potential breaks is paramount, especially as VDFs become integrated into critical financial and governance infrastructure.

## 1.5.2 5.2 The Trusted Setup Problem

For RSA-based VDFs (the most hardware-efficient construction), the generation of the public modulus  $N = pq$  represents a critical point of vulnerability – the **trusted setup**. If the primes  $p$  and  $q$  are known to *any* party, that party possesses a **trapdoor** enabling them to compute VDF outputs  $y = g^{(2^T)} \bmod N$  instantly using Euler's theorem ( $y = g^{\{ (2^T \bmod \phi(N)) \}} \bmod N$ ), completely breaking sequentiality. This creates a single point of failure antithetical to the decentralization ethos VDFs often serve.

- **The Trapdoor Risk:**
- **Backdoors and Collusion:** A malicious entity generating  $N$  could retain  $p$  and  $q$ . They could then compute VDF outputs arbitrarily fast, allowing them to manipulate randomness beacons, win leader elections unfairly, or break time-lock encryption. Even if the initial generator is honest,  $p$  and  $q$  could be stolen later via hacking or coercion.

- **“Nothing-Up-My-Sleeve” Failure:** Simple attempts to generate  $N$  transparently (e.g., using digits of  $\pi$ ) are insecure, as anyone can reverse-engineer  $p$  and  $q$  if the generation process is deterministic. The process *must* involve secrecy.
- **Example - The RSA-2048 MPC Ceremony (Ethereum/Filecoin):** Recognizing the severity of this risk, the Ethereum Foundation and Protocol Labs (Filecoin) initiated a massive **Multi-Party Computation (MPC) ceremony** in 2018 to collaboratively generate an RSA modulus  $N = pq$  where *no single party*, and ideally *no small coalition*, learns the factorization. Participants (over 20 globally distributed entities/individuals) contributed random data used to generate candidate primes. Using sophisticated MPC protocols (like [GG18](#) and [Lin17](#)), they multiplied their contributions such that only the final product  $N$  was revealed, while the individual factors remained secret. The ceremony concluded in May 2020, producing the 2048-bit modulus  $N$  intended for potential use in Ethereum’s VDF-based randomness beacon. While a landmark achievement in practical MPC, it wasn’t without controversy:
- **Complexity & Cost:** The ceremony was logistically complex, expensive, and required significant expertise to participate securely.
- **Trust in Participants:** While MPC minimizes trust, the security relies on the assumption that not too many participants collude. A sufficiently large coalition could reconstruct  $p$  and  $q$ . The “tribes of trust” problem remains.
- **Parameter Obsolescence:** Concerns were raised that 2048 bits might become insecure within the lifetime of Ethereum 2.0, necessitating a future, even larger ceremony.
- **Mitigations: Multi-Party Computation (MPC) Ceremonies:**

MPC is the primary tool to mitigate the trusted setup risk. The goal is “**trapdoor-free**” setup where the probability of any coalition learning the factors is negligible.

- **Threshold Security:** MPC protocols can be designed with a threshold  $t$ : the factorization remains secret as long as fewer than  $t$  participants collude. Choosing a large, diverse set of participants increases  $t$  practically.
- **Verifiability:** Participants can cryptographically verify that their contribution was correctly incorporated and that the final  $N$  is a product of two primes, without learning what those primes are.
- **Challenges:** MPC ceremonies are complex, vulnerable to implementation bugs, require secure execution environments for participants, and can be difficult to audit fully. The security guarantee is probabilistic and relies on the honesty of a majority (or supermajority) of participants.
- **The Class Group Advantage:**

This is where **class groups of imaginary quadratic fields** shine. Their public parameters consist solely of a discriminant  $-D$ , which can be generated *transparently*:

1. Choose a public, random seed  $S$  (e.g., the hash of a Bitcoin block header, the digits of  $\pi$ , or a common string).
2. Use a cryptographic hash function  $H$  to derive a candidate discriminant  $-D = H(S)$  or iterate  $H(S || \text{counter})$  until a suitable fundamental discriminant is found.
3. Publish  $-D$ .

Since no secrets are involved in the group generation, **there is no trapdoor**. Anyone can verify the discriminant was generated correctly from the public seed. This eliminates the trusted setup risk entirely, making class groups the preferred choice for deployments prioritizing maximal decentralization and verifiability, such as **Chia Network’s Proofs of Space and Time** consensus. The trade-off, as discussed, is slower arithmetic.

The trusted setup problem is a stark reminder that cryptographic security extends beyond abstract mathematics into the messy realm of implementation, procedure, and human coordination. While MPC offers a powerful mitigation for RSA-based VDFs, it introduces its own complexities and residual trust assumptions. Class groups provide an elegant, trust-minimized alternative, albeit at a computational cost. The choice between them hinges on the application’s tolerance for setup complexity versus its need for raw performance.

### 1.5.3 5.3 Hardware Advantages and the ASIC Threat

Perhaps the most fundamental and often misunderstood limitation of VDFs is their inherent vulnerability to **faster hardware**. Unlike cryptographic assumptions, which are mathematical conjectures, Moore’s Law and specialized hardware design are relentless physical realities. **Sequentiality guarantees a minimum number of computational steps  $T$ , not a minimum wall-clock time**. This distinction has profound implications for security and decentralization.

- **The Inevitability of Hardware Speedup:**

The evaluation time for a VDF with parameter  $T$  is:

$$t_{\text{eval}} = T * t_{\text{step}}$$

where  $t_{\text{step}}$  is the time per sequential operation (e.g., one modular squaring in  $\mathbb{Z}_N^*$  or one ideal squaring in  $\text{Cl}(-D)$ ). While  $T$  is fixed by the protocol,  $t_{\text{step}}$  is determined by the hardware. Faster hardware directly reduces  $t_{\text{eval}}$ :

- **Commodity Hardware:** Initial implementations use CPUs (slow  $t_{\text{step}}$ ).

- **GPUs:** Offer moderate speedups for modular arithmetic (RSA), but are less effective for complex class group operations.
- **FPGAs (Field-Programmable Gate Arrays):** Provide significant speedups (e.g., 5-10x over CPU) by implementing highly optimized squaring circuits. Projects like the [Ethereum Foundation's FPGA VDF implementation](#) demonstrated this potential early on.
- **ASICs (Application-Specific Integrated Circuits):** Represent the pinnacle of optimization, offering orders of magnitude speedup. Dedicated circuits eliminate general-purpose CPU overhead, optimize data paths, and run at higher clock speeds. **Supranational's VDF ASIC** designs (developed partly for the Ethereum competition) demonstrated the feasibility, targeting massive speedups for RSA modular squaring.
- **The ASIC Reality:**
  - **Performance Gains:** Well-designed ASICs can reduce  $t_{\text{step}}$  by 100x or more compared to high-end CPUs. For a fixed  $T$ , this directly reduces the real-world delay  $t_{\text{eval}}$  by the same factor. An ASIC completing  $T$  steps in 1 second implies an honest CPU user might take 100 seconds, violating the protocol's intended timing.
  - **Centralization Pressure:** Designing and fabricating ASICs requires significant capital and expertise. This risks centralizing VDF computation power in the hands of a few entities (e.g., large mining pools, specialized hardware manufacturers), mirroring the centralization concerns of Bitcoin's PoW mining. These entities could gain disproportionate influence:
  - **Randomness Beacons:** Faster computation could allow predicting or influencing beacon outputs slightly earlier than others, potentially enabling exploitative strategies (e.g., in DeFi).
  - **Leader Election (PoS):** ASIC-equipped validators could consistently win leader elections by completing the mandatory VDF faster, leading to validator centralization.
  - **MEV Mitigation:** If VDFs enforce a delay for fair transaction ordering, entities with faster VDF hardware could gain an advantage in the subsequent ordering process.
- **The Ethereum Foundation's RSA ASIC Challenge:** Recognizing this threat early, the Ethereum Foundation launched a [VDF ASIC competition](#) (2018-2019). The explicit goal was to **democratize access** by fostering open-source, efficient ASIC designs. Teams from EPFL, Supranational, and others submitted designs. While the competition spurred innovation (e.g., [EPFL's open-source 1GHz modular multiplier](#)), it also starkly illustrated the performance gulf between ASICs and general hardware, potentially accelerating the centralization it aimed to prevent. Ethereum's subsequent decision to defer VDF integration into its beacon chain was partly influenced by these hardware centralization concerns.
- **Mitigation Strategies and the "ASIC-Resistance" Debate:**

- **Conservative Parameter Setting ( $\mathbf{T}$ ):** The primary defense is setting  $\mathbf{T}$  very high based on *anticipated* future hardware capabilities, not current ones. This creates a buffer. However, this is inherently speculative and wasteful in the present, as honest users suffer unnecessarily long delays until hardware catches up. It also makes the system sluggish.
- **Hardware Diversity:** Choosing VDF constructions where the core operation (e.g., class group arithmetic) is less amenable to massive ASIC speedups could slow centralization. While class group squaring *will* see ASIC gains, the gains might be less extreme than for simpler modular arithmetic. However, this is a delaying tactic, not a solution.
- **Frequent Parameter Updates:** Periodically increasing  $\mathbf{T}$  (or changing the GUO parameters) to counter hardware advances. This adds complexity and coordination overhead.
- **Is “ASIC-Resistance” Desirable or Possible?** A key philosophical debate exists:
- **Pro-Resistance:** Argues that minimizing ASIC speedups preserves decentralization and fairness. Hash-based VDFs or complex operations target this, though with weaker sequentiality (Section 4.4).
- **Anti-Resistance / Pro-ASIC:** Argues that ASICs are inevitable and efficient. Open-source ASIC designs and commoditization (like those encouraged by Ethereum’s challenge) can make them accessible, potentially *improving* decentralization compared to a landscape dominated by clandestine, optimized ASICs. Furthermore, ASICs represent sunk cost, creating economic barriers against Sybil attacks similar to PoW. The focus should be on managing access, not futile resistance.
- **Economic Mechanisms:** Designing tokenomics where VDF computation is rewarded, but mechanisms exist to tax or redistribute rewards if centralization thresholds are exceeded. This is complex and game-theoretically challenging.

The hardware advantage dilemma underscores that VDFs anchor trust in *physical computation time*, which is inherently tied to engineering progress. There is no cryptographic magic that can freeze hardware development. Managing this reality involves a combination of prudent parameter choices, potential architectural diversity, economic design, and acceptance that some degree of specialization and associated centralization pressure is unavoidable for high-performance VDFs.

#### 1.5.4 5.4 Amortization and Precomputation Attacks

VDFs guarantee sequentiality *for a single evaluation on a specific input  $x$* . Clever adversaries can exploit parallelism *across multiple inputs* or leverage predictable inputs to gain an advantage, chipping away at the enforced delay. These amortization and precomputation attacks highlight that sequentiality is fundamentally input-bound.

- **Amortization Attacks: Parallelism Across Inputs:**

- **The Vulnerability:** While an adversary cannot compute  $\text{Eval}(\text{pp}, x)$  for a *single*  $x$  faster than  $\sim T$  sequential time, they can compute  $\text{Eval}(\text{pp}, x_1), \text{Eval}(\text{pp}, x_2), \dots, \text{Eval}(\text{pp}, x_k)$  concurrently on  $k$  *different* inputs  $x_i$  using  $k$  parallel processors. The *average* time per VDF output approaches  $T / k$  as  $k$  increases. This violates the *perception* of a fixed delay per output if the outputs are used independently.
- **Impact on Applications:**
- **Timed Releases / Time-Lock Encryption:** If an adversary wants to decrypt  $k$  messages encrypted for time  $T$ , they can compute all  $k$  decryptions in parallel, finishing in time  $\sim T$  (if they have  $k$  processors), rather than  $k \cdot T$ . This negates the intended release schedule per message.
- **Spam Prevention:** If a VDF is used as a spam deterrent (one VDF proof per email), a spammer with a botnet ( $k$  machines) can generate  $k$  proofs in time  $\sim T$ , sending  $k$  emails in the time it should take to send one. This defeats the purpose.
- **Randomness Beacons & Leader Election:** These typically rely on a *single* VDF output per epoch derived from a *single* unpredictable seed  $x$ . Amortization doesn't directly help an adversary control *this specific* output faster. However, if an adversary seeks to influence *some* leader election within a window, they could precompute VDFs for many potential future seeds, hoping one matches (see Precomputation below).
- **Mitigation:** The core defense is ensuring that the critical application relies on a **single VDF output derived from a single, unpredictable input  $x$  per time period**. For applications requiring many independent delays (like spam prevention), VDFs are generally unsuitable; traditional Proof-of-Work, despite its parallelizability and energy cost, might be more appropriate. Chaining VDF inputs (using output  $y_i$  as input  $x_{i+1}$ ) can force sequentiality across outputs but adds complexity and latency.
- **Precomputation Attacks: Exploiting Predictable Inputs:**
- **The Vulnerability:** If an adversary can predict (or influence) a *future* input  $x$  to the VDF *before* the computation needs to start, they can begin computing  $\text{Eval}(\text{pp}, x)$  in advance. When  $x$  is officially revealed, they may have already completed the computation or be far ahead, negating the intended delay.
- **Impact on Applications:**
- **Randomness Beacons (RANDAO + VDF):** Recall the setup: participants commit to random seeds  $s_i$ , then reveal them sequentially. The final seed  $x$  is the hash of all revealed  $s_i$ , fed into a VDF. If an adversary is the *last* to reveal ( $s_k$ ), they know all previous seeds  $s_1, \dots, s_{k-1}$  *before* revealing their own. They can compute  $x' = H(s_1, \dots, s_{k-1}, s_{k'})$  for *many* candidate  $s_{k'}$  values. For each candidate  $x'$ , they start computing  $\text{Eval}(\text{pp}, x')$  in parallel on many machines. When they finally reveal a chosen  $s_{k'}$ , they select the candidate output  $y'$

corresponding to the  $x'$  whose VDF computation is *most favorable* (e.g., biases the final randomness). While they still must wait  $\sim T$  time *after* choosing  $s_k'$ , the massive parallel precomputation across candidates allows them to effectively choose  $s_k'$  to bias the outcome, partially negating the VDF's benefit. This is a **limited** form of bias compared to no VDF, but not perfect neutrality. The VDF delay  $T$  limits the number of candidates they can explore ( $\text{number\_of\_candidates} = \text{parallel\_machines} * (T / t_{\text{step\_per\_candidate}})$ ), but determined adversaries with vast resources can still exert influence.

- **Leader Election:** If the input to the VDF for slot  $n+1$  is predictable from slot  $n$  (e.g., derived from the blockchain state), an adversary could start computing the VDF early, gaining an advantage over honest validators who start at the official slot time.
- **Mitigation:**
- **Unpredictable Inputs:** This is the most crucial defense. Ensure the input  $x$  to the critical VDF cannot be predicted significantly before the computation must start. In RANDAO+VDF, using a commit-reveal scheme with a *short* reveal phase relative to  $T$  limits the adversary's window for parallel precomputation across candidates. Chia uses the previous VDF output as part of the input for the next, enhancing unpredictability.
- **Input Commitment and Binding:** Require participants to commit to their contribution (e.g.,  $s_k$ ) *before* seeing others, and enforce that they reveal the committed value. This prevents them from adaptively choosing  $s_k'$  based on others' reveals, but precomputation on the *known* commitments is still possible if the final combination  $x$  is computable from them.
- **Sufficiently Long Delay ( $T$ ):** Setting  $T$  large enough makes the advantage gained by precomputation (e.g., the number of candidates an adversary can explore) negligible for the desired security level. This requires careful calculation based on estimates of an adversary's maximum parallel power.

Amortization and precomputation attacks reveal that VDFs are not a panacea. Their sequentiality guarantee is tightly bound to the uniqueness and unpredictability of the input. Applications must be carefully designed to minimize opportunities for parallelism across evaluations and to ensure inputs cannot be gamed or predicted prematurely. The VDF delay  $T$  must be calibrated not just against raw hardware speed, but also against an adversary's potential for massive parallel precomputation when inputs are partially controllable.

The security landscape of VDFs is complex and dynamic. Cryptanalysis probes the mathematical foundations, hardware advancements erode the wall-clock delay, trusted setups introduce procedural risks, and clever attackers exploit parallelism across inputs or time. These limitations do not negate the value of VDFs but define the parameters within which they can be reliably deployed. Robust VDF systems demand conservative parameter choices, trust-minimized setups (preferably class groups), vigilant monitoring for cryptanalytic breaks, and application designs that mitigate amortization and precomputation. Having scrutinized the potential pitfalls, we now turn to the vibrant domain where VDFs are actively shaping technology: their



**Core Applications and Use Cases** in blockchain randomness, consensus enhancement, and beyond, exploring how these powerful primitives are being harnessed despite the challenges.

---

## 1.6 Section 6: VDFs in Action: Core Applications and Use Cases

The rigorous mathematical foundations and intricate security considerations explored in previous sections are not abstract intellectual exercises—they are the bedrock upon which Verifiable Delay Functions (VDFs) deliver transformative capabilities in real-world systems. Having navigated the theoretical depths and confronted practical vulnerabilities, we now witness VDFs emerge from the cryptographic laboratory into the vibrant arena of deployment. Their unique ability to generate trustless, publicly verifiable proof of sequential time elapsed unlocks solutions to fundamental challenges in decentralized coordination, fairness, and security. This section illuminates the diverse and impactful applications where VDFs are not merely theoretical constructs but active, critical components, primarily within the blockchain ecosystem but extending into broader domains of digital trust.

The security challenges outlined in Section 5—cryptanalytic threats, trusted setup dilemmas, the relentless march of hardware acceleration, and the nuances of amortization—underscore that VDF deployment demands careful design. Yet, these challenges are navigated because VDFs solve problems that are otherwise intractable in trust-minimized environments. From ensuring unbiased randomness in multi-billion dollar protocols to anchoring Proof-of-Stake consensus in physical time and enabling sustainable alternatives to energy-intensive mining, VDFs are proving their worth as foundational cryptographic primitives for the decentralized age.

### 1.6.1 6.1 Randomness Beacons: Unpredictable and Unbiasable

**The Problem:** Generating public, unpredictable, and unbiased randomness in decentralized systems is notoriously difficult. This “randomness beacon” is crucial for applications like:

- **Proof-of-Stake (PoS) Leader Election:** Randomly selecting which validator proposes the next block.
- **Sharding:** Randomly assigning validators to shards.
- **Lotteries and Gaming:** Fair on-chain games and prize distribution.
- **Governance:** Random sampling for decentralized autonomous organization (DAO) committees.

Naïve approaches are vulnerable:

- **Single Oracle:** A trusted party is a single point of failure and corruption.



- **Block Hashes:** Miners/validators can manipulate their block content to influence the hash, making it predictable and biasable.
- **Commit-Reveal Schemes:** Participants commit to random numbers, then reveal them. The last revealer sees all others first and can choose their own number to manipulate the final output (`last-revealer bias`). Simple commit-reveal is highly vulnerable.

**How VDFs Solve It:** VDFs act as a cryptographic “delay mixer,” neutralizing last-revealer bias and ensuring unpredictability. The canonical solution is the **RANDAO + VDF hybrid beacon**, prominently adopted in **Ethereum 2.0’s (now the Ethereum consensus layer) design**:

1. **Commit Phase (RANDAO):** In each epoch (e.g., ~6.4 minutes), each validator  $i$  commits to a random seed  $s_i$  by publishing a hash  $H(s_i)$ .
2. **Reveal Phase:** Validators reveal their  $s_i$ . The beacon input  $x$  is constructed as the hash of all revealed seeds:  $x = H(s_1 || s_2 || \dots || s_k)$ .
3. **VDF Delay:** The input  $x$  is fed into a VDF set with a significant delay parameter  $T$  (e.g., targeting 1-10 minutes). The VDF computes  $y = \text{Eval}(pp, x)$  and outputs  $(y, \pi)$ .
4. **Output:** The final randomness is derived from  $y$  (often simply  $y$  itself or a hash of it).

#### Why it Works:

- **Neutralizing Last-Revealer Bias:** The last validator to reveal  $s_k$  knows all previous seeds  $s_1, \dots, s_{\{k-1\}}$  and thus can compute  $x' = H(s_1, \dots, s_{\{k-1\}}, s_k')$  for many candidate  $s_k'$  values. *However*, they cannot compute  $y' = \text{Eval}(pp, x')$  instantly. The VDF imposes a mandatory, sequential delay  $T$ . They can start computing  $\text{Eval}(pp, x')$  for many candidates in parallel, but the number of candidates they can explore is limited by  $(\text{parallel\_processors} * T) / t_{\text{step}}$ . With  $T$  set sufficiently long and  $t_{\text{step}}$  determined by the fastest known hardware (ASICs), the number of candidates an adversary can meaningfully influence becomes negligible for practical security (e.g.,  $2^{30}$  candidates might be feasible for a powerful adversary, but influencing a 256-bit output requires finding a bias among  $2^{256}$  possibilities – astronomically harder). They are forced to commit to  $s_k'$  before knowing which  $y'$  it will produce.
- **Unpredictability:** Until the VDF computation completes  $T$  sequential steps after the last reveal, *no one* knows  $y$ . The output is unpredictable until the delay has genuinely elapsed.
- **Public Verifiability:** Anyone can quickly `Verify(pp, x, y, π)` to confirm  $y$  is the correct output of  $T$  steps on the agreed input  $x$ .

#### Case Studies:

- **Ethereum 2.0:** While Ethereum’s initial Phase 0 beacon chain launch used a simpler RANDAO with attester shuffling (due to VDF hardware complexity concerns), the **long-term roadmap explicitly includes VDF integration** as the ultimate solution for unbiased randomness. The Ethereum Foundation invested heavily in VDF research, ASIC development, and MPC ceremonies for RSA modulus generation. The plan is to incorporate VDFs once robust, decentralized hardware execution is feasible. This demonstrates the perceived criticality of VDFs for the protocol’s security and fairness.
- **Chia Network:** Chia uses VDFs (specifically Wesolowski’s scheme with class groups) as a core component of its timelord infrastructure. While Chia’s primary randomness comes from its Proofs of Space, VDFs are used to finalize blocks and provide a source of verifiable delay, contributing to the overall unpredictability and security of its consensus.
- **Dfinity / Internet Computer:** Dfinity employs a threshold BLS signature scheme combined with a non-interactive distributed key generation (NI-DKG) protocol for its randomness beacon. While not strictly VDF-based, its design shares the goal of unbiased randomness and leverages cryptographic delays in a related manner. VDFs were seriously considered during its design phase.
- **Comparison:** Pure BLS-based beacons (using threshold signatures) offer fast, bias-resistant randomness but require complex key management and potentially weaker liveness guarantees. Commit-reveal schemes are simpler but vulnerable without VDFs. RANDAO+VDF offers a compelling balance: relatively simple participant interaction, strong unbiasedness guarantees contingent on VDF security and parameter tuning, and public verifiability.

The VDF-based randomness beacon exemplifies the power of sequential time as a trust anchor. By enforcing an unavoidable computational delay between the commitment to a seed and the revelation of the final output, VDFs create a window where manipulation becomes computationally infeasible, fostering fairness in critical decentralized processes.

### 1.6.2 6.2 Enhancing Proof-of-Stake (PoS) Security

**The Problem:** Pure Proof-of-Stake (PoS) consensus mechanisms, while energy-efficient, face unique security challenges compared to Proof-of-Work (PoW):

- **Nothing at Stake:** In the event of a blockchain fork, a rational validator has no direct cost (like burned electricity in PoW) to vote on *multiple* forks simultaneously to maximize potential rewards. This can hinder consensus finality and make chain reorganizations easier.
- **Long-Range Attacks:** An attacker who compromises or acquires old validator private keys (even for coins that have since been spent or slashed) could theoretically create an alternative blockchain history branching from a point far in the past. Since creating blocks in PoS has negligible computational cost compared to PoW, they could “outpace” the honest chain from that old point if they control enough

stake. Defenses like “weak subjectivity” require new nodes to trust recent checkpoints, reintroducing an element of trust.

- **Grinding Attacks:** A validator selected to propose a block might subtly manipulate the block’s content (e.g., transaction ordering or including specific transactions) to influence the seed for the *next* leader election, potentially increasing their chances of being selected again.

**How VDFs Solve It - Proof-of-Sequential-Time (PoST):** VDFs introduce a **physical time dimension** into PoS by acting as a “rate limiter” or a mandatory “cooldown” period for critical actions:

1. **Leader Election Cooldown:** When a validator is pseudo-randomly selected (e.g., via the RANDAO+VDF beacon) to propose a block for slot  $n$ , they must first compute a VDF proof *specific to that slot* before they can broadcast the block. The VDF input  $x_n$  is derived from the current beacon state and the slot number. The VDF delay  $T$  is set to be a significant fraction of the slot time (e.g., 50-80%).
2. **Block Finalization Delay:** Validators might be required to compute a VDF after attesting to a block before that attestation is considered fully finalized, adding a time cost to supporting forks.

### Why it Works:

- **Mitigating Nothing at Stake & Long-Range Attacks:** Adding a mandatory VDF computation *per block proposal* imposes a **real-world time cost** on creating blocks. An attacker attempting to build a long-range fork or support multiple forks simultaneously must perform the sequential VDF computations *for each block* on each fork. This significantly slows down their ability to produce conflicting chains compared to the honest chain, as parallelization doesn’t help within one VDF instance. Creating  $k$  competing blocks for the same slot would require computing  $k$  independent VDFs, taking roughly  $k * T$  time, making rapid chain reorganization or long fork creation infeasible within the protocol’s finality window. The VDF acts as a “proof-of-patience,” forcing attackers to wait.
- **Deterring Grinding Attacks:** Attempting to manipulate the next leader seed by grinding through different block variations would require recomputing the VDF for each variation. The sequential delay  $T$  makes grinding over more than a trivial number of variations computationally impractical within the slot time.
- **Fairer Leader Scheduling:** The VDF delay ensures that even if a validator knows they are selected well in advance (e.g., via predictable RANDAO), they cannot act *instantly*. They must wait the fixed VDF computation time, giving other validators and network participants time to prepare.

### Case Studies:

- **Ethereum 2.0 (Proposer Boost):** While full VDF integration for leader election was deferred, Ethereum incorporates the *concept* of enforced delay through its “proposer boost” mechanism. A selected proposer has a limited time window to broadcast their block. If they fail, the slot is skipped. While not cryptographically enforced via VDF, this creates a similar economic pressure. VDFs remain a potential future upgrade for stronger guarantees.
- **Filecoin:** Filecoin’s Expected Consensus (EC) uses VDFs as part of its leader election process. A miner’s power (storage proven) determines their probability of election. Upon being elected, they must compute a VDF proof before broadcasting their block. This directly implements the “VDF cooldown” to mitigate grinding and nothing-at-stake risks. Filecoin utilizes Wesolowski VDFs.
- **Mina Protocol (Ouroboros Samisika):** Mina’s succinct blockchain design uses a variant of Ouroboros PoS. While its primary innovation is recursive zk-SNARKs for constant-sized blockchain proofs, its consensus mechanism incorporates VDFs to add a time delay element, enhancing security against certain attacks within its unique architecture.
- **Comparison:** Alternative PoS security mechanisms include **slashing** (penalizing validators provably signing conflicting blocks, addressing nothing-at-stake directly but requiring complex detection) and **long unbonding periods** (delaying stake withdrawal to allow time to detect and slash misbehavior, mitigating long-range attacks but reducing capital efficiency). VDFs provide a complementary, physics-based defense layer.

By anchoring validator actions to the passage of sequential computation time, VDFs fortify PoS consensus against its inherent economic-game-theoretic vulnerabilities. They transform stake from a purely virtual asset into one bound by the physical constraints of computation time, enhancing the protocol’s resilience.

### 1.6.3 6.3 Proofs of Space-Time (PoST) and Sustainable Consensus

**The Problem:** Proof-of-Work (PoW) provides robust security but at an enormous and unsustainable environmental cost due to its massive energy consumption. Pure Proof-of-Stake (PoS) solves the energy issue but relies solely on economic incentives, lacking a direct physical resource cost that some argue is fundamental for robust decentralization (the “nothing at stake” problem being one manifestation). How can we achieve secure, decentralized consensus that leverages physical resources but avoids PoW’s energy waste?

**The Solution - Proofs of Space-Time (PoST):** PoST combines two primitives:

1. **Proof-of-Space (PoS):** A prover demonstrates they have allocated a specific amount of disk space  $N$  bytes by storing data (often in a “plot” file generated through a complex, one-way process). They can quickly generate proofs that they store specific data chunks upon challenge. Examples include Chia’s plots and Filecoin’s sectors.
2. **Verifiable Delay Function (VDF):** The VDF enforces that the prover must have *stored the data for a minimum duration* – the **Time** component.

## How VDFs Enable Proofs of Space-Time:

1. **Initialization:** A miner (“farmer”) generates a PoS commitment by plotting data onto their storage drives, binding it to their public key. This is expensive and slow but done once.
2. **Challenge:** The network broadcasts a challenge  $c$  (e.g., derived from the blockchain tip).
3. **Space Proof Generation:** The farmer uses their stored plots to generate a proof  $\pi_{\text{space}}$  demonstrating they store data relevant to challenge  $c$ . This is computationally easy.
4. **VDF Delay:** Crucially, the farmer *cannot* instantly broadcast  $\pi_{\text{space}}$ . They must first compute a VDF:  $(y, \pi_{\text{vdf}}) = \text{Eval}(pp, x)$ , where the input  $x$  incorporates  $\pi_{\text{space}}$  and  $c$ . The VDF delay  $T$  is significant (minutes to hours).
5. **Block Proposal:** The farmer broadcasts the block containing  $c$ ,  $\pi_{\text{space}}$ ,  $y$ , and  $\pi_{\text{vdf}}$ .
6. **Verification:** The network verifies  $\pi_{\text{space}}$  is valid for  $c$  and then verifies  $\pi_{\text{vdf}}$  proves  $y$  was computed from  $x = H(c, \pi_{\text{space}})$  after  $T$  sequential steps.

## Why it Works - The Role of the VDF:

- **Preventing Instant Proof Generation (Grinding/Sybil):** Without the VDF, a farmer could generate  $\pi_{\text{space}}$  on demand instantly upon seeing  $c$ . This opens attacks:
- **Grinding:** The farmer could discard unfavorable  $\pi_{\text{space}}$  outputs and recompute new ones until they get one that gives them an advantage (e.g., winning leader election or producing a favorable block). The VDF delay  $T$  makes grinding over more than one or a few  $\pi_{\text{space}}$  outputs per challenge infeasible within the block time.
- **Sybil Attacks / Outsourcing:** An attacker could rent massive amounts of *computation* to generate  $\pi_{\text{space}}$  proofs on the fly without actually storing the data long-term, or outsource proof generation instantly. The VDF forces a mandatory time delay *after* the challenge is seen, making such instant computation useless. The prover *must* have the data stored and ready *before* the challenge to compute  $\pi_{\text{space}}$  fast enough to then feed it into the VDF and finish within the required timeframe. The VDF ensures the proof demonstrates “space held over time.”
- **Enforcing Sequentiality:** The VDF ensures that even if an adversary has the data stored, they cannot bypass the time delay  $T$  using parallelism or faster hardware for *this specific proof instance*. They must wait the sequential time, proving continuous commitment.
- **Sustainability:** PoST leverages underutilized disk space rather than burning vast amounts of energy. While plotting consumes energy (a one-time cost), ongoing farming is extremely energy-efficient, primarily involving disk reads and VDF computation (which, while energy-consuming per instance, is vastly less than continuous PoW hashing).

## Case Study: Chia Network

Chia is the flagship implementation of Proofs of Space-Time. Its consensus relies heavily on VDFs:

- **Timelords:** Specialized nodes run VDF evaluators (“Timelords”). Farmers generate PoSpace proofs ( $\pi_{\text{space}}$ ) in response to network challenges. They send  $\pi_{\text{space}}$  to a Timelord.
- **VDF Computation:** The Timelord computes  $(y, \pi_{\text{vdf}}) = \text{Eval}(pp, x)$  where  $x = H(\text{challenge}, \pi_{\text{space}})$ . Chia uses Wesolowski VDFs with **class groups of imaginary quadratic fields** ( $\text{Cl}(-D)$ ), avoiding any trusted setup.
- **Block Creation:** The Timelord returns  $y$  and  $\pi_{\text{vdf}}$  to the farmer, who assembles and broadcasts the block.
- **Security & Fairness:** The VDF delay  $T$  (configurable, often minutes) prevents farmers from grinding PoSpace proofs and ensures that creating a block requires provable elapsed time after seeing the challenge. Chia’s open-source [VDF implementation](#) is a key part of its ecosystem.
- **Performance:** Class group arithmetic is slower than RSA, but Chia optimizes it heavily. The focus is on the ratio of space proven over time, not ultra-low latency.

PoST, powered by VDFs, offers a compelling vision for sustainable blockchain consensus. By replacing energy burn with the allocation and *persistence* of storage space over verifiable time, it leverages a different physical resource while maintaining strong cryptographic security guarantees against grinding and Sybil attacks.

### 1.6.4 6.4 Mitigating Miner Extractable Value (MEV)

**The Problem:** Miner Extractable Value (MEV) arises when block producers (miners in PoW, validators/proposers in PoS) can reorder, include, or exclude transactions within their blocks to extract additional profit beyond standard block rewards and fees. Common forms include:

- **Front-running:** Seeing a pending lucrative transaction (e.g., a large DEX trade) and placing one’s own transaction ahead of it to profit from the anticipated price impact.
- **Back-running:** Placing a transaction immediately after a known event (e.g., an oracle price update) to capitalize on it.
- **Sandwich Attacks:** Placing transactions both before and after a victim’s large trade to profit from the induced price movement.

MEV leads to inefficiencies (users overpaying), centralization (specialized “searchers” and block builders dominate extraction), and potential censorship.

**How VDFs Can Help - Enforcing Fair Ordering Delays:** VDFs introduce a mandatory delay between the moment transactions are ordered (or commitments to them are made) and the moment the block is built, allowing decentralized actors time to detect and potentially counter manipulative ordering.

### 1. Commit-Delay-Build Paradigm:

- **Propose Order:** A designated entity (could be a decentralized committee, or even the current block proposer) proposes an *ordering* of transactions (or transaction hashes/commitments) for the next block. This proposal is published immediately.
- **VDF Delay:** A VDF computation is initiated using the proposed ordering as input  $x$ . This imposes a fixed delay  $T$ .
- **Build Block:** *After* the VDF delay  $T$  has elapsed (and the VDF output  $y$  is published and verified), the block *builder* (who could be the same entity or a different one) constructs the actual block by including the transactions in the *previously committed order*.
- **Verification:** The builder publishes the block and the VDF proof  $\pi$ . Verifiers check the block's transactions match the committed order and that  $\pi$  verifies  $y$  was computed from the commitment after  $T$  steps.

### Why it Works:

- **Prevents Last-Second Manipulation:** The builder cannot change the order after seeing the VDF input  $x$  (the committed order). They are locked in by the commitment. The VDF delay  $T$  ensures that any attempt to propose a new, manipulative order would be detected because the VDF output for the *correct* order has already been in progress for time  $T$ .
- **Enables Fair Response:** During the delay  $T$ , network participants (users, other builders, decentralized watchdogs) can scrutinize the proposed order. If they detect front-running or other manipulation, they can potentially react:
- **Counter-Trades:** Users could submit counter-trades to mitigate the impact of detected front-running.
- **Challenging Orders:** In more advanced designs, participants might have a mechanism to challenge maliciously ordered commitments during the delay period, forcing a fallback ordering mechanism.
- **Reputation Systems:** Consistently proposing unfair orders can damage the proposer's reputation.
- **Reduces Time Pressure for Builders:** By separating ordering from execution, builders have the full  $T$  time to optimize block *execution* (e.g., gas efficiency) within the fixed order, rather than frantically trying to both order and build optimally instantly.

### Examples and Initiatives:



- **“TimeBoost” in MEV-Boost / PBS (Proposer-Builder Separation):** While not using a full VDF, Ethereum’s MEV-Boost market incorporates the concept of a “time boost” relay. Builders submit bids (blocks with ordering) to relays. Proposers select the highest bid. A “time boost” auction allows builders to pay extra for their bid to be revealed to the proposer slightly *later* than others, giving them a last-mover advantage to potentially incorporate the latest transactions. While controversial, this highlights the role of *timing* in MEV extraction. A VDF-enforced delay provides a more structured, less gameable approach than ad-hoc time boosts.
- **Theoretical Proposals:** Several research papers explicitly propose VDF-based fair ordering, such as **Aequitas** and **Themis**. These designs leverage VDFs to enforce the commit-delay-build pipeline, often using threshold cryptography for decentralized ordering committees.
- **Challenges:** Implementing practical, decentralized, and efficient commit-delay-build protocols is complex. VDF delays add latency to block production. Defining “fair” ordering objectively is difficult. Malicious proposers might still try to propose unfair orders, relying on the difficulty of detection within  $T$ . However, VDFs provide a crucial *enforcement mechanism* for the delay, making manipulation more costly and detectable.

VDFs offer a promising cryptographic tool in the battle against MEV. By introducing a verifiable, mandatory delay between transaction ordering commitment and block building, they create a window for transparency and response, potentially fostering fairer and more efficient decentralized markets.

### 1.6.5 6.5 Beyond Blockchain: Timed Releases, Auctions, and More

While blockchain applications have driven much VDF development, their utility extends to classic problems in cryptography and distributed systems requiring verifiable delays or proofs of sequential work.

- **Timed Releases and “Crypto Time Capsules”:**
- **Problem:** Securely encrypt a message so it can only be decrypted after a specific future date, without relying on a trusted party to release the key. Rivest’s original time-lock puzzle (1996) pioneered this concept but lacked verifiability.
- **VDF Solution:** Use the VDF output  $y = \text{Eval}(pp, x)$  as the key (or part of the key) to encrypt the message  $M$ :  $C = \text{Encrypt}(y, M)$ . Publish  $pp, x, C$ . To decrypt, compute  $y$  after time  $T$  and then  $M = \text{Decrypt}(y, C)$ . Crucially, anyone can *verify* a claimed  $y'$  and proof  $\pi$  using  $\text{Verify}(pp, x, y', \pi)$  to confirm it’s the genuine key before attempting decryption. This prevents spoofing.
- **Use Cases:** Embargoed document releases (e.g., news, financial reports, wills), delayed activation of software features or licenses, phased decryption of sensitive archives.

- **Enhancement:** Combine with threshold cryptography to require multiple parties to contribute to  $x$ , ensuring no single party can decrypt early.
- **Preventing Auction Sniping:**
- **Problem:** In online auctions, “sniping” occurs when bidders wait until the last possible moment to submit a bid, preventing others from responding. A trusted auctioneer can enforce hard deadlines, but decentralized auctions lack this.
- **VDF Solution:** Require bidders to submit a *commitment* to their bid  $b$  (e.g.,  $H(b, \text{nonce})$ ) well before the deadline. When the deadline passes, bidders reveal  $b$  and  $\text{nonce}$ . Crucially, they must also submit a VDF proof computed on their commitment:  $(y, \pi) = \text{Eval}(pp, H(\text{commitment}))$ . The VDF delay  $T$  starts *after* the commitment phase ends. The highest valid bid (with valid commitment opening and VDF proof) wins.
- **Why it Works:** A sniper cannot wait until the last second to decide their bid. To have their bid considered, they must have committed to it early enough that their VDF computation (started after commitment) finishes within a reasonable time after the deadline. This deters last-second bids by making them impossible to validate in time if started too late. The VDF proof verifies they committed early enough to start the computation.
- **Resource Pricing and Spam Prevention Revisited:**
- **Problem:** Deterring Sybil attacks and spam (e.g., email, API access, comment posting) by imposing a mandatory cost per action. Proof-of-Work (like Hashcash) is commonly used but wastes energy and is highly parallelizable by botnets.
- **VDF Solution:** Require a valid VDF proof  $(y, \pi)$  for a service request. The input  $x$  could be derived from the request content and a nonce. The VDF delay  $T$  imposes a mandatory time cost per request.
- **Advantages over PoW:** Resistant to parallelization by botnets; the cost per request is fixed in *time* rather than energy (though energy is still consumed). An adversary with many machines can make many requests, but each request still takes a minimum wall-clock time  $T$  per machine. This might be more user-friendly than unpredictable PoW solve times.
- **Challenges:** Requires efficient verification and potentially client-side VDF computation libraries. May not be suitable for very high-frequency requests due to the inherent delay. Hash-based VDFs (Sloth) are sometimes considered here due to simplicity, despite weaker sequentiality.
- **Secure E-Voting:**
- **Problem:** Preventing last-minute coercion or vote buying in electronic voting. If voters can change their vote until the very last second, coercers can demand proof of the final vote.

- **VDF Solution:** Voters submit a commitment to their vote early. To finalize their vote, they must compute and submit a VDF proof based on that commitment *after* a set deadline. This creates a “point of no return” enforced by sequential computation time. Changing one’s mind after committing would require recomputing the VDF, which is impossible within the remaining time if  $T$  is set appropriately relative to the commitment deadline. Verifiability ensures the final vote corresponds to the early commitment.
- **Proof of Sequential Work for Non-Blockchain Apps:** Any application needing proof that a client performed a minimum amount of sequential computation (not just work) could leverage VDFs, such as:
- **Rate Limiting Access Tokens:** Granting API access tokens only after solving a VDF, ensuring clients cannot rapidly acquire vast numbers of tokens via parallel requests.
- **Delayed Key Derivation:** Deriving decryption keys for archival data only after a VDF delay, ensuring keys aren’t exposed too soon after encryption.

VDFs transcend the blockchain domain, offering a general-purpose mechanism to enforce and verify the passage of sequential computation time in decentralized settings. From ensuring fair auctions and secure document releases to combating spam and enhancing e-voting, the ability to create trustless delays opens avenues for fairer, more secure, and more robust digital interactions. The journey of VDFs, from conceptual necessity to mathematical formalization, efficient construction, and diverse application, underscores their emergence as a foundational pillar of modern cryptography, shaping the future of decentralized trust.

The exploration of these core applications demonstrates the tangible impact VDFs are having and are poised to have. However, harnessing this power requires translating elegant mathematical schemes into efficient, secure, and reliable software and hardware systems. This brings us to the critical **Implementation Landscape: Hardware, Software, and Optimization**, where the rubber meets the road in deploying VDFs at scale.

---

## 1.7 Section 7: Implementation Landscape: Hardware, Software, and Optimization

The journey of Verifiable Delay Functions—from cryptographic abstraction to deployed infrastructure—reaches its critical juncture in the realm of implementation. While Sections 5 and 6 explored the security boundaries and transformative applications of VDFs, the practical realization of these concepts hinges on overcoming formidable engineering challenges. Translating the elegant mathematics of repeated squaring in Groups of Unknown Order (GUOs) into efficient, reliable, and scalable systems demands specialized hardware, optimized software, and rigorous standardization. **This section delves into the trenches of VDF implementation, where nanoseconds per modular squaring determine economic viability, open-source**

**libraries become pillars of trust, and custom silicon reshapes the landscape of decentralized timekeeping.** Here, the theoretical guarantees of sequentiality confront the realities of transistor physics, thermal constraints, and the relentless pursuit of performance.

The security analyses and application landscapes previously discussed set non-negotiable requirements: VDFs must deliver verifiable delays ranging from seconds to hours, with verification times measured in milliseconds, all while resisting cryptanalytic and hardware-based attacks. Meeting these demands, especially for high-value blockchain consensus and randomness beacons, necessitates pushing the boundaries of computational efficiency. The implementation landscape is thus a dynamic interplay between algorithmic refinement, hardware acceleration, and collaborative software ecosystems, all striving to make verifiable delay a practical, trustworthy primitive in global-scale systems.

### 1.7.1 7.1 The Computational Core: Modular Exponentiation and Beyond

At the heart of the dominant GUO-based VDFs (Pietrzak, Wesolowski) lies a computationally intensive primitive: **iterated modular exponentiation**, specifically **repeated modular squaring** ( $y = g^{(2^T)}$  mod  $N$  for RSA,  $[g]^{2^T}$  for class groups). Optimizing this core operation is paramount, as it consumes the vast majority of the evaluation time  $T$ .

- **The Bottleneck: Why Squaring Dominates:**

Each squaring operation within the  $T$ -step chain is fundamentally a large integer multiplication followed by a reduction modulo the group modulus ( $N$  for RSA, the discriminant-derived modulus for class groups). For security-relevant parameters (e.g., RSA-2048, RSA-3072, RSA-4096; class group discriminants of 700-1000+ bits), these operations involve manipulating integers with hundreds or thousands of bits. Performing  $T$  such operations sequentially (where  $T$  can be  $10^9$  or more for minute-scale delays) creates an immense computational burden. Reducing the time per squaring ( $t_{\text{step}}$ ) directly reduces the wall-clock delay  $t_{\text{eval}} = T * t_{\text{step}}$ .

- **Optimizing Modular Arithmetic:**

Efficient algorithms are crucial:

1. **Montgomery Multiplication:** The gold standard for modular reduction. It transforms operands into a residue system where reduction modulo  $N$  can be performed using only shifts, additions, and multiplications, avoiding expensive division operations. Precomputation based on  $N$  allows for highly efficient reduction steps. Its dominance is near-universal in high-performance implementations.
2. **Barrett Reduction:** An alternative method using precomputed reciprocal approximations of  $N$  to estimate quotients, followed by correction steps. While theoretically comparable in complexity to Montgomery, it often involves more instructions and is less commonly used in cutting-edge VDF implementations due to Montgomery's superior performance and simpler correction logic.

3. **Karatsuba and Toom-Cook Multiplication:** For multiplying the large integers before reduction, these algorithms reduce the computational complexity below the naive  $O(n^2)$  for  $n$ -bit numbers. Karatsuba ( $O(n^{\log_2 3}) \approx O(n^{1.585})$ ) is commonly used for intermediate sizes, while Toom-Cook (e.g., Toom-3,  $O(n^{\log_3 5}) \approx O(n^{1.465})$ ) can be faster for very large numbers (thousands of bits), though with higher overhead. For the sizes used in VDFs (up to  $\sim 4096$  bits), highly optimized schoolbook or Comba multiplication combined with Montgomery reduction often outperforms more complex algorithms due to lower constant factors and better cache locality.
4. **Vectorization (SIMD):** Exploiting Single Instruction, Multiple Data (SIMD) units in modern CPUs (e.g., AVX2, AVX-512 on x86; NEON on ARM) can significantly speed up the inner loops of multiplication and Montgomery reduction. Processing multiple 32-bit or 64-bit limbs in parallel within wide registers (256-bit or 512-bit) reduces instruction count and improves throughput. Libraries like GMP (GNU Multiple Precision Arithmetic Library) leverage SIMD heavily.

- **The Class Group Challenge:**

Arithmetic in **ideal class groups of imaginary quadratic fields** ( $\mathbb{C}l(-D)$ ) is inherently more complex than RSA modular arithmetic:

- **Representation:** Ideals are represented by integer triples  $(a, b, c)$  satisfying  $b^2 - 4ac = -D$ , where  $a > 0$  and  $\gcd(a, b, c) = 1$ .

- **Operations:** Ideal multiplication involves:

1. Computing the product of the ideals' forms.
2. Computing the greatest common divisor (GCD) of related values.
3. **Reduction:** Finding the equivalent "reduced" ideal (minimizing  $a$ ) using a specialized algorithm akin to the Euclidean algorithm. This reduction step is computationally expensive and lacks the straightforward parallelizability of modular multiplication.

- **Optimizations:** Implementations (like Chia's) rely heavily on:

- **Nucomp and Nudupl:** Efficient algorithms for composition and reduction of binary quadratic forms, optimized for the case of negative discriminants. These minimize the number of GCD steps required during reduction.
- **Assembly Optimization:** Critical inner loops (GCD, partial modular inverses) are often hand-optimized in assembly language to minimize overhead.
- **Avoiding Full Reduction:** Some intermediate steps might use partially reduced forms where possible, deferring full reduction until necessary.

- **Performance Gap:** Despite intense optimization, class group squaring remains significantly slower than RSA modular squaring at equivalent security levels (estimates range from 10x to 100x slower). This is the primary trade-off for gaining transparent setup.
- **Beyond Squaring: Wesolowski’s Proof Generation:** In Wesolowski’s scheme, proof generation involves an expensive operation: computing  $\pi = (y * g^{-r})^{1/\ell} \bmod N$  (or equivalent in class groups). This requires:
  - Computing  $g^{-r} \bmod N$  (fast, as  $r$  is small).
  - Computing the modular inverse of  $(y * g^{-r})$  raised to the power  $1/\ell$  modulo  $N$  – essentially finding a modular  $\ell$ -th root. This typically involves solving the discrete logarithm in the subgroup of order  $\ell$  (using Pollard’s Rho or similar,  $O(\sqrt{\ell})$  time) or leveraging the Adleman-Manders-Miller algorithm, which requires knowing the factorization of  $\ell - 1$  (often feasible for prime  $\ell$ ). This step, while independent of  $T$ , can be a bottleneck for high-frequency VDF evaluations due to its  $O(\sqrt{\ell})$  complexity. Optimizations focus on efficient implementations of these root-finding algorithms.

The relentless optimization of modular and class group arithmetic forms the foundation of practical VDFs. Every cycle shaved off a squaring operation translates directly into shorter real-world delays or reduced hardware costs. However, for delays measured in minutes or hours at scale, even the fastest CPUs struggle, necessitating specialized hardware.

### 1.7.2 7.2 Hardware Acceleration: ASICs and FPGAs

Achieving practical VDF delays ( $T$  seconds) for real-world deployment, especially in high-throughput scenarios like blockchain consensus, often requires moving beyond general-purpose CPUs. The inherently parallel nature of modular arithmetic operations makes them exceptionally well-suited for **hardware acceleration** using Field-Programmable Gate Arrays (FPGAs) and Application-Specific Integrated Circuits (ASICs).

- **The Imperative for Hardware:**
- **Performance Demands:** A delay  $T$  of 1 minute requires completing  $T$  sequential squarings within 60 seconds. For RSA-3072,  $t_{\text{step}}$  on a high-end CPU might be  $\sim 1$  microsecond, requiring  $T = 60,000,000$  steps. Achieving a  $t_{\text{step}}$  of 10 nanoseconds would reduce  $T$  to 6,000,000 steps for the same wall-clock time, but more importantly, it allows achieving much *shorter* wall-clock times for the same  $T$  (e.g., 0.6 seconds instead of 60 seconds). For class groups, hardware acceleration is even more critical due to their higher  $t_{\text{step}}$ .
- **Energy Efficiency:** Dedicated hardware can perform the specific computation (modular squaring) orders of magnitude more efficiently (in terms of operations per Joule) than a general-purpose CPU executing complex instruction streams.

- **Deterministic Latency:** Hardware designs offer precise control over timing, crucial for meeting strict slot times in consensus protocols.
- **FPGAs: Flexibility in Silicon:**
  - **Technology:** FPGAs consist of programmable logic blocks, DSP slices, and memory blocks interconnected by a configurable fabric. Developers describe hardware circuits using Hardware Description Languages (HDLs) like VHDL or Verilog, which are synthesized and loaded onto the FPGA.
  - **Advantages:** Faster time-to-market and lower NRE (Non-Recurring Engineering) costs than ASICs. Designs can be updated in the field. Suitable for prototyping and initial deployment.
  - **VDF Implementations:** The **Ethereum Foundation** spearheaded early FPGA development for RSA-based VDFs. Their efforts, documented on GitHub, focused on highly pipelined Montgomery multipliers targeting Xilinx Ultrascale+ FPGAs. Key optimizations included:
    - **Deep Pipelining:** Breaking the squaring operation into dozens or hundreds of stages, allowing a new squaring operation to start every few clock cycles (high throughput).
    - **High Clock Frequencies:** Pushing FPGA clock rates to 300-500+ MHz through careful timing closure.
    - **Efficient Memory Access:** Designing caches and data paths to feed the computational units without bottlenecks.
  - **Results:** Demonstrated  $t_{\text{step}}$  in the range of **10-50 nanoseconds** for RSA-2048/3072, representing **10-50x speedups** over optimized CPU software.
  - **Challenges:** FPGAs are typically 3-5x slower and less energy-efficient than equivalent ASICs. Resource constraints limit the size of the multiplier that can be implemented (e.g., supporting RSA-4096 requires more resources than RSA-2048).
- **ASICs: The Pinnacle of Performance:**
  - **Technology:** ASICs are custom-designed integrated circuits fabricated for a specific purpose. They offer the ultimate in performance, power efficiency, and area density.
  - **Advantages:** Potential for **100-1000x speedup** and **10-100x better energy efficiency** compared to CPUs. Ultimate performance for mass deployment.
  - **Disadvantages:** Extremely high NRE costs (millions of dollars for design, verification, mask creation, and fabrication). Long development cycles (12-24+ months). Fixed functionality; bugs require respins.
- **Major Projects:**



- **Supranational:** A leader in cryptographic hardware acceleration, Supranational developed highly optimized ASIC designs for both RSA and BLS signatures, targeting Ethereum’s needs. Their RSA VDF core focused on:
  - **Massive Parallelism:** Employing many modular multiplier units working on different limbs of the large integers simultaneously.
  - **Sophisticated Pipelining:** Multi-stage pipelines operating at high frequencies (targeting >1 GHz).
  - **Advanced Process Nodes:** Targeting cutting-edge nodes (e.g., 7nm, 5nm) for maximum performance and efficiency.
  - **Benchmarks:** Projected  $t_{\text{step}}$  well **below 10 nanoseconds** for RSA-3072, enabling VDF evaluations in seconds even for very large  $T$ .
  - **Ethereum Foundation’s VDF Alliance / Competition:** To foster open, accessible hardware and mitigate centralization risks, the Ethereum Foundation launched a [VDF Alliance](#) and ran a [VDF ASIC Competition](#) (2018-2019). Teams from **EPFL** (École polytechnique fédérale de Lausanne), **Supranational**, and **Synopsys** participated. EPFL’s entry, [PipeZK](#), focused on a scalable, open-source modular multiplier design capable of 1 GHz operation. While no single design was mass-produced, the competition advanced open-source VDF hardware IP and highlighted the performance potential.
  - **Class Group ASICs:** While less mature than RSA ASICs, designing hardware accelerators for class group operations is an active research area. The challenge lies in efficiently implementing the complex reduction algorithms in hardware. FPGA prototypes are a likely first step.
- **Performance Comparison Landscape:**

The table below provides *illustrative* performance figures for different hardware platforms targeting RSA-3072 squaring (lower  $t_{\text{step}}$  is better). Actual numbers depend heavily on specific implementations, clock speeds, and process nodes.

Hardware Platform | Approx.  $t_{\text{step}}$  | Notes | Speedup vs. CPU |

:—————- | :—————- | :—————- | :—————- |

High-End CPU (AVX2) | 500 - 1000 ns | e.g., Optimized GMP or dedicated C/Rust | 1x (Baseline) |

High-End CPU (AVX-512) | 300 - 700 ns | Exploiting wider SIMD | ~1.5x |

High-End FPGA (Ultrascale+) | 10 - 50 ns | e.g., Ethereum Foundation design | 10-50x |

ASIC (Modern Node) | 1 - 10 ns | e.g., Supranational / EPFL competition targets | 50-500x+ |

The relentless drive towards specialized hardware underscores the performance demands of practical VDF deployment. While FPGAs offer a flexible stepping stone, ASICs represent the inevitable frontier for achieving the shortest possible delays and enabling their use in latency-sensitive consensus mechanisms. However, this hardware arms race also raises centralization concerns, necessitating robust software ecosystems and open standards.

### 1.7.3 7.3 Software Implementations and Libraries

Bridging the gap between cryptographic theory, hardware acceleration, and application integration falls to software libraries. Robust, efficient, and well-audited open-source implementations are essential for adoption, security, and fostering a diverse ecosystem.

- **Key Libraries and Projects:**

1. **Chia VDF (Python/C++):** The reference implementation for Chia Network’s class group-based VDFs. Features:

- Core performance-critical operations (ideal multiplication, reduction) in optimized C++.
- Python bindings for high-level integration and testing.
- Implements both Wesolowski proof generation/verification and the core repeated squaring.
- Focus on correctness and security for Chia’s production blockchain.

- **GitHub:** <https://github.com/Chia-Network/vdf>

2. **Filecoin rust-fil-proofs (Rust):** Filecoin’s implementation of its Proof-of-Spacetime and associated VDFs (Wesolowski scheme). Written in Rust for performance and safety. Integrated tightly with Filecoin’s Proof-of-Replication (PoRep) and consensus mechanisms. **GitHub:** <https://github.com/filecoin-project/rust-fil-proofs> (VDF components within).

3. **Ethereum Research VDF (Python/C):** Reference and research code from the Ethereum Foundation. Includes implementations of Pietrzak and Wesolowski VDFs for RSA groups, often used as a benchmark and research testbed. Features MPC code for RSA modulus generation. **GitHub:** <https://github.com/ethereum/research/tree/master/vdf>

4. **VDF Alliance Competition Entries (Various):** The open-source submissions from the Ethereum VDF competition provide valuable implementations, particularly EPFL’s PipeZK (C++/HDL) and Supranational’s designs, showcasing high-performance approaches. **GitHub:** <https://github.com/facebookincubator/vdf-competition>

5. **Cryptography Libraries w/ VDF Support:** Emerging integrations into broader cryptographic libraries:

- **OpenSSL Engine:** Proof-of-concept engines integrating VDF computation (e.g., via OpenSSL’s engine interface) exist, though not yet mainstream.
- **libsodium / libsecp Consideration:** While not currently including VDFs, future versions of major crypto libraries might incorporate VDF primitives as standardization progresses.

- **Language Choices and Rationale:**
- **C/C++:** Dominates performance-critical sections (modular arithmetic, class group reduction) due to low-level hardware access, mature compilers, and minimal runtime overhead. Essential for squeezing out maximum CPU/FPGA performance.
- **Rust:** Increasingly popular (e.g., Filecoin) due to its strong memory safety guarantees, performance comparable to C/C++, and excellent concurrency support. Reduces the risk of critical vulnerabilities like buffer overflows common in C/C++.
- **Python / Go:** Often used for higher-level bindings, testing frameworks, prototyping, and integration layers due to developer productivity and rich ecosystems. The core heavy computation is still offloaded to C/Rust modules.
- **Integration Challenges:**

Integrating VDFs into complex systems like blockchain nodes presents hurdles:

- **Hardware Abstraction:** Supporting both CPU fallback and hardware acceleration (FPGA/ASIC) requires clean abstraction layers. Standards like **PCIe** or **OpenCL** might be used, but custom interfaces are common.
- **Concurrency and Resource Management:** VDF evaluation is long-running and blocking. Efficiently managing evaluation threads/processes, especially when multiple VDFs might be needed concurrently (e.g., in PoST systems), is crucial. Interruptibility for block production deadlines adds complexity.
- **Proof Handling:** Efficient serialization, deserialization, and network transmission of proofs ( $\pi$ ) are important, especially for Pietrzak's  $O(\log T)$ -sized proofs. Wesolowski's constant-size proofs have an advantage here.
- **Dependency Management:** VDF libraries often have complex dependencies (GMP, custom hardware drivers). Packaging and distribution, especially for end-user blockchain clients, can be challenging. Statically linked binaries or containerization are common solutions.

The availability of high-quality, open-source implementations fosters trust, enables independent verification, and lowers the barrier to entry for projects seeking to leverage VDFs. These libraries are the workhorses translating cryptographic promises into operational reality.

#### 1.7.4 7.4 Optimizations for Verification and Proof Generation

While evaluation (squaring) dominates the time cost, optimizing verification and proof generation is vital for system scalability, responsiveness, and usability, especially on resource-constrained devices.

- **Verification Optimizations:**
- **Wesolowski’s Edge:** Wesolowski’s scheme inherently offers near-optimal verification:  $O(\lambda)$  time (a few milliseconds) and constant proof size (one group element). The primary optimizations focus on:
- **Fast Modular Exponentiation:** Efficiently computing  $\pi^\ell \bmod N$  and  $g^r \bmod N$  using standard exponentiation algorithms (e.g., fixed-window, sliding-window) optimized for the small exponent sizes ( $|\ell|, |r| \approx 128\text{-}256$  bits). Leveraging SIMD or hardware acceleration is less critical here but still beneficial for very high throughput verification.
- **Batch Verification:** If multiple VDF outputs  $(x_i, y_i, \pi_i)$  need verification (e.g., verifying multiple blocks or timelord outputs), techniques can combine checks. For Wesolowski, a simple approach is verifying each independently but sharing the cost of setting up modular contexts. More advanced batching using techniques like [Pippenger’s algorithm](#) for multi-exponentiation might offer marginal gains but is often overkill given the already low per-proof cost.
- **Pietrzak’s Verification:** Pietrzak’s  $O(\log T)$  verification involves recursively checking  $O(\log T)$  group elements and equations. Optimizations include:
- **Parallel Recursion:** The recursive verification of left and right sub-proofs  $(\pi_L, \pi_R)$  at each level can often be performed concurrently.
- **Multi-Exponentiation:** Combining the exponentiations needed at each recursive step (e.g., checking  $\mu = g^{(2^{\lceil T/2 \rceil})}$  involves  $g$  and  $\mu$ ) using fast multi-exponentiation algorithms like Pippenger’s or Bos-Coster can significantly reduce the total number of group operations.
- **Proof Size Compression:** While the proof is  $O(\log T)$  elements, techniques exist to compress the representation of intermediate points  $(\mu_i)$ , especially in class groups where elements have a compact canonical form.
- **Proof Generation Optimizations (Wesolowski):**

The  $O(\sqrt{\ell})$  complexity of generating  $\pi = (y * g^{-r})^{\{1/\ell\}}$  can be a bottleneck, particularly for class groups or high-frequency evaluations.

- **Optimized Root Finding:** Highly tuned implementations of Pollard’s Rho or Adleman-Manders-Miller for the specific prime  $\ell$  derived from the challenge. Precomputation based on the factorization of  $\ell - 1$  (if known) can accelerate Adleman-Manders-Miller.
- **Parallelization Potential:** While the core root-finding algorithm is sequential, exploring multiple potential solutions or parallelizing parts of Pollard’s Rho might offer limited gains. The inherently sequential nature limits speedups.
- **ASIC/FPGA Offload:** For extreme throughput needs, the proof generation step itself could be accelerated in hardware, though the variable prime  $\ell$  makes this more complex than fixed squaring.

- **Proof Generation Optimizations (Pietrzak):**

Pietrzak’s proof generation (`Eval`) naturally computes the intermediate points ( $\mu$ ) needed for the proof. The main overhead is managing the recursive proof structure. Optimizations include:

- **Parallel Sub-Proof Generation:** Once a midpoint  $\mu$  is computed, the proofs for the left half ( $\pi_L$  for  $g \rightarrow \mu$ ) and right half ( $\pi_R$  for  $\mu \rightarrow y$ ) can be generated concurrently on separate CPU cores or machines.
- **Efficient Storage/Recomputation:** Storing intermediate squaring states allows efficient recomputation for sub-proofs without restarting the entire chain. Trading memory for computation time.

Optimizing the “fast” parts of VDFs ensures the system remains responsive and scalable. Wesolowski’s verification efficiency makes it particularly attractive for networks with many lightweight participants, while Pietrzak’s simpler proof generation offers advantages in high-throughput proving scenarios.

### 1.7.5 7.5 Benchmarking, Testing, and Standardization Efforts

Establishing reliable performance metrics, ensuring correctness, and defining interoperability standards are crucial for the maturation, adoption, and security of VDF technology.

- **Benchmarking: Establishing Fair Comparisons:**

Measuring VDF performance requires standardized methodologies focusing on:

- **Delay Time ( $t_{eval}$ ) vs.  $T$ :** The core metric: How long does `Eval` ( $pp, x$ ) take to produce  $(y, \pi)$  for a given security level ( $\lambda$ ), delay parameter  $T$ , and hardware platform? Must be measured on *representative* hardware and averaged over many runs.
- **Verification Time ( $t_{verify}$ ):** Time for `Verify` ( $pp, x, y, \pi$ ) to run. Critical for client performance.
- **Proof Size:** Size of  $\pi$  in bytes. Impacts network transmission and storage.
- **Proof Generation Time (for Wesolowski):** Time to generate  $\pi$  *after*  $y$  is known. Relevant for high-frequency VDFs.
- **Energy Consumption:** Power usage during `Eval` and `Verify`. Important for sustainability assessments.
- **Challenges:** Reproducibility across different hardware environments, isolating VDF computation from system noise, accurately measuring wall-clock time for long-running `Eval`. Projects like Chia and Ethereum publish detailed benchmarks for their implementations.

- **Testing and Verification: Ensuring Correctness and Security:**

Rigorous testing is paramount, given VDFs' critical role in consensus:

- **Unit Tests:** Comprehensive tests for individual components (modular arithmetic, class group operations, proof generation, verification).
- **Functional Tests:** End-to-end tests verifying correct output  $y$  and proof  $\pi$  for known inputs and parameters.
- **Property-Based Testing:** Using frameworks to generate random inputs and parameters, checking that properties like `Verify(Eval(...)) == true` and uniqueness hold.
- **Differential Testing:** Running multiple independent implementations (e.g., Chia VDF vs. a research implementation) on the same input and comparing outputs to detect discrepancies.
- **Fuzz Testing:** Feeding malformed or adversarial inputs to uncover crashes or vulnerabilities.
- **Formal Verification:** Applying mathematical methods to prove correctness of algorithms (especially complex ones like class group reduction or Wesolowski proof generation) and implementations. While challenging, efforts are emerging, particularly for critical components. The [HACL\\*](#) project's ethos (formally verified crypto) is relevant here.
- **Test Vectors:** Publicly available sets of precomputed  $(pp, x, T, y, \pi)$  for specific schemes and parameters, allowing implementers to verify correctness. Essential for interoperability.
- **Standardization: Building Consensus and Interoperability:**

Standardization fosters trust, enables interoperability between different implementations, and provides clear security guidelines.

- **IETF VDF Working Group:** The Internet Engineering Task Force established a [VDF Working Group](#) to develop standards. Key areas include:
- **Algorithm Specifications:** Precise, unambiguous definitions of VDF schemes (e.g., Pietrzak RSA, Wesolowski RSA, Wesolowski Class Group), including parameter generation, evaluation, and verification steps. Drafts like [draft-irtf-cfrg-vdf-01](#) are under development.
- **Serialization Formats:** Standard formats for public parameters  $(pp)$ , inputs  $(x)$ , outputs  $(y)$ , and proofs  $(\pi)$  to enable cross-platform compatibility.
- **Security Considerations:** Guidance on parameter selection (modulus size, discriminant size, challenge prime size) based on security level  $(\lambda)$  and delay  $T$ , considering known attacks and hardware advancements. Addressing trusted setup requirements for RSA groups.

- **Test Vectors:** Defining standard test vectors.
- **Industry Consortia:** Groups like the [VDF Alliance](#) (historically driven by Ethereum) play a role in fostering collaboration, sharing research, and promoting best practices, complementing formal standardization.

The ongoing work in benchmarking, testing, and standardization transforms VDFs from research curiosities into robust, deployable cryptographic infrastructure. By establishing common ground, rigorous validation, and performance baselines, these efforts pave the way for wider adoption and integration into the fabric of decentralized systems. As VDFs mature from prototypes to production-grade components, their implementation landscape will continue to evolve, driven by the dual engines of hardware innovation and collaborative software refinement.

The intricate dance between specialized silicon, optimized algorithms, and rigorous software engineering brings the power of verifiable delay into tangible reality. However, the proliferation of VDFs, particularly within high-stakes economic systems like blockchain, inevitably triggers broader questions about power dynamics, environmental sustainability, and equitable access. Having mastered the technical implementation, we must now confront the **Societal and Economic Implications** of binding decentralized trust to the physics of sequential computation.

---

## 1.8 Section 8: Societal and Economic Implications

The relentless drive for hardware acceleration and software optimization, detailed in Section 7, transforms Verifiable Delay Functions from cryptographic abstractions into operational infrastructure with profound societal consequences. As VDFs become embedded in blockchain consensus, randomness generation, and decentralized coordination mechanisms, they inevitably reshape power structures, resource allocation, and economic incentives. **This section examines the tectonic shifts triggered by the fusion of verifiable time and decentralized systems—where the physics of sequential computation collide with human governance, environmental sustainability, and global equity.** The very mechanisms designed to decentralize trust now face scrutiny over potential centralization vectors, the environmental footprint of specialized hardware, novel economic models emerging around “time-as-a-resource,” and geopolitical battles over cryptographic sovereignty. Beyond technical specifications, we confront the uncomfortable question: *Who controls the clocks in a decentralized world, and at what cost to society?*

The implementation landscape revealed a fundamental tension: VDFs require significant computational resources to enforce meaningful delays, creating barriers to entry that contradict the egalitarian ideals of decentralization. This paradox lies at the heart of their societal impact, forcing a reevaluation of what “decentralization” truly means when physical hardware and capital concentration become prerequisites for participation. Simultaneously, VDFs offer a tantalizing vision of sustainability by replacing energy-intensive mining with



proofs of persistent storage over time—yet this shift generates its own environmental controversies and e-waste dilemmas. Economically, VDFs birth new markets for provable sequential computation and novel financial instruments tied to verifiable delays, while geopolitically, they become tools for both censorship resistance and state control. Understanding these implications is crucial as VDFs evolve from cryptographic curiosities into pillars of digital infrastructure.

### 1.8.1 8.1 Decentralization Revisited: Power Dynamics and Access

The foundational promise of VDFs—to enable trust without central authorities—confronts a harsh reality: *specialized hardware creates inherent centralization pressure*. This tension manifests in three critical dimensions:

- **The ASIC Centralization Dilemma:** As explored in Sections 5.3 and 7.2, VDF evaluation demands high-performance hardware (FPGAs/ASICs) for practical deployment. This creates a barrier far higher than that of typical Proof-of-Stake (PoS) validation:
- **Capital Concentration:** Designing and fabricating cutting-edge ASICs costs millions of dollars, favoring well-funded entities (established mining pools, semiconductor giants, venture-backed startups). **Supranational’s VDF ASIC project**, while technologically impressive, exemplified this dynamic—only players with deep pockets could compete. This risks replicating Bitcoin’s ASIC mining centralization, where a handful of pools control hash power. In Ethereum’s planned VDF beacon, entities controlling fast ASICs could gain outsized influence by consistently being first to compute outputs, subtly biasing downstream processes like leader election or MEV extraction windows.
- **Geographic Disparities:** Access to cheap electricity, favorable regulations, and semiconductor supply chains is uneven. Regions with subsidized industrial power (e.g., parts of China, Kazakhstan) or lax environmental oversight could become VDF computation hubs, concentrating influence geographically—a stark contrast to the vision of globally distributed, permissionless participation. The **Ethereum Foundation’s ASIC competition** aimed to democratize access through open-source designs like EPFL’s PipeZK, but fabrication costs and chip shortages still limit real-world decentralization.
- **Governance Vulnerability:** Concentrated VDF computation power creates a single point of failure for governance attacks. Entities controlling significant VDF capacity could collude to stall randomness beacons (by delaying outputs), disrupt leader elections, or censor transactions by manipulating the timing of MEV mitigation protocols. Unlike PoS slashing, where malicious acts are cryptographically provable, timing-based manipulation via controlled VDF delays can be subtle and hard to penalize.
- **The Class Group Advantage and Its Limits:** While class group VDFs (used in **Chia Network**) eliminate the trusted setup risk of RSA groups, they don’t solve hardware centralization. Class group arithmetic is *slower* than RSA (Section 7.1), making hardware acceleration *more* critical for competitive performance. Chia’s “Timelords” (nodes running VDF evaluators) are theoretically permissionless, but in practice:

- **Early Centralization:** At Chia’s launch, the Chia Network company operated most Timelords to ensure chain stability, highlighting the bootstrap problem. While community Timelords emerged, performance disparities between optimized setups and consumer hardware create a hierarchy of influence.
- **Resource Asymmetry:** Entities with custom FPGA/ASIC implementations for class group reduction (an active research area) will outperform CPU-based nodes, potentially dominating block finalization and earning disproportionate rewards. This mirrors the centralization in Filecoin’s VDF-enhanced Proof-of-Spacetime, where large storage providers invest in hardware-accelerated VDF proofs.
- **Comparative Decentralization Models:** VDFs force a reevaluation of decentralization trade-offs:
- **PoW (Bitcoin):** Centralized at the *computation* layer (ASIC farms), but open at the validation layer (any node can verify blocks). Energy-intensive.
- **Pure PoS (e.g., Cardano, early ETH2):** Low hardware barriers for validation, but potential centralization via stake concentration and delegated voting. Minimal energy use.
- **PoST with VDFs (Chia, Filecoin):** Aims for decentralization via distributed *storage*, but risks centralization at the *time-verification* layer (VDF computation). Moderate energy use (mostly from storage plotting/farming).
- **VDF-PoS Hybrids (e.g., planned ETH2):** Combines stake-based validation with VDF-based rate-limiting. Decentralization depends on mitigating VDF hardware centralization via open designs and broad participation.

**The verdict:** VDFs do not inherently decentralize power; they shift the locus of potential centralization from stake concentration (PoS) or hash power (PoW) to control over *sequential computation resources*. True decentralization requires proactive measures: open-source hardware (like PipeZK), accessible fabrication, protocols that reward distributed VDF participation, and vigilance against geographic or capital-based consolidation. Without these, VDFs risk creating a new cryptocrat class—the owners of time.

## 1.8.2 8.2 Environmental Footprint: Energy vs. Storage vs. Time

VDFs emerged partly as a response to Bitcoin’s staggering energy consumption. Yet, their environmental impact is complex and contested, involving direct energy use, embodied energy in hardware, and electronic waste:

- **Energy Consumption: Orders of Magnitude Difference:** VDFs are fundamentally more energy-efficient than PoW, but not negligible:
- **VDF ASICs vs. PoW ASICs:** A Bitcoin mining ASIC (e.g., Bitmain S19 XP, 140 TH/s) consumes ~3-5 kW continuously. A high-performance VDF ASIC (e.g., targeting RSA-3072 at 1 ns/squaring)

might consume 100-500 W while active. Crucially, VDFs only run when needed (e.g., per block proposal), not 24/7. **Filecoin estimates** its VDF-based leader election uses >10,000x less energy than equivalent PoW. Even with millions of VDF instances globally, their aggregate energy use would pale compared to Bitcoin's ~150 TWh/year.

- **The Storage Factor (PoST):** Chia's PoST model shifts environmental impact from energy to storage hardware. Plotting terabytes of data for Proof-of-Space is energy-intensive (one-time burst), but farming (ongoing verification) is efficient (~0.1-1 W/TB for HDDs). However, the 2021 Chia boom caused:
- **HDD/SSD Shortages:** Soaring demand for high-capacity drives, disrupting supply chains.
- **E-Waste Surge:** Low-end SSDs used for plotting failed rapidly under intensive write cycles (some within weeks). Millions of TBs of write endurance were consumed, generating premature e-waste. While Chia shifted plotting to HDDs, the incident highlighted the environmental cost of storage-based consensus bootstrapping.
- **Embodied Energy:** The environmental cost of *manufacturing* hardware matters. ASICs and high-end SSDs/HDDs require significant resources (silicon, rare earth metals, water). A lifecycle analysis must compare:
- PoW: High operational energy + frequent ASIC turnover (1-2 years).
- VDF ASICs: Lower operational energy + longer lifespan (5+ years) if designed for protocol stability.
- PoST: Low operational energy + high upfront storage cost + drive replacement cycles (3-5 years).
- **Sustainability Claims vs. Reality:** Proponents tout PoST/VDF systems as "green" alternatives. The reality is nuanced:
- **Relative Efficiency:** Yes, PoST/VDF systems consume vastly less *operational* energy than PoW. Filecoin and Chia networks likely use less than 0.1% of Bitcoin's energy.
- **Absolute Impact:** However, large-scale PoST farming still consumes terawatts annually. VDF ASICs add to global semiconductor demand. The embodied energy of manufacturing millions of storage drives and ASICs is substantial.
- **E-Waste Legacy:** The accelerated obsolescence of SSDs during Chia's launch is a cautionary tale. While HDD farming is sustainable, plotting still favors SSDs/NVMe for speed. Responsible recycling is essential but often lacking.
- **The "Time" Resource:** VDFs uniquely tie energy consumption directly to *time*:  $\text{Energy} \approx T \times \text{Power}$ . This creates a linear, predictable relationship unlike PoW's difficulty-driven arms race. Protocols can tune  $T$  for optimal energy-security trade-offs.
- **Mitigation Strategies:** Sustainable VDF deployment requires:

- **Hardware Longevity:** Designing ASICs/FPGAs for protocol stability (e.g., supporting multiple modulus sizes) to avoid obsolescence.
- **Renewable Integration:** Locating VDF computation in regions with surplus renewable energy.
- **Storage Efficiency:** Developing less write-intensive plotting algorithms for PoST.
- **Circular Economy:** Mandating recycling programs for VDF hardware and storage drives within blockchain ecosystems.

VDFs offer a path towards dramatically more sustainable blockchain infrastructure, but they are not zero-impact. A holistic view encompassing operational energy, manufacturing, and e-waste is essential to avoid greenwashing and ensure genuine environmental responsibility.

### 1.8.3 8.3 Economic Models and Incentives

VDFs introduce “provable sequential time” as a new economic primitive, spawning novel incentive structures, markets, and financial instruments:

- **Tokenomics of VDF-Based Systems:** How value flows in ecosystems using VDFs:
- **PoST Block Rewards (Chia):** Farmers (storage providers) earn block rewards for providing storage *and* contributing to the VDF input. Timelords (VDF evaluators) are essential but often uncompensated directly in Chia’s model, leading to concerns about under-provisioning. Solutions include:
- **Implicit Rewards:** Timelords win blocks they finalize, incentivizing participation.
- **Protocol Fees:** Dedicate a portion of transaction fees to active Timelords.
- **Staking Requirements:** Require Timelords to stake tokens, earning rewards proportional to uptime and correct computation.
- **VDF Services in PoS (e.g., ETH2):** If VDFs are used for leader election or randomness, specialized “VDF Provers” could emerge as a service. Validators might outsource VDF computation to these provers for a fee, creating a market. Protocols could mandate staking for provers to ensure accountability.
- **MEV Mitigation Markets:** In VDF-enforced fair ordering protocols (Section 6.4), entities bidding for the right to propose the initial transaction order (“Orderers”) might pay fees to VDF provers who enforce the delay, creating a layered market structure.
- **The Computation Marketplace:** A potential future development is a decentralized marketplace for VDF computation:

- **Provers and Verifiers:** Users needing VDF outputs (e.g., for time-locked encryption, secure auctions) pay “Provers” (hardware operators) to compute  $(y, \pi)$ . Smart contracts could escrow payments and release funds upon successful verification.
- **Reputation Systems:** Provers build reputation based on speed, reliability, and correctness. Higher reputation commands premium fees.
- **Bundling and Batching:** Provers could batch multiple VDF requests, amortizing hardware costs. Protocols like TrueBit or DECO could facilitate trustless off-chain computation with on-chain verification.
- **Example - Chainlink VDF Service:** Oracle networks could integrate VDFs, offering verifiable delay as an on-demand service for smart contracts needing timed releases or unbiased randomness.
- **Novel Financial Primitives:** VDFs enable financial instruments anchored to verifiable time:
- **Time-Locked Vesting:** Tokens or assets automatically unlock after a VDF-proven delay, replacing centralized vesting contracts. DAOs could use this for contributor compensation.
- **Verifiable Delayed Options:** Financial derivatives that can only be exercised after a publicly verifiable delay, preventing front-running.
- **Decentralized Lotteries:** Participants commit funds. A winner is selected via VDF-based randomness *after* a mandatory delay, preventing last-second manipulation. Fees fund the VDF provers.
- **Anti-Sybil Bonds:** Users post collateral locked by a VDF. Spamming requires forfeiting the bond *and* waiting the delay period, raising the attack cost significantly.
- **Cost of Attack Economics:** VDFs alter the security economics of blockchains:
- **Randomness Beacons:** Biasing a RANDAO+VDF beacon requires influencing the final revealer *and* controlling enough parallel VDF capacity to explore a significant fraction of candidate seeds within  $T$ . The cost scales with the number of parallel instances needed (hardware cost + energy) and the value at stake. A sufficiently large  $T$  makes this economically irrational.
- **PoS Security:** Attacking a VDF-enhanced PoS chain requires not just acquiring stake ( $\geq 33\%$  for liveness attacks) but also controlling enough VDF hardware to produce blocks faster than the honest network across multiple forks (Nothing at Stake mitigation). The hardware cost adds a significant physical barrier beyond pure capital.
- **PoST Security:** Dominating Chia’s consensus requires controlling  $>50\%$  of netSPACE *and* sufficient VDF capacity to outpace honest Timelords. The storage cost dominates, but VDF capacity ensures attacks can’t be executed instantly.

VDFs transform time from an abstract concept into a quantifiable, tradable resource with inherent economic value. This creates opportunities for innovative markets and financial products while adding robust physical

cost layers to blockchain security models. However, it also risks commodifying access to “trustworthy time,” potentially excluding those without capital for specialized hardware.

#### 1.8.4 8.4 Accessibility, Open Source, and Geopolitics

The global deployment of VDFs intersects with issues of equitable access, the critical role of open collaboration, and the geopolitical struggle for technological control:

- **Open Source as a Trust Anchor:** Given VDFs’ role in critical infrastructure (e.g., Ethereum’s beacon chain, Chia’s consensus), transparency is paramount:
- **Auditability:** Open-source implementations (Chia VDF, Filecoin rust-fil-proofs, Ethereum Research VDF) allow independent security audits, fostering trust. Closed-source VDF hardware or software would be anathema to decentralization ideals and a security risk.
- **Mitigating Centralization:** Open-source hardware designs (like EPFL’s PipeZK from the Ethereum competition) lower barriers to ASIC manufacturing, preventing monopolies. The **IETF VDF standardization effort** relies on open specifications and public scrutiny.
- **Reproducible Builds:** Ensuring that deployed binaries match open-source code is crucial for protocols like Filecoin and Ethereum, where VDF correctness underpins consensus safety. This combats backdoors and supply-chain attacks.
- **Accessibility and the Digital Divide:** The hardware demands of VDFs risk exacerbating existing inequalities:
- **Global Participation Barriers:** Individuals in regions with limited capital, unreliable electricity, or restricted semiconductor imports cannot participate as VDF provers or competitive Timelords. This biases influence towards wealthy nations and corporations, undermining the global decentralization vision.
- **Verification Asymmetry:** While VDF *evaluation* requires hardware, *verification* (especially Wesolowski) is lightweight. This allows anyone to *use* VDF outputs (e.g., verify beacon randomness) without expensive hardware, preserving some access. However, influence over *generation* remains concentrated.
- **Protocol Design Choices:** Systems prioritizing decentralization must favor constructions with lower hardware barriers (e.g., class groups with CPU-friendly parameters) or explicitly subsidize distributed participation.
- **Geopolitical Dimensions:** VDF technology intersects with state power and control:
- **Resource Nationalism:** Countries rich in renewable energy (Iceland, Paraguay) or semiconductor manufacturing (Taiwan, South Korea) could become VDF computation hubs, wielding influence over decentralized networks. This mirrors Bitcoin mining geopolitics but on a smaller scale.

- **Regulatory Crackdowns:** Governments could ban VDF ASIC imports/operation (like China’s Bitcoin mining ban), restrict access to class group parameter generation tools, or mandate backdoors in hardware. The “**Nothing Up My Sleeve**” property of class groups provides some resistance, as parameters are generated transparently.
- **Censorship Resistance vs. State Control:** VDFs enhance censorship resistance in applications like unbiased randomness for voting or timed document releases. Conversely, states might deploy VDFs in centralized systems for “sovereign” timekeeping or controlled information release, leveraging the technology for authority rather than decentralization. Iran’s exploration of national blockchain infrastructure highlights this dual-use potential.
- **Sanctions Evasion?:** The ability to prove computation time without trusted parties could hypothetically be misused in covert communication or coordination, though no significant cases are known.
- **The Open-Source Imperative:** Countering centralization and geopolitical risks hinges on robust open ecosystems:
- **Public Goods Funding:** Mechanisms like Ethereum’s Protocol Guild or Gitcoin Grants are crucial for funding ongoing development and security audits of open-source VDF software and hardware designs.
- **Decentralized Manufacturing:** Initiatives like Open Compute Project (OCP) adapting to crypto hardware could foster geographically distributed, open ASIC production.
- **Knowledge Sharing:** Academic collaboration and open publications (e.g., via IACR) remain vital for advancing VDF security and efficiency globally.

The societal promise of VDFs—decentralized, trustless coordination—can only be realized through vigilant commitment to openness, equitable access, and resistance to centralized control, whether corporate or state-sponsored. The technology itself is neutral; its impact depends on how humanity chooses to deploy it.

The societal and economic ripples from VDF deployment underscore that cryptography is never merely technical. It reshapes power, resources, and access on a global scale. While offering solutions to decentralization’s trust problems, VDFs introduce new challenges around hardware equity, environmental trade-offs, and economic concentration. Navigating these complexities requires ongoing research and innovation. This leads us to the **Frontiers of Research and Open Problems**, where cryptographers strive to build VDFs that are more secure, efficient, decentralized, and resilient against future threats like quantum computers.

---

## 1.9 Section 9: Frontiers of Research and Open Problems

The societal and economic implications explored in Section 8 reveal a profound truth: Verifiable Delay Functions have evolved from cryptographic curiosities into infrastructure with planetary-scale consequences.



Yet this maturation reveals new frontiers where fundamental limitations persist and paradigm-shifting innovations beckon. **The quest for trust in decentralized time now confronts the quantum computing revolution, adapts to evolving threat models, and reimagines sequential computation itself through distributed protocols and novel mathematical structures.** This section charts the bleeding edge of VDF research—where cryptographers battle quantum adversaries, dismantle trusted setups through multi-party computation, and wrestle with the stubborn incompatibility of parallelism and sequentiality. From isogeny labyrinths to massively parallel proof systems, the race to perfect verifiable delay continues to redefine the boundaries of decentralized trust.

The urgency of these challenges is magnified by real-world deployments. Ethereum’s beacon chain postponement of VDF integration underscores the fragility of RSA-based schemes in a quantum-threat landscape. Chia’s class groups, while trustless, face hardware centralization pressures that demand distributed solutions. Filecoin’s VDF-accelerated leader election highlights the need for adaptive security as adversaries evolve. As VDFs become embedded in billion-dollar ecosystems, theoretical breakthroughs transform into critical infrastructure upgrades. This section illuminates five seismic shifts defining VDFs’ next decade: the quantum migration, adaptive and continuous time models, distributed computation frontiers, proof system revolutions, and enduring mathematical enigmas.

### 1.9.1 9.1 Post-Quantum Secure VDFs

The sword of Damocles hanging over modern cryptography—Shor’s algorithm—renders current VDF constructions obsolete in a quantum future. RSA and class groups rely on integer factorization and discrete logarithms, problems quantum computers solve in polynomial time. **When sufficiently large quantum processors emerge, attackers could compute  $g^{(2^T)}$  in seconds by recovering group orders via Shor, nullifying sequentiality guarantees.** Google’s 2019 quantum supremacy demonstration (Sycamore) and IBM’s 2023 433-qubit Osprey chip signal rapid progress, making post-quantum VDFs an existential priority. Four approaches dominate research:

- **Isogeny-Based VDFs: Navigating Elliptic Curve Labyrinths**

Supersingular isogeny Diffie-Hellman (SIDH) and commutative variants like CSIDH leverage the computational hardness of navigating isogeny graphs—maps between elliptic curves. The core intuition: computing an isogeny chain  $\phi_1 \circ \phi_2 \circ \dots \circ \phi_T$  requires sequential curve operations resistant to quantum speedups.

- **CSIDH (Commutative SIDH):** Castryck, Lange, Martindale, Panny, and Renes’ 2018 construction enables public-key operations with inherent sequentiality. **VDF potential:** Repeated application of class group actions forces sequential walks through isogeny space. Challenges include massive parameter sizes (~10KB public keys) and slow evaluation (~100ms/step vs. RSA’s ns/step). The 2022 [SQIsign](#) breakthrough by De Feo, Kohel, Leroux, Petit, and Wesolansky improved signing efficiency but VDF adaptations remain theoretical.

- **Obstacles:** Trusted setup requirements for supersingular curves, vulnerability to [subexponential classical attacks](#), and lack of efficient proof systems compatible with Wesolowski/Pietrzak verification. The [2023 attack](#) on SIDH further underscores fragility.
- **Lattice-Based VDFs: The Shortest Path to Sequentiality**

Leveraging the hardness of Shortest Vector Problem (SVP) or Learning With Errors (LWE), lattice schemes promise quantum resistance and versatility. Pietrzak’s 2018 proposal for “Simple VDFs” using incrementally verifiable computation (IVC) with lattice-based SNARKs sparked interest, but bottlenecks persist:

- **Proof Size Explosion:** STARKs for lattice VDFs generate proofs scaling linearly with  $T$  (e.g., 1GB for  $T=2^{30}$ ), negating efficient verification. Döttling et al.’s [2019 construction](#) reduced assumptions but retained impractical overheads.
- **Sequentiality Gap:** Lattice reduction exhibits *some* parallelism (e.g., BKZ algorithm blocks), violating strict sequentiality. Alwen et al.’s [2021 work](#) on “memory-hard” lattices aims to enforce sequentiality via large-state computations but increases hardware costs.
- **Hash-Based VDFs: Depth-Robust Graphs to the Rescue?**

Sloth (Lenstra and Wesolowski, 2015) demonstrated weak sequentiality via repeated modular square roots, but parallel attacks limited security. Current research focuses on **depth-robust graphs**—combinatorial structures where any shortcut requires massive memory.

- **Balloon Hashing & Beyond:** Alwen and Serbinenko’s 2015 STOC work formalized graph-based sequential functions. [Bünz et al. \(2020\)](#) adapted this to VDFs using Merkle tree accumulation over depth-robust graphs.
- **Limitations:** Proof sizes remain large ( $O(T)$  without SNARKs), and quantum Grover searches could halve sequential time. The [MinRoot VDF](#) by Boneh et al. improved efficiency but still lacks robust proofs against specialized hardware.
- **The Verdict:** No post-quantum VDF candidate matches RSA/class group efficiency. Isogenies lead for plausibly quantum-sequential operations, lattices offer versatile frameworks with SNARK compatibility, and hash-based approaches provide minimal assumptions at high overhead. Ethereum’s PQC-VDF working group prioritizes isogeny-lattice hybrids, while NIST’s PQC project influences standardization.

## 1.9.2 9.2 Continuous VDFs and Adaptive Security

Current VDFs suffer from rigidity: the delay parameter  $T$  is fixed at setup. Adversaries who observe public parameters  $pp$  can precompute attacks tailored to  $T$ . **Continuous VDFs (cVDFs)** solve this by enabling evaluation over arbitrary intervals, while **adaptive security** prevents adversaries from choosing  $T$  maliciously after seeing  $pp$ .

- **The cVDF Vision:** A function where:
  - $\text{Eval}(\text{pp}, x, T)$  outputs  $y_T$  and proof  $\pi_T$  for any  $T > 0$
  - Computing  $y_T$  requires  $\approx T$  sequential steps from scratch
  - Computing  $y_{T+k}$  from  $y_T$  requires  $\approx k$  steps (incremental efficiency)
  - Verification of  $y_T$  remains efficient.
- **Döttling et al.’s Breakthrough (CRYPTO 2020):** First formalized cVDFs using:
  1. A **sequential function**  $F$  (e.g., repeated squaring)
  2. A **verifiable computation scheme** (SNARK) proving  $y_T = F^{\{T\}}(x)$
- **Trade-offs:** SNARKs introduce trusted setups (for SNARKs), large proving times, and recursive proof composition overhead. Their construction remained theoretical, with  $T=2^{30}$  requiring hours of proving on a supercomputer.
- **Adaptive Root Assumption Limitations:** Wesolowski’s uniqueness relies on the Adaptive Root Assumption (ARA), which assumes adversaries cannot find  $\ell$ -th roots for adaptively chosen primes  $\ell$ . Continuous VDFs require security against adversaries choosing  $T$  (and thus implicitly  $\ell$ ) *after* seeing  $\text{pp}$ . Current ARA formulations don’t guarantee this, leaving a security gap.
- **Practical Progress:** [Ephemeral VDFs](#) (Gorbunov et al., 2023) enable “on-demand”  $T$  selection via cryptographic commitments, but require new hardness assumptions. Class groups offer hope—their transparent setup resists parameter-targeting attacks better than RSA.

### 1.9.3 9.3 Distributed VDFs and MPC Protocols

VDFs’ sequential core resists parallelization, but **distributed computation** could enhance security and decentralization. By splitting computation across parties, we mitigate hardware centralization risks and eliminate trusted setups.

- **Threshold VDFs: Sharing the Sequential Burden**

Boneh et al.’s [2019 proposal](#) introduced threshold VDFs for class groups:

- $n$  parties hold shares of the initial state  $g$
- Each sequentially computes their share of  $g^{(2^T)}$  via MPC-friendly operations
- Reconstruction requires  $t+1$  parties to combine shares into  $y$

- **Innovation:** Parties compute *independently* after setup, avoiding interactive MPC during evaluation. The 2022 [Fiat-Shamir VDF](#) by Ganesh et al. reduced rounds using non-interactive proofs.
- **MPC for Trusted Setup Perfection:** RSA modulus generation via MPC (Ethereum’s RSA-2048 ceremony) remains vulnerable to covert attacks if  $>t$  parties collude. Cutting-edge research focuses on:
- **One-Round MPC:** [Lin17](#) and [GG18](#) protocols reduced interaction but require reliable broadcast. [Groth’s 2023 SIMD-based MPC](#) allows single-round modulus generation with abort security.
- **Transparent Setup via Class Groups:** Eliminates MPC for setup but shifts burden to distributed VDF evaluation.
- **The Parallelism Paradox:** Can sequential computation be *securely distributed*? While parties compute independently in threshold schemes, the computation itself remains sequential per party. True parallel evaluation is impossible without breaking sequentiality. Current solutions trade decentralization for efficiency: each party computes a full VDF, but output requires threshold reconstruction. The quest continues for protocols where  $n$  parties compute in  $T/n$  time—a theoretical impossibility for strictly sequential functions.

#### 1.9.4 9.4 Improved Constructions and Proof Systems

Beyond quantum and distribution, VDFs face efficiency frontiers. Wesolowski’s fast verification battles slow proof generation; Pietrzak’s simplicity fights logarithmic proof sizes. Next-generation constructions blend algebraic innovation with proof system synergies.

- **SNARK/STARK-VDF Hybrids: Zero-Knowledge Time Proofs**

Combining VDFs with succinct arguments creates “verifiable verifiable delay functions”:

- **VeeDo (Ethereum Research):** Uses STARKs to prove correct Wesolowski VDF execution. Benefits: constant-time verification and post-quantum security (STARKs). Costs: 100x slower proof generation and 1MB+ proofs.
- **Nova (Spartan) + VDFs:** Kothapalli et al.’s [2022 Nova](#) enables incremental SNARKs. Applied to Pietrzak VDFs, it allows continuous proving with sublinear costs. Challenge: trusted setup for SNARK recursion.
- **Beyond GUOs: Hyperelliptic Curves and Beyond**

Groups of Unknown Order (GUOs) underpin most VDFs, but new structures emerge:

- **Jacobians of Hyperelliptic Curves:** Offer GUO properties with smaller parameters than class groups. [Castryck and Decru’s 2023 attack](#) on SIDH dampened enthusiasm, but Jacobians remain viable for VDFs.
- **Pairing-Free Isogenies:** SIKE’s cryptanalysis stalled isogeny-based crypto, but CSIDH derivatives like [CSURF](#) provide efficient, sequential isogeny walks without pairings.
- **Endomorphism Rings:** [SeaSign](#) constructs leverage endomorphism ring computations, conjectured to require exponential time. VDF adaptations remain untapped.
- **Proof Generation Revolution:**
- **Wesolowski Optimization:** [Shamir’s trick](#) for simultaneous exponentiation accelerates  $\pi = (y * g^{\{-r\}})^{\{1/\ell\}}$  computation. Class group implementations (Chia) use optimized Tonelli-Shanks variants.
- **Pietrzak Parallelism:** [Döttling et al.’s 2020 work](#) parallelizes proof generation across subtree ranges, reducing wall-clock time for large  $T$ .

### 1.9.5 9.5 Major Unsolved Problems

Despite astonishing progress, foundational challenges endure:

1. **Falsifiable Assumptions:** All efficient VDFs rely on non-falsifiable assumptions (e.g., Adaptive Root Assumption). “*Prove that breaking VDF sequentiality implies factoring RSA integers*” remains elusive. Constructing VDFs from falsifiable assumptions (e.g., LWE hardness) would revolutionize security arguments.
2. **Perfect Uniqueness:** Current VDFs offer computational uniqueness—finding two valid outputs is hard but not impossible. Information-theoretic uniqueness (where only one  $y$  exists per  $x$ ) may require radical paradigms, like embedding VDFs in error-correcting code lattices.
3. **The Parallelism Gap:** In practice, parallel attacks *do* reduce effective delay. For class groups, 1000 GPUs might cut  $T$  by 10% via simultaneous ideal reductions. Minimizing this gap—ideally to near-zero—demands algebraic structures with inherent sequential bottlenecks. The 2023 [Gentry competition](#) seeks such “parallelism-resistant” functions.
4. **Optimality Proofs:** What is the minimal computational overhead for a VDF? Pietrzak’s scheme requires  $2T$  squarings for evaluation and proof; Wesolowski needs  $T + O(\sqrt{\ell})$ . Theoretical lower bounds are unknown. A grand challenge: “*Build a VDF where evaluation is  $(1+\varepsilon)T$  steps and verification is  $O(1)$ , for  $\varepsilon \rightarrow 0$ .*”
5. **Energy-Latency Equivalence:** Can VDFs be constructed where energy consumption—not wall-clock time—is the sequential resource? Physical proposals using optical computing or superconducting loops exist but lack cryptographic formalization.

---

These frontiers represent not merely academic puzzles but the next evolutionary leap for decentralized systems. Solving them would forge VDFs immune to quantum collapse, adaptable to dynamic environments, distributable across global networks, and verifiable with near-zero overhead. As researchers navigate isogeny mazes, refine multi-party protocols, and battle the parallelism paradox, they lay foundations for a future where trust in time transcends centralized authorities and even the limits of classical computation. Yet the ultimate trajectory of VDFs extends beyond cryptography into the realms of philosophy, governance, and human coordination—a synthesis we explore in our concluding section.

---

## 1.10 Section 10: Synthesis and Future Trajectories

The journey through Verifiable Delay Functions—from their conceptual origins in Rivest’s time-lock puzzles to their implementation as specialized ASICs enforcing blockchain consensus—reveals a profound evolution. What began as a cryptographic curiosity addressing the “tyranny of parallelism” has matured into a foundational primitive reshaping decentralized systems. As we stand at this inflection point, VDFs represent more than just algorithms; they embody a paradigm shift in how humanity coordinates trust across digital networks without central authorities. The societal implications explored in Section 8—centralization risks, environmental trade-offs, and economic transformations—underscore that VDFs have transcended theoretical computer science to become infrastructure with planetary consequences. Yet, as research frontiers like post-quantum security and distributed VDFs advance (Section 9), their ultimate trajectory remains unwritten. This synthesis examines VDFs as both technological artifacts and social catalysts while charting their potential to underpin humanity’s next era of digital collaboration.

### 1.10.1 10.1 VDFs as Foundational Cryptographic Primitives

VDFs occupy a unique niche in the cryptographic pantheon, distinguished by their ability to *temporalize trust*. Unlike traditional primitives that secure *data* (encryption) or *identity* (signatures), VDFs secure *process*—specifically, the irreversible passage of sequential computation time. This capability fills a critical gap in decentralized systems:

- **Complementary to Zero-Knowledge Proofs (ZKPs):** While ZKPs (like zk-SNARKs) verify computational integrity *without revealing inputs*, VDFs enforce computational *duration* without shortcuts. Ethereum’s planned integration of VDFs with ZKPs illustrates their synergy: VDFs provide unbiased randomness for ZKP-based rollups (e.g., StarkNet), while ZKPs could verify VDF correctness in constant time. The Mina Protocol’s recursive zk-SNARKs combined with VDF timestamps exemplify this convergence.

- **Beyond Proof-of-Work:** VDFs solve PoW’s fatal flaw—parallelizability—by replacing *work* with *sequential time*. This transforms the security model: whereas Bitcoin miners compete via energy expenditure, Chia’s Timelords cooperate by *waiting* through mandatory VDF delays. The shift from “proof-of-burned-energy” to “proof-of-elapsed-time” represents a cryptographic evolution as significant as the move from symmetric to asymmetric encryption.
- **The Time-Authority Paradigm:** Historically, societies relied on centralized timekeepers (from sundials to NTP servers). VDFs enable *decentralized time authorities*—network participants collectively verifying elapsed intervals without trusted third parties. Filecoin’s leader election VDFs or Ethereum’s RANDAO+VDF beacon function as algorithmic sundials, their shadows lengthening at a rate governed by modular squarings rather than Earth’s rotation.

The 2018 Boneh et al. formalization crystallized VDFs as a cryptographic class alongside commitments and oblivious transfer. Their mathematical elegance—sequentiality enforced by groups of unknown order—belies revolutionary implications: **time itself becomes a verifiable public good.**

### 1.10.2 10.2 Impact Assessment: Promises Fulfilled and Challenges Ahead

#### Successes in Target Applications:

- **Randomness Beacons:** Ethereum’s beacon chain delay in deploying VDFs (due to hardware concerns) paradoxically proves their necessity. Without VDFs, RANDAO remains vulnerable to last-revealer bias—a flaw exploited in minor incidents like the 2020 **Medalla testnet incident**, where validators manipulated timing to influence outcomes. Projects like **Drand** (a production VDF-based beacon) have provided unbiased randomness for Filecoin and Polkadot since 2019, demonstrating robustness at scale.
- **Proof-of-Stake Security:** Filecoin’s leader election VDFs have mitigated “grinding attacks” since mainnet launch (2020), forcing validators to wait through mandatory delays before proposing blocks. The absence of catastrophic forks in Filecoin—compared to early PoS chains like Tezos—validates VDFs as a “rate-limiting” solution to Nothing-at-Stake vulnerabilities.
- **Sustainable Consensus:** Chia Network’s “Proofs of Space-Time” has operated since 2021 with an estimated 0.36 TWh/year energy consumption—0.2% of Bitcoin’s footprint. While criticized for SSD wear during plotting, its VDF-enforced storage proofs have created the first viable large-scale alternative to energy-intensive mining.

#### Persistent Challenges:

- **Hardware Centralization:** The **Supranational ASIC project** highlighted the risk: entities controlling custom VDF hardware could dominate Ethereum’s beacon chain. Open-source designs like



**EPFL’s PipeZK** mitigate but don’t eliminate this; true decentralization requires commoditized VDF chips accessible to hobbyists—a goal yet unrealized.

- **Quantum Vulnerability:** Shor’s algorithm threatens all GUO-based VDFs. Ethereum’s cautious VDF rollout reflects this existential risk. While lattice-based constructions exist theoretically (e.g., Pietrzak’s IVC proposal), their verification overhead (1GB proofs for  $T=2^{30}$ ) renders them impractical. The 2023 **NIST PQC finalist BIKE** offers hope for hybrid approaches, but production-ready post-quantum VDFs remain years away.
- **Economic Barriers:** Running competitive Timelords in Chia requires ~\$10,000 FPGA setups—far costlier than staking in PoS chains. This creates a participation asymmetry contradicting Web3’s egalitarian ideals.

### Lessons Learned:

1. **Trusted setups are toxic:** Ethereum’s RSA-2048 MPC ceremony (2020) was a logistical marvel but proved so complex it delayed VDF adoption. Class groups’ transparent setup (Chia) emerged as the superior model.
2. **Hardware defines security:** The 2021 revelation that **custom FPGAs** could compute Chia VDFs 15x faster than CPUs forced recalibration of network parameters. Physical infrastructure, not mathematics, now dictates attack costs.
3. **Delay is contextual:** A “1-minute VDF” provides different security in Ethereum (high-value MEV) versus a timestamping service (low-value docs). Parameter tuning must reflect economic stakes.

### 1.10.3 10.3 The Broader Vision: VDFs in the Fabric of Future Systems

Beyond current blockchain applications, VDFs enable architectures where verifiable delay orchestrates human-machine collaboration:

- **Decentralized Identity & Reputation:** Imagine a “VDF-sealed” identity credential where users prove continuous engagement. A DAO could require members to accumulate VDF proofs over time, preventing sybil attacks more elegantly than token-weighted voting. Microsoft’s **ION DID** project explores similar time-based attestations.
- **Anti-Collusion Supply Chains:** In diamond certification, a VDF could enforce a mandatory “consideration period” between conflict-free audits and issuance, preventing last-minute bribery. De Beers’ Tracr blockchain may adopt such mechanics to harden ethical sourcing.
- **Democratic Governance:** Quadratic voting systems could integrate VDFs to prevent snapshot manipulation. By enforcing a fixed delay between vote commitment and tallying (as in **Aragon V2** prototypes), adversaries lose the ability to pivot based on early results.

- **Proofs of Useful Time:** Current VDFs “waste” computation on modular squaring. Future iterations could usefully sequence:
- Protein folding simulations (Folding@home + VDFs)
- Climate modeling increments
- **AI training checkpoints**

The 2023 **FoldingCoin** initiative demonstrates early steps—rewarding contributors whose Folding@home work is VDF-sequenced to prevent cheating.

- **Temporal Contracts:** Smart contracts with VDF-enforced delays could automate:
- Inheritance releases (funds unlock after 1 year of proven identity activity)
- Graduated intellectual property licensing
- **Dynamic carbon credits** that depreciate verifiably over time

The common thread: **VDFs as temporal glue**, binding digital actions to irreversible intervals in ways previously requiring notaries, escrow, or regulatory oversight.

#### 1.10.4 10.4 Ethical Considerations and Responsible Development

As VDFs mature, ethical imperatives emerge:

- **Equitable Access:** The **Open Compute Project’s Crypto Working Group** must prioritize low-cost VDF hardware designs. Protocols should incentivize distributed Timelord networks—perhaps via “VDF staking pools” where small holders delegate resources, akin to Rocket Pool for Ethereum validators.
- **Environmental Stewardship:** While PoST reduces energy use, the e-waste from specialized hardware remains problematic. A **VDF Hardware Sustainability Standard** should mandate:
- 5-year minimum lifespans for ASICs
- Modular upgradability (e.g., socketed RSA modulus chips)
- Takeback programs like Fairphone’s, funded by protocol fees

Chia’s transition from SSD-plotting to HDDs in 2021 averted a e-waste crisis, setting a precedent.

- **Transparency vs. Obscurity:** Closed-source VDF hardware risks backdoors. The **Linux Foundation’s CHIPS Alliance** should host open-source VDF implementations (RISC-V cores for modular reduction). Regulatory attempts to mandate “lawful access” (e.g., the EU’s Chat Control proposal) must exempt VDFs, as backdoors would nullify uniqueness guarantees.
- **Geopolitical Equity:** VDF computation hubs must avoid Bitcoin mining’s concentration in authoritarian states. The **Internet Archive’s decentralized Timelord initiative**—spreading nodes across libraries globally—offers a model for anti-fragile, jurisdictionally diverse deployment.

### 1.10.5 10.5 Concluding Thoughts: The Enduring Quest for Trust in Time

Human civilization has always sought objective timekeeping—from Babylonian water clocks to Harrison’s marine chronometers. VDFs represent the digital era’s solution to this ancient problem: **cryptographic clocks** whose ticks are verifiable squarings in class groups rather than pendulum swings. Their invention responds to a fundamental need—coordinating trust across adversarial networks where physical time cannot be observed directly.

The journey from Rivest’s 1996 time-lock puzzle to Chia’s live Timelord network embodies cryptography’s evolution from academic exercise to societal infrastructure. Like public-key encryption in the 1990s, VDFs face skepticism today (“Why not just use NTP?”). Yet their value shines where central timekeepers fail: mitigating MEV extraction, securing \$40B+ in staked Ethereum, or enabling Filecoin’s 19-exabyte decentralized storage network.

Challenges remain—quantum vulnerability looms, hardware centralization threatens, and energy debates persist. But the trajectory is clear. As research advances in:

- **Threshold VDFs** (Gorbunov et al.)
- **Continuous VDFs** (Döttling’s cVDFs)
- **SNARK-wrapped VDFs** (VeeDo)

the primitive grows more resilient.

In 2164, when Encyclopedia Galactica’s physical edition is opened after a 200-year VDF-sealed time capsule, readers may marvel that humanity’s first steps toward cosmic coordination began with modular exponentiation in groups of unknown order. Verifiable Delay Functions, once a cryptographic curiosity, will have become the silent heartbeat of decentralized civilization—proof that in the digital realm, time itself can be forged into a tool for collective trust.

The quest continues.